


OFFERED BY KLAYM AI

The AI SEO Brain



10+ years of SEO packed into one complete Notion system

The strategy, execution systems, templates, workflows, prompts, audits, and playbooks required to run modern SEO in the age of AI.



THE SEO OPERATING SYSTEM

10 operating areas · 15 databases · 25 prompts · 16 SOPs



DASHBOARDS · SOPs · PROMPT LIBRARY · SCORECARDS · WEEKLY RHYTHM

Become the Brand AI Recommends

Join the waitlist at <https://klaym.ai>

SEO is not dead. The interface changed.

Search was never really about ten blue links. It was about a person with a problem trying to find someone who could solve it. That has not changed. What changed is the surface the answer arrives on — and how much of the decision is made before anyone visits a website.

Your buyer now searches across Google, AI Overviews, Google AI Mode, ChatGPT, Claude, Perplexity, Gemini, Reddit, YouTube, LinkedIn, review sites, marketplaces, directories, and private communities. Often several of these in one afternoon, for one decision.

The fundamentals still hold. Relevance, authority, trust, technical accessibility, and content that actually answers the question are as load-bearing as they have ever been. What sits on top of them is new: discoverability across surfaces you do not control, extractability so a system can quote you safely, entity clarity so a model knows who you are, structured facts that stay consistent everywhere you appear, and citation worthiness — the thing that decides whether you are named or merely summarized.

Modern SEO optimizes for two readers at once: a person deciding whether to trust you, and a machine deciding whether to repeat you.

WHERE BUYERS ACTUALLY SEARCH NOW

CLASSIC SEARCH	AI ANSWER SURFACES	COMMUNITY & SOCIAL	MARKETPLACES / DIRECTORIES
- Google web results - Bing - Image and video search - News and Discover	- Google AI Overviews - Google AI Mode - ChatGPT - Claude - Perplexity - Gemini	- Reddit - YouTube - LinkedIn - Forums and Slack - Review sites	- App and software directories - Industry associations - Comparison sites - Maps and local listings

One brand, four surface families, one underlying requirement: be findable, understandable, and credible on each.

WHAT THIS PLAYBOOK WILL NOT PROMISE

No action described here guarantees a ranking, a click, a citation, an AI recommendation, traffic, or revenue. Search systems and AI answer engines change their behaviour without notice and do not publish their weighting. What this playbook gives you is a system that improves the probability of being found, understood, trusted, cited, and recommended — and, just as usefully, a way to tell whether it is working.

What is inside the Brain

Thirty-three sections: ten operating areas, thirty-two database specifications, eighty-seven prompts, forty-two SOPs, fifty-five templates, a worked example, ten workflows, and the build instructions.

32

DATABASES

87

CLAUDE PROMPTS

42

SOPS

55

TEMPLATES

00	SEO Is Not Dead. The Interface Changed.	2
01	What This Is: The Complete Blueprint	5
02	The Workspace Architecture	7
03	The Notion Database Standard	17
04	SEO Foundations	18
05	Market, Audience, and Search Demand	22
06	Keyword Research	26
07	Content Strategy	30
08	On-Page SEO and Answer Readiness	33
09	Technical SEO	35
10	Internal Linking and Information Architecture	37
11	Link Building, Digital PR, and Authority	39
12	Local SEO	42
13	LLM and AI SEO	45
14	The SEO Brain Command Center	50
15	The Twenty Technical Modules	58
16	The AI Visibility Operating Map	60
17	How the System Changes by Business Type	62
18	The Thirty-Two Databases	63
19	The Claude Prompt Library	114
20	The SOP Library	147
21	The Template Library	161
22	Worked Example: Northstar Analytics	164
23	Ten Cross-Functional Workflows	169
24	How to Build It in Notion	173
25	The Agency and Multi-Project Model	176
26	The SEO Operating Dashboard	177
27	The Weekly SEO Operating Rhythm	181
28	Prompt Library: The Core Twenty-Five	183
29	SOP Library: The Core Sixteen	208

30	SEO Scorecards	216
31	The 30-Day Implementation Roadmap	219
32	Final Readiness Checklist	221
33	Build Your SEO Brain	222

HOW TO BUILD THIS IN NOTION

Each of the ten operating areas becomes a page. Each database described in Section 11 becomes a Notion database with the properties listed there. The SOPs in Section 14 become a single SOP database with one entry per procedure, and the prompt library in Section 13 becomes a prompt database you filter by task. The dashboard in Section 11 is the home page that links them together. Nothing here requires a paid tier or a template purchase.

SECTION 01 · WHAT THIS IS

The complete blueprint

This is not a guide to building an SEO workspace. It is the workspace itself, written out — every page, database, property, relation, view, formula, template, SOP, prompt, and example required to reproduce it.

This playbook contains the complete architecture and content of The AI SEO Brain. Each folder below represents a page or operating area in the Notion workspace. Each database includes its properties, relationships, views, formulas, templates, workflows, and example records.

WHAT YOU ARE RECEIVING, STATED HONESTLY

This is a PDF blueprint and operating manual — a complete, implementation-ready specification you rebuild in your own Notion workspace. It is not a live Notion template and no duplication link is included. Nothing here connects to Google Search Console, your crawler, or any AI engine automatically; every integration described is something you would build or operate manually. Rebuild time is roughly six to ten hours for the full system, or about ninety minutes for the minimum viable version described in Section 18.

The asset inventory

Everything in this document, counted honestly. No asset appears twice under a different name.

ASSET CATEGORY	INCLUDED	WHERE	WHAT IT LETS YOU DO
Workspace pages and folders	17 top-level + 8 operating sub-pages = 25	Section 02	Reproduce the full Brain structure in Notion
Notion databases fully specified	32	Sections 04 and 05	Build the data layer with correct types and relations
Database views specified	96 (three or more per database)	Inside each database spec	See only what a given decision needs
Formulas and rollups	38 (24 formulas, 14 rollups)	Inside each database spec	Calculate scores and coverage without manual work
Claude prompts	87	Section 12	Run every recurring task at a consistent standard
Standard operating procedures	42	Section 13	Execute the same way regardless of who is doing it
Templates, briefs, and trackers	52	Section 14	Never start a document from a blank page
Scorecards	9	Section 15	Diagnose where to look next, consistently
End-to-end workflows	10	Section 16	See how the parts connect across operating areas
Populated fictional examples	60+ across every important database	Throughout, labelled as fictional	See what a completed record actually looks like
Dashboard blocks specified	6 blocks, 22 view configurations	Section 03	Open one page and know what needs attention
Vertical adaptations	5 business types	Section 17	Adapt fields and priorities to your model

What changed in this edition

The first edition was a framework. This one is the workspace, written out.

AREA	FIRST EDITION	THIS EDITION
Workspace pages	Described in prose	17 pages, each with owner, inputs, outputs, views, and a definition of done

AREA	FIRST EDITION	THIS EDITION
Databases	15, with property lists	32, with types, defaults, required flags, relations, rollups, formulas, views, workflows, and examples
Prompts	25	87, each with when-not-to-use, failure conditions, and the next prompt in sequence
SOPs	16, with steps and gates	42, each with tools, preconditions, failure conditions, handoff, and an example run
Templates	Referenced inside sections	55, each with fields, completion criterion, and destination database
Command Center	Eleven views listed	6 blocks, 28 view configurations with filters, sorts, and the decision each supports
Technical SEO	A checklist	20 modules with evidence requirements and classic failure modes
AI visibility	A method and five measures	35 components, 10 measures, and what each cannot tell you
Worked examples	None	A 16-step end-to-end walkthrough plus example records throughout, all fictional
Workflows	Implied	10 cross-functional workflows with owners, gates, and learning captured
Build guidance	A 30-day plan	A 10-step build order, a database dependency chain, and 3 deployment modes
Multi-client	Not covered	A full model with separation rules and leakage prevention

WHAT IS DELIBERATELY ABSENT

No guaranteed rankings, traffic, citations, or AI recommendations. No invented statistics, benchmarks, or research findings. No claim that any scoring model is scientifically validated — every score here is judgement made consistent and reviewable. No manipulative link tactics. Every company, metric, and record shown as an example is fictional and labelled as such.

SECTION 02 · ARCHITECTURE

The workspace architecture

Seventeen top-level pages. The ten operating areas from the original framework are preserved and expanded; seven supporting areas hold the command center, the libraries, and the record of how the system learns.

The seventeen top-level pages

#	PAGE	PURPOSE	OWNER	REVIEW CADENCE
00	Command Center	The single page the team opens each morning; surfaces what needs attention now.	SEO lead	Daily glance, weekly working session
01	Strategy & Governance	Holds the strategy, the charter, ownership, constraints, and the decision log.	SEO lead	Quarterly, with a monthly check
02	Market & Audience Intelligence	The demand evidence layer — what customers actually say, in their words.	Content lead	Quarterly deep pass, monthly additions
03	Keyword Intelligence	One scored, mapped keyword list the whole team works from.	SEO specialist	Weekly status, quarterly re-scoring
04	Topic & Content System	The production and maintenance engine: clusters, briefs, pipeline, refresh queue.	Content lead	Weekly for status, monthly for decay
05	On-Page & Answer Readiness	Page-level quality: intent match, structure, extractability, citation readiness.	Editor	Per piece, plus a quarterly sample audit
06	Technical SEO	Twenty modules covering everything between a crawler and a rendered page.	Technical SEO owner	Monthly crawl, weekly backlog review
07	Internal Linking & Site Architecture	Where authority and readers are directed. The cheapest lever in the system.	SEO specialist	Monthly
08	Authority, Links & Digital PR	Earning references: assets, campaigns, outreach, reclamation, monitoring.	Marketing or PR owner	Weekly outreach review, monthly link review
09	Local SEO	Per-location consistency, profiles, reviews, citations, and local content.	Local SEO owner	Monthly per location
10	LLM & AI Visibility	The largest operating area: prompt universe, runs, mentions, citations, sources.	SEO lead	Monthly, identical wording every run
11	Measurement & Reporting	Where numbers become decisions rather than status updates.	SEO lead	Weekly and monthly
12	Prompt Library	Every Claude prompt in the system, filterable by task and category.	SEO lead	Quarterly review of what is actually being used
13	SOP Library	Every recurring procedure, with owners, triggers, gates, and escalation.	SEO lead	Quarterly
14	Templates	Fifty-two reusable documents: briefs, audits, records, reports, plans.	Whoever owns the task	Quarterly
15	Experiments & Learning	One variable at a time, with the decision rule written before the result.	SEO lead	Weekly check, evaluation on the committed date
16	Archive & Decision Log	What was decided, why, by whom, and what happened afterwards.	SEO lead	Monthly review, quarterly retrospective

HOW THIS MAPS TO THE ORIGINAL TEN AREAS

Operating areas 01 through 10 are the same ten from the original framework, renamed only where the expanded scope required it — Market & Audience Intelligence, Keyword Intelligence, Topic & Content System, On-Page & Answer Readiness, and LLM & AI Visibility. Areas 00 and 11 through 16 are the operating layer that makes the other ten repeatable: the command center, measurement, the three libraries, experiments, and the decision log.

Full page specifications

Each of the seventeen pages, with everything needed to build and operate it.

00 Command Center

WORKSPACE PAGE

Purpose	The single page the team opens each morning; surfaces what needs attention now.
When to use it	Daily, and at the start of every working session.
Owner	SEO lead
Inputs	Every database in the workspace, filtered
Outputs	Priorities, risks, decisions awaiting approval
Related databases	All 32 databases
Related SOPs	SOP-02, SOP-15
Related prompts	PR-01, PR-70
Dashboard views	Executive Snapshot · This Week · Opportunities · Risk Center · Performance · Decisions
Review cadence	Daily glance, weekly working session

QUALITY GATE

No block shows more than 10 rows; every row has an owner

DEFINITION OF DONE

A person can see what to do next in under 60 seconds

01 Strategy & Governance

WORKSPACE PAGE

Purpose	Holds the strategy, the charter, ownership, constraints, and the decision log.
When to use it	Quarterly planning, and any time a direction is questioned.
Owner	SEO lead
Inputs	Business goals, performance history, maturity assessment
Outputs	Versioned strategy, three bets, decisions
Related databases	Business Goals, SEO Strategy, Decisions & Change Log, Constraints
Related SOPs	SOP-01, SOP-30, SOP-32
Related prompts	PR-01 to PR-06
Dashboard views	Current strategy · Bets and signals · Open decisions · Constraint register
Review cadence	Quarterly, with a monthly check

QUALITY GATE

Every bet has a measurable signal and a named owner

DEFINITION OF DONE

A new hire could read one page and know what we are trying to win

02 Market & Audience Intelligence

WORKSPACE PAGE

Purpose	The demand evidence layer — what customers actually say, in their words.
When to use it	Before keyword or content work, and quarterly as language shifts.
Owner	Content lead
Inputs	Calls, reviews, tickets, interviews, community threads
Outputs	Problems, quotes, questions, demand records
Related databases	Customer Problems, Voice of Customer, Search Demand, Question Bank
Related SOPs	SOP-03, SOP-04, SOP-05
Related prompts	PR-07 to PR-14
Dashboard views	Validated problems · Question gaps · Emerging demand · Vocabulary gaps
Review cadence	Quarterly deep pass, monthly additions

QUALITY GATE

Every record is verbatim and carries a source and a date

DEFINITION OF DONE

Every topic in the pipeline traces to a customer statement

03 Keyword Intelligence

WORKSPACE PAGE

Purpose	One scored, mapped keyword list the whole team works from.
When to use it	Quarterly, and when entering a new topic or market.
Owner	SEO specialist
Inputs	Demand records, tool exports, current rankings
Outputs	Scored keywords mapped to owning pages
Related databases	Keyword Master, Query Patterns, Declined Keywords
Related SOPs	SOP-06, SOP-07, SOP-08, SOP-09
Related prompts	PR-15 to PR-24
Dashboard views	Build order · Unmapped · Cannibalization risks · Declined log
Review cadence	Weekly status, quarterly re-scoring

QUALITY GATE

No keyword maps to two pages; every score records its criteria

DEFINITION OF DONE

A writer can see what to write next without asking anyone

04 Topic & Content System

WORKSPACE PAGE

Purpose	The production and maintenance engine: clusters, briefs, pipeline, refresh queue.
When to use it	Continuously — this is where most weekly work happens.
Owner	Content lead
Inputs	Keywords, topics, demand evidence, performance data
Outputs	Briefs, drafts, published pages, refresh decisions
Related databases	Topics, Topic Clusters, Content, Content Briefs, Refresh Queue
Related SOPs	SOP-10 to SOP-14
Related prompts	PR-25 to PR-34
Dashboard views	Pipeline · Refresh queue · Decay watch · Cluster coverage
Review cadence	Weekly for status, monthly for decay

QUALITY GATE

Nothing is commissioned without a differentiated angle

DEFINITION OF DONE

Every published page has a purpose, an owner, and a next action

05 On-Page & Answer Readiness

WORKSPACE PAGE

Purpose	Page-level quality: intent match, structure, extractability, citation readiness.
When to use it	Before every publication, and on any page being refreshed.
Owner	Editor
Inputs	Drafts, briefs, on-page and extractability standards
Outputs	QA verdicts, rewrites, metadata
Related databases	Content, Content Briefs, Brand Facts
Related SOPs	SOP-13, SOP-15
Related prompts	PR-35 to PR-42
Dashboard views	Awaiting QA · Failed QA · Extractability gaps · Metadata queue
Review cadence	Per piece, plus a quarterly sample audit

QUALITY GATE

Unsourced claims are an automatic return, not a note for later

DEFINITION OF DONE

A system could quote any key passage without introducing an error

06 Technical SEO

WORKSPACE PAGE

Purpose	Twenty modules covering everything between a crawler and a rendered page.
When to use it	Monthly audit, after every release, and immediately on any alert.
Owner	Technical SEO owner
Inputs	Crawls, Search Console, log files, release notes
Outputs	Scored issues, tickets, verified fixes
Related databases	Technical Issues, Releases & Verification, Site Architecture
Related SOPs	SOP-16 to SOP-21
Related prompts	PR-43 to PR-50
Dashboard views	Backlog by priority · Critical open · Awaiting verification · Release watch
Review cadence	Monthly crawl, weekly backlog review

QUALITY GATE

No issue closes until verified in production, with a date

DEFINITION OF DONE

Every known issue has a score, an owner, and a ticket

07 Internal Linking & Site Architecture

WORKSPACE PAGE

Purpose	Where authority and readers are directed. The cheapest lever in the system.
When to use it	Monthly, and after every batch of publishing.
Owner	SEO specialist
Inputs	Crawl export, link counts, commercial importance
Outputs	Implemented links, architecture decisions
Related databases	Internal Link Opportunities, Site Architecture, Content
Related SOPs	SOP-22, SOP-23
Related prompts	PR-51 to PR-55
Dashboard views	Under-linked money pages · Orphans · Implementation queue · Verified
Review cadence	Monthly

QUALITY GATE

Every implemented link records a reader-benefit rationale

DEFINITION OF DONE

Commercial pages are among the best-linked pages on the site

08 Authority, Links & Digital PR

WORKSPACE PAGE

Purpose	Earning references: assets, campaigns, outreach, reclamation, monitoring.
When to use it	Continuously, campaign by campaign.
Owner	Marketing or PR owner
Inputs	Linkable assets, prospect research, mention monitoring
Outputs	Coverage, citations, reclaimed links
Related databases	Backlinks, Authority Targets, Digital PR Campaigns
Related SOPs	SOP-24 to SOP-28
Related prompts	PR-56 to PR-62
Dashboard views	Live campaigns · Qualified targets · Unlinked mentions · Reclamation queue
Review cadence	Weekly outreach review, monthly link review

QUALITY GATE

Nothing pursued that would need to be hidden if published

DEFINITION OF DONE

Every outcome is logged, including the refusals and the silence

09 Local SEO

WORKSPACE PAGE

Purpose	Per-location consistency, profiles, reviews, citations, and local content.
When to use it	Monthly per location, and immediately on any NAP change.
Owner	Local SEO owner
Inputs	Canonical NAP records, profile exports, citation audits
Outputs	Corrections, responses, local pages
Related databases	Local Locations, Local Citations, Content
Related SOPs	SOP-29 to SOP-32
Related prompts	PR-63 to PR-66
Dashboard views	Inconsistencies · Unclaimed listings · Reviews awaiting response · Audit due
Review cadence	Monthly per location

QUALITY GATE

Every difference is quoted exactly, with its source

DEFINITION OF DONE

One canonical NAP record generates every listing and page

10 LLM & AI Visibility

WORKSPACE PAGE

Purpose	The largest operating area: prompt universe, runs, mentions, citations, sources.
When to use it	Monthly full run, weekly on high-priority prompts.
Owner	SEO lead
Inputs	Prompt universe, brand facts, captured answers
Outputs	Scorecards, citation gaps, correction actions
Related databases	Brand Facts, Prompt Universe, Prompt Runs, AI Answers, Citations & Sources, AI Competitors, AI Opportunities, Historical Scorecards
Related SOPs	SOP-33 to SOP-38
Related prompts	PR-67 to PR-78
Dashboard views	Latest run · Absent prompts · Factual errors · Citation gaps · Trend
Review cadence	Monthly, identical wording every run

QUALITY GATE

Prompt wording never changes between runs; errors escalate immediately

DEFINITION OF DONE

You can show a twelve-month trend for five defined measures

11 Measurement & Reporting

WORKSPACE PAGE

Purpose	Where numbers become decisions rather than status updates.
When to use it	Weekly and monthly, on a fixed slot.
Owner	SEO lead
Inputs	All performance data, with sources and as-of dates
Outputs	Reports, decisions, system changes
Related databases	Reports, Business Goals, Decisions & Change Log
Related SOPs	SOP-39, SOP-40
Related prompts	PR-79 to PR-84
Dashboard views	Latest report · Goal progress · Decisions outstanding · Trend views
Review cadence	Weekly and monthly

QUALITY GATE

Every figure carries a source and an as-of date; nothing is forecast

DEFINITION OF DONE

Every report ends in three owned decisions and a system change

12 Prompt Library

WORKSPACE PAGE

Purpose	Every Claude prompt in the system, filterable by task and category.
When to use it	Whenever a task in any operating area begins.
Owner	SEO lead
Inputs	The prompt specifications in this playbook
Outputs	Consistent, reviewable AI output
Related databases	Prompt Library
Related SOPs	All SOPs reference it
Related prompts	All 84 prompts
Dashboard views	By category · By difficulty · Most used · Recently revised
Review cadence	Quarterly review of what is actually being used

QUALITY GATE

Every prompt names its inputs, its failure conditions, and its reviewer

DEFINITION OF DONE

A new team member can run any workflow from the library alone

13 SOP Library

WORKSPACE PAGE

Purpose	Every recurring procedure, with owners, triggers, gates, and escalation.
When to use it	On the trigger defined in each SOP.
Owner	SEO lead
Inputs	The SOP specifications in this playbook
Outputs	Consistent execution across people
Related databases	SOP Library
Related SOPs	All 40 SOPs
Related prompts	Referenced per SOP
Dashboard views	By cadence · By owner · Due this week · Never run
Review cadence	Quarterly

QUALITY GATE

Every SOP has a named owner and a written escalation route

DEFINITION OF DONE

The work happens the same way regardless of who does it

14 Templates

WORKSPACE PAGE

Purpose	Fifty-two reusable documents: briefs, audits, records, reports, plans.
When to use it	Whenever a document is created — never start from a blank page.
Owner	Whoever owns the task
Inputs	The template specifications in this playbook
Outputs	Consistent, complete documents
Related databases	Feeds every database
Related SOPs	Referenced per template
Related prompts	Referenced per template
Dashboard views	By operating area · By type · Most used
Review cadence	Quarterly

QUALITY GATE

Every template names its destination database and its completion criterion

DEFINITION OF DONE

Nobody in the team writes a document from scratch

15 Experiments & Learning

WORKSPACE PAGE

Purpose	One variable at a time, with the decision rule written before the result.
When to use it	When the team disagrees about what would work.
Owner	SEO lead
Inputs	Hypothesis, baseline with dates, available volume
Outputs	Adopt, reject, or extend decisions
Related databases	Experiments, Reports
Related SOPs	SOP-41
Related prompts	PR-85, PR-86
Dashboard views	Running · Awaiting evaluation · Concluded · Adopted changes
Review cadence	Weekly check, evaluation on the committed date

QUALITY GATE

The decision rule predates the result; one variable only

DEFINITION OF DONE

Every change to the system traces to an experiment or a decision

16 Archive & Decision Log

Purpose	What was decided, why, by whom, and what happened afterwards.
When to use it	On every material decision, and on every archive action.
Owner	SEO lead
Inputs	Decisions, their evidence, and their outcomes
Outputs	An auditable record of how the system evolved
Related databases	Decisions & Change Log, Archive
Related SOPs	SOP-42
Related prompts	PR-87
Dashboard views	Recent decisions · Reversed decisions · Due for review · Archived assets
Review cadence	Monthly review, quarterly retrospective

QUALITY GATE

Every decision records its evidence and its review date

DEFINITION OF DONE

You can explain any past decision without relying on memory

The Notion database standard

Every one of the thirty-two databases in this system follows the same specification. Once you can read one, you can build all of them.

SECTION	WHAT IT SPECIFIES	WHY IT IS THERE
Header	Purpose, owner, who uses it, inputs, outputs	So nobody has to guess what the database is for or who maintains it
Properties	Name, Notion type, allowed values, required flag, default, definition, example	So the database can be built exactly, without interpretation
Relations	Destination database, cardinality, purpose	Relations are where most Notion builds go wrong — they are specified explicitly
Rollups	Source relation, property, calculation, why it matters	So the numbers on a card come from data rather than from manual entry
Formulas	The exact expression where Notion supports it, or the method where it does not	So scores calculate themselves and stay comparable
Views	Name, type, filters, sort, grouping, intended user, decision supported	A view exists to support one decision — if it does not, it is clutter
Template	A database template with pre-filled sections and a completion checklist	So new records start complete rather than empty
Workflow	Status values, allowed transitions, blocking conditions, definition of done	So a record cannot advance past a stage it has not satisfied
Cadence, archive, gate	Review frequency, archive rule, quality gate	So the database stays useful instead of becoming a graveyard
Example records	Three fictional records for the most important databases	So you can see what a completed record actually looks like

Property conventions used throughout

Applied consistently across all thirty-two databases.

CONVENTION	RULE	REASON
Required flag	R = required before the record leaves Draft status	Prevents half-complete records entering a workflow
Unknown values	Left empty, never estimated or filled with a placeholder number	An estimated value becomes an apparent fact within a month
Dates	Every measurement carries an as-of date property	A number without a date cannot be compared to anything
Sources	Every factual property has a paired Source property	Provenance is what makes a record reviewable later
Owner	Person type, never text — one named human per record	Text owners drift; a Person property can be filtered and notified
Scores	Number properties feeding a Formula, never a manually entered total	Manual totals stop matching their components within weeks
Status	Select with a fixed set, never multi-select	A record in two states cannot be reported on
Archive	A Checkbox plus an Archived date, never deletion	Deleted records take their reasoning with them

BUILD THE PROPERTIES BEFORE THE RELATIONS


Notion will let you create a relation to a database that does not exist yet, and the resulting mess is the single most common reason people abandon a workspace build halfway through. Create all thirty-two databases with their properties first, then add relations in the order given in Section 18, then rollups, then formulas. That order is not a preference — it is a dependency chain.

SECTION 01 · FOUNDATIONS

SEO foundations

The principles that hold regardless of channel, algorithm, or which assistant your buyer happens to be using. Everything later in this playbook assumes these are in place.

THE CHAIN EVERY SEO DECISION HANGS FROM



If a piece of work cannot be traced back to a business goal along this chain, it is activity rather than strategy.

The twelve things that still decide everything

PRINCIPLE	WHAT IT MEANS IN PRACTICE	HOW YOU KNOW IT IS FAILING
Search intent	The page answers the question actually being asked	High impressions, low clicks, high bounce, no conversions
User problems	You write about problems, not about your product features	Content that only makes sense to someone who already knows you
Relevance	The page is clearly about one thing, thoroughly	Pages that mention ten topics and rank for none
Authority	Other credible sources reference you on this subject	Good content that never outranks thinner competitors
Trust	Named authors, real credentials, verifiable claims, clear ownership	No author, no about page, no way to tell who is behind it
Technical accessibility	Systems can reach, render, and read the page	Pages missing from the index or rendered empty without JavaScript
Information architecture	A logical structure a person could draw from memory	Important pages four or five clicks from the homepage
Crawlability	Crawl budget goes to pages that matter	Log files showing crawlers spending time on filters and pagination
Indexability	The right pages are indexable and the rest are not	Thin, duplicate, or parameter pages competing with your real ones
Content quality	Original substance a competitor cannot copy in an afternoon	Content that summarizes what already ranks and adds nothing
Brand demand	People search for you by name	All traffic non-branded, all of it fragile
Conversion alignment	The page leads somewhere commercially sensible	Traffic that grows while pipeline does not

THE SEO MATURITY ASSESSMENT

LEVEL 1	Reactive SEO happens when someone remembers or a problem appears	
LEVEL 2	Tactical Keywords and content exist, but nothing connects to a goal	
LEVEL 3	Structured A documented strategy, a backlog, and a review cadence	
LEVEL 4	Systematic One workspace, defined SOPs, owners, and quality gates	
LEVEL 5	Compounding Every cycle updates the system that produced it	

Most teams sit at Level 2 and believe they are at Level 3. The difference is whether the strategy exists as a document anyone could follow.

The SEO maturity assessment

Score each dimension 1 to 5 against the levels above. Your maturity is your lowest score, not your average.

DIMENSION	THE QUESTION TO ASK HONESTLY
Strategy	Could a new hire read one document and know what we are trying to win?
Demand understanding	Do we know what buyers ask, in their words, with evidence?
Keyword system	Is there one list everyone works from, with a scoring rule?
Content system	Is there a brief, a standard, and a refresh cadence?
Technical health	Do we know our issues, ranked, with owners and tickets?
Internal linking	Do we know which money pages have no internal links?
Authority	Do we know who cites us and why, and where the gaps are?
AI visibility	Do we know what AI systems say about us today, tracked over time?
Measurement	Can we connect an SEO action to a business metric?
Governance	Is there an owner, a cadence, and a place decisions are recorded?

The one-page SEO strategy template

If your strategy does not fit here, it is a task list rather than a strategy.

TEMPLATE — ONE-PAGE SEO STRATEGY

SEO Strategy — [company] — [quarter]

BUSINESS GOAL: [the commercial outcome SEO is serving this quarter]

WHY SEO: [why this channel, versus the alternatives]

AUDIENCE: [who has the problem — segment, role, situation]

THEIR SEARCH BEHAVIOUR: [where they look and what they type or ask]

WHERE WE COMPETE: [topics and surfaces we intend to win]

WHERE WE DO NOT: [explicit — this list matters more than the one above]

OUR UNFAIR ADVANTAGE: [first-party data, expertise, product, customers, reach]

THIS QUARTER'S THREE BETS:

1. [bet] → success signal: [measurable] → owner: [name]

2. [bet] → success signal: [measurable] → owner: [name]

3. [bet] → success signal: [measurable] → owner: [name]

CONSTRAINTS: [budget, people, dev capacity, publishing capacity]

REVIEW DATE: [when] VERSION: [date] — [owner]

Business objective to SEO metric mapping

Every metric you report should trace back to something the business already cares about.

BUSINESS OBJECTIVE	SEO OUTCOME THAT SERVES IT	METRIC TO TRACK	LEADING INDICATOR
Grow qualified pipeline	Visibility on high-intent commercial queries	Demo or trial starts from organic	Rankings and clicks on BOFU pages
Lower acquisition cost	Compounding non-paid discovery	Organic share of new customers	Non-branded impression growth
Enter a new segment	Coverage of that segment's vocabulary	Organic sessions from segment-specific pages	Published cluster coverage
Build brand demand	Being named as an option, everywhere	Branded search volume	Mentions and citations across surfaces
Shorten the sales cycle	Comparison and objection content that pre-answers	Content-assisted opportunities	Views on comparison and pricing pages
Reduce support load	Documentation and FAQ discoverability	Self-serve resolution rate	Search visibility on support queries
Defend against a competitor	Presence on comparison and alternative queries	Win rate on competitive deals	Rankings on "[competitor] alternative"
Be recommended by AI systems	Findable, understandable, credible, citable content	Citation and mention rate in tracked prompts	Entity consistency and source coverage

The modern SEO principles checklist

Run this quarterly. Anything unchecked becomes a backlog item with an owner.

- ☐ A written strategy exists and names what we will not pursue
- ☐ Every priority page has a stated intent and a stated conversion path
- ☐ Content is written from customer problems, in customer language
- ☐ Authors are named and their credentials are verifiable
- ☐ Claims on the site are sourced or removed
- ☐ Every important page is reachable in three clicks or fewer
- ☐ Crawl and index status is checked on a cadence, not on a hunch
- ☐ Structured data is present, valid, and honest about the page
- ☐ Key facts about the brand are identical everywhere they appear
- ☐ Content is extractable — a system could quote one passage safely
- ☐ Internal links point at commercially important pages
- ☐ Measurement connects at least one SEO metric to a business objective

THE ORDER MATTERS MORE THAN THE LIST

Teams routinely work on authority and AI visibility while their own site is partly unindexable and their money pages have no internal links. Foundations first is not conservatism — it is the only sequence where later work compounds instead of leaking.

Market, audience, and search demand

Keyword tools tell you what has already been typed enough times to be measured. This section is about understanding the problem before you accept a tool's framing of it.

Almost every weak SEO programme starts the same way: someone opens a keyword tool, exports a list, and begins. The list is real, but it is a record of the past, phrased in the vocabulary that already ranks. It contains no emerging demand, none of the language your customers use in private, and nothing about the questions people now ask an assistant instead of a search box.

Demand research comes first. It costs a week and changes everything downstream.

Six sources of demand evidence, ranked by usefulness

SOURCE	WHAT IT GIVES YOU	HOW TO MINE IT	STRENGTH
Sales and support calls	Exact objections and problem language	Pull transcripts, extract verbatim phrases by stage	Highest — real buyers, real stakes
Customer interviews	The trigger that started the search	Ask what they typed and where they looked first	Very high — small samples, deep signal
Reviews and communities	Unfiltered complaints and comparisons	Read category reviews and relevant threads, log recurring themes	High — public and current
Site search and internal queries	What people cannot find on your own site	Export site search logs and support ticket subjects	High — zero-cost, under-used
Keyword tools	Volume and difficulty for existing demand	Use to size and validate, not to originate	Medium — historical by nature
AI assistants	How questions are being phrased conversationally	Ask the assistant what people ask about the category	Medium — directional, needs validation

Workflow: mining customer language

Two hours, once a quarter. It produces the vocabulary every other section uses.

#	STEP	OUTPUT
1	Gather every transcript, review, ticket, and interview note from the quarter	A single labelled corpus with source and date on each item
2	Extract verbatim problem statements — no paraphrasing, no cleaning up	A raw quote list with role, segment, and source
3	Group by underlying problem, not by similar wording	Problem clusters with a count of distinct people per cluster
4	Note the trigger that preceded each search	A trigger list you can later map to content timing
5	Flag every phrase your team uses that customers never do	A vocabulary gap list — usually uncomfortable reading
6	Flag every phrase customers use that your site never does	The highest-value input to keyword research

Separating real demand from vanity topics

Signs of real demand

- Multiple customers described it unprompted
- It appears in lost-deal reasons or support volume
- It maps to a decision someone has to make
- People search it with commercial or comparison intent
- You have something specific to say that others do not

Signs of a vanity topic

- High volume, no relationship to what you sell
- Interesting to your team, not to a buyer
- Every competitor has covered it identically
- Traffic would arrive and leave with nothing to do
- You would be summarizing rather than contributing

Zero-click and AI-answer opportunities

Some queries will be answered without a visit. That is not automatically a loss.

QUERY TYPE	LIKELY OUTCOME	WHAT TO DO ABOUT IT
Simple definition or fact	Answered in place, no click	Cover it for citation and entity clarity, not for traffic
How-to with a clear procedure	Partially answered, some clicks for detail	Make the summary extractable and the depth worth the visit
Comparison between named options	Frequently synthesized by AI answers	Publish honest, structured comparisons that are safe to quote
Pricing and cost questions	Answered from third-party sources if not from you	Publish real pricing structure so the answer comes from you
Product or vendor selection	Answered as a shortlist of recommendations	This is where AI visibility work in Section 10 pays
Deep, situation-specific problems	Rarely answered fully without a click	The highest-value content to own — depth is defensible

Six databases for the demand layer

Build these in Notion. Every later section reads from them.

TOPICS, KEYWORDS, CONTENT Customer Problem Database		
PROPERTY	TYPE	WHAT IT HOLDS
Problem	Title	The problem stated as a customer would state it
Verbatim quotes	Text	Exact wording, one per line, never paraphrased
Source	Select	Call / interview / review / ticket / community
Source date	Date	When it was said — freshness matters
Segment	Multi-select	Which segment or role raised it
Distinct people	Number	How many separate people, not how many mentions
Buyer stage	Select	Problem-aware / solution-aware / vendor-aware
Evidence strength	Select	Strong / moderate / single data point
Linked topics	Relation	Topics database
Status	Select	New / validated / covered / declined
PURPOSE Every problem a customer has described, with evidence and frequency.		REVIEW Hold the raw demand evidence everything else is built from.
		RELATED Quarterly

KEYWORDS, TOPICS, AI VISIBILITY Search Demand Database		
PROPERTY	TYPE	WHAT IT HOLDS
Query or question	Title	As typed or as asked
Demand type	Select	Established / emerging / conversational / zero-click
Evidence	Text	Where you saw it, with the source
Volume	Number	If measurable — leave empty rather than guessing
Intent	Select	Informational / comparison / commercial / navigational
Linked problem	Relation	Customer Problem Database
Surface	Multi-select	Google / AI answers / community / video / marketplace
Status	Select	Unvalidated / validated / mapped / rejected
PURPOSE What people search, with volume and evidence, before it becomes a keyword list.		REVIEW Bridge raw problems and the keyword system without losing the evidence.
		RELATED Monthly

CONTENT, AI VISIBILITY Question Bank		
PROPERTY	TYPE	WHAT IT HOLDS
Question	Title	Phrased exactly as asked
Asked by	Select	Role or segment
Stage	Select	Where in the journey it appears
Answered on	Relation	Content database — blank means a gap
Answer is extractable	Checkbox	Can a system quote our answer safely?
Priority	Select	High / medium / low
Source	Text	Call, community, assistant, or support ticket

PURPOSE Every question a buyer asks, in the form they ask it.

REVIEW Feed FAQ content, on-page coverage, and the AI prompt universe.

RELATED Monthly

The demand validation checklist

Before a topic enters the content pipeline, it should pass all six.

- ☐ At least three distinct people described this problem, or the evidence is labelled thin
- ☐ The problem connects to a decision that leads somewhere commercially
- ☐ We have something specific to say that is not already published elsewhere
- ☐ The vocabulary comes from customers, not from our internal naming
- ☐ We know which surfaces this question gets asked on
- ☐ We know what a good answer looks like and whether it can be extracted safely

Keyword research

A keyword list is not a strategy — but without one, every content decision is reopened from scratch. This is the system that turns demand evidence into a prioritized, mapped, defensible list.

The sixteen keyword types, and what each is for

TYPE	EXAMPLE SHAPE (PLACEHOLDER)	WHAT IT IS GOOD FOR
Seed	[category]	Starting expansion — rarely a target itself
Head term	[category] software	Long-term authority, hard to win early
Long-tail	how to [job] when [constraint]	Early wins and precise intent
Commercial	best [category] for [segment]	Directly precedes a purchase decision
Informational	what is [concept]	Topical coverage and entity clarity
Comparison	[brand A] vs [brand B]	High intent, often AI-synthesized
Alternative	[competitor] alternative	Buyers already in market and dissatisfied
Review	[brand] review	Trust-stage queries you rarely own outright
Branded	[your brand] [feature]	Demand you created — protect it
Local	[service] in [city]	Geographic intent, covered in Section 09
Question	why does [problem] happen	FAQ, extractability, AI answers
Problem-aware	[symptom] keeps happening	Earliest useful stage to reach someone
Solution-aware	how to [approach] [problem]	They know the category, not the vendor
Purchase-intent	[category] pricing	Closest to revenue, usually under-covered
Conversational	what should I use to [job] if [constraint]	How queries arrive in assistants
Jobs-based	tool to [specific job]	Maps directly to what your product does

The keyword scoring model

Ten criteria, scored 0–5. Transparent, adjustable, and openly imperfect.

CRITERION	SCORE 0	SCORE 5	WEIGHT
Business relevance	Unrelated to what we sell	Directly precedes a purchase decision	×3
Search demand	No measurable or evidenced demand	Strong, sustained, evidenced demand	×2
Intent match	Intent conflicts with our page type	Intent matches exactly what we offer	×3
Ranking difficulty	Dominated by far stronger domains	Weak, outdated results we can beat	×2
Conversion potential	Traffic with nowhere to go	Clear next step and a real offer	×3
Current authority	No topical presence at all	We already rank nearby and are trusted here	×1
Content gap	Well covered by everyone including us	A real gap we can fill distinctively	×2
AI answer potential	Cannot be summarized safely	Extractable, factual, quotable	×2
Brand fit	Off-positioning or off-audience	Reinforces exactly what we stand for	×1
Strategic value	Isolated, leads nowhere	Anchors a cluster or defends a position	×2

THE PRIORITIZATION FORMULA

PRIORITY SCORE = sum of (criterion score x weight)

Maximum possible = 105

Suggested bands:

80-105 Build now – cluster anchor or revenue page

60-79 Build this quarter

40-59 Build when capacity allows, or fold into an existing page

20-39 Monitor only

0-19 Decline and record the reason

ADJUST THE WEIGHTS. They are a starting point reflecting a typical B2B situation, not a scientific model. If conversion matters more than volume in your business, raise it. If you are pre-authority, raise ranking difficulty. Record every weight change with a date so scores stay comparable.

BE HONEST ABOUT WHAT A SCORE IS

This model makes decisions consistent and reviewable. It does not make them correct. Two competent SEOs will score the same keyword differently, and the weights are judgement expressed as numbers. Its value is that it forces the judgement to be written down, so it can be argued with and improved — not that it produces objective truth.

The search intent matrix

Map each keyword to a cell. A mismatch between cell and page type is the most common reason good content does not rank.

	INFORMATIONAL	COMPARISON	COMMERCIAL	NAVIGATIONAL
Problem-aware	Explainer, guide, glossary	Approach vs approach	Rare — usually too early	Rare
Solution-aware	How-to, framework, template	Category vs category	Category landing page	Category directory
Vendor-aware	Deep guide with product context	Brand vs brand, alternatives	Product page, pricing, demo	Brand and login queries
Customer	Documentation, advanced use	Upgrade and tier comparison	Expansion and add-ons	Support and account

TOPICS, CONTENT, AI VISIBILITY

Keyword Master Database

PROPERTY	TYPE	WHAT IT HOLDS
Keyword	Title	The query as searched
Type	Select	One of the sixteen types above
Intent	Select	Informational / comparison / commercial / navigational
Stage	Select	TOFU / MOFU / BOFU
Branded	Checkbox	Branded or non-branded
Volume	Number	Leave blank rather than estimating
Difficulty	Number	Whatever your tool reports, with the tool named
Criterion scores	Number ×10	One property per scoring criterion
Priority score	Formula	The weighted total
Target page	Relation	Content database — one keyword, one owning page
Cluster	Relation	Topics database
Current position	Number	Updated on your tracking cadence
Owner	Person	Who is accountable for this query
Status	Select	Unmapped / mapped / published / ranking / declined

PURPOSE One list everyone works from, scored and mapped.

REVIEW The single source of truth for what the company is trying to be found for.

RELATED Weekly for status, quarterly for scoring

Keyword-to-page mapping, and the cannibalization audit

One primary keyword, one owning page. Everything else is a supporting variation.

#	AUDIT STEP	WHAT YOU ARE LOOKING FOR
1	Export every keyword with its ranking URL over the last 90 days	Queries where the ranking URL changes back and forth
2	Group by keyword and count distinct ranking URLs	Two or more of your pages competing for one query
3	For each conflict, identify the intended owner	Which page should win, given intent and conversion path
4	Decide: consolidate, differentiate, de-optimize, or redirect	A decision per conflict, recorded with a reason
5	Check internal links point at the intended owner	Links pulling authority towards the page you decided to demote
6	Re-check after 60 days	Whether the intended page now owns the query consistently

THE MOST COMMON KEYWORD MISTAKE

Mapping many similar keywords to many similar pages, in the belief that more pages means more coverage. It produces a set of pages that each rank weakly, compete with each other, split internal link equity, and are individually too thin to be cited by anything. One strong page beats four adequate ones on every surface discussed in this playbook.

Content strategy

Publishing is not a strategy. This section is about deciding what to build, in what order, how deep, and — the decision most teams never make — what to stop maintaining.

The hub-and-spoke system

One pillar, six to twelve supporting pieces, linked in both directions. This is what produces topical authority rather than a pile of posts.

LAYER	WHAT IT IS	WHAT IT DOES	HOW MANY
Pillar page	The comprehensive page on the core topic	Ranks for the head term, holds the cluster together	One per cluster
Supporting content	Depth on one sub-question each	Ranks for long-tail, links up to the pillar	6–12 per cluster
Commercial page	The product, category, or pricing page	The destination the cluster points at	1–2 per cluster
Comparison content	Honest vs and alternative pages	Captures vendor-aware intent and is frequently cited	2–5 per cluster
Reference assets	Glossary, templates, tools, data	Earns links, gets cited, defends the cluster	As the topic warrants

The create / update / merge / redirect / remove decision

Run this on every existing page once a quarter. Most teams have more to fix than to write.

IF THE PAGE...	AND...	THEN	WHY
Ranks and converts	Nothing has changed in the topic	Leave it alone	Do not touch what is working
Ranks but does not convert	Intent is commercial	Fix the conversion path	The traffic is right, the destination is wrong
Ranked and is now sliding	Competitors have gone deeper	Update against what overtook it	Refreshing beats writing new almost every time
Gets impressions, no clicks	Position 4–10	Rewrite title and meta	This is a copy problem, not a ranking problem
Overlaps another page	Both rank weakly for one query	Merge into the stronger URL	One strong page beats two adequate ones
Has no traffic and no links	Live for over a year	Remove and redirect	Thin pages dilute crawl and topical clarity
Has links but no traffic	The topic is no longer relevant	Redirect to the closest live page	Preserve the equity, drop the page
Is factually out of date	It is still cited or linked	Update urgently	Outdated cited content is a credibility risk on AI surfaces

The content brief template

If a writer has to ask what the page is for, the brief was incomplete.

TEMPLATE — CONTENT BRIEF

Content Brief — [working title]

PRIMARY KEYWORD: [one] INTENT: [informational / comparison / commercial]

STAGE: [TOFU / MOFU / BOFU] CLUSTER: [which pillar this supports]

READER: [who, and what they already know]

THE DECISION THIS SUPPORTS: [what they are trying to decide]

AFTER READING THEY SHOULD BE ABLE TO: [specific capability]

ANGLE: [what we say that the current top results do not]

OUR EVIDENCE: [first-party data, customer quotes, product experience, expertise]

WHAT WE WILL NOT CLAIM: [anything we cannot support]

REQUIRED COVERAGE: [the questions the page must answer, in order]

EXTRACTABLE ANSWERS: [which passages must stand alone if quoted]

STRUCTURE: [H2s and H3s]

INTERNAL LINKS IN: [pages that must link to this]

INTERNAL LINKS OUT: [pages this must link to — include the commercial page]

FORMAT: [article / comparison / template / tool / data study]

LENGTH GUIDE: [approximate — coverage decides, not word count]

AUTHOR: [named, with the credential that makes them credible here]

REVIEWER: [named] DUE: [date]

KEYWORDS, TOPICS, INTERNAL LINKS, REPORTING Content Database

PROPERTY	TYPE	WHAT IT HOLDS
Title	Title	The page
URL	URL	Live URL, or blank if unpublished
Primary keyword	Relation	Keyword Master Database — exactly one
Cluster	Relation	Topics database
Stage	Select	TOFU / MOFU / BOFU
Format	Select	Article / comparison / template / tool / study / glossary
Author	Person	Named, with a credential shown on the page
Published	Date	First publication date
Last updated	Date	Drives the refresh queue
Performance	Select	Growing / stable / declining / never landed
Next action	Select	None / update / merge / redirect / remove
Extractable	Checkbox	Contains a self-contained quotable passage
Owner	Person	Accountable for the next action
Status	Select	Brief / drafting / review / published / archived

PURPOSE Every page, its purpose, its performance, and its next action.

REVIEW The production and maintenance backbone of the whole system.

RELATED Weekly for status, quarterly for decay

The content scoring model

Score before commissioning. Six criteria, 0–5 each, weighted.

CRITERION	WHAT A 5 LOOKS LIKE	WEIGHT
Business proximity	The reader is one step from a commercial decision	×3
Evidence we hold	First-party data or expertise nobody else can replicate	×3

CRITERION	WHAT A 5 LOOKS LIKE	WEIGHT
Demand evidence	Multiple validated sources show people want this	×2
Competitive gap	A real gap — current results are thin or outdated	×2
Cluster contribution	It strengthens a pillar we are actively building	×2
Citation potential	Structured, factual, and safe for a system to quote	×2

THE REFRESH-FIRST RULE

Before commissioning anything new, check the decay monitor. A page that once ranked and is now sliding will almost always recover faster, and for less effort, than a new page will climb. Most content teams have this backwards, because writing is more satisfying than editing.

On-page SEO

Making one page as easy as possible to understand — for a person skimming it, a crawler parsing it, and a model deciding whether it can be quoted without risk.

On-page work has quietly changed job. It used to be about signalling relevance to a ranking algorithm. It is now equally about being extractable: can a system lift one passage from this page, present it as an answer, and attribute it to you without introducing an error?

A page that is genuinely clear performs well at both jobs. A page written to game one usually fails the other.

The on-page checklist

Run before publication. Nothing here is optional on a commercially important page.

ELEMENT	THE STANDARD	THE COMMON FAILURE
Title	Primary keyword, real benefit, roughly 60 characters	Keyword stuffing, or a title that describes nothing
Meta description	Accurate summary that earns the click	Auto-generated, duplicated, or unrelated to the page
URL	Short, readable, stable, no dates or parameters	Deep nesting and IDs that make the URL meaningless
H1	One per page, matching the search intent	Multiple H1s, or an H1 that differs from the promise
Heading structure	A hierarchy a person could outline from	Headings chosen for styling rather than structure
Introduction	Answers the question in the first two sentences	Three paragraphs of context before the answer
Intent match	The page type matches what the query wants	A blog post targeting a query that wants a product page
Semantic coverage	Related concepts covered because they belong	Keyword variations sprinkled without substance
Internal links	Contextual, descriptive anchors, pointing at money pages	"Click here", or no links to anything commercial
External references	Sources cited where a claim needs support	Unsupported statistics with no origin
Images and alt text	Compressed, described accurately	Alt text stuffed with keywords or left empty
Tables and lists	Used where structure genuinely aids comprehension	Walls of prose that should have been a table
Definitions	Key terms defined once, clearly, in one place	Terms used before they are explained
FAQs	Real questions from the Question Bank, answered directly	Invented questions padded for schema markup
CTA	One clear next step matched to the reader's stage	Four competing CTAs, or none at all

The extractability checklist

This is the part most on-page guidance still omits. It decides whether you are cited or merely summarized.

- ☐ The direct answer appears within the first 100 words, before any preamble
- ☐ Each key answer is self-contained — it makes sense lifted out of the page
- ☐ No answer depends on a previous paragraph to be understood
- ☐ Facts are stated plainly, not implied through narrative
- ☐ Numbers carry their unit, their period, and their source
- ☐ Definitions are single sentences that could stand alone
- ☐ Lists and steps are genuinely parallel and complete
- ☐ Tables have real headers rather than styled first rows
- ☐ Comparisons state the criterion before the verdict
- ☐ Nothing important lives only inside an image or a video
- ☐ The page renders fully without JavaScript execution
- ☐ Structured data matches what a human sees on the page
- ☐ The author, the publication date, and the last update are visible
- ☐ Claims that cannot be supported have been removed rather than softened

The AI citation readiness checklist

A narrower question: could a system quote this page and be confident it is correct?

REQUIREMENT	WHY IT MATTERS	HOW TO CHECK
One unambiguous claim per passage	Ambiguity makes a passage unsafe to quote	Read one paragraph alone — is it still true and clear?
Facts consistent across your own site	Self-contradiction undermines everything else	Search your domain for the same fact stated differently
Facts consistent with third parties	Corroboration carries weight	Check your directory listings, profiles, and about page
Dated and attributed	Freshness and authorship affect trust	Visible publication date, update date, and named author
Sources for anything non-obvious	An unsourced claim is risky to repeat	Every statistic links to its origin
No overclaiming	Superlatives make a page less quotable, not more	Remove "best", "leading", "#1" unless externally verified

Technical SEO

Everything above this section assumes machines can reach, render, and read your pages. This layer decides whether that assumption holds — and it is where fixes are cheapest per unit of effect.

The prioritization model

Technical backlogs fail because everything looks urgent. Score every issue the same way.

ISSUE PRIORITIZATION FORMULA

$PRIORITY = (Impact \times Affected\ Pages \times Business\ Importance \times Confidence) / Effort$

Impact 1-5 How much does fixing this change discoverability?

Affected pages 1-5 1 = a handful, 5 = sitewide

Business importance 1-5 Are the affected pages commercially important?

Confidence 1-5 How sure are we this is actually the cause?

Effort 1-5 1 = minutes, 5 = a project

Confidence is the term teams omit, and it is the one that prevents a quarter of engineering time being spent on a hypothesis. If confidence is below 3, the next action is a diagnostic, not a fix.

The severity matrix

SEVERITY	EXAMPLES	RESPONSE
Critical — site level	Robots.txt blocking, sitewide noindex, site down, certificate failure	Immediate incident response, fixed within hours
High — section level	A template unindexable, broken canonicals, key content rendered only by JS	Next sprint, with a named engineering owner
Medium — page level	Slow templates, redirect chains, orphan money pages, missing structured data	Scheduled into the backlog by priority score
Low — cosmetic	Minor markup warnings, image sizing, non-critical alt text	Batched, or fixed alongside other work

The technical audit checklist

Twenty-six checks across four groups. Run quarterly, and after any release that touches templates.

Crawl and index

- Robots.txt permits what should be crawled
- XML sitemaps current, containing only indexable URLs
- Index coverage matches the page count you intend
- Canonicals self-referencing or deliberate
- No redirect chains longer than one hop
- Status codes correct — no soft 404s
- Orphan pages identified, then linked or removed
- Crawl budget not consumed by filters and pagination

Structure and duplication

- Site architecture reachable in three clicks
- Faceted navigation controlled, not indexable by default
- Duplicate content consolidated or canonicalized
- Parameter handling documented
- Breadcrumbs present and marked up

Rendering and performance

- Key content renders without JavaScript
- Core Web Vitals measured on field data, not lab only
- Mobile usability verified on real devices
- Images compressed and correctly sized
- Server response time stable under load
- No render-blocking resources on critical templates
- HTTPS enforced everywhere, no mixed content
- Pagination handled consistently

International, markup, and monitoring

- Hreflang complete, reciprocal, and valid where used
- Structured data present, valid, matching visible content
- Broken internal and external links identified
- Log files reviewed for where crawlers actually spend time
- Security headers and certificate validity checked

The SEO dev-ticket template

Engineering does not reject SEO tickets because they disagree. They reject them because the ticket does not say what to change.

TEMPLATE — SEO DEV TICKET

TITLE: [what changes, in one line]

PROBLEM: [what is happening now, with the evidence]

EVIDENCE: [crawl export, log sample, screenshot, URL list – attached]

AFFECTED URLS: [count, plus a sample of 10 and the full list attached]

EXPECTED BEHAVIOUR: [what should happen instead, specifically]

ACCEPTANCE CRITERIA:

- [testable statement a QA engineer could verify]
- [testable statement]
- [testable statement]

PRIORITY SCORE: [from the formula] = impact [n] x pages [n] x business [n] x confidence [n] / effort [n]

BUSINESS IMPACT: [which pages, why they matter commercially]

OUT OF SCOPE: [explicitly, to prevent scope creep]

HOW WE WILL VERIFY: [what we check after deploy, and when]

ROLLBACK: [what to do if it makes things worse]

REQUESTED BY: [name] ENGINEERING OWNER: [name] TARGET: [sprint]

CONTENT, REPORTING, EXPERIMENTS

Technical Issues Database

PROPERTY	TYPE	WHAT IT HOLDS
Issue	Title	One line describing what is wrong
Severity	Select	Critical / high / medium / low
Type	Select	Crawl / index / render / performance / markup / duplication
Affected URLs	Number	Count, with the list attached
Evidence	Files	Crawl export, log sample, screenshot
Impact, Pages, Business, Confidence, Effort	Number ×5	The five scoring inputs
Priority score	Formula	The calculated result
Owner	Person	Engineering owner, not the SEO who found it
Ticket	URL	Link to the engineering system
Status	Select	Found / diagnosed / ticketed / in progress / deployed / verified
Verified date	Date	When the fix was confirmed working

PURPOSE Every known technical issue, scored, owned, and tracked to verification.	REVIEW Turn technical findings into engineering work that actually ships.	RELATED Weekly
--	---	----------------

Migration and relaunch checklist

Most catastrophic traffic losses are migrations that skipped one of these.

- ☐ Full crawl of the current site archived before anything changes
- ☐ Complete URL inventory with current traffic, rankings, and backlinks
- ☐ Redirect map covering every URL, one hop, tested on staging
- ☐ Canonical, hreflang, and structured data ported and validated
- ☐ Robots.txt and sitemaps prepared for the new structure
- ☐ Analytics and Search Console configured before launch, not after
- ☐ Internal links updated to final URLs, not pointing through redirects
- ☐ Staging environment blocked from indexing, and verified as blocked
- ☐ Launch window chosen to avoid the highest-traffic period
- ☐ Post-launch crawl within 24 hours, and again at 7 and 30 days
- ☐ Rollback plan written, tested, and owned by a named person
- ☐ Rankings and index coverage baselined for comparison

Internal linking and information architecture

The cheapest and most under-used lever in SEO. You already own every page involved, no outreach is required, and the effect is usually visible within weeks.

Internal linking does three jobs at once. It distributes authority towards the pages that matter commercially. It tells search systems and models how your topics relate. And it moves a reader from a question towards a decision. Most sites do the first two by accident and the third not at all.

The internal linking audit

Six checks. The first alone usually finds work worth a quarter of content production.

#	CHECK	WHAT YOU ARE LOOKING FOR	TYPICAL FINDING
1	Money pages with few internal links	Commercially important pages nothing points at	The pricing page has fewer links than a blog post
2	Orphan pages	Pages reachable only from the sitemap	Older content nobody linked to after publication
3	Click depth	Important pages more than three clicks from home	A whole category buried under pagination
4	Anchor text distribution	Descriptive anchors versus generic ones	"Learn more" pointing at your priority page
5	Blog-to-product links	Whether informational content leads anywhere	Hundreds of posts with no commercial link
6	Link decay	Links lost in redesigns and template changes	A navigation change that quietly removed a hub

The link priority score

Which internal link opportunity to build first.

LINK PRIORITY SCORE

$\text{LINK PRIORITY} = \text{Target page value} \times \text{Topical relevance} \times \text{Source page strength}$

Target page value 1-5 How commercially important is the destination?

Topical relevance 1-5 Would a reader genuinely want this link here?

Source page strength 1-5 Does the linking page have traffic and links itself?

Build in descending order. Stop when relevance drops below 3 – an irrelevant internal link helps nobody and dilutes the ones that matter.

RULE: never add a link a reader would not want to click. Internal linking done for the algorithm alone is visible to readers and to reviewers, and it degrades the page it sits on.

Anchor text guidelines

SITUATION	DO	DO NOT
Linking to a commercial page	Use a descriptive phrase matching the page topic	Use "click here", "this page", or a bare URL
Linking to a pillar	Use the topic name as it appears in the H1	Use exact-match keyword text on every instance
Multiple links to one page	Vary the anchor naturally across pages	Repeat identical exact-match anchors sitewide
Linking within a sentence	Link the phrase describing the destination	Link a whole sentence, or a single word like "here"
Navigation and footers	Keep them stable and meaningful	Stuff the footer with keyword anchors to everything

CONTENT, KEYWORDS, TOPICS

Internal Links Database

PROPERTY	TYPE	WHAT IT HOLDS
Opportunity	Title	Source page → target page
Source page	Relation	Content database
Target page	Relation	Content database
Suggested anchor	Text	The phrase to use
Target value, Relevance, Source strength	Number ×3	The three scoring inputs
Priority	Formula	The calculated score
Rationale	Text	Why a reader would want this link
Owner	Person	Who implements it
Status	Select	Identified / approved / implemented / verified

PURPOSE	Every planned and completed internal link opportunity.	REVIEW	Make internal linking a tracked workstream rather than an afterthought.	RELATED	Weekly
---------	--	--------	---	---------	--------

WHERE TO START ON A SITE YOU HAVE JUST INHERITED

Export every URL with its internal inbound link count and its commercial importance. Sort by importance descending and link count ascending. The top twenty rows of that list are almost always the highest-return SEO work available to you, and none of it requires writing anything new or asking anyone for a link.

Link building, digital PR, and authority

Links are one expression of authority, not the whole of it. This section covers the six kinds of authority that matter now, and how to build each without doing anything you would have to hide.

Six kinds of authority, and why the distinction matters

TYPE	WHAT IT IS	HOW IT IS BUILT	WHERE IT SHOWS UP
Link acquisition	Other sites linking to yours	Assets worth referencing, plus outreach	Classic ranking signals
Brand authority	People search for you by name	Consistent presence, product quality, and reputation	Branded search volume, direct traffic
Entity authority	Systems know who you are and what you do	Consistent facts everywhere, structured data, corroboration	Knowledge panels, AI answers
Topical authority	Depth and coverage on one subject	Clusters, not scattered posts	Rankings across a topic, not one query
Citation authority	You are quoted as a source	Original data, clear facts, extractable passages	AI answers, journalism, other sites
Trust signals	Evidence you are who you claim to be	Named authors, credentials, reviews, transparency	Everywhere — it gates all of the above

Linkable asset types, ranked by durability

The asset determines the outreach. Weak assets cannot be rescued by good outreach.

ASSET	WHY PEOPLE REFERENCE IT	EFFORT	DURABILITY
Original research from first-party data	Nobody else has the numbers	High	Years — and frequently cited by AI systems
Free tool or calculator	It solves a problem repeatedly	High	Years, with maintenance
Definitive template or framework	People use it in their own work	Medium	Years
Expert commentary on a live development	Journalists need a named source quickly	Low, but time-sensitive	Weeks — value is in speed
Comprehensive glossary	It becomes the reference for a term	Medium	Years, if maintained
Case study with real numbers	Proof others cite when arguing a point	Medium	Years, if the numbers are real and permitted
Curated industry resource	It saves someone an afternoon	Medium, plus upkeep	Decays fast without maintenance

The digital PR campaign template

TEMPLATE — DIGITAL PR CAMPAIGN

Digital PR Campaign — [name]

THE ASSET: [what exists that is worth referencing]

THE STORY: [the one sentence a journalist would write]

WHY NOW: [the timing hook — a season, a change, a live debate]

WHY US: [what makes us the credible source for this specifically]

THE DATA: [what we can share, from where, with what permission]

WHAT WE CANNOT CLAIM: [limits of the data — state these before outreach]

METHODOLOGY: [published alongside — this is what makes it citable]

TARGET PUBLICATIONS: [type, not just names — trade, national, niche, newsletter]

TARGET JOURNALISTS: [who covers this beat, with a link to recent work]

WHAT THEY GET: [the angle for their audience, not for ours]

ASSETS PREPARED: [data file, charts, quotes, spokesperson bio, images]

SPOKESPERSON: [named, with the credential that matters here]

OUTREACH SEQUENCE: [first contact, follow-up, close — max three touches]

NOT DOING: [no paid placements, no link exchanges, no undisclosed sponsorship]

SUCCESS SIGNALS: [coverage, citations, mentions, referral traffic]

OWNER: [name] WINDOW: [dates]

Outreach template library

Four situations, four openings. Each under 120 words in practice.

SITUATION	THE OPENING MOVE	THE ASK
Journalist, live story	Reference the specific piece they just published	Offer a named source and a data point, no link request
Resource page	Name the page and what your asset adds that is missing	Suggest the addition, and accept a no
Unlinked mention	Thank them for the mention, note the correct URL	Ask whether they would be willing to link it
Original research pitch	Lead with the finding, not with your company	Offer the data, the methodology, and an interview

WHAT THIS PLAYBOOK WILL NOT HELP YOU DO

Buying links, private blog networks, link exchanges, undisclosed sponsored placements, automated comment or forum links, scraped guest-post networks, and anything that requires hiding the arrangement. These violate search engine guidelines, carry real penalty risk, and — increasingly relevant — do nothing for entity or citation authority, which is where the growth now is. If a tactic only works while nobody knows about it, it is not in this playbook.

CONTENT, AI VISIBILITY, REPORTING Backlinks and Authority Database

PROPERTY	TYPE	WHAT IT HOLDS
Target	Title	The site, publication, or resource page
Type	Select	Journalist / resource page / unlinked mention / partner / directory
Asset	Relation	Content database — what we are offering
Relevance	Number	1–5, topical fit with our subject
Authority indicator	Number	Whatever metric you use, with the tool named
Contact	Text	Role and public contact route only
Status	Select	Identified / contacted / responded / earned / declined
Outcome	Select	Link / mention / citation / coverage / none
Date	Date	When it landed
Owner	Person	Who is running it

PURPOSE Every link opportunity, mention, and earned reference.

REVIEW Track authority as a workstream with real outcomes, not just link counts.

RELATED Weekly

The toxic-link review checklist

Most sites need this far less often than tools suggest. Review before you disavow anything.

- ☐ The link genuinely came from a manipulative source, not just a low-metric one
- ☐ It is part of a pattern, not an isolated scrape or aggregator
- ☐ You did not build it yourself, or it was built by a previous agency
- ☐ Removal was attempted first, and documented
- ☐ There is a plausible connection to an actual ranking problem
- ☐ The decision and its reasoning are recorded before any disavow file is filed

Local SEO

For anyone who serves a place: a service business, a multi-location brand, or a company whose buyers care where you are. Local visibility runs on different signals and rewards consistency above almost everything else.

The local SEO checklist

Five groups. NAP consistency is the one that quietly breaks everything else.

Profile and listings

- Google Business Profile claimed and fully completed
- Primary category correct, secondary categories deliberate
- Hours accurate, including holidays
- Services and products listed with descriptions
- Photos current and genuinely representative
- NAP identical across every listing, character for character
- Core citations claimed and consistent
- Duplicate listings identified and merged

Site and content

- One landing page per location, with unique content
- Service-area pages that describe real service, not just a place name
- Location schema present and matching the profile
- Embedded map and directions on each location page
- Local-intent keywords covered in real language
- Local content that would interest a local reader
- Internal links from the site to each location page
- Mobile experience verified — most local search is mobile

The local landing page template

TEMPLATE — LOCAL LANDING PAGE

[Service] in [Location]

H1: [Service] in [Location]

INTRO: What we do here, for whom, and what makes this location specific.
Two sentences. No stuffed place names.

WHAT WE OFFER HERE: [services actually available at this location]

WHO WE SERVE HERE: [the local customer, described specifically]

LOCAL CONTEXT: [something true about operating in this place — regulations, conditions, timing, the local situation]

TEAM: [who works here, named where possible]

ADDRESS AND HOURS: [exactly matching the Business Profile, character for character]

MAP AND DIRECTIONS: [embedded, plus written directions from a landmark]

CONTACT: [local number where one exists]

PROOF: [reviews or work from this location — never invented, never borrowed from another location]

FAQS: [real questions this location actually receives]

CTA: [the one action — call, book, visit, request]

SCHEMA: LocalBusiness with address, geo, hours, and telephone matching the profile exactly.

Review response prompts and rules

REVIEW TYPE	RESPOND BY	NEVER
Positive and specific	Thanking them and referencing what they mentioned	Copy-pasting the same response to every review
Positive but generic	A short, warm, human reply	Adding keywords or a sales pitch
Negative and fair	Acknowledging it, stating what you are doing, moving offline	Arguing, or disputing the facts in public
Negative and inaccurate	Correcting one fact politely, offering to resolve directly	A long public rebuttal

REVIEW TYPE	RESPOND BY	NEVER
Suspected fake	Reporting through the platform, responding neutrally meanwhile	Accusing the reviewer publicly

ON REVIEW GENERATION

Ask customers for reviews. Do not offer anything in exchange, do not filter who gets asked based on how happy they seem, do not write reviews for anyone, and do not use services that supply them. Beyond the platform-policy risk, incentivized and gated reviews are increasingly detectable and they poison the exact signal you were trying to build.

CONTENT, KEYWORDS, REPORTING

Local SEO Database

PROPERTY	TYPE	WHAT IT HOLDS
Location	Title	Branch, city, or service area
NAP	Text	The canonical name, address, and phone — the single source of truth
Profile status	Select	Unclaimed / claimed / optimized / verified
Primary category	Text	As set in the Business Profile
Landing page	Relation	Content database
Citations	Number	Claimed and consistent count
Citation issues	Text	Where NAP does not match, with the source
Reviews	Number	Count and average, with the as-of date
Local keywords	Relation	Keyword Master Database
Owner	Person	Who maintains this location
Last audited	Date	Drives the review cadence

PURPOSE	Every location, its profile, its listings, and its performance.	REVIEW	Keep every location consistent — inconsistency is the main local ranking problem.	RELATED	Monthly per location
---------	---	--------	---	---------	----------------------

MULTI-LOCATION: THE ONE RULE THAT SCALES

Maintain one canonical NAP record per location and generate every listing and page from it. The moment addresses are edited independently in three systems, they drift — and NAP inconsistency across the web is the single most common cause of weak local visibility for businesses that are otherwise doing everything right.

LLM and AI SEO

The newest surface, and the one with the least reliable folklore around it. What follows is what is actually observable, what you can legitimately influence, and how to measure whether any of it is working.

AI answer engines do not rank pages the way a search engine does. They assemble an answer, and in doing so they decide which sources to draw on and which to name. That decision runs on things you can influence: whether you can be found at all, whether your facts are unambiguous, whether other credible sources say the same thing, whether your content can be extracted safely, and whether you are an obvious answer to a specific question.

Nobody outside these companies knows the weighting, and it changes. What follows is a system for improving the inputs and measuring the outputs — not a formula for being recommended.

THE AI VISIBILITY PROGRESSION

01 FINDABLE	Crawlable, indexable, discoverable on every surface your buyer uses
02 UNDERSTANDABLE	Clear entities, structured facts, unambiguous language and markup
03 CREDIBLE	Corroborated by sources other than you, consistently, over time
04 CITABLE	Extractable in self-contained passages a system can quote safely
05 RECOMMENDABLE	The obvious answer for a specific query, from a source with standing

Each stage depends on the one before it. Nothing here guarantees a recommendation — it changes the probability that one is possible.

What each stage actually requires

STAGE	WHAT INCREASES THE PROBABILITY	HOW YOU CHECK IT
Findable	Indexable pages, presence on third-party sources these systems draw on, coverage of the question in public places	Ask the assistant a question you should obviously appear for
Understandable	Consistent entity facts, unambiguous naming, structured data, clear category language, an about page that states plainly what you do	Ask an assistant to describe your company and check every fact
Credible	Corroboration from sources other than you, named expertise, consistent facts across the web, real reviews and independent coverage	Ask what the assistant knows about your reputation and sources
Citable	Self-contained passages, dated and attributed content, sourced claims, no overclaiming, facts stated plainly rather than implied	Check whether answers quote you or paraphrase without attribution
Recommendable	Being a genuinely good answer to a specific question, with standing in that topic and third-party validation to support it	Track recommendation prompts over time and record who is named

Building the prompt universe

You cannot track AI visibility without deciding what to track. This is the foundational step and most teams skip it.

#	STEP	OUTPUT
1	List the decisions your buyer makes where you should be a candidate	A decision list, not a keyword list
2	For each decision, write how someone would ask an assistant about it	Conversational prompts in natural phrasing
3	Add category prompts — "best X for Y", "X alternatives", "how to choose X"	The prompts where recommendations actually happen
4	Add brand prompts — "what is [brand]", "is [brand] good for [use case]"	Prompts that test understanding and sentiment
5	Add competitor prompts — "[competitor] vs", "alternatives to [competitor]"	Where you should appear but may not
6	Add problem prompts in customer language, from the Question Bank	Early-stage prompts where a first impression forms
7	Cap the set at what you will genuinely re-run on a cadence	Usually 30–60 prompts. Tracking 300 once is worth less than 40 monthly

The AI visibility audit

Run the prompt universe across the assistants your buyers use. Record the same fields every time.

WHAT TO RECORD	WHY IT MATTERS
The exact prompt used	Small wording changes produce different answers — comparability requires consistency
Which assistant and when	Behaviour differs between systems and changes over time
Whether you were mentioned	The base visibility question
Whether you were cited with a link	Mention and citation are different outcomes
Which page was cited	Tells you what kind of content actually gets used
Which competitors were named	Your share of the answer, and who owns the rest
What was said about you	Accuracy matters as much as presence
Any factual error	The highest-priority finding in the whole audit
Which sources the answer drew on	Shows where to build presence next

The brand fact consistency audit

Systems weigh corroboration. If your own facts disagree with each other, none of them are safe to repeat.

TEMPLATE — BRAND FACT SHEET

Brand Fact Sheet — the single source of truth

LEGAL NAME: [exact] TRADING NAME: [exact]

FOUNDED: [year] HEADQUARTERS: [city, country]

CATEGORY: [how we describe what we are, in one phrase]

WHAT WE DO: [one sentence, used verbatim everywhere]

WHO WE SERVE: [one sentence]

KEY PRODUCTS: [names, spelled and capitalized consistently]

PRICING MODEL: [structure, not necessarily figures]

FOUNDERS AND LEADERSHIP: [names and roles as publicly stated]

NOTABLE FACTS: [only what is publicly verifiable]

CHECK THESE PLACES MATCH IT, EXACTLY:

☐ Homepage and about page ☐ Structured data on the site

☐ LinkedIn company page ☐ Crunchbase or equivalent

☐ Business Profile listings ☐ Directory and marketplace profiles

☐ Press releases and coverage ☐ Founder and team social profiles

☐ Review platform profiles ☐ Partner and integration pages

RULE: one canonical wording per fact. Where a source is wrong, correct it at the source rather than adding a third version.

The citation-gap and source-authority analysis

When an assistant answers a question in your category, where does it get the answer?

#	STEP	WHAT IT TELLS YOU
1	Run your category prompts and record every source cited	The set of places these systems actually draw on
2	Group sources by type — publisher, directory, community, vendor, review site	Where the category’s credibility currently lives
3	Check which of those sources mention you, and how	Your presence in the pool the answer is built from
4	Identify sources that cover competitors but not you	The citation gap — usually the most actionable finding
5	Assess which gaps are legitimately closable	Editorial coverage, directory listings, community presence, review platforms
6	Rank by effort and by how often the source appears	A prioritized authority backlog that feeds Section 08

CONTENT, KEYWORDS, BACKLINKS, REPORTING

AI Visibility Database

PROPERTY	TYPE	WHAT IT HOLDS
Prompt	Title	The exact wording, never edited between runs
Prompt type	Select	Category / brand / competitor / problem / comparison
Assistant	Select	Which system was used
Run date	Date	Every run is a separate entry
Mentioned	Checkbox	Were we named at all
Cited	Checkbox	Was a link or source attributed to us
Page cited	Relation	Content database
Competitors named	Multi-select	Who else appeared
What was said	Text	The relevant passage, copied verbatim
Factual errors	Text	Anything incorrect — highest priority
Sources drawn on	Text	Where the answer came from
Action	Relation	Opportunities database
Owner	Person	Who runs this prompt set

PURPOSE

Every tracked prompt, every run, and what came back.

REVIEW

Turn AI visibility from an impression into a tracked, comparable measurement.

RELATED

Monthly, same prompts, same wording

The AI visibility scorecard

Five measures. Track the trend rather than the absolute number — the absolute number is not comparable to anyone else’s.

MEASURE	HOW TO CALCULATE IT	WHAT TO WATCH
Mention rate	Prompts where you appear / total prompts run	Trend across months, at fixed prompt wording
Citation rate	Prompts where you are cited with a source / total prompts	The gap between mention and citation — that gap is extractability
Share of answer	Your mentions / total brand mentions across the set	Who owns the category answer, and whether that is shifting
Accuracy rate	Answers about you with zero factual errors / answers about you	Anything below full accuracy is an entity consistency problem
Recommendation presence	Recommendation prompts where you are among those named	The hardest measure to move, and the one that matters commercially

THE HONEST STATEMENT ABOUT ALL OF THIS

No action described in this section guarantees that an AI system will find, cite, or recommend your brand. These systems are opaque, they change without notice, and two runs of the same prompt on the same day can differ. What you can do is remove the reasons a system would not use you: be indexable, be factually consistent, be corroborated by others, be extractable, and be genuinely good on a specific question. Then measure, on a fixed cadence, and let the trend tell you whether it is working.

The entity consistency checklist

- ☐ One canonical brand fact sheet exists and has an owner
- ☐ Company description is identical on site, LinkedIn, and every directory
- ☐ Product names are spelled and capitalized the same way everywhere
- ☐ Organization structured data is present and matches the about page
- ☐ Founder and leadership names match across all public profiles
- ☐ Category description is consistent — you are one thing, not four
- ☐ Outdated descriptions from old positioning have been corrected at source
- ☐ Third-party profiles are claimed rather than abandoned
- ☐ Where a source states something wrong, a correction has been requested
- ☐ The fact sheet is reviewed quarterly and versioned

The SEO Brain Command Center

One page. Six blocks. Twenty-eight view configurations. This is what the team opens each morning and what leadership opens each month.

The Command Center holds no data of its own. Every block is a linked view of a database specified in Section 05, filtered to answer one question. That is the discipline that keeps it useful: if a block does not answer a question somebody actually asks, it is removed rather than tolerated.

EX Executive Snapshot

DAILY GLANCE

BLOCK	SOURCE DATABASE	FILTER	SORT / GROUP	VISIBLE PROPERTIES	F O R
Current priorities	SEO Strategy	Period is this quarter	Sort by Bet order	Bet, Signal, Owner	L e a d e r s h i p
Critical issues	Technical Issues	Severity is Critical and Status is not Verified	Sort by Priority score desc	Issue, Module, Owner, Days open	E v e r y o n e
Content in production	Content	Status is Drafting or Review	Group by Status	Title, Author, Due	C o n t e n t l e a d
Pages declining	Content	Performance is Declining	Sort by Age desc	Title, Age, Next action	C o n t e n t l e a d
AI visibility changes	Historical Scorecards	Month is this or last month	Sort by Month desc	Mention rate, Citation rate, Coverage	L e a d e r s h i p
Authority opportunities	Authority Targets	Status is Qualified	Sort by Relevance desc	Target, Asset, Owner	O u t r e a c h w n e r
Decisions awaiting approval	Decisions & Change Log	Verdict is empty and Review date <= today	Sort by Review date	Decision, Evidence, Decided by	S E O l e a d

TH This Week

MONDAY PLANNING, DAILY CHECK

BLOCK	SOURCE DATABASE	FILTER	SORT / GROUP	VISIBLE PROPERTIES	F O R
Tasks due	All databases with a Due property	Due within 7 days and Status is not Done	Group by Owner	Task, Owner, Due, Priority	E v e r y o n e
Blocked	All databases	Status is Review or Waiting	Group by blocking person	Item, Blocker, Days waiting	S E O l e a d
Linked to a goal	All databases	Business goal is not empty and Due within 7 days	Group by Business goal	Item, Goal, Owner	L e a d e r s h i p
SOPs due	SOP Library	Next due within 7 days	Sort by Next due	SOP, Owner, Cadence, Last run	E v e r y o n e

OP Opportunities

WEEKLY WORKING SESSION

BLOCK	SOURCE DATABASE	FILTER	SORT / GROUP	VISIBLE PROPERTIES	F O R
Quick wins	AI Visibility Opportunities + Internal Link Opportunities	Effort is Hours and Status is not Done	Sort by Priority desc	Item, Area, Effort, Owner	E v e r y o n e
Content gaps	Question Bank	Answered on is empty	Sort by Frequency desc	Question, Frequency, Priority	C o n t e n t l e a d
Technical opportunities	Technical Issues	Action is Fix and Status is Diagnosed	Sort by Priority score desc	Issue, Score, Effort	T e c h n i c a l S E O o w n e r
Internal-link opportunities	Internal Link Opportunities	Status is Approved	Sort by Priority desc	Source, Target, Anchor, Priority	E d i t o r
Authority opportunities	AI Citations & Sources	Gap is Yes and Closable is Yes	Sort by Citation frequency desc	Source, Route, Owner	O u t r e a c h w o n e r
AI visibility opportunities	AI Visibility Opportunities	Priority is Now	Sort by Frequency desc	Opportunity, Gap type, Area, Owner	S E O l e a d

BLOCK	SOURCE DATABASE	FILTER	SORT / GROUP	VISIBLE PROPERTIES	F O R
Indexation incidents	Technical Issues	Module is Indexation and Severity is Critical or High	Sort by Created desc	Issue, Affected URLs, Owner	T e c h n i c a l S E O w n e r
Content decay	Content Refresh Queue	Completed is empty and Priority is High	Sort by Decline started	Page, Cause type, Decision	C o n t e n t l e a d
Unsupported claims	Content	Status is Review and QA flag is Unsourced claim	Sort by Due	Title, Author, Reviewer	E d i t o r
Brand-fact inconsistencies	Brand Facts	Status is Conflicted	Sort by Fact type	Fact, Canonical value, Conflict	M a r k e t i n g w n e r
Inaccurate AI answers	AI Prompt Runs	Factual error is true	Sort by Run date desc	Prompt, Assistant, Error detail	M a r k e t i n g w n e r

BLOCK	SOURCE DATABASE	FILTER	SORT / GROUP	VISIBLE PROPERTIES	F O R
Lost citations	Backlinks	Lost is true and Last seen within 90 days	Sort by Last seen desc	Source, Target, Reclamation action	O u t r e a c h w n e r
Local inconsistencies	Local Citations	Matches canonical is false	Group by Location	Directory, Difference, Status	L o c a l S E O w n e r

PE Performance

WEEKLY AND MONTHLY

BLOCK	SOURCE DATABASE	FILTER	SORT / GROUP	VISIBLE PROPERTIES	F O R
Business goals	Business Goals	Status is Active	Sort by Period	Goal, Primary metric, Baseline, Owner	L e a d e r s h i p
Search visibility	Keyword Master	Status is Ranking	Group by Stage	Keyword, Position, Position date	S E O l e a d
Qualified organic acquisition	Reports	Period is latest	Sort by Period desc	Metric, Value, Source, As-of date	L e a d e r s h i p
AI mention and citation trend	Historical Scorecards	All	Sort by Month asc	Month, Mention rate, Citation rate, Share	L e a d e r s h i p
Content performance	Content	Status is Published	Group by Performance	Title, Performance, Age, Next action	C o n t e n t l e a d
Technical health	Technical Issues	Status is not Verified	Group by Severity	Count by severity	T e c h n i c a l S E O w n e r

BLOCK	SOURCE DATABASE	FILTER	SORT / GROUP	VISIBLE PROPERTIES	F O R
Authority growth	Backlinks	First seen within 90 days	Group by Type	Source, Type, Campaign	M a r k e t i n g l e a d

DE Decisions

WEEKLY

BLOCK	SOURCE DATABASE	FILTER	SORT / GROUP	VISIBLE PROPERTIES	F O R
Decision required	Decisions & Change Log	Status is Recorded and Verdict is empty	Sort by Review date	Decision, Evidence, Decided by, Review date	S E O l e a d
Evidence attached	Decisions & Change Log	Evidence is empty	Sort by Date desc	Decision, Decided by	S E O l e a d
Deadline passed	Decisions & Change Log	Review date < today and Verdict is empty	Sort by Review date	Decision, Review date, Owner	S E O l e a d
Consequence of no decision	Decisions & Change Log	Status is Under review	Sort by Review date	Decision, Context, Expected outcome	L e a d e r s h i p

BUILD THE RISK CENTER FIRST

Of the six blocks, the Risk Center is the one that changes behaviour fastest. It surfaces the seven things that quietly cost the most — indexation incidents, decay, unsupported claims, brand-fact conflicts, inaccurate AI answers, lost citations, and local drift — none of which anyone goes looking for voluntarily. Build it in week one, even if the rest of the Command Center waits.

SECTION 09 · TECHNICAL MODULES

The twenty technical modules

Technical SEO is not one discipline. It is twenty, each with its own checks, evidence, and failure mode. Every issue record in Database 16 is tagged to one of these.

#	MODULE	WHAT IT CHECKS	EVIDENCE REQUIRED	TYPICAL SEVERITY	CLASSIC FAILURE
0 1	Crawling	Robots.txt directives, crawl rate, crawl budget distribution, blocked resources	Crawl export, log file sample	High	Crawlers spending budget on filters instead of money pages
0 2	Indexation	Index coverage, noindex directives, canonical conflicts, thin-page exclusion	Search Console coverage report, crawl export	Critical	A template carrying an accidental noindex after a release
0 3	Rendering and JavaScript	Server-rendered content, hydration, lazy-loaded copy, blocked scripts	Rendered vs raw HTML comparison	Critical	Pillar content only present after JS execution
0 4	Canonicals and duplication	Self-referencing canonicals, cross-domain canonicals, parameter duplicates	Crawl export with canonical column	High	Canonicals pointing at a redirected URL
0 5	Redirects and status codes	Chains, loops, soft 404s, 302s used permanently	Crawl export, server log sample	High	A migration leaving three-hop chains
0 6	XML sitemaps	Coverage, freshness, only indexable URLs, correct lastmod	Sitemap file, index report	Medium	Sitemaps containing noindexed or redirected URLs
0 7	Robots directives	Meta robots, X-Robots-Tag, conflicts between them	Header inspection, crawl export	High	A header directive contradicting the meta tag
0 8	Site architecture	Click depth, hub coverage, orphan pages, navigation stability	Crawl depth report	High	Money pages at depth five behind pagination
0 9	Core Web Vitals	Field data on key templates, layout stability, interaction delay	Field data export by template	Medium	Optimising lab scores while field data is unchanged
1 0	Mobile	Real-device rendering, tap targets, viewport, parity of content	Device testing notes, screenshots	High	Content hidden on mobile that ranks on desktop
1 1	Images	Compression, dimensions, formats, alt text accuracy, lazy loading	Crawl export with image data	Low	Alt text stuffed with keywords rather than describing the image
1 2	Structured data	Validity, match with visible content, required properties	Validator output, page source	Medium	FAQ markup for questions not visible on the page
1 3	International SEO	Hreflang completeness and reciprocity, locale detection, canonical interaction	Hreflang audit export	High	Hreflang tags that are not reciprocal
1 4	Faceted navigation	Indexable combinations, crawl traps, parameter handling	Crawl export, parameter list	Critical	Thousands of facet URLs indexable by default
1 5	Pagination	Consistent handling, canonical treatment, content on page two onward	Crawl export of paginated sets	Medium	Page two canonicalised to page one, hiding content
1 6	Log-file analysis	Where crawlers actually spend time, response codes served to them	Raw server logs, filtered by user agent	High	Discovering half the crawl budget goes to a legacy directory
1 7	Migrations	URL inventory, redirect map, parity checks, staged verification	Pre-migration crawl, redirect map	Critical	A redirect map missing the long tail
1 8	Monitoring and incident response	Alerting thresholds, escalation route, first-hour checklist	Alert configuration, incident log	Critical	No alert on a sitewide index drop
1 9	Technical tickets	Testable acceptance criteria, scope boundaries, rollback plan	The ticket itself	Medium	Tickets rejected because the criteria are not testable
2 0	Release verification	Pre-deploy staging check, post-deploy crawl, 7-day and 30-day index checks	Release record, post-deploy crawl	High	A regression found six weeks after the deploy

How a module becomes a fix

The same eight steps regardless of which module the issue came from.

STEP	WHAT HAPPENS	RECORDED WHERE
1 · Detect	The module check finds something	Technical Issues, status Found
2 · Evidence	Crawl export, log sample, or screenshot attached	Evidence property
3 · Score	Impact, pages, business, confidence, effort entered	Five score properties
4 · Route	Confidence below 3 becomes a diagnostic, not a ticket	Action formula
5 · Ticket	Dev ticket written with testable acceptance criteria	Ticket URL
6 · Deploy	Engineering ships it in a named release	Releases & Verification
7 · Verify	Production check within 48 hours	Verified date
8 · Watch	7-day and 30-day index checks on the release	Release record

THE FAILURE THIS SEQUENCE EXISTS TO PREVENT

An SEO finds an issue, describes it in a message, engineering interprets it differently, ships something adjacent, and nobody checks afterwards. Six weeks later the issue is still open and both sides believe the other dropped it. Steps 2, 5, and 7 are the ones that prevent it — evidence, testable criteria, and a production verification with a date.

The AI visibility operating area in full

Thirty-five components across six groups. Section 12 of the original framework covered the method; this is the full map, with every part located in a database, an SOP, or a prompt.

Foundations

- AI visibility foundations
- What can and cannot be influenced
- Why mention and citation are different outcomes
- The limits of every measure in this section

Prompt design

- Prompt-universe creation
- Prompt taxonomy across six families
- Intent distribution across the set
- Branded and non-branded balance
- Recommendation prompts
- Comparison prompts
- Alternative prompts
- Informational prompts
- Purchase-intent prompts
- Prompt localization by market
- Prompt versioning and the wording lock

Tracking

- Engine tracking across assistants
- Mention tracking
- Citation tracking
- Answer coverage
- Share of AI voice
- Sentiment classification
- Position within the answer
- Competitor recommendations
- Sources cited
- Monthly tracking discipline
- Historical comparison

Inputs you control

- Brand facts and the canonical sheet
- Entity consistency across sources
- Content extractability
- Source authority
- Third-party validation
- Freshness and dating
- Structured facts and markup

Gap analysis

- Citation gaps

- Recommendation gaps
- Factual errors
- Source concentration
- Competitor overlap
- Opportunity prioritization

Output

- AI-ready content briefs
- Executive reporting
- What to report and what to withhold

The ten AI visibility measures, and what each cannot tell you

Every measure here is comparable only to your own previous months, at identical prompt wording and assistant set.

MEASURE	METHOD	WHAT IT CANNOT TELL YOU
Mention rate	Runs where we are named ÷ total runs	Whether the mention was favourable, or whether anyone saw it
Citation rate	Runs where we are cited with attribution ÷ total runs	Whether the cited page was the one you wanted used
Answer coverage	Prompts with a mention on any assistant ÷ prompts run	How consistently you appear — one assistant of three still counts
Share of AI voice	Our mentions ÷ all brand mentions in the set	Market share. It measures a set you defined, nothing more
Sentiment distribution	Manual classification of each answer about us	Anything reliable at small n — classification is a judgement
Recommendation rate	Recommendation prompts naming us ÷ recommendation prompts	Why you were or were not recommended
Competitor overlap	Prompts where both we and a competitor appear	Whether the comparison favoured either party
Source concentration	Distinct sources cited ÷ total citations observed	Whether those sources are reachable for you
Factual accuracy	Error-free answers about us ÷ answers about us	Errors on prompts you did not include in the set
Trend change	This month minus last month, per measure	Causation. A change may follow your work or something the system changed

THE HONEST FRAMING FOR EVERY NUMBER IN THIS AREA

These systems are opaque, they change without notice, and two runs of the same prompt on the same day can differ. Every measure above is a description of what you observed in a set you chose, not a measurement of a market. Report the trend and the method together, never the number alone, and never compare your figures to a published benchmark — no comparable source exists.

SECTION 17 · VERTICALS

How the system changes by business type

The architecture is the same for every business. What changes is which fields matter, which keyword types earn priority, and where the technical risk concentrates.

FIELD OR DECISION	B2B SAAS	ECOMMERCE	AGENCY	LOCAL SERVICE	PUBLISHER
Primary conversion	Trial or demo start	Add to basket	Enquiry form	Call or booking	Subscription or ad view
Highest-value keyword type	Comparison and alternatives	Product and category	Service plus outcome	Service plus location	Question and topic
Cluster anchor	Use-case pillar	Category page	Service page	Location page	Topic hub
Content depth driver	Buyer sophistication	Product attribute coverage	Proof of outcomes	Local specificity	Editorial completeness
Technical priority	Rendering and index coverage	Faceted navigation and duplication	Speed and lead capture	Local schema and mobile	Crawl budget and pagination
Internal link priority	Blog to feature and pricing	Category to product	Case study to service	Content to location page	Article to hub
Authority route	Original data and expert commentary	Product data and buying guides	Client outcomes and thought leadership	Local partnerships and press	Original reporting and syndication
AI prompt emphasis	Category and comparison prompts	Product recommendation prompts	Service selection prompts	Local recommendation prompts	Topical authority prompts
Reporting focus	Pipeline contribution	Revenue by category	Enquiry quality	Calls and bookings by location	Sessions and subscriptions

What each type should build first

BUSINESS TYPE	BUILD FIRST	WHY	WHAT TO DEFER
B2B SaaS	Keyword Master, Content, AI Prompt Universe	Comparison and category prompts are where your buyers form shortlists	Local databases entirely
Ecommerce	Technical Issues, Site Architecture, Internal Links	Faceted navigation and duplication cost more than content ever will	Digital PR until the technical layer is stable
Agency	SOP Library, Prompt Library, Client relation on every database	Repeatability across clients is the product	Deep per-client AI tracking until the retainer justifies it
Local service	Local Locations, Local Citations, Question Bank	NAP consistency predicts local visibility more than anything else	Topic clusters beyond the core services
Publisher	Topics, Topic Clusters, Content, Site Architecture	Topical authority and crawl efficiency are the whole game	Commercial-page linking, which barely applies

ONE SYSTEM, FIVE EMPHASES

Do not build five versions of the Brain. Build one, and use the Client or Segment property to filter it. The differences above are about sequencing and priority, not about structure — a local service business and a B2B SaaS company need the same Keyword Master, they just fill different columns of it first.

The thirty-two databases

Every database in the Brain, specified to the property. Build them in the order given in Section 18 — the relations depend on it.

Databases 01 to 16 cover strategy, demand, keywords, and content. Databases 17 to 32 cover technical, architecture, authority, local, AI visibility, and the operating layer. Eight of the thirty-two carry full formula, rollout, and example-record specifications; the remainder carry properties, relations, views, and workflow, which is everything needed to build them.

THE FICTIONAL REFERENCE COMPANY

Every example record in this document belongs to Northstar Analytics, an invented B2B analytics SaaS company selling to data teams at mid-market software businesses. Northstar does not exist. Every figure, URL, keyword, and result shown for it is a placeholder written to show the shape of a completed record, not a claim about anything real.

DB 01 Business Goals

Purpose	Hold the commercial objectives SEO is serving, so every downstream record can trace to one.
Primary owner	Business owner or founder
Used by	Strategy, Reports, Content, Keywords
Inputs	Company plan, revenue targets, quarterly objectives
Outputs	The goals every SEO metric maps back to

PROPERTIES

PROPERTY	TYPE	VALUES / DEFAULT	R	DEFINITION	EXAMPLE
Goal	Title	Free text	R	The commercial objective in business language	Grow qualified trial starts
Period	Select	Q1-Q4 / Annual	R	The window it applies to	Q3
Owner	Person	One person	R	Who is accountable in the business	Head of Growth
SEO outcome	Text	Free text	R	What SEO must produce to serve it	Visibility on high-intent comparison queries
Primary metric	Text	Free text	R	The measurable business number	Trial starts from organic
Leading indicator	Text	Free text	·	What moves first	Impressions on BOFU pages
Baseline	Number	Empty if unknown	·	The starting figure	38
Baseline date	Date	—	·	When the baseline was taken	2026-07-01
Status	Select	Active / Met / Missed / Retired	R	Current state	Active

RELATIONS

RELATION	DESTINATION	CARDINALITY	PURPOSE
Strategy bets	SEO Strategy	One-to-many	Show which bets serve this goal
Content	Content	One-to-many	Trace published pages to a commercial goal
Reports	Reports	One-to-many	Show progress against the goal over time

VIEWS

VIEW	TYPE	FILTER	SORT / GROUP	FOR	DECISION IT SUPPORTS
Active goals	Table	Status is Active	Sort by Period	Everyone	What are we actually trying to achieve this quarter
Goal progress	Board	Status is Active	Group by Owner	Leadership	Who owns what, and is anything unowned
Retired	Table	Status is Met or Missed	Sort by Period descending	SEO lead	What did we set out to do last year and what happened

WORKFLOW

Status values	Draft / Active / Met / Missed / Retired
Transitions	Draft to Active on approval; Active to Met or Missed at period end only
Blocking conditions	Cannot become Active without a primary metric and a named owner
Definition of done	The goal has a metric, an owner, a baseline with a date, and a period

TEMPLATE, CADENCE, ARCHIVE, GATE

Database template	Goal statement, metric definition, baseline capture instructions, linked bets checklist
Review cadence	Quarterly
Archive rule	Retire rather than delete; keep for year-on-year comparison
Quality gate	No goal is Active without a measurable primary metric

DB 02

SEO Strategy

Purpose	The versioned strategy record: the bets, their signals, and the not-doing list.
Primary owner	SEO lead
Used by	Everyone; referenced by every operating area
Inputs	Business goals, maturity assessment, performance history
Outputs	A one-page strategy anyone can act from

PROPERTIES

PROPERTY	TYPE	VALUES / DEFAULT	R	DEFINITION	EXAMPLE
Version	Title	e.g. 2026-Q3-v2	R	Dated version identifier	2026-Q3-v2
Period	Select	Quarter	R	The quarter it governs	Q3 2026
Business goal	Relation	Business Goals	R	What it serves	Grow qualified trials
Bet 1 / 2 / 3	Text ×3	Free text	R	The three commitments	Own the comparison cluster
Signal 1 / 2 / 3	Text ×3	Free text	R	How we will know each worked	Rankings on 6 comparison queries
Not doing	Text	Free text	R	Explicitly out of scope this quarter	No new TOFU cluster
Constraints	Text	Free text	R	Capacity, budget, dev time	2 pieces/month, no dev sprint
Maturity score	Number	1–5	·	From the maturity assessment	2
Weakest dimension	Text	Free text	·	Where to focus	Measurement
Approved by	Person	One person	R	Who signed it	Head of Growth
Approved date	Date	—	R	When	2026-06-24

RELATIONS

RELATION	DESTINATION	CARDINALITY	PURPOSE
Business goal	Business Goals	Many-to-one	Anchor the strategy in a commercial objective
Decisions	Decisions & Change Log	One-to-many	Show what changed and why
Experiments	Experiments	One-to-many	Link bets to the tests that inform them

VIEWS

VIEW	TYPE	FILTER	SORT / GROUP	FOR	DECISION IT SUPPORTS
Current	Table	Period is this quarter	Sort by Approved date descending	Everyone	What are we doing and not doing right now
History	Table	All	Sort by Period descending	SEO lead	How has the strategy changed, and did the bets pay
Bets and signals	Table	Period is this quarter	Show Bet and Signal properties only	Leadership	Are the commitments measurable

WORKFLOW

Status values	Draft / Approved / Superseded
Transitions	Draft to Approved on sign-off; Approved to Superseded when the next version is approved
Blocking conditions	Cannot be Approved without three bets, three signals, and a not-doing list
Definition of done	Three bets each with a measurable signal, a not-doing list, and an approver

TEMPLATE, CADENCE, ARCHIVE, GATE

Database template	Strategy one-pager structure with prompts for each section and a completion checklist
Review cadence	Quarterly
Archive rule	Supersede, never delete — the history is the value
Quality gate	Every bet has a signal that could be shown to be false

DB 03

Customer Problems

Purpose	The evidence layer: what customers say their problem is, in their own words.
Primary owner	Content lead
Used by	Keywords, Topics, Content, AI Visibility
Inputs	Call transcripts, reviews, tickets, interview notes
Outputs	Validated problems with distinct-person counts

PROPERTIES

PROPERTY	TYPE	VALUES / DEFAULT	R	DEFINITION	EXAMPLE
Problem	Title	Free text	R	Stated as a customer would state it	Reporting takes three days a month
Verbatim quotes	Text	One per line	R	Exact wording, never paraphrased	"I lose the first week of every month to this"
Source	Select	Call / Interview / Review / Ticket / Community	R	Where it came from	Call
Source date	Date	—	R	When it was said	2026-05-14
Segment	Multi-select	Your segments	·	Who raised it	Mid-market SaaS
Role	Text	Free text	·	The role that raised it	Data lead
Distinct people	Number	Integer	R	Separate people, not mentions	7
Buyer stage	Select	Problem-aware / Solution-aware / Vendor-aware	R	Where in the journey	Problem-aware
Evidence strength	Formula	Derived	R	Strong / moderate / single	Strong
Status	Select	New / Validated / Covered / Declined	R	Current state	Validated

RELATIONS

RELATION	DESTINATION	CARDINALITY	PURPOSE
Search demand	Search Demand	One-to-many	Link the problem to how it is searched
Questions	Question Bank	One-to-many	Link to the questions it generates
Topics	Topics	Many-to-many	Show which clusters address it
Voice of customer	Voice of Customer	One-to-many	Attach the full quote records

ROLLUPS

ROLLUP	VIA RELATION	PROPERTY	CALCULATION	WHY IT MATTERS
Content coverage	Topics	Published count	Sum	Shows whether a validated problem is actually addressed

FORMULAS

EVIDENCE STRENGTH

```
if(prop("Distinct people") >= 5, "Strong",
  if(prop("Distinct people") >= 3, "Moderate", "Single data point"))
```

Thresholds are a judgement call, not a statistical rule. Raise them if your sample sizes are large; lower them only if you are prepared to act on thin evidence.

VIEWS

VIEW	TYPE	FILTER	SORT / GROUP	FOR	DECISION IT SUPPORTS
Validated problems	Table	Status is Validated	Sort by Distinct people descending	Content lead	Which problems are worth building content around
Uncovered	Table	Status is Validated and Topics is empty	Sort by Distinct people descending	Content lead	Which validated problems have no content addressing them
By segment	Board	Status is not Declined	Group by Segment	Strategy	Does one segment have problems we systematically ignore
Recent evidence	Table	Source date within 90 days	Sort by Source date descending	Everyone	What are customers saying right now

WORKFLOW

Status values	New / Validated / Covered / Declined
Transitions	New to Validated at 3+ distinct people; Validated to Covered when a page publishes
Blocking conditions	Cannot be Validated without at least one verbatim quote and a source date
Definition of done	Verbatim quotes present, source and date recorded, distinct-person count entered

TEMPLATE, CADENCE, ARCHIVE, GATE

Database template	Problem record with quote field, source picker, distinct-person prompt, and a validation checklist
Review cadence	Quarterly deep pass, additions weekly
Archive rule	Declined records stay with the reason; nothing is deleted
Quality gate	No record without at least one verbatim quote

EXAMPLE RECORDS — ALL FICTIONAL

PROBLEM	DISTINCT PEOPLE	SOURCE / DATE	STAGE	STATUS
Monthly reporting takes three days	7	Call · 2026-05-14	Problem-aware	Validated
Cannot tell which dashboards are used	4	Interview · 2026-04-02	Solution-aware	Validated
Migration from legacy tool feels risky	2	Review · 2026-06-11	Vendor-aware	New

Fictional illustrative records. Northstar Analytics is an invented company; every figure shown is a placeholder.

DB 04

Voice of Customer

Purpose	Individual verbatim quotes, retrievable by theme, role, and stage for use in copy.
Primary owner	Content lead
Used by	Content, On-page, Authority, AI Visibility
Inputs	Transcripts, reviews, community threads
Outputs	A searchable language bank

PROPERTIES

PROPERTY	TYPE	VALUES / DEFAULT	R	DEFINITION	EXAMPLE
Quote	Title	Verbatim	R	Exact wording, never cleaned up	"I lose the first week of every month to this"
Theme	Multi-select	Your themes	R	What it is about	Reporting effort
Role	Text	Free text	R	Who said it	Data lead
Segment	Select	Your segments	·	Their segment	Mid-market SaaS
Stage	Select	Problem / Solution / Vendor-aware	·	Journey stage	Problem-aware
Source	Text	Free text	R	Where exactly	Discovery call, account 412
Date	Date	—	R	When	2026-05-14
Permission to publish	Checkbox	Default false	R	Cleared for external use	false
Used in	Relation	Content	·	Where it has been used	Reporting guide

RELATIONS

RELATION	DESTINATION	CARDINALITY	PURPOSE
Problem	Customer Problems	Many-to-one	Group quotes under the problem they evidence
Content	Content	Many-to-many	Track where a quote has been used

VIEWS

VIEW	TYPE	FILTER	SORT / GROUP	FOR	DECISION IT SUPPORTS
Publishable bank	Gallery	Permission to publish is true	Group by Theme	Writers	Which exact phrases can appear in copy
Unused	Table	Used in is empty and Permission is true	Sort by Date descending	Writers	Strong language nobody has used yet
By role	Board	All	Group by Role	Content lead	How does each buying role describe this differently

WORKFLOW

Status values	Captured / Cleared / Published / Retired
Transitions	Captured to Cleared when permission is confirmed; Cleared to Published on first use
Blocking conditions	Cannot be Cleared without a recorded permission decision
Definition of done	Verbatim quote, role, source, date, and a permission decision

TEMPLATE, CADENCE, ARCHIVE, GATE

Database template	Quote record with a permission checklist and theme picker
Review cadence	Monthly
Archive rule	Retire quotes older than 24 months rather than deleting them
Quality gate	No quote is published without a recorded permission decision

DB 05 Search Demand

Purpose	What people search and ask, with evidence, before it becomes a keyword list.
Primary owner	SEO specialist
Used by	Keywords, Topics, AI Visibility
Inputs	Problem records, tool exports, assistant queries, site search logs
Outputs	Validated demand records with intent and surface

PROPERTIES

PROPERTY	TYPE	VALUES / DEFAULT	R	DEFINITION	EXAMPLE
Query	Title	As typed or asked	R	The actual phrasing	best analytics tool for data teams
Demand type	Select	Established / Emerging / Conversational / Zero-click	R	What kind of demand	Established
Evidence	Text	Free text	R	Where you saw it, with source	GSC export 2026-06
Volume	Number	Empty if unknown	-	Never estimated	390
Volume source	Text	Tool name	-	Which tool reported it	Tool X, 2026-06
Intent	Select	Informational / Comparison / Commercial / Navigational	R	What the searcher wants	Commercial
Surface	Multi-select	Google / AI answers / Community / Video / Marketplace	R	Where it is asked	Google, AI answers
Status	Select	Unvalidated / Validated / Mapped / Rejected	R	State	Validated

RELATIONS

RELATION	DESTINATION	CARDINALITY	PURPOSE
Problem	Customer Problems	Many-to-one	Tie demand to evidenced problems
Keyword	Keyword Master	One-to-one	Promote validated demand into the keyword list
Prompt	Prompt Universe	One-to-many	Conversational demand becomes an AI prompt

VIEWS

VIEW	TYPE	FILTER	SORT / GROUP	FOR	DECISION IT SUPPORTS
Ready to promote	Table	Status is Validated and Keyword is empty	Sort by Volume descending	SEO specialist	What should enter the keyword list next
Emerging	Table	Demand type is Emerging	Sort by Evidence date descending	Strategy	What is new that tools have not caught up with
Conversational	Table	Demand type is Conversational	Group by Intent	AI visibility owner	What should enter the prompt universe

WORKFLOW

Status values	Unvalidated / Validated / Mapped / Rejected
Transitions	Unvalidated to Validated with evidence; Validated to Mapped when a keyword record exists
Blocking conditions	Cannot be Validated without an evidence entry naming a source
Definition of done	Evidence recorded, intent classified, surface identified

TEMPLATE, CADENCE, ARCHIVE, GATE

Database template	Demand record with evidence prompt and intent classification guidance
Review cadence	Monthly
Archive rule	Rejected records keep their reason
Quality gate	No volume figure without a named source and date

DB 06 Question Bank

Purpose	Every question a buyer asks, in the form they ask it, and whether you answer it.
Primary owner	Content lead
Used by	Content, On-page, AI Visibility
Inputs	Calls, community threads, support tickets, assistant queries
Outputs	A gap list that drives FAQ, on-page, and prompt work

PROPERTIES

PROPERTY	TYPE	VALUES / DEFAULT	R	DEFINITION	EXAMPLE
Question	Title	As asked	R	Exact phrasing	How long does implementation take?
Asked by	Select	Role or segment	R	Who asks it	Data lead
Stage	Select	Problem / Solution / Vendor-aware / Customer	R	Journey stage	Vendor-aware
Frequency	Number	Integer	·	Distinct people who asked	11
Answered on	Relation	Content	·	Blank means a gap	Implementation guide
Answer extractable	Checkbox	Default false	R	Could a system quote it safely	true
Answer location	Text	Free text	·	Where on the page	First H2, para 1
Priority	Select	High / Medium / Low	R	Commercial proximity	High
Source	Text	Free text	R	Where it was heard	Sales calls, Q2

RELATIONS

RELATION	DESTINATION	CARDINALITY	PURPOSE
Problem	Customer Problems	Many-to-one	Group questions under their problem
Content	Content	Many-to-one	Show where it is answered
Prompt	Prompt Universe	One-to-one	High-frequency questions become tracked prompts

VIEWS

VIEW	TYPE	FILTER	SORT / GROUP	FOR	DECISION IT SUPPORTS
Unanswered	Table	Answered on is empty	Sort by Frequency descending	Content lead	What are we not answering that people keep asking
Answered but not extractable	Table	Answered on is not empty and Extractable is false	Sort by Priority	Editor	Which answers cannot safely be quoted
By stage	Board	All	Group by Stage	Content lead	Is our coverage skewed to one part of the journey

WORKFLOW

Status values	Captured / Answered / Extractable / Retired
Transitions	Captured to Answered when a page covers it; Answered to Extractable on QA pass
Blocking conditions	Cannot reach Extractable without the QA checklist in Section 07 being run
Definition of done	Question captured verbatim, answered on a named page, and quotable in isolation

TEMPLATE, CADENCE, ARCHIVE, GATE

Database template	Question record with source prompt, frequency counter, and extractability checklist
Review cadence	Monthly
Archive rule	Retire questions no longer asked, with a date
Quality gate	Answer extractable is only checked after a human runs the extractability checklist

EXAMPLE RECORDS — ALL FICTIONAL

QUESTION	FREQUENCY	ANSWERED ON	EXTRACTABLE
How long does implementation take?	11	Implementation guide	Yes
Do you support our warehouse?	8	— (gap)	No
What happens to data if we leave?	5	Security page	No

Fictional illustrative records. Northstar Analytics is an invented company; every figure shown is a placeholder.

DB 07

Keyword Master

Purpose	The single source of truth for what the company is trying to be found for.
Primary owner	SEO specialist
Used by	Content, Topics, Reporting, AI Visibility
Inputs	Search demand records, tool exports, current rankings
Outputs	A scored, ranked, mapped keyword list

PROPERTIES

PROPERTY	TYPE	VALUES / DEFAULT	R	DEFINITION	EXAMPLE
Keyword	Title	As searched	R	The query	analytics tool for data teams
Type	Select	16 types (Section 08)	R	Query family	Commercial
Intent	Select	Informational / Comparison / Commercial / Navigational	R	What it wants	Commercial
Stage	Select	TOFU / MOFU / BOFU	R	Funnel position	BOFU
Branded	Checkbox	Default false	R	Branded or not	false
Volume	Number	Empty if unknown	·	Never estimated	390
Difficulty	Number	0-100	·	With the tool named	42
Tool	Text	Tool name	·	Which tool reported it	Tool X
S1 Business relevance	Number	0-5	R	Proximity to a purchase	5
S2 Search demand	Number	0-5	R	Evidenced demand	3
S3 Intent match	Number	0-5	R	Fits a page we can offer	5
S4 Difficulty (inverted)	Number	0-5	R	5 = easy to win	2
S5 Conversion potential	Number	0-5	R	Clear next step exists	5
S6 Current authority	Number	0-5	R	Do we already rank nearby	2
S7 Content gap	Number	0-5	R	Is there a real opening	4
S8 AI answer potential	Number	0-5	R	Extractable and factual	4
S9 Brand fit	Number	0-5	R	On-positioning	5
S10 Strategic value	Number	0-5	R	Anchors a cluster	4
Priority score	Formula	Derived	R	Weighted total out of 105	87
Band	Formula	Derived	R	Build now / this quarter / later / monitor / decline	Build now
Target page	Relation	Content	·	Exactly one owning page	Comparison hub
Cluster	Relation	Topic Clusters	·	Which cluster	Tool selection
Current position	Number	Integer	·	Latest tracked position	14
Position date	Date	—	·	As-of date	2026-08-01
Owner	Person	One person	R	Accountable	SEO specialist
Status	Select	Unmapped / Mapped / Published / Ranking / Declined	R	State	Mapped

RELATIONS

RELATION	DESTINATION	CARDINALITY	PURPOSE
Target page	Content	Many-to-one	One keyword, one owning page — enforced by review
Cluster	Topic Clusters	Many-to-one	Roll keyword coverage up to the cluster
Demand record	Search Demand	One-to-one	Preserve the evidence trail
Cannibalization conflicts	Content	Many-to-many	Record pages competing for this query

ROLLUPS

ROLLUP	VIA RELATION	PROPERTY	CALCULATION	WHY IT MATTERS
Cluster keyword count	Cluster	Keyword	Count	Shows whether a cluster has enough query coverage to be worth building
Cluster average score	Cluster	Priority score	Average	Tells you whether a cluster is commercially worth the build
Page keyword count	Target page	Keyword	Count	Anything above 1 primary keyword per page is a cannibalization risk

FORMULAS

PRIORITY SCORE

```
prop("S1 Business relevance") * 3 + prop("S2 Search demand") * 2
+ prop("S3 Intent match") * 3 + prop("S4 Difficulty (inverted)") * 2
+ prop("S5 Conversion potential") * 3 + prop("S6 Current authority") * 1
+ prop("S7 Content gap") * 2 + prop("S8 AI answer potential") * 2
+ prop("S9 Brand fit") * 1 + prop("S10 Strategic value") * 2
```

Maximum 105. The weights encode a typical B2B situation and are meant to be edited. Record every weight change with a date, or scores stop being comparable across quarters.

BAND

```
if(prop("Priority score") >= 80, "Build now",
  if(prop("Priority score") >= 60, "This quarter",
    if(prop("Priority score") >= 40, "When capacity allows",
      if(prop("Priority score") >= 20, "Monitor", "Decline"))))
```

Bands are thresholds on a judgement-based score. They make the build order consistent between people; they do not make it objectively correct.

VIEWS

VIEW	TYPE	FILTER	SORT / GROUP	FOR	DECISION IT SUPPORTS
Build order	Table	Status is Mapped and Band is Build now or This quarter	Sort by Priority score descending	Content lead	What do we write next
Unmapped	Table	Target page is empty and Band is not Decline	Sort by Priority score descending	SEO specialist	Which scored keywords have no page
Cannibalization risk	Table	Page keyword count > 1	Group by Target page	SEO specialist	Where are two keywords fighting over one page
BOFU coverage	Table	Stage is BOFU	Group by Status	Leadership	Are we covering the queries closest to revenue
Declined log	Table	Status is Declined	Sort by Modified descending	SEO lead	What did we refuse, and why

WORKFLOW

Status values	Unmapped / Mapped / Published / Ranking / Declined
Transitions	Unmapped to Mapped when a target page is set; Mapped to Published on publication; Published to Ranking when a position is recorded
Blocking conditions	Cannot be Mapped without all ten criterion scores entered; cannot be Published if the target page already owns another primary keyword
Definition of done	Ten scores entered, one target page assigned, cluster set, owner named

TEMPLATE, CADENCE, ARCHIVE, GATE

Database template	Keyword record pre-loaded with the ten scoring criteria, their definitions, and a mapping checklist
Review cadence	Weekly for status and positions, quarterly for re-scoring
Archive rule	Declined keywords stay with the decline reason; never delete
Quality gate	No keyword reaches Mapped with an unscored criterion

EXAMPLE RECORDS — ALL FICTIONAL

KEYWORD	INTENT / STAGE	SCORE / BAND	TARGET PAGE	STATUS
analytics tool for data teams	Commercial · BOFU	87 · Build now	Comparison hub	Mapped
how to reduce reporting time	Informational · TOFU	54 · When capacity allows	Reporting guide	Published
[competitor] alternative	Comparison · BOFU	79 · This quarter	— (unmapped)	Unmapped

Fictional illustrative records. Northstar Analytics is an invented company; every figure shown is a placeholder.

DB 08

Query Patterns

Purpose	The sixteen query types, their shapes, and which page type each requires.
Primary owner	SEO specialist
Used by	Keywords, Content, AI Visibility
Inputs	Keyword research, SERP observation
Outputs	A consistent classification everyone applies the same way

PROPERTIES

PROPERTY	TYPE	VALUES / DEFAULT	R	DEFINITION	EXAMPLE
Pattern	Title	Free text	R	The query family	Alternative
Shape	Text	Placeholder form	R	How it is phrased	[competitor] alternative
Intent	Select	Four intents	R	What it wants	Comparison
Stage	Select	TOFU / MOFU / BOFU	R	Funnel position	BOFU
Required page type	Text	Free text	R	What must be built	Alternatives page
AI answer likelihood	Select	High / Medium / Low	R	How often it is synthesized rather than clicked	High
Notes	Text	Free text	·	Observed behaviour	Frequently cited in AI answers

RELATIONS

RELATION	DESTINATION	CARDINALITY	PURPOSE
Keywords	Keyword Master	One-to-many	Classify every keyword by pattern

VIEWS

VIEW	TYPE	FILTER	SORT / GROUP	FOR	DECISION IT SUPPORTS
All patterns	Table	All	Group by Stage	SEO specialist	Which query families are we covering and which are we ignoring
High AI likelihood	Table	AI answer likelihood is High	Sort by Stage	AI visibility owner	Where should extractability work be concentrated
BOFU patterns	Table	Stage is BOFU	Sort by Intent	Content lead	Are the revenue-adjacent patterns covered

WORKFLOW

Status values	Reference / Retired
Transitions	Retired when a pattern stops appearing in your market
Blocking conditions	None — this is a reference database
Definition of done	All sixteen patterns present with a required page type

TEMPLATE, CADENCE, ARCHIVE, GATE

Database template	Pattern record with shape, intent, and page-type guidance
Review cadence	Quarterly
Archive rule	Retire rather than delete
Quality gate	Every pattern names the page type it requires

DB 09 Declined Keywords

Purpose	What you refused to pursue, and why — so the same debate does not reopen.
Primary owner	SEO lead
Used by	Keywords, Strategy
Inputs	Rejected keyword records and their reasoning
Outputs	A defensible refusal log

PROPERTIES

PROPERTY	TYPE	VALUES / DEFAULT	R	DEFINITION	EXAMPLE
Keyword	Title	As searched	R	The query declined	free analytics dashboard
Reason	Select	Off-ICP / No conversion path / Cannot win / Off-positioning / Capacity	R	Why it was refused	No conversion path
Detail	Text	Free text	R	The specific reasoning	Attracts free-tier only
Declined by	Person	One person	R	Who made the call	SEO lead
Date	Date	—	R	When	2026-06-30
Revisit if	Text	Free text	·	The condition that would change it	If we launch a self-serve tier
Revisit date	Date	—	·	When to reconsider	2027-01-15

RELATIONS

RELATION	DESTINATION	CARDINALITY	PURPOSE
Keyword	Keyword Master	One-to-one	Preserve the link to the original record

VIEWS

VIEW	TYPE	FILTER	SORT / GROUP	FOR	DECISION IT SUPPORTS
All declines	Table	All	Group by Reason	SEO lead	What have we consistently refused, and is the pattern still right
Due for revisit	Table	Revisit date is on or before today	Sort by Revisit date	SEO lead	Has anything changed that would reopen a decision
Recent	Table	Date within 90 days	Sort by Date descending	Everyone	What did we just decide not to do

WORKFLOW

Status values	Declined / Revisiting / Reinstated
Transitions	Declined to Revisiting on the revisit date; Revisiting to Reinstated only with new evidence
Blocking conditions	Cannot be Reinstated without stating what changed
Definition of done	A reason, a detail, a person, and a date

TEMPLATE, CADENCE, ARCHIVE, GATE

Database template	Decline record with reason picker and a revisit-condition prompt
Review cadence	Quarterly
Archive rule	Never delete — this log exists to prevent repeated debates
Quality gate	Every decline names a person and a date

DB 10

Topics

Purpose	The subject areas the company intends to be known for.
Primary owner	Content lead
Used by	Content, Keywords, AI Visibility
Inputs	Problems, demand, keyword clusters
Outputs	A prioritized topic list with coverage tracking

PROPERTIES

PROPERTY	TYPE	VALUES / DEFAULT	R	DEFINITION	EXAMPLE
Topic	Title	Customer language	R	The subject	Analytics tool selection
Belief	Text	Free text	R	What we argue about it	Selection should start from the reporting job, not the feature list
Audience	Select	Your segments	R	Who it serves	Data leads
Business outcome	Relation	Business Goals	R	What it supports	Qualified trials
Our evidence	Text	Free text	R	What we hold that others do not	Anonymized product usage patterns
Priority	Select	This quarter / Next / Later / Declined	R	Build order	This quarter
Owner	Person	One person	R	Accountable	Content lead
Authority status	Select	None / Building / Established	R	Where we stand	Building

RELATIONS

RELATION	DESTINATION	CARDINALITY	PURPOSE
Clusters	Topic Clusters	One-to-many	A topic may need more than one cluster
Problems	Customer Problems	Many-to-many	Which evidenced problems it addresses
Business goal	Business Goals	Many-to-one	Trace the topic to a commercial objective

ROLLUPS

ROLLUP	VIA RELATION	PROPERTY	CALCULATION	WHY IT MATTERS
Published pieces	Clusters	Published count	Sum	Shows real coverage rather than intent

VIEWS

VIEW	TYPE	FILTER	SORT / GROUP	FOR	DECISION IT SUPPORTS
This quarter	Board	Priority is This quarter	Group by Authority status	Content lead	What are we building authority on right now
No evidence	Table	Our evidence is empty	Sort by Priority	Strategy	Which topics would we be summarizing rather than contributing to
Established	Table	Authority status is Established	Sort by Topic	Leadership	What are we actually known for

WORKFLOW

Status values	Proposed / Active / Established / Declined
Transitions	Proposed to Active on prioritization; Active to Established when the cluster is complete and ranking
Blocking conditions	Cannot become Active without a stated belief and our own evidence
Definition of done	A belief, our evidence, an audience, an owner, and a linked business goal

TEMPLATE, CADENCE, ARCHIVE, GATE

Database template	Topic record with belief prompt, evidence prompt, and cluster planning checklist
Review cadence	Monthly
Archive rule	Declined topics keep the reason
Quality gate	No topic is Active if we have nothing of our own to say about it

DB 11

Topic Clusters

Purpose	The hub-and-spoke plan for each topic: pillar, spokes, commercial destination.
Primary owner	Content lead
Used by	Content, Keywords, Internal Links
Inputs	Topic records, mapped keywords
Outputs	A build plan with coverage tracking

PROPERTIES

PROPERTY	TYPE	VALUES / DEFAULT	R	DEFINITION	EXAMPLE
Cluster	Title	Free text	R	The cluster name	Tool selection cluster
Topic	Relation	Topics	R	Parent topic	Analytics tool selection
Pillar page	Relation	Content	R	The anchor	How to choose an analytics tool
Spokes planned	Number	Integer	R	Target count	9
Spokes published	Rollup	Derived	R	Actual count	4
Coverage	Formula	Derived	R	Percentage complete	44%
Commercial destination	Relation	Content	R	Where the cluster points	Pricing page
Comparison pages	Number	Integer	-	Vs and alternative pages planned	3
Owner	Person	One person	R	Accountable	Content lead
Status	Select	Planned / Building / Complete / Maintaining	R	State	Building

RELATIONS

RELATION	DESTINATION	CARDINALITY	PURPOSE
Topic	Topics	Many-to-one	Parent topic
Content	Content	One-to-many	All pieces in the cluster
Keywords	Keyword Master	One-to-many	All queries the cluster targets
Internal links	Internal Link Opportunities	One-to-many	Interlinking within the cluster

ROLLUPS

ROLLUP	VIA RELATION	PROPERTY	CALCULATION	WHY IT MATTERS
Spokes published	Content	Status	Count where Published	Coverage is the only honest measure of whether a cluster exists yet
Average position	Keywords	Current position	Average	Cluster-level performance rather than single-query noise

FORMULAS

COVERAGE

<pre>if(prop("Spokes planned") = 0, 0, round(prop("Spokes published") / prop("Spokes planned") * 100)) + "%"</pre>
--

A blunt completion measure. It says nothing about quality — a cluster at 100% coverage with thin pages is worse than one at 50% with strong ones.

VIEWS

VIEW	TYPE	FILTER	SORT / GROUP	FOR	DECISION IT SUPPORTS
Building	Board	Status is Building	Group by Owner	Content lead	Which clusters are in flight and who owns each
Incomplete	Table	Coverage below 60% and Status is Building	Sort by Coverage ascending	Content lead	Which clusters are stalled half-built
Missing commercial link	Table	Commercial destination is empty	Sort by Topic	SEO lead	Which clusters lead nowhere commercially

WORKFLOW

Status values	Planned / Building / Complete / Maintaining
Transitions	Planned to Building on first brief; Building to Complete at planned spoke count; Complete to Maintaining after the first refresh cycle
Blocking conditions	Cannot reach Complete without a commercial destination and a published pillar
Definition of done	Pillar published, planned spokes published, commercial destination linked, interlinking verified

TEMPLATE, CADENCE, ARCHIVE, GATE

Database template	Cluster plan with pillar brief, spoke list, and an interlinking checklist
Review cadence	Monthly
Archive rule	Maintaining clusters stay live; nothing is archived while ranking
Quality gate	No cluster is Complete without a linked commercial destination

DB 12

Content

Purpose	Every page, its purpose, its performance, and its next action.
Primary owner	Content lead
Used by	Everyone; the most referenced database in the system
Inputs	Briefs, drafts, performance data, refresh decisions
Outputs	A live inventory of what exists and what it needs

PROPERTIES

PROPERTY	TYPE	VALUES / DEFAULT	R	DEFINITION	EXAMPLE
Title	Title	Free text	R	The page	How to choose an analytics tool
URL	URL	Absolute	·	Blank if unpublished	https://example.com/choose
Primary keyword	Relation	Keyword Master	R	Exactly one	analytics tool selection
Cluster	Relation	Topic Clusters	R	Parent cluster	Tool selection
Role	Select	Pillar / Spoke / Comparison / Commercial / Reference	R	Its job in the cluster	Pillar
Stage	Select	TOFU / MOFU / BOFU	R	Funnel position	MOFU
Format	Select	Article / Comparison / Template / Tool / Study / Glossary	R	Content type	Article
Author	Person	One person	R	Named on the page	Head of Data
Credential shown	Checkbox	Default false	R	Is expertise visible	true
Published	Date	—	·	First publication	2026-03-11
Last updated	Date	—	·	Drives the refresh queue	2026-07-02
Age (days)	Formula	Derived	R	Since last update	35
Performance	Select	Growing / Stable / Declining / Never landed	R	Trend	Growing
Next action	Select	None / Update / Merge / Redirect / Remove	R	The decision	None
Extractable	Checkbox	Default false	R	Passed the extractability checklist	true
AI cited	Checkbox	Default false	·	Has appeared in a tracked AI answer	true
Owner	Person	One person	R	Accountable for the next action	Content lead
Status	Select	Brief / Drafting / Review / Published / Archived	R	State	Published

RELATIONS

RELATION	DESTINATION	CARDINALITY	PURPOSE
Primary keyword	Keyword Master	One-to-one	One page, one primary query
Cluster	Topic Clusters	Many-to-one	Roll coverage up to the cluster
Brief	Content Briefs	One-to-one	The commissioning document
Questions answered	Question Bank	One-to-many	What this page covers
Internal links in	Internal Link Opportunities	One-to-many	What points here
AI citations	AI Citations & Sources	One-to-many	Where this page has been cited
Refresh entries	Content Refresh Queue	One-to-many	Its refresh history

ROLLUPS

ROLLUP	VIA RELATION	PROPERTY	CALCULATION	WHY IT MATTERS
Inbound internal links	Internal links in	Status	Count where Implemented	The single most actionable number on a commercial page
Questions covered	Questions answered	Question	Count	Shows depth of coverage rather than word count
AI citation count	AI citations	Run date	Count	Which pages AI systems actually use

FORMULAS

AGE (DAYS)

```
if(empty(prop("Last updated")), toNumber(""),
    dateBetween(now(), prop("Last updated"), "days"))
```

Age alone does not mean a page needs updating. Use it with Performance — an old page that is still growing should be left alone.

VIEWS

VIEW	TYPE	FILTER	SORT / GROUP	FOR	DECISION IT SUPPORTS
Pipeline	Board	Status is Brief, Drafting, or Review	Group by Status	Content lead	What is in production and where is it stuck
Refresh queue	Table	Performance is Declining or Next action is Update	Sort by Age descending	Content lead	What needs updating before we write anything new
Under-linked money pages	Table	Role is Commercial and Inbound internal links < 5	Sort by Inbound internal links ascending	SEO specialist	Which commercial pages nothing points at
Not extractable	Table	Status is Published and Extractable is false	Sort by Stage	Editor	Which published pages cannot safely be quoted
AI cited	Table	AI cited is true	Group by Format	AI visibility owner	What kind of content do AI systems actually use
Never landed	Table	Performance is Never landed and Age > 180	Sort by Published	Content lead	What should be merged, redirected, or removed

WORKFLOW

Status values	Brief / Drafting / Review / Published / Archived
Transitions	Brief to Drafting on assignment; Drafting to Review on submission; Review to Published only on a QA pass; anything to Archived with a redirect recorded
Blocking conditions	Cannot reach Published without a QA pass and a named author with a visible credential; cannot be Archived without a redirect destination
Definition of done	Published with an author, a primary keyword, a cluster, a commercial link, and an extractability verdict

TEMPLATE, CADENCE, ARCHIVE, GATE

Database template	Content record pre-linked to its brief, with the QA checklist and the publishing handoff steps
Review cadence	Weekly for status, monthly for decay
Archive rule	Archive with a redirect destination; never delete a URL record
Quality gate	No page publishes with an unsourced claim or an unnamed author

EXAMPLE RECORDS — ALL FICTIONAL

TITLE	ROLE / STAGE	PERFORMANCE	NEXT ACTION	STATUS
How to choose an analytics tool	Pillar · MOFU	Growing	None	Published
Reporting time benchmark guide	Spoke · TOFU	Declining	Update	Published
[Competitor] vs Northstar	Comparison · BOFU	—	—	Drafting

Fictional illustrative records. Northstar Analytics is an invented company; every figure shown is a placeholder.

DB 13 Content Briefs

Purpose The commissioning document — no page is written without one.

Primary owner Content lead

Used by Writers, editors, SEO

Inputs Keyword, intent, evidence, SERP analysis, cluster plan

Outputs A brief a writer can work from without asking questions

PROPERTIES

PROPERTY	TYPE	VALUES / DEFAULT	R	DEFINITION	EXAMPLE
Brief	Title	Working title	R	The piece	Choosing an analytics tool
Primary keyword	Relation	Keyword Master	R	One query	analytics tool selection
Reader	Text	Free text	R	Who, and what they know	Data lead, evaluating options
Decision supported	Text	Free text	R	What they are deciding	Which tool to shortlist
Angle	Text	Free text	R	What we say that others do not	Start from the reporting job, not the feature grid
Our evidence	Text	Free text	R	What we hold	Anonymized usage patterns
Will not claim	Text	Free text	R	Explicit limits	No time-saving percentages
Required coverage	Text	Questions in order	R	What it must answer	5 questions listed
Extractable passages	Text	Free text	R	What must stand alone	The definition and the criteria list
Links in / out	Text	Free text	R	Internal linking plan	From 3 spokes; to pricing
Author	Person	One person	R	Who writes it	Head of Data
Reviewer	Person	One person	R	Who approves it	Editor
Due	Date	—	R	Deadline	2026-08-22
Status	Select	Draft / Approved / In production / Delivered	R	State	Approved

RELATIONS

RELATION	DESTINATION	CARDINALITY	PURPOSE
Content	Content	One-to-one	The page this brief produces
Keyword	Keyword Master	Many-to-one	The query it targets
Questions	Question Bank	Many-to-many	What it must answer

VIEWS

VIEW	TYPE	FILTER	SORT / GROUP	FOR	DECISION IT SUPPORTS
Awaiting approval	Table	Status is Draft	Sort by Due	Content lead	What is waiting on me before writing can start
In production	Table	Status is In production	Group by Author	Content lead	Who is writing what, and when is it due
No differentiating angle	Table	Angle is empty	Sort by Due	Content lead	Which briefs should not have been commissioned

WORKFLOW

Status values Draft / Approved / In production / Delivered

Transitions Draft to Approved on editor sign-off; Approved to In production on assignment

Blocking conditions Cannot be Approved without an angle, our evidence, and a will-not-claim list

Definition of done Angle differentiated, coverage listed as questions, extractable passages named, links specified

TEMPLATE, CADENCE, ARCHIVE, GATE

Database template Full brief template with every field pre-prompted and a completion checklist

Review cadence Per piece

Archive rule Delivered briefs stay attached to their page permanently

Quality gate No brief is approved without a differentiated angle

DB 14 Content Refresh Queue

Purpose	Declining and ageing pages, with a decided action for each.
Primary owner	Content lead
Used by	Writers, SEO specialist
Inputs	Performance history, current SERPs, cannibalization checks
Outputs	A prioritized refresh backlog with causes identified

PROPERTIES

PROPERTY	TYPE	VALUES / DEFAULT	R	DEFINITION	EXAMPLE
Page	Relation	Content	R	The declining page	Reporting time guide
Decline started	Date	—	R	When it began	2026-05-01
Evidence	Text	Free text	R	The performance data	Clicks -38% over 90d, GSC
Cause type	Select	Content / Technical / Intent drift / Cannibalization / Competitive	R	Diagnosed cause	Competitive
What overtook it	Text	Free text	·	What now ranks	3 newer guides with data
Decision	Select	Update / Merge / Redirect / Remove / Investigate	R	The action	Update
Specific changes	Text	Free text	R	What exactly to change	Add the usage data section, rewrite intro
Priority	Select	High / Medium / Low	R	By business importance	High
Owner	Person	One person	R	Who does it	Content lead
Completed	Date	—	·	When it shipped	
Recovered	Checkbox	Default false	·	Verified 60 days later	false

RELATIONS

RELATION	DESTINATION	CARDINALITY	PURPOSE
Page	Content	Many-to-one	The page being refreshed
Keyword	Keyword Master	Many-to-one	The query that declined

VIEWS

VIEW	TYPE	FILTER	SORT / GROUP	FOR	DECISION IT SUPPORTS
This month	Table	Completed is empty	Sort by Priority then Decline started	Content lead	What to refresh before writing anything new
By cause	Board	Completed is empty	Group by Cause type	SEO lead	Is the decline pattern content, technical, or competitive
Awaiting verification	Table	Completed is not empty and Recovered is false	Sort by Completed	Content lead	Did the refresh actually work

WORKFLOW

Status values	Identified / Diagnosed / Scheduled / Shipped / Verified
Transitions	Identified to Diagnosed when a cause type is set; Shipped to Verified 60 days after
Blocking conditions	Cannot be Scheduled without a cause type and specific changes listed
Definition of done	Cause diagnosed, specific changes listed, shipped, and verified after 60 days

TEMPLATE, CADENCE, ARCHIVE, GATE

Database template	Refresh record with a cause-diagnosis checklist and a change list prompt
Review cadence	Monthly
Archive rule	Verified entries stay as history for pattern analysis
Quality gate	No refresh is scheduled with "improve the content" as the change list

DB 15

Site Architecture

Purpose	The intended structure: hubs, depth, navigation, and where each template lives.
Primary owner	SEO specialist
Used by	Technical, Internal linking, Content
Inputs	Crawl exports, navigation audit, template inventory
Outputs	A documented intended architecture to audit against

PROPERTIES

PROPERTY	TYPE	VALUES / DEFAULT	R	DEFINITION	EXAMPLE
Node	Title	Free text	R	Section, hub, or template	Comparison hub
Type	Select	Hub / Category / Template / Utility	R	What it is	Hub
Intended depth	Number	Integer	R	Clicks from home	2
Actual depth	Number	Integer	·	From the last crawl	3
Depth gap	Formula	Derived	R	Actual minus intended	1
Parent	Relation	Site Architecture	·	Its parent node	Resources
In main nav	Checkbox	Default false	R	Linked from navigation	false
Pages	Number	Integer	·	How many pages sit under it	14
Owner	Person	One person	R	Accountable	SEO specialist
Status	Select	Intended / Live / Deprecated	R	State	Live

RELATIONS

RELATION	DESTINATION	CARDINALITY	PURPOSE
Parent	Site Architecture	Many-to-one	Self-relation building the tree
Content	Content	One-to-many	Pages within this node
Technical issues	Technical Issues	One-to-many	Issues affecting this section

FORMULAS

DEPTH GAP

<pre>if(empty(prop("Actual depth")), toNumber(""), prop("Actual depth") - prop("Intended depth"))</pre>

A positive gap means the page is deeper than intended. Anything above 1 on a commercially important node is worth a navigation conversation.

VIEWS

VIEW	TYPE	FILTER	SORT / GROUP	FOR	DECISION IT SUPPORTS
Too deep	Table	Depth gap > 0	Sort by Depth gap descending	SEO specialist	Which sections are buried deeper than intended
Tree	Table	All	Group by Parent	Everyone	What is the intended structure
Not in nav	Table	Type is Hub and In main nav is false	Sort by Pages descending	SEO lead	Which hubs have no navigational support

WORKFLOW

Status values	Intended / Live / Deprecated
Transitions	Intended to Live once built; Live to Deprecated on removal, with redirects recorded
Blocking conditions	Cannot be Deprecated without a redirect plan
Definition of done	Intended depth recorded, actual depth measured, gap reviewed

TEMPLATE, CADENCE, ARCHIVE, GATE

Database template	Node record with depth targets and a navigation checklist
Review cadence	Quarterly, and after any navigation change
Archive rule	Deprecated nodes stay with their redirect record
Quality gate	Every commercially important node has an intended depth of 3 or less

DB 16 Technical Issues

Purpose	Every known technical issue, scored, owned, ticketed, and tracked to verification.
Primary owner	Technical SEO owner
Used by	Engineering, SEO, Reporting
Inputs	Crawls, Search Console, log files, alerts
Outputs	Scored issues with tickets and verified fixes

PROPERTIES

PROPERTY	TYPE	VALUES / DEFAULT	R	DEFINITION	EXAMPLE
Issue	Title	One line	R	What is wrong	Facet URLs indexable
Module	Select	20 modules (Section 09)	R	Which technical area	Faceted navigation
Severity	Select	Critical / High / Medium / Low	R	Scope of impact	High
Affected URLs	Number	Integer	R	Count	2400
Evidence	Files	Crawl export, log, screenshot	R	Proof it exists	crawl-2026-08.csv
Impact	Number	1-5	R	Effect of fixing it	4
Pages	Number	1-5	R	1 = a handful, 5 = sitewide	4
Business	Number	1-5	R	Commercial importance of affected pages	3
Confidence	Number	1-5	R	How sure we are of the diagnosis	4
Effort	Number	1-5	R	1 = minutes, 5 = a project	2
Priority score	Formula	Derived	R	The calculated result	96
Action	Formula	Derived	R	Fix or diagnose first	Fix
Engineering owner	Person	One person	·	Who implements	Platform engineer
Ticket	URL	Absolute	·	Link to the tracker	https://tracker/ENG-4412
Deployed	Date	—	·	When it shipped	
Verified	Date	—	·	When it was confirmed working	
Status	Select	Found / Diagnosed / Ticketed / In progress / Deployed / Verified	R	State	Ticketed

RELATIONS

RELATION	DESTINATION	CARDINALITY	PURPOSE
Architecture node	Site Architecture	Many-to-one	Which section it affects
Release	Technical Releases & Verification	Many-to-one	Which deploy contained the fix
Content affected	Content	Many-to-many	Which pages are impacted

ROLLUPS

ROLLUP	VIA RELATION	PROPERTY	CALCULATION	WHY IT MATTERS
Open critical count	Architecture node	Severity	Count where Critical	The one number that should always be zero
Days open	Release	Deployed	Date range	Shows whether findings are becoming fixes or a backlog

FORMULAS

PRIORITY SCORE

```
if(prop("Effort") = 0, 0,
  round(prop("Impact") * prop("Pages") * prop("Business")
    * prop("Confidence") / prop("Effort")))
```

Confidence is the term most teams omit, and the one that prevents a sprint being spent on a hypothesis. Maximum is 625 at effort 1; treat the number as a ranking device, not a magnitude.

ACTION

```
if(prop("Confidence") < 3, "Diagnose first", "Fix")
```

A low-confidence issue is a research task, not an engineering ticket. Sending it to a sprint is how SEO loses engineering credibility.

IEWS

VIEW	TYPE	FILTER	SORT / GROUP	FOR	DECISION IT SUPPORTS
Backlog	Table	Status is not Verified	Sort by Priority score descending	Technical SEO owner	What should engineering pick up next
Critical open	Table	Severity is Critical and Status is not Verified	Sort by Created	Everyone	Is anything on fire right now
Diagnose first	Table	Action is Diagnose first	Sort by Priority score descending	Technical SEO owner	Which issues need evidence before a ticket
Awaiting verification	Table	Deployed is not empty and Verified is empty	Sort by Deployed	Technical SEO owner	What shipped that we have not confirmed
By module	Board	Status is not Verified	Group by Module	Technical SEO owner	Which technical area generates the most problems

WORKFLOW

Status values	Found / Diagnosed / Ticketed / In progress / Deployed / Verified
Transitions	Found to Diagnosed with evidence and confidence; Diagnosed to Ticketed only if Action is Fix; Deployed to Verified only after production confirmation
Blocking conditions	Cannot be Ticketed with confidence below 3; cannot be Verified without a production check
Definition of done	Evidence attached, scored, ticketed, deployed, and verified in production with a date

TEMPLATE, CADENCE, ARCHIVE, GATE

Database template	Issue record pre-loaded with the five scoring inputs, the evidence prompt, and the dev-ticket structure
Review cadence	Weekly backlog review, monthly full audit
Archive rule	Verified issues stay as history; recurrence patterns matter
Quality gate	No issue closes without a production verification date

EXAMPLE RECORDS — ALL FICTIONAL

ISSUE	MODULE	SCORE / ACTION	STATUS
Facet URLs indexable	Faceted navigation	96 · Fix	Ticketed
Pillar renders empty without JS	Rendering	150 · Fix	Verified
Suspected crawl budget waste on filters	Crawling	24 · Diagnose first	Found

Fictional illustrative records. Northstar Analytics is an invented company; every figure shown is a placeholder.

DB 17

Technical Releases & Verification

Purpose	Every deploy that could affect search, and whether it was checked afterwards.
Primary owner	Technical SEO owner
Used by	Engineering, SEO
Inputs	Release notes, deploy schedule, post-deploy crawls
Outputs	A verification record per release

PROPERTIES

PROPERTY	TYPE	VALUES / DEFAULT	R	DEFINITION	EXAMPLE
Release	Title	Free text	R	Name or version	Template refresh v4.2
Deploy date	Date	—	R	When it shipped	2026-07-18
Templates touched	Multi-select	Your templates	R	What changed	Article, Comparison
SEO reviewed pre-deploy	Checkbox	Default false	R	Was staging checked	true
Post-deploy crawl	Date	—	·	When the crawl ran	2026-07-19
Index check 7d	Checkbox	Default false	R	Checked after a week	true
Index check 30d	Checkbox	Default false	R	Checked after a month	false
Issues found	Relation	Technical Issues	·	Anything it broke	Canonical loss
Verdict	Select	Clean / Issues found / Rolled back	·	Outcome	Issues found
Owner	Person	One person	R	Who verified	Technical SEO owner

RELATIONS

RELATION	DESTINATION	CARDINALITY	PURPOSE
Issues found	Technical Issues	One-to-many	Link regressions to their release
Architecture	Site Architecture	Many-to-many	Which sections were affected

VIEWS

VIEW	TYPE	FILTER	SORT / GROUP	FOR	DECISION IT SUPPORTS
Unverified	Table	Deploy date is not empty and Index check 7d is false	Sort by Deploy date	Technical SEO owner	Which releases have not been checked
Regressions	Table	Verdict is Issues found	Sort by Deploy date descending	SEO lead	Which releases break things, and is there a pattern
Release log	Table	All	Sort by Deploy date descending	Everyone	What changed and when — the first thing to check on any anomaly

WORKFLOW

Status values	Scheduled / Deployed / Verifying / Closed
Transitions	Scheduled to Deployed on ship; Deployed to Verifying on the post-deploy crawl; Verifying to Closed after the 30-day check
Blocking conditions	Cannot be Closed without both the 7-day and 30-day index checks
Definition of done	Pre-deploy review recorded, post-deploy crawl run, both index checks completed

TEMPLATE, CADENCE, ARCHIVE, GATE

Database template	Release record with the pre-deploy and post-deploy checklists
Review cadence	Per release, plus a monthly review of the log
Archive rule	Keep permanently — the release log is the first place to look on any anomaly
Quality gate	No release closes without a 30-day index check

DB 18 Internal Link Opportunities

Purpose Every planned and implemented internal link, scored and verified.

Primary owner SEO specialist

Used by Content, Editors

Inputs Crawl export, link counts, commercial importance

Outputs A prioritized, verified link implementation queue

PROPERTIES

PROPERTY	TYPE	VALUES / DEFAULT	R	DEFINITION	EXAMPLE
Opportunity	Title	Source to target	R	The link	Reporting guide → Pricing
Source page	Relation	Content	R	Where the link lives	Reporting guide
Target page	Relation	Content	R	Where it points	Pricing
Suggested anchor	Text	Free text	R	Descriptive phrase	how our pricing works
Reader rationale	Text	Free text	R	Why a reader would click	They ask about cost at this point
Target value	Number	1–5	R	Commercial importance of destination	5
Relevance	Number	1–5	R	Would a reader want it here	4
Source strength	Number	1–5	R	Traffic and links of the source	3
Priority	Formula	Derived	R	Product of the three	60
Owner	Person	One person	R	Who implements	Editor
Implemented	Date	—	·	When	
Verified	Checkbox	Default false	·	Confirmed live	false
Status	Select	Identified / Approved / Implemented / Verified / Rejected	R	State	Approved

RELATIONS

RELATION	DESTINATION	CARDINALITY	PURPOSE
Source page	Content	Many-to-one	The linking page
Target page	Content	Many-to-one	The destination — drives the inbound rollup
Cluster	Topic Clusters	Many-to-one	Interlinking within a cluster

FORMULAS

PRIORITY

```
prop("Target value") * prop("Relevance") * prop("Source strength")
```

Maximum 125. Discard anything scoring below 3 on Relevance regardless of the total — an irrelevant internal link helps nobody and degrades the page it sits on.

VIEWS

VIEW	TYPE	FILTER	SORT / GROUP	FOR	DECISION IT SUPPORTS
Implementation queue	Table	Status is Approved	Sort by Priority descending	Editor	Which links to build this week
Low relevance	Table	Relevance < 3	Sort by Priority descending	SEO specialist	Which suggestions should be rejected rather than built
Unverified	Table	Implemented is not empty and Verified is false	Sort by Implemented	Editor	Which links were added but never confirmed live
By target	Board	Status is Verified	Group by Target page	SEO specialist	Which pages are actually gaining internal support

WORKFLOW

Status values Identified / Approved / Implemented / Verified / Rejected

Transitions Identified to Approved on editor review; Implemented to Verified on a live check

Blocking conditions Cannot be Approved with Relevance below 3

Definition of done Anchor written, rationale recorded, implemented, and verified live

TEMPLATE, CADENCE, ARCHIVE, GATE

Database template	Opportunity record with the three scoring inputs and a rationale prompt
Review cadence	Weekly
Archive rule	Rejected opportunities keep their reason
Quality gate	No link is approved without a reader-benefit rationale

DB 19

Backlinks

Purpose	Links and mentions earned, monitored for loss and for reclamation opportunities.
Primary owner	Marketing or PR owner
Used by	SEO, Authority campaigns
Inputs	Backlink exports, mention monitoring, campaign outcomes
Outputs	A monitored link profile with reclamation actions

PROPERTIES

PROPERTY	TYPE	VALUES / DEFAULT	R	DEFINITION	EXAMPLE
Source	Title	Domain or page	R	Who linked or mentioned	industryjournal.example
Type	Select	Link / Unlinked mention / Citation / Coverage	R	What it is	Unlinked mention
Target	Relation	Content	·	Which page it points at	Reporting benchmark study
Anchor	Text	Free text	·	Anchor text used	reporting benchmark study
Relevance	Number	1–5	R	Topical fit	5
Authority indicator	Number	Tool metric	·	With the tool named	58
Tool	Text	Tool name	·	Which tool	Tool Y
First seen	Date	—	R	When found	2026-06-02
Last seen	Date	—	R	Most recent check	2026-08-01
Lost	Checkbox	Default false	·	No longer present	false
Reclamation action	Select	None / Request link / Request correction / Monitor	·	What to do	Request link
Campaign	Relation	Digital PR Campaigns	·	Which campaign earned it	Reporting benchmark

RELATIONS

RELATION	DESTINATION	CARDINALITY	PURPOSE
Target	Content	Many-to-one	Which asset earned it
Campaign	Digital PR Campaigns	Many-to-one	Attribute outcomes to campaigns
Authority target	Authority Targets	One-to-one	Close the loop on an outreach target

ROLLUPS

ROLLUP	VIA RELATION	PROPERTY	CALCULATION	WHY IT MATTERS
Campaign link count	Campaign	Source	Count	The only honest measure of whether a campaign worked

VIEWS

VIEW	TYPE	FILTER	SORT / GROUP	FOR	DECISION IT SUPPORTS
Unlinked mentions	Table	Type is Unlinked mention and Reclamation action is Request link	Sort by Authority indicator descending	Outreach owner	Which mentions could become links with one email
Lost links	Table	Lost is true	Sort by Last seen descending	Outreach owner	What disappeared, and is it worth reclaiming
By campaign	Board	Campaign is not empty	Group by Campaign	Marketing lead	Which campaigns actually earned coverage

WORKFLOW

Status values	Found / Actioned / Earned / Lost
Transitions	Found to Actioned on outreach; Actioned to Earned on a live link
Blocking conditions	None — this is a monitoring database
Definition of done	Source, type, relevance, and dates recorded; reclamation action decided

TEMPLATE, CADENCE, ARCHIVE, GATE

Database template	Link record with a reclamation decision prompt
Review cadence	Weekly for new and lost, monthly for full review
Archive rule	Lost links stay with their last-seen date
Quality gate	No authority figure without the tool that produced it

DB 20 Authority Targets

Purpose Qualified outreach prospects, ranked by relevance rather than by metric.

Primary owner Outreach owner

Used by PR campaigns, SEO

Inputs Prospect research, journalist beats, resource pages

Outputs A qualified, prioritized outreach list

PROPERTIES

PROPERTY	TYPE	VALUES / DEFAULT	R	DEFINITION	EXAMPLE
Target	Title	Publication or page	R	Who we are approaching	Data Weekly — tools roundup
Type	Select	Journalist / Resource page / Partner / Directory / Community	R	What kind	Resource page
Relevance	Number	1-5	R	Topical fit — this leads, not metrics	5
Specific reason	Text	Free text	R	Why they would care	Their roundup omits the selection-criteria angle
Asset offered	Relation	Content	R	What we are offering	Reporting benchmark study
Contact route	Text	Public route only	R	How to reach them	Public editorial address
Touches sent	Number	Max 2	R	Contact count	1
Outcome	Select	No response / Declined / Link / Mention / Coverage	·	What happened	
Owner	Person	One person	R	Who runs it	Outreach owner
Status	Select	Identified / Qualified / Contacted / Closed	R	State	Contacted

RELATIONS

RELATION	DESTINATION	CARDINALITY	PURPOSE
Asset	Content	Many-to-one	What is being offered
Campaign	Digital PR Campaigns	Many-to-one	Which campaign it belongs to
Backlink	Backlinks	One-to-one	The outcome record if it lands

VIEWS

VIEW	TYPE	FILTER	SORT / GROUP	FOR	DECISION IT SUPPORTS
Ready to contact	Table	Status is Qualified	Sort by Relevance descending	Outreach owner	Who to approach next
Follow-up due	Table	Touches sent = 1 and Outcome is empty	Sort by Modified	Outreach owner	Who needs one follow-up, and only one
Closed outcomes	Board	Status is Closed	Group by Outcome	Marketing lead	What actually converts, including the refusals

WORKFLOW

Status values Identified / Qualified / Contacted / Closed

Transitions Identified to Qualified on a relevance check; Contacted to Closed after two touches or an outcome

Blocking conditions Cannot be Contacted without a specific reason and an asset

Definition of done Relevance assessed, specific reason written, asset linked, outcome recorded either way

TEMPLATE, CADENCE, ARCHIVE, GATE

Database template Target record with a qualification checklist and personalization prompt

Review cadence Weekly during a campaign

Archive rule Closed targets stay with their outcome, including refusals

Quality gate Maximum two touches; every outcome is logged, including silence

DB 21

Digital PR Campaigns

Purpose	Each authority campaign: the asset, the story, the targets, and the outcomes.
Primary owner	Marketing or PR owner
Used by	SEO, Content, Leadership
Inputs	Linkable assets, methodology, permissions, target research
Outputs	Campaign records with honest outcome reporting

PROPERTIES

PROPERTY	TYPE	VALUES / DEFAULT	R	DEFINITION	EXAMPLE
Campaign	Title	Free text	R	Name	Reporting benchmark 2026
Asset	Relation	Content	R	What is being pitched	Reporting benchmark study
Story	Text	Free text	R	The headline a journalist would write	Data teams spend a fifth of the month on reporting
Timing hook	Text	Free text	R	Why now	Annual planning season
Methodology published	Checkbox	Default false	R	Is the method visible	true
Stated limits	Text	Free text	R	What the data cannot support	Self-reported, 140 responses, one sector
Permission cleared	Checkbox	Default false	R	Data owner signed off	true
Spokesperson	Person	One person	R	Named, with credentials	Head of Data
Window	Date range	—	R	Campaign dates	2026-09-01 to 2026-09-26
Targets contacted	Rollup	Derived	·	Count	34
Outcomes earned	Rollup	Derived	·	Links and coverage	6
Status	Select	Planning / Live / Closed	R	State	Planning

RELATIONS

RELATION	DESTINATION	CARDINALITY	PURPOSE
Asset	Content	Many-to-one	The linkable asset
Targets	Authority Targets	One-to-many	Everyone approached
Backlinks	Backlinks	One-to-many	What was earned

ROLLUPS

ROLLUP	VIA RELATION	PROPERTY	CALCULATION	WHY IT MATTERS
Targets contacted	Targets	Target	Count	Effort, so outcomes can be read against it
Outcomes earned	Backlinks	Source	Count	The honest result — including zero, which is a valid outcome

VIEWS

VIEW	TYPE	FILTER	SORT / GROUP	FOR	DECISION IT SUPPORTS
Live campaigns	Table	Status is Live	Sort by Window	Marketing lead	What is running right now
Outcome review	Table	Status is Closed	Sort by Outcomes earned descending	Leadership	Which kinds of asset actually earn coverage
Not cleared	Table	Permission cleared is false and Status is not Closed	Sort by Window	Marketing lead	What cannot legally go out yet

WORKFLOW

Status values	Planning / Live / Closed
Transitions	Planning to Live only when permission is cleared and methodology is published
Blocking conditions	Cannot go Live without permission cleared, methodology published, and stated limits written
Definition of done	Asset published with methodology, targets contacted, every outcome logged

TEMPLATE, CADENCE, ARCHIVE, GATE

Database template	Campaign record with the full campaign template and a permission checklist
Review cadence	Weekly while live, retrospective at close
Archive rule	Closed campaigns stay permanently for pattern analysis
Quality gate	No pitch claim exceeds what the published methodology supports

DB 22

Local Locations

Purpose	The canonical record for each location — the single source every listing derives from.
Primary owner	Local SEO owner
Used by	Local citations, Content, Reporting
Inputs	Business records, profile exports, audits
Outputs	One authoritative NAP per location

PROPERTIES

PROPERTY	TYPE	VALUES / DEFAULT	R	DEFINITION	EXAMPLE
Location	Title	Free text	R	Branch or service area	Manchester office
Canonical name	Text	Exact	R	The one true name	Northstar Analytics Ltd
Canonical address	Text	Exact, character for character	R	The one true address	14 Example Street, Manchester, M1 1AA
Canonical phone	Phone	Exact format	R	The one true number	+44 161 000 0000
Primary category	Text	As set in the profile	R	Main category	Software company
Secondary categories	Multi-select	Your list	·	Deliberate additions	Data analytics service
Profile status	Select	Unclaimed / Claimed / Optimized / Verified	R	State	Optimized
Landing page	Relation	Content	R	Its page	Manchester office
Reviews count	Number	Integer	·	With an as-of date	46
Reviews as-of	Date	—	·	When counted	2026-08-01
Citation consistency	Rollup	Derived	R	Percentage matching	86%
Owner	Person	One person	R	Who maintains it	Local SEO owner
Last audited	Date	—	R	Drives the cadence	2026-08-01

RELATIONS

RELATION	DESTINATION	CARDINALITY	PURPOSE
Citations	Local Citations	One-to-many	Every listing derived from this record
Landing page	Content	One-to-one	The location page
Keywords	Keyword Master	One-to-many	Local queries for this location

ROLLUPS

ROLLUP	VIA RELATION	PROPERTY	CALCULATION	WHY IT MATTERS
Citation consistency	Citations	Matches canonical	Percent checked	The single number that predicts local visibility problems

VIEWS

VIEW	TYPE	FILTER	SORT / GROUP	FOR	DECISION IT SUPPORTS
Audit due	Table	Last audited more than 30 days ago	Sort by Last audited	Local SEO owner	Which locations need auditing this week
Inconsistent	Table	Citation consistency below 95%	Sort ascending	Local SEO owner	Where is NAP drifting
Unclaimed	Table	Profile status is Unclaimed	Sort by Location	Local SEO owner	Which profiles are not under our control

WORKFLOW

Status values	Setup / Active / Under review / Closed
Transitions	Setup to Active when the profile is claimed and the page is live
Blocking conditions	Cannot be Active without a canonical NAP and a live landing page
Definition of done	Canonical NAP recorded, profile claimed, page live, citations audited

TEMPLATE, CADENCE, ARCHIVE, GATE

Database template	Location record with the canonical NAP block and the audit checklist
Review cadence	Monthly per location
Archive rule	Closed locations keep their record with a closure date and redirects
Quality gate	Every listing derives from this record — never edited independently

DB 23 Local Citations

Purpose	Every listing, compared character by character against the canonical record.
Primary owner	Local SEO owner
Used by	Local locations
Inputs	Directory exports, manual checks
Outputs	An inconsistency list with corrections assigned

PROPERTIES

PROPERTY	TYPE	VALUES / DEFAULT	R	DEFINITION	EXAMPLE
Directory	Title	Free text	R	Where the listing is	Industry directory A
Location	Relation	Local Locations	R	Which location	Manchester office
Name as listed	Text	Exactly as shown	R	Verbatim	Northstar Analytics
Address as listed	Text	Exactly as shown	R	Verbatim	14 Example St, Manchester M1 1AA
Phone as listed	Text	Exactly as shown	R	Verbatim	0161 000 0000
Matches canonical	Checkbox	Default false	R	Character-for-character match	false
Difference	Text	Free text	·	What exactly differs	"St" vs "Street"; no +44
Claimed	Checkbox	Default false	R	Under our control	true
Correction requested	Date	—	·	When we asked	2026-08-02
Corrected	Date	—	·	When it was fixed	
Status	Select	Found / Inconsistent / Correction requested / Correct / Duplicate	R	State	Inconsistent

RELATIONS

RELATION	DESTINATION	CARDINALITY	PURPOSE
Location	Local Locations	Many-to-one	Feeds the consistency rollout

VIEWS

VIEW	TYPE	FILTER	SORT / GROUP	FOR	DECISION IT SUPPORTS
Inconsistent	Table	Matches canonical is false	Group by Location	Local SEO owner	What to correct, and where
Awaiting correction	Table	Correction requested is not empty and Corrected is empty	Sort by Correction requested	Local SEO owner	Who has not responded yet
Duplicates	Table	Status is Duplicate	Group by Location	Local SEO owner	Which listings need merging

WORKFLOW

Status values	Found / Inconsistent / Correction requested / Correct / Duplicate
Transitions	Found to Correct on a match; Inconsistent to Correction requested on outreach
Blocking conditions	Cannot be marked Correct without a character-level comparison
Definition of done	Listing compared exactly, difference quoted, correction requested and tracked

TEMPLATE, CADENCE, ARCHIVE, GATE

Database template	Citation record with the exact-comparison fields prefilled from the canonical record
Review cadence	Monthly
Archive rule	Corrected entries stay as history
Quality gate	Differences are quoted exactly; never described in summary

DB 24

Brand Facts

Purpose	The canonical statement of every fact about the brand, and where each must match.
Primary owner	Marketing owner
Used by	AI visibility, Content, On-page, Local
Inputs	Site copy, profiles, listings, press, structured data
Outputs	One canonical version of every fact, and a conflict list

PROPERTIES

PROPERTY	TYPE	VALUES / DEFAULT	R	DEFINITION	EXAMPLE
Fact	Title	Free text	R	What the fact is	Category description
Canonical value	Text	Exact wording	R	The one true version	Analytics platform for data teams
Fact type	Select	Identity / Product / People / Commercial / History	R	Grouping	Identity
Must match on	Multi-select	Site / LinkedIn / Directories / Structured data / Press / Reviews	R	Where it must appear identically	Site, LinkedIn, Directories
Verified sources	Number	Integer	·	How many checked and matching	4
Conflicting sources	Text	Free text	·	Where it differs, quoted	Directory B: "BI reporting tool"
Publicly verifiable	Checkbox	Default false	R	Can a third party confirm it	true
Owner	Person	One person	R	Who maintains it	Marketing owner
Last audited	Date	—	R	When checked	2026-07-15
Version	Text	Free text	R	Version identifier	v3
Status	Select	Canonical / Conflicted / Under revision / Retired	R	State	Conflicted

RELATIONS

RELATION	DESTINATION	CARDINALITY	PURPOSE
AI answers	AI Answers	One-to-many	Answers checked against this fact
Content	Content	Many-to-many	Pages that state this fact
Citations	Local Citations	Many-to-many	Listings that must match

ROLLUPS

ROLLUP	VIA RELATION	PROPERTY	CALCULATION	WHY IT MATTERS
Conflict count	AI answers	Factual error	Count where true	How often AI systems get this fact wrong — the highest-priority signal in the system

VIEWS

VIEW	TYPE	FILTER	SORT / GROUP	FOR	DECISION IT SUPPORTS
Conflicted	Table	Status is Conflicted	Sort by Fact type	Marketing owner	Which facts about us are inconsistent in public
Audit due	Table	Last audited more than 90 days ago	Sort by Last audited	Marketing owner	What has not been checked this quarter
Unverifiable	Table	Publicly verifiable is false	Sort by Fact type	SEO lead	Which facts cannot be corroborated by anyone but us
By type	Board	Status is Canonical	Group by Fact type	Everyone	The fact sheet, as a single reference

WORKFLOW

Status values	Canonical / Conflicted / Under revision / Retired
Transitions	Canonical to Conflicted when an audit finds a difference; Conflicted to Canonical only after every source is corrected
Blocking conditions	Cannot return to Canonical while a conflicting source is unresolved and un-escalated
Definition of done	One canonical value, every required surface checked and matching, versioned with a date

TEMPLATE, CADENCE, ARCHIVE, GATE

Database template	Fact record with the must-match checklist and a conflict-quoting prompt
Review cadence	Quarterly, and immediately on any AI factual error
Archive rule	Retired facts keep their history — old descriptions persist in third-party sources
Quality gate	A third version is never introduced; corrections happen at source

EXAMPLE RECORDS — ALL FICTIONAL

FACT	CANONICAL VALUE	CONFLICT FOUND	STATUS
Category description	Analytics platform for data teams	Directory B says "BI reporting tool"	Conflicted
Founded	2019	None across 6 sources	Canonical
Headquarters	Manchester, United Kingdom	Profile C still lists a 2021 address	Conflicted

Fictional illustrative records. Northstar Analytics is an invented company; every figure shown is a placeholder.

DB 25 AI Prompt Universe

Purpose	The fixed set of prompts tracked every month, with wording that never changes.
Primary owner	SEO lead
Used by	AI visibility, Content, Strategy
Inputs	ICP decisions, question bank, competitor list
Outputs	A stable, comparable tracking set

PROPERTIES

PROPERTY	TYPE	VALUES / DEFAULT	R	DEFINITION	EXAMPLE
Prompt	Title	Exact wording	R	Never edited after first run	What is the best analytics tool for a small data team?
Prompt ID	Text	PU-001 format	R	Stable identifier	PU-014
Type	Select	Category / Brand / Competitor / Comparison / Problem / Purchase-intent	R	Prompt family	Category
Intent	Select	Informational / Comparison / Commercial	R	What it wants	Commercial
Locale	Select	Your markets	R	Language and market	UK English
Version	Number	Integer	R	Increments only on a deliberate rewrite	1
Superseded by	Relation	AI Prompt Universe	·	If rewritten, the new record	
Good result looks like	Text	Free text	R	Definition of success	Named among 3–5 options with a link
Priority	Select	Core / Secondary	R	Core prompts run monthly without fail	Core
Source	Select	Question bank / Decision list / Competitor set / Manual	R	Where it came from	Decision list
Active	Checkbox	Default true	R	In the current run set	true

RELATIONS

RELATION	DESTINATION	CARDINALITY	PURPOSE
Runs	AI Prompt Runs	One-to-many	Every execution of this prompt
Question	Question Bank	One-to-one	The customer question it derives from
Opportunities	AI Visibility Opportunities	One-to-many	Gaps this prompt revealed

ROLLUPS

ROLLUP	VIA RELATION	PROPERTY	CALCULATION	WHY IT MATTERS
Run count	Runs	Run date	Count	A prompt with fewer than three runs cannot show a trend
Mention rate	Runs	Mentioned	Percent checked	The core visibility measure for this prompt over time

VIEWS

VIEW	TYPE	FILTER	SORT / GROUP	FOR	DECISION IT SUPPORTS
Core set	Table	Active is true and Priority is Core	Sort by Prompt ID	AI visibility owner	What must be run this month
By type	Board	Active is true	Group by Type	SEO lead	Is the set balanced across the six prompt families
Never mentioned	Table	Mention rate = 0% and Run count >= 3	Sort by Type	SEO lead	Where are we consistently absent
Superseded	Table	Active is false	Sort by Version descending	AI visibility owner	Which prompts were rewritten, and when the trend breaks

WORKFLOW

Status values	Drafted / Active / Superseded / Retired
Transitions	Drafted to Active on the first run; Active to Superseded only when a deliberate rewrite creates a new record
Blocking conditions	Wording cannot be edited on an Active prompt — a change creates a new record at version+1
Definition of done	Wording fixed, type and locale set, good-result definition written, first run recorded

TEMPLATE, CADENCE, ARCHIVE, GATE

Database template	Prompt record with the wording-lock warning and the good-result prompt
Review cadence	Monthly run, quarterly set review
Archive rule	Superseded prompts stay permanently — they anchor the historical trend
Quality gate	Wording never changes on an active prompt; comparability is the entire point

EXAMPLE RECORDS — ALL FICTIONAL

PROMPT ID / TYPE	PROMPT WORDING	GOOD RESULT	RUNS / MENTION RATE
PU-014 · Category	What is the best analytics tool for a small data team?	Named among 3–5 options with a link	6 · 33%
PU-021 · Brand	What does Northstar Analytics do?	Accurate category and no factual errors	6 · 100%
PU-033 · Comparison	Northstar Analytics vs [competitor] — which is better?	Both named, our differentiator stated correctly	6 · 50%

Fictional illustrative records. Northstar Analytics is an invented company; every figure shown is a placeholder.

DB 26

AI Prompt Runs

Purpose	One record per prompt per assistant per month — the raw tracking data.
Primary owner	AI visibility owner
Used by	Reporting, Content, Strategy
Inputs	The prompt universe and the assistants your buyers use
Outputs	A comparable dated dataset

PROPERTIES

PROPERTY	TYPE	VALUES / DEFAULT	R	DEFINITION	EXAMPLE
Run	Title	Auto: prompt ID + assistant + date	R	Identifier	PU-014 · Assistant A · 2026-08
Prompt	Relation	AI Prompt Universe	R	Which prompt	PU-014
Assistant	Select	Your tracked systems	R	Which system	Assistant A
Run date	Date	—	R	When	2026-08-04
Run by	Person	One person	R	Who ran it	AI visibility owner
Mentioned	Checkbox	Default false	R	Were we named	false
Cited	Checkbox	Default false	R	Linked or attributed	false
Position in answer	Number	Integer or empty	·	Order named, if named	
Sentiment	Select	Positive / Neutral / Negative / Not mentioned	R	Tone about us	Not mentioned
Answer verbatim	Text	Full text	R	Copied exactly	[full answer pasted]
Factual error	Checkbox	Default false	R	Anything wrong about us	false
Error detail	Text	Free text	·	What was wrong	
Competitors named	Multi-select	Your competitor list	·	Who else appeared	Competitor A, Competitor B

RELATIONS

RELATION	DESTINATION	CARDINALITY	PURPOSE
Prompt	AI Prompt Universe	Many-to-one	Feeds the mention-rate rollup
Citations	AI Citations & Sources	One-to-many	Sources this answer drew on
Competitors	AI Competitors	Many-to-many	Track who is named where
Brand facts checked	Brand Facts	Many-to-many	Which facts this answer touched

VIEWS

VIEW	TYPE	FILTER	SORT / GROUP	FOR	DECISION IT SUPPORTS
Latest run	Table	Run date is this month	Group by Assistant	AI visibility owner	What came back this month
Errors	Table	Factual error is true	Sort by Run date descending	Marketing owner	What is being said about us that is wrong — the highest priority in the system
Absent	Table	Mentioned is false	Group by Prompt type	SEO lead	Where are we missing, and is there a pattern
Mentioned not cited	Table	Mentioned is true and Cited is false	Sort by Run date descending	Editor	Where are we named without attribution — usually an extractability gap
Trend	Table	All	Group by Run month	Leadership	Is visibility moving in either direction

WORKFLOW

Status values	Captured / Analysed / Actioned
Transitions	Captured to Analysed once errors and gaps are logged; Analysed to Actioned when opportunities are created
Blocking conditions	Cannot be Analysed without the verbatim answer pasted
Definition of done	Verbatim answer captured, mention and citation recorded separately, errors flagged

TEMPLATE, CADENCE, ARCHIVE, GATE

Database template	Run record with the capture checklist and a verbatim paste field
Review cadence	Monthly
Archive rule	Keep every run permanently — the trend is the asset
Quality gate	Prompt wording must match the universe record exactly, or the run is void

DB 27

AI Citations & Sources

Purpose	Which sources AI answers draw on, and whether we are among them.
Primary owner	AI visibility owner
Used by	Content, Authority, Strategy
Inputs	Captured answers and the sources they name
Outputs	A citation-gap list that feeds authority work

PROPERTIES

PROPERTY	TYPE	VALUES / DEFAULT	R	DEFINITION	EXAMPLE
Source	Title	Domain or page	R	What was cited	industryjournal.example
Source type	Select	Publisher / Directory / Community / Vendor / Review site / Own site	R	Category	Publisher
Cited in run	Relation	AI Prompt Runs	R	Which answer	PU-014 · A · 2026-08
Our page	Relation	Content	·	If it was us	
Mentions us	Checkbox	Default false	R	Does this source cover us	false
Mentions competitor	Multi-select	Competitor list	·	Who it covers	Competitor A
Gap	Formula	Derived	R	Covers competitors but not us	Yes
Closable	Select	Yes / Difficult / No	·	Can we realistically appear here	Yes
Route	Text	Free text	·	How we would appear	Editorial pitch with the benchmark study
Owner	Person	One person	·	Who would pursue it	Outreach owner

RELATIONS

RELATION	DESTINATION	CARDINALITY	PURPOSE
Run	AI Prompt Runs	Many-to-one	Which answer cited it
Our page	Content	Many-to-one	Which of our pages was used
Authority target	Authority Targets	One-to-one	Promote a gap into an outreach target

ROLLUPS

ROLLUP	VIA RELATION	PROPERTY	CALCULATION	WHY IT MATTERS
Citation frequency	Run	Source	Count	How often a source appears across your prompt set — the ones that appear repeatedly are where authority work pays

FORMULAS

GAP

if(and(not prop("Mentions us"), not empty(prop("Mentions competitor"))),
"Yes", "No")

A gap means a source AI systems draw on covers your competitors and not you. It is the most actionable finding the AI visibility work produces.

VIEWS

VIEW	TYPE	FILTER	SORT / GROUP	FOR	DECISION IT SUPPORTS
Citation gaps	Table	Gap is Yes and Closable is Yes	Sort by Citation frequency descending	Outreach owner	Which sources should we be in, ranked by how often they are used
Our pages cited	Table	Our page is not empty	Group by Our page	Content lead	What kind of our content gets used
Source landscape	Board	All	Group by Source type	Strategy	Where does credibility in this category actually live

WORKFLOW

Status values	Observed / Assessed / Targeted / Closed
Transitions	Observed to Assessed on a closability judgement; Assessed to Targeted when an outreach record is created
Blocking conditions	Cannot be Targeted without a stated route
Definition of done	Source recorded, gap calculated, closability assessed, route written where closable

TEMPLATE, CADENCE, ARCHIVE, GATE

Database template	Source record with a closability prompt and route field
Review cadence	Monthly after each run
Archive rule	Keep permanently — the source landscape shifts slowly
Quality gate	Gaps are assessed for closability before they enter an outreach list

DB 28

AI Competitors

Purpose	Who AI systems name in your category, how often, and what they are credited with.
Primary owner	SEO lead
Used by	Strategy, Content, Authority
Inputs	Captured answers, competitor public presence
Outputs	A share-of-answer picture with closable differences

PROPERTIES

PROPERTY	TYPE	VALUES / DEFAULT	R	DEFINITION	EXAMPLE
Competitor	Title	Name	R	Who	Competitor A
Times named	Rollup	Derived	R	Across the prompt set	18
Share of answer	Formula	Derived	R	Their share of all brand mentions	31%
Reasons given	Text	Free text	·	What answers credit them with	Cited for pricing transparency
Public presence	Text	Free text	·	What we observe	Listed in 4 directories we are absent from
Closable difference	Text	Free text	·	What we could legitimately match	Publish a comparable pricing page
Effort	Select	Hours / Days / Weeks	·	To close it	Days
Last reviewed	Date	—	R	When assessed	2026-08-05

RELATIONS

RELATION	DESTINATION	CARDINALITY	PURPOSE
Runs	AI Prompt Runs	Many-to-many	Which answers named them
Opportunities	AI Visibility Opportunities	One-to-many	Gaps derived from them

ROLLUPS

ROLLUP	VIA RELATION	PROPERTY	CALCULATION	WHY IT MATTERS
Times named	Runs	Competitors named	Count	Raw presence across your tracked set

FORMULAS

SHARE OF ANSWER

```
if(prop("Total brand mentions") = 0, "0%",
  round(prop("Times named") / prop("Total brand mentions") * 100) + "%")
```

Total brand mentions is a rollup across all runs. This measures share of a set you defined — it is not a market share figure and should never be presented as one.

VIEWS

VIEW	TYPE	FILTER	SORT / GROUP	FOR	DECISION IT SUPPORTS
Share of answer	Table	All	Sort by Times named descending	Leadership	Who owns the category answer today
Closable differences	Table	Closable difference is not empty	Sort by Effort	SEO lead	What could we legitimately match, and how quickly
Stale	Table	Last reviewed more than 90 days ago	Sort by Last reviewed	SEO lead	Whose presence have we not re-checked

WORKFLOW

Status values	Tracked / Assessed / Archived
Transitions	Tracked to Assessed after a presence comparison; Archived when they stop appearing
Blocking conditions	Cannot state a closable difference without observed presence evidence
Definition of done	Presence observed, differences described factually, closability assessed

TEMPLATE, CADENCE, ARCHIVE, GATE

Database template	Competitor record with an observation prompt and a no-speculation reminder
Review cadence	Quarterly
Archive rule	Archived competitors keep their history
Quality gate	Only observable differences are recorded — never speculation about ranking factors

DB 29 AI Visibility Opportunities

Purpose Every gap the AI tracking reveals, prioritized and assigned.

Primary owner SEO lead

Used by Content, Authority, On-page, Marketing

Inputs Prompt runs, citation gaps, competitor comparisons, factual errors

Outputs An actionable, prioritized AI visibility backlog

PROPERTIES

PROPERTY	TYPE	VALUES / DEFAULT	R	DEFINITION	EXAMPLE
Opportunity	Title	Free text	R	What could be done	Publish a pricing structure page
Gap type	Select	Absence / Citation gap / Factual error / Extractability / Source gap	R	What kind of gap	Absence
Evidence	Relation	AI Prompt Runs	R	Which runs showed it	PU-014 ×3
Frequency	Number	Integer	R	How many runs revealed it	3
Proposed action	Text	Free text	R	The specific work	Publish pricing structure with extractable summary
Area	Select	Content / On-page / Authority / Brand facts / Technical	R	Who does it	Content
Effort	Select	Hours / Days / Weeks	R	Size	Days
Priority	Select	Now / This quarter / Later / Declined	R	When	Now
Owner	Person	One person	R	Accountable	Content lead
Status	Select	Identified / Scheduled / In progress / Done / Declined	R	State	Scheduled

RELATIONS

RELATION	DESTINATION	CARDINALITY	PURPOSE
Runs	AI Prompt Runs	Many-to-many	The evidence
Content	Content	Many-to-one	The page that would close it
Authority target	Authority Targets	One-to-one	If it needs outreach

VIEWS

VIEW	TYPE	FILTER	SORT / GROUP	FOR	DECISION IT SUPPORTS
Errors first	Table	Gap type is Factual error	Sort by Frequency descending	Marketing owner	What is wrong about us and being repeated
Now	Table	Priority is Now and Status is not Done	Sort by Frequency descending	Everyone	What are we doing about AI visibility this month
By area	Board	Status is not Done	Group by Area	SEO lead	Where does the work actually sit
Quick wins	Table	Effort is Hours and Priority is not Declined	Sort by Frequency descending	SEO lead	What could be closed this week

WORKFLOW

Status values Identified / Scheduled / In progress / Done / Declined

Transitions Identified to Scheduled with an owner and a proposed action; Done when the work ships

Blocking conditions Factual errors cannot be Declined — they escalate to the brand facts owner

Definition of done Evidence linked, action specified, owner named, shipped and confirmed in a later run

TEMPLATE, CADENCE, ARCHIVE, GATE

Database template	Opportunity record pre-linked to its evidence runs
Review cadence	Monthly after each tracking run
Archive rule	Declined entries keep their reason
Quality gate	Every opportunity links to the runs that evidenced it

DB 30 Historical Scorecards

Purpose	The monthly snapshot of the five AI visibility measures, so a trend exists.
Primary owner	AI visibility owner
Used by	Leadership, Reporting
Inputs	Completed prompt runs for the month
Outputs	A comparable twelve-month trend

PROPERTIES

PROPERTY	TYPE	VALUES / DEFAULT	R	DEFINITION	EXAMPLE
Month	Title	YYYY-MM	R	The period	2026-08
Prompts run	Number	Integer	R	Set size that month	40
Assistants	Number	Integer	R	How many systems	3
Mention rate	Formula	Derived	R	Mentioned / total runs	38%
Citation rate	Formula	Derived	R	Cited / total runs	14%
Answer coverage	Formula	Derived	R	Prompts mentioned on at least one assistant	52%
Share of AI voice	Formula	Derived	R	Our mentions / all brand mentions	19%
Accuracy rate	Formula	Derived	R	Error-free answers about us	92%
Sentiment split	Text	Free text	·	Positive / neutral / negative counts	6 / 9 / 0
Notes	Text	Free text	·	What changed and possible reasons, labelled as hypotheses	Citation rate up after extractability pass — hypothesis, not proven

RELATIONS

RELATION	DESTINATION	CARDINALITY	PURPOSE
Runs	AI Prompt Runs	One-to-many	The month's underlying data

FORMULAS

MENTION RATE

```
if(prop("Total runs") = 0, "0%",  
  round(prop("Mentioned count") / prop("Total runs") * 100) + "%")
```

Comparable only against your own previous months, at identical prompt wording and the same assistant set. It is not comparable to any published figure or to another company.

ANSWER COVERAGE

```
if(prop("Prompts run") = 0, "0%",  
  round(prop("Prompts with a mention") / prop("Prompts run") * 100) + "%")
```

Coverage counts prompts, not runs — a prompt where you appear on one assistant out of three still counts as covered. Report it alongside mention rate, never instead of it.

SHARE OF AI VOICE

```
if(prop("All brand mentions") = 0, "0%",  
  round(prop("Our mentions") / prop("All brand mentions") * 100) + "%")
```

A share of the brands named within your own defined prompt set. Changing the prompt set changes the number. It is not market share and must never be described as such.

VIEWS

VIEW	TYPE	FILTER	SORT / GROUP	FOR	DECISION IT SUPPORTS
Trend	Table	All	Sort by Month ascending	Leadership	Are the five measures moving, and in which direction
Latest	Table	Month is this month	Show all measures	Everyone	Where do we stand right now
Set changes	Table	Prompts run differs from previous month	Sort by Month	SEO lead	Which months are not comparable because the set changed

WORKFLOW

Status values	Draft / Final / Superseded
Transitions	Draft to Final once every core prompt has been run; never edited after Final
Blocking conditions	Cannot be Final if any core prompt was skipped that month
Definition of done	All core prompts run, five measures calculated, set size recorded

TEMPLATE, CADENCE, ARCHIVE, GATE

Database template	Scorecard record with the five formulas and a comparability note
Review cadence	Monthly
Archive rule	Keep permanently — this is the only historical record
Quality gate	A month with a changed prompt set is flagged as not comparable

EXAMPLE RECORDS — ALL FICTIONAL

MONTH	MENTION / CITATION	COVERAGE / SHARE	ACCURACY
2026-06	31% / 9%	45% / 15%	86%
2026-07	34% / 11%	48% / 17%	92%
2026-08	38% / 14%	52% / 19%	92%

Fictional illustrative records. Northstar Analytics is an invented company; every figure shown is a placeholder.

DB 31 Experiments

Purpose One variable, one baseline, one decision rule written before the result.

Primary owner SEO lead

Used by Everyone

Inputs A hypothesis, a dated baseline, available volume

Outputs Attributable change rather than coincidence

PROPERTIES

PROPERTY	TYPE	VALUES / DEFAULT	R	DEFINITION	EXAMPLE
Experiment	Title	Free text	R	One variable, stated	Direct-answer block above the fold
Hypothesis	Text	Free text	R	What we believe and why	Extractable answers raise citation rate
Baseline	Text	Figure, source, date	R	Before anything changed	Citation rate 11%, 40 prompts, 2026-07
Variable changed	Text	Free text	R	Exactly one	Added a 40-word answer block to 12 pages
Held constant	Text	Free text	R	Everything else	No new pages, no prompt changes, no link work
Minimum sample	Text	Free text	R	Before judging	3 monthly runs
Decision rule	Text	Written in advance	R	What result means what	Adopt if citation rate rises 5pp over 3 runs
Evaluation date	Date	—	R	Committed in advance	2026-11-05
Result	Text	Figure with sample	·	What happened	
Decision	Select	Adopt / Reject / Extend / Inconclusive	·	The call	
Owner	Person	One person	R	Who runs it	SEO lead

RELATIONS

RELATION	DESTINATION	CARDINALITY	PURPOSE
Related content	Content	Many-to-many	Pages involved
Decision record	Decisions & Change Log	One-to-one	The resulting decision
Strategy bet	SEO Strategy	Many-to-one	Which bet it informs

VIEWS

VIEW	TYPE	FILTER	SORT / GROUP	FOR	DECISION IT SUPPORTS
Running	Table	Decision is empty	Sort by Evaluation date	SEO lead	What is in flight and when is each judged
Due now	Table	Evaluation date on or before today and Decision is empty	Sort by Evaluation date	SEO lead	What must be decided this week
Concluded	Board	Decision is not empty	Group by Decision	Leadership	What have we actually learned, including the rejections

WORKFLOW

Status values Designed / Running / Evaluating / Concluded

Transitions Designed to Running on the start date; Running to Evaluating on the committed date only

Blocking conditions Cannot start without a baseline, a minimum sample, and a decision rule

Definition of done One variable, dated baseline, decision rule predating the result, outcome recorded either way

TEMPLATE, CADENCE, ARCHIVE, GATE

Database template Experiment record enforcing baseline capture before the variable field unlocks

Review cadence Weekly check, evaluation on the committed date

Archive rule Concluded experiments stay permanently, including inconclusive ones

Quality gate The decision rule is written before the experiment starts, without exception

DB 32 Decisions & Change Log

Purpose	What was decided, on what evidence, by whom, and what happened afterwards.
Primary owner	SEO lead
Used by	Everyone
Inputs	Reports, experiments, strategy reviews, incidents
Outputs	An auditable record of how the system evolved

PROPERTIES

PROPERTY	TYPE	VALUES / DEFAULT	R	DEFINITION	EXAMPLE
Decision	Title	Free text	R	What was decided	Stop pursuing the free-tier keyword set
Date	Date	—	R	When	2026-07-31
Decided by	Person	One person	R	Who made the call	SEO lead
Context	Text	Free text	R	The situation at the time	Traffic growing, trials flat for two quarters
Evidence	Text	Free text	R	What informed it	Segment analysis: free-tier pages, 0 trials in 90d
Options considered	Text	Free text	R	What else was on the table	Retarget the pages, or gate the content
Expected outcome	Text	Falsifiable	R	What we predict	Trial rate per session rises without traffic loss
Review date	Date	—	R	When to check	2026-10-31
Actual outcome	Text	Free text	·	What happened	
Verdict	Select	Correct / Partly correct / Wrong / Too early	·	On review	
Affected	Multi-select	Operating areas	R	What it changed	Keywords, Content
Reversed	Checkbox	Default false	·	Was it later undone	false

RELATIONS

RELATION	DESTINATION	CARDINALITY	PURPOSE
Strategy	SEO Strategy	Many-to-one	Which strategy version it belongs to
Experiment	Experiments	One-to-one	If it came from a test
Report	Reports	Many-to-one	Which report surfaced it

VIEWS

VIEW	TYPE	FILTER	SORT / GROUP	FOR	DECISION IT SUPPORTS
Due for review	Table	Review date on or before today and Verdict is empty	Sort by Review date	SEO lead	Which past decisions should now be checked
Recent	Table	Date within 90 days	Sort by Date descending	Everyone	What have we decided lately
Wrong calls	Table	Verdict is Wrong or Partly correct	Sort by Date descending	SEO lead	Where does our judgement go wrong, and is there a pattern
Reversed	Table	Reversed is true	Sort by Date descending	Leadership	What did we undo, and why

WORKFLOW

Status values	Recorded / Under review / Verdict given
Transitions	Recorded to Under review on the review date; Verdict given once the outcome is known
Blocking conditions	Cannot be Recorded without evidence and a falsifiable expected outcome
Definition of done	Evidence recorded, prediction falsifiable, review date set, verdict given on review

TEMPLATE, CADENCE, ARCHIVE, GATE

Database template	Decision record with an evidence prompt and a falsifiable-prediction check
Review cadence	Monthly review, quarterly retrospective
Archive rule	Never delete — the log exists precisely so old reasoning survives
Quality gate	Every decision records a prediction specific enough to be shown wrong

EXAMPLE RECORDS — ALL FICTIONAL

DECISION	EVIDENCE	PREDICTION	REVIEW / VERDICT
Stop pursuing free-tier keywords	Segment analysis: 0 trials in 90d	Trial rate per session rises	2026-10-31 · pending
Move comparison pages up one level	Depth gap 2 on 6 pages	Comparison pages gain internal links	2026-09-15 · Correct
Delay the glossary build	Capacity: 2 pieces/month	No effect on cluster coverage	2026-12-01 · Too early

Fictional illustrative records. Northstar Analytics is an invented company; every figure shown is a placeholder.

Databases 33 and beyond: the libraries

The Prompt Library and SOP Library are databases too. Their schemas are simple and their content is Sections 12 and 13.

DATABASE	PROPERTIES	KEY VIEWS
Prompt Library	Prompt ID · Name · Category · Difficulty · When to use · When not to use · Required inputs · Databases used · Prompt body · Expected output · Quality checklist · Failure conditions · Next prompt · Times used · Last revised	By category · By difficulty · Most used · Never used · Recently revised
SOP Library	SOP ID · Name · Purpose · Owner · Trigger · Inputs · Tools · Preconditions · Steps · Output · Quality gate · Failure conditions · Escalation · Handoff · Cadence · Related database · Related prompt · Last run · Next due	Due this week · By owner · By cadence · Never run · Overdue
Template Library	Template ID · Name · Operating area · Fields · Completion criterion · Related SOP · Related prompt · Destination database · Example link	By operating area · By type · Most used
Reports	Period · Metrics with sources and as-of dates · What changed · Hypotheses · Three decisions · System changes · Author	Latest · Trend · Decisions outstanding

WHY THE LIBRARIES ARE DATABASES AND NOT PAGES

A prompt on a page gets copied, edited, and forked until three versions exist and nobody knows which is current. A prompt in a database has one record, a version, a usage count, and a last-revised date. The same argument applies to SOPs and templates — and the Never Used view on each is how you find out what you built that nobody wanted.

The Claude prompt library

Eighty-seven prompts across fourteen categories. Prompts PR-01 to PR-25 appear in full form later in this document; PR-26 onward use the same ten-part standard in a denser layout.

CATEGORY	PROMPTS	RANGE
Strategy	9	PR-01 to PR-06, PR-26 to PR-30
Audience and demand	10	PR-07 to PR-14, PR-31 to PR-38
Keywords	11	PR-15 to PR-24, PR-39 to PR-45
Content planning	12	PR-46 to PR-57
Content creation	6	PR-25 to PR-34 (creation subset)
On-page	8	PR-35 to PR-42
Technical SEO	8	PR-43 to PR-50
Internal linking	5	PR-51 to PR-55
Authority and PR	7	PR-56 to PR-62
Local SEO	4	PR-63 to PR-66
LLM and AI SEO	12	PR-67 to PR-78
Measurement	3	PR-79 to PR-81
Reporting	3	PR-82 to PR-84
QA and governance	3	PR-85 to PR-87

THE RULE THAT APPLIES TO ALL EIGHTY-SEVEN

Every output is a draft. Nothing here goes live, gets sent, or changes a page without a person reviewing it. Where a prompt says it will name a gap rather than fill it, that behaviour is the point — an answer that invents a benchmark is worse than no answer, because it looks like research.

PR-26 Stakeholder alignment interview

STRATEGY · INTERMEDIATE

When to use	Before a quarterly strategy review, to gather what each stakeholder actually needs.
When NOT to use	Do not use it as a substitute for talking to them — it prepares the conversation, it does not replace it.
Required inputs	Stakeholder roles, their functions, the business goal under discussion.
Databases used	Business Goals, SEO Strategy

COPY-PASTE PROMPT

ROLE: You are preparing a stakeholder interview for an SEO strategy review.

OBJECTIVE: Surface what each stakeholder needs from search, in their terms.

CONTEXT: Stakeholders: [ROLES AND FUNCTIONS]. Business goal: [GOAL].
Current SEO position: [PASTE].

CONSTRAINTS: One question set per role, six questions maximum each. No leading questions. Do not assume what any function cares about — the questions exist to find out.

PROCESS: 1 For each role, identify what they are measured on. 2 Write questions that surface their actual need rather than their stated request. 3 Include one question that could reveal a conflict with another function. 4 Note what a useful answer sounds like.

OUTPUT: A question set per role with the metric each maps to, plus the three conflicts most likely to surface.

QUALITY CHECKS: No leading questions; every question maps to something the role is measured on.

MISSING INFORMATION: If a role's metrics are unknown, ask for them first.

HUMAN REVIEW: The SEO lead runs the interviews personally.

Expected output	A question set per role plus the likely conflicts between functions.
Quality checklist	Questions map to what each role is measured on; none are leading.
Failure conditions	Stops if stakeholder metrics are unknown.
Next prompt	PR-27

PR-27 Business goal to SEO metric mapping

STRATEGY · FOUNDATIONAL

When to use	When a business objective needs translating into something SEO can be held to.
When NOT to use	Do not use it to invent a metric you cannot actually measure with the data you have.
Required inputs	The business objective, available data sources, current baselines.
Databases used	Business Goals, Reports

COPY-PASTE PROMPT

ROLE: You are mapping business objectives to measurable SEO outcomes.

OBJECTIVE: Produce a metric chain from a commercial goal to a leading indicator.

CONTEXT: Business objective: [GOAL]. What we can measure today: [DATA SOURCES].
Current baselines with dates: [PASTE].

CONSTRAINTS: Only propose metrics we can actually measure with the listed sources. Do not propose vanity metrics. Every metric needs a baseline or a plan to capture one.

PROCESS: 1 Restate the objective. 2 Identify the SEO outcome that serves it. 3 Propose one primary metric and one leading indicator. 4 State where the data comes from. 5 Flag anything unmeasurable today.

OUTPUT: A metric chain table with source, baseline, and as-of date per metric.

QUALITY CHECKS: Every metric has a named source; unmeasurable items are flagged rather than substituted.

MISSING INFORMATION: If no baseline exists, produce a baseline capture plan instead of a metric.

HUMAN REVIEW: The business owner confirms the metric is one they recognize.

Expected output	A metric chain with sources, baselines, and flagged gaps.
Quality checklist	Measurable with listed sources; baselines present or planned.
Failure conditions	Produces a capture plan when no baseline exists.
Next prompt	PR-28

PR-28 SEO operating charter

STRATEGY · INTERMEDIATE

When to use	Once, at setup, and revised when the team or the remit changes.
When NOT to use	Do not use it to assign work — it defines decision rights, not tasks.
Required inputs	Team structure, current owners, escalation practice.
Databases used	SEO Strategy, SOP Library

COPY-PASTE PROMPT

ROLE: You are writing an SEO operating charter for a team.

OBJECTIVE: Define who decides what, who is consulted, and what escalates.

CONTEXT: Team: [ROLES AND NAMES]. Current pain points in decision-making: [DESCRIBE]. Escalation practice today: [DESCRIBE].

CONSTRAINTS: Name real roles. Every decision type gets exactly one decider. Do not create committees. Where a decision genuinely needs two functions, say which one breaks a tie.

PROCESS: 1 List the recurring decision types. 2 Assign a single decider to each. 3 Identify who must be consulted. 4 Define what escalates, to whom, and within what time. 5 Flag any decision currently made by nobody.

OUTPUT: A decision-rights table, an escalation route, and the unowned decisions.

QUALITY CHECKS: One decider per decision type; escalation names a person and a time.

MISSING INFORMATION: Flag decision types where no owner exists rather than assigning one arbitrarily.

HUMAN REVIEW: The business owner approves decision rights before circulation.

Expected output	A decision-rights table, escalation routes, and unowned decisions named.
Quality checklist	One decider per decision; escalation names a person and a window.
Failure conditions	Flags unowned decisions rather than assigning them.
Next prompt	PR-29

PR-29 Constraint register

STRATEGY · FOUNDATIONAL

When to use	At the start of any planning cycle, before commitments are made.
When NOT to use	Do not use it to justify inaction — it defines what a realistic plan looks like.
Required inputs	Capacity, budget, dev availability, publishing rate, approval bottlenecks.
Databases used	SEO Strategy

COPY-PASTE PROMPT

ROLE: You are documenting the constraints an SEO plan must fit inside.

OBJECTIVE: Make the real limits explicit before the plan is written.

CONTEXT: People and hours: [DESCRIBE]. Publishing capacity: [PIECES/MONTH]. Dev capacity: [DESCRIBE]. Approval steps and their typical duration: [DESCRIBE]. Budget: [DESCRIBE].

CONSTRAINTS: Record actual observed capacity, not stated intent. Where capacity is unknown, say what to measure for four weeks.

PROCESS: 1 List each constraint with its observed value. 2 Identify which one binds hardest. 3 State what a plan fitting these constraints looks like in volume terms. 4 Identify the constraint most cheaply relaxed.

OUTPUT: A constraint register, the binding constraint named, and the realistic monthly output it implies.

QUALITY CHECKS: Values are observed rather than aspirational; one binding constraint is named.

MISSING INFORMATION: Propose a four-week measurement rather than estimating.

HUMAN REVIEW: The team confirms the observed capacity is accurate.

Expected output	A constraint register with the binding constraint and realistic output volume.
Quality checklist	Observed values only; one binding constraint identified.
Failure conditions	Proposes measurement rather than estimating unknown capacity.
Next prompt	PR-30

PR-30 Quarterly strategy review

STRATEGY · ADVANCED

When to use	Two weeks before each quarter, using the accumulated weekly reports.
When NOT to use	Do not run it without the previous quarter's reports — the reconstruction is the point.
Required inputs	All weekly reports, bet outcomes, where effort actually went.
Databases used	SEO Strategy, Reports, Decisions & Change Log

COPY-PASTE PROMPT

ROLE: You are running a quarterly SEO strategy review.

OBJECTIVE: Compare the stated strategy with what actually happened, and revise.

CONTEXT: Stated strategy: [PASTE]. Weekly reports: [PASTE ALL]. Where effort went: [DESCRIBE]. Bet outcomes against their signals: [PASTE].

CONSTRAINTS: Reconstruct the strategy from where effort went before comparing. Name the gap plainly. Do not forecast. Fit the next plan to the constraint register.

PROCESS: 1 Reconstruct actual strategy from effort. 2 Compare with stated. 3 Assess each bet against its signal — informative or inconclusive. 4 Propose three bets for next quarter. 5 Write the not-doing list.

OUTPUT: The stated-versus-actual gap, bet verdicts, three new bets with signals and owners, and the not-doing list.

QUALITY CHECKS: The gap is named explicitly; every new bet has a falsifiable signal.

MISSING INFORMATION: If reports are missing, say which quarters cannot be assessed.

HUMAN REVIEW: The business owner approves the revised strategy and versions it.

Expected output	A stated-versus-actual gap, bet verdicts, and three new owned bets.
Quality checklist	The gap is named; every bet signal could be shown false.
Failure conditions	Names unassessable periods rather than inferring them.
Next prompt	PR-31

PR-31 Call mining for demand evidence

AUDIENCE AND DEMAND · FOUNDATIONAL

When to use	After any batch of sales or support calls with transcripts available.
When NOT to use	Do not use it on summarized notes — paraphrase destroys the language bank.
Required inputs	Verbatim transcripts, labelled by source, role, and date.
Databases used	Customer Problems, Voice of Customer, Question Bank

COPY-PASTE PROMPT

ROLE: You are mining call transcripts for demand evidence.

OBJECTIVE: Extract verbatim problems, questions, and objections with sources.

CONTEXT: Transcripts: [PASTE, each labelled with role, date, and call type]. Our ICP: [DESCRIBE].

CONSTRAINTS: Verbatim only — never smooth a quote. Count distinct people, not mentions. Mark anything raised by one person as a single data point.

PROCESS: 1 Extract every problem statement verbatim with its source. 2 Extract every question asked, in the form asked. 3 Extract objections separately. 4 Group problems by underlying cause. 5 Note the trigger where stated.

OUTPUT: Three tables — problems with distinct-person counts, questions with frequency, objections with context — each row carrying its source and date.

QUALITY CHECKS: All quotes verbatim; distinct-person counts, not mention counts.

MISSING INFORMATION: If transcripts are summaries, stop and say the output would be unreliable.

HUMAN REVIEW: The content lead validates before records enter the database.

Expected output	Problem, question, and objection tables with verbatim quotes and sources.
Quality checklist	Verbatim only; distinct-person counts; single mentions flagged.
Failure conditions	Stops on summarized rather than verbatim input.
Next prompt	PR-32

PR-32 Review mining

AUDIENCE AND DEMAND · FOUNDATIONAL

When to use	Quarterly, on category reviews rather than only your own.
When NOT to use	Do not use it as a satisfaction measure for your own product — that is a different job.
Required inputs	Review text with sources and dates, for you and comparable products.
Databases used	Customer Problems, Voice of Customer

COPY-PASTE PROMPT

ROLE: You are analysing category review text for demand evidence.

OBJECTIVE: Find what buyers praise and complain about across the category.

CONTEXT: Review text: [PASTE with source and date]. Our category: [DESCRIBE].

CONSTRAINTS: Separate product complaints from onboarding, support, pricing, and expectation complaints — they need different responses. Rank by frequency, never by how strongly a review is worded. Cite the source of each pattern.

PROCESS: 1 Categorize every complaint by underlying cause. 2 Count by category. 3 Do the same for praise. 4 Identify which complaints we could credibly address. 5 Extract the exact language buyers use.

OUTPUT: Ranked complaint and praise tables with counts and sources, the addressable complaints, and the language to reuse.

QUALITY CHECKS: Frequency drives ranking; complaint types separated; single reviews are not treated as trends.

MISSING INFORMATION: If review dates are absent, note that recency cannot be assessed.

HUMAN REVIEW: Marketing approves any language before it enters copy.

Expected output	Ranked complaint and praise tables with sources and addressable items.
Quality checklist	Frequency-ranked; categories separated; no single-review trends.
Failure conditions	Flags when recency cannot be assessed.
Next prompt	PR-33

PR-33 Community research

AUDIENCE AND DEMAND · INTERMEDIATE

When to use	When you need to know how a problem is discussed where you have no presence.
When NOT to use	Do not use it to harvest personal information or to plan covert participation.
Required inputs	Thread content you have gathered, with the community and date noted.
Databases used	Customer Problems, Question Bank, Search Demand

COPY-PASTE PROMPT

ROLE: You are analysing community discussion for demand signals.

OBJECTIVE: Understand how the problem is discussed by people who are not our customers.

CONTEXT: Threads: [PASTE with community and date]. Our category: [DESCRIBE].

CONSTRAINTS: Record no personal information about any individual. Do not propose covert participation, undisclosed promotion, or astroturfing. Distinguish a recurring theme from one vocal participant.

PROCESS: 1 Extract the problems described, verbatim. 2 Note what solutions are recommended and by whom in role terms. 3 Identify vocabulary differences from our own. 4 Flag recurring unanswered questions. 5 Note where our category is discussed without us.

OUTPUT: A theme table with frequency, the vocabulary gaps, the unanswered questions, and where we are absent from the conversation.

QUALITY CHECKS: No personal data recorded; themes distinguished from single voices; no covert tactics proposed.

MISSING INFORMATION: Note which communities were not covered.

HUMAN REVIEW: Any participation plan is approved by marketing and disclosed.

Expected output	A theme table, vocabulary gaps, unanswered questions, and absence signals.
Quality checklist	No personal data; no covert tactics; themes separated from single voices.
Failure conditions	Names uncovered communities as a limit on the analysis.
Next prompt	PR-34

PR-34 Jobs-to-be-done mapping

AUDIENCE AND DEMAND · INTERMEDIATE

When to use	When feature requests and content topics keep drifting apart.
When NOT to use	Do not use it without interview evidence — inventing emotional jobs is worse than having none.
Required inputs	Interview notes, support themes, call summaries with sources.
Databases used	Customer Problems, Topics

COPY-PASTE PROMPT

ROLE: You are building a jobs-to-be-done map from evidence.

OBJECTIVE: Identify what customers hire the product to do, and what blocks them.

CONTEXT: Evidence: [PASTE notes, transcripts, tickets with sources]. Our offer: [DESCRIBE].

CONSTRAINTS: Ground every job in a cited line of evidence. Do not invent emotional or social jobs nobody expressed. Mark inference clearly.

PROCESS: 1 Identify functional jobs with their evidence. 2 Identify emotional and social jobs only where expressed. 3 For each, note the trigger situation, the current alternative, the definition of progress, and the switching anxiety. 4 Rank by evidence weight.

OUTPUT: A jobs table with trigger, alternative, progress definition, anxiety, and the cited evidence per row.

QUALITY CHECKS: Every job cites a line; unexpressed jobs are omitted rather than inferred.

MISSING INFORMATION: If only functional evidence exists, say so and omit the other categories.

HUMAN REVIEW: The content lead validates before it drives topic selection.

Expected output	A jobs table with triggers, alternatives, anxieties, and cited evidence.
Quality checklist	Every job cites evidence; nothing inferred is presented as expressed.
Failure conditions	Omits categories with no evidence rather than inventing them.
Next prompt	PR-35

PR-35 Emerging demand detection

AUDIENCE AND DEMAND · ADVANCED

When to use	Monthly, to find demand before keyword tools register it.
When NOT to use	Do not use it to justify a topic your ICP has never mentioned.
Required inputs	Recent community, support, and assistant-query evidence with dates.
Databases used	Search Demand, Question Bank, Topics

COPY-PASTE PROMPT

ROLE: You are detecting emerging demand ahead of tool data.

OBJECTIVE: Identify questions gaining frequency that tools do not yet show.

CONTEXT: Recent evidence: [PASTE with dates]. Previous cycle's emerging list: [PASTE]. Our ICP: [DESCRIBE].

CONSTRAINTS: Emerging means rising frequency in recent evidence, not a topic that sounds novel. Compare against the previous cycle. Do not treat a single new phrasing as a trend.

PROCESS: 1 Extract questions and phrasings from the last 90 days. 2 Compare with the previous cycle. 3 Identify what is genuinely new or rising. 4 Check whether tools show volume — absence supports emergence. 5 Assess ICP relevance.

OUTPUT: An emerging demand table with evidence, cycle-on-cycle change, tool volume status, and ICP relevance.

QUALITY CHECKS: Comparison against a previous cycle is present; single instances are not promoted.

MISSING INFORMATION: Without a previous cycle, state that emergence cannot be assessed and record a baseline instead.

HUMAN REVIEW: The SEO lead decides what enters the keyword list.

Expected output	An emerging demand table with cycle comparison and relevance assessment.
Quality checklist	Cycle comparison present; novelty distinguished from rising frequency.
Failure conditions	Records a baseline rather than claiming emergence on the first run.
Next prompt	PR-36

PR-36 Pain-to-query mapping

AUDIENCE AND DEMAND · INTERMEDIATE

When to use	Between demand research and keyword research, to preserve the evidence trail.
When NOT to use	Do not use it to attach a keyword to a pain nobody actually described.
Required inputs	Validated customer problems, the question bank, keyword candidates.
Databases used	Customer Problems, Search Demand, Keyword Master

COPY-PASTE PROMPT

ROLE: You are mapping evidenced customer problems to the queries they generate.

OBJECTIVE: Ensure every keyword traces to a stated problem.

CONTEXT: Validated problems: [PASTE with distinct-person counts]. Question bank: [PASTE]. Keyword candidates: [PASTE].

CONSTRAINTS: A keyword maps to a problem only if the connection is evident from customer language. Keywords with no traceable problem are flagged, not forced.

PROCESS: 1 For each problem, list the queries someone with it would type or ask. 2 Match against the keyword candidates. 3 Flag keywords with no problem behind them. 4 Flag problems with no keyword coverage. 5 Note conversational phrasings separately for the prompt universe.

OUTPUT: A mapping table, the orphan keywords, the uncovered problems, and the conversational variants.

QUALITY CHECKS: Every mapping is defensible from customer language; orphans are flagged rather than assigned.

MISSING INFORMATION: If problems are unvalidated, say the mapping is provisional.

HUMAN REVIEW: The SEO specialist confirms before keywords are scored.

Expected output	A mapping table with orphan keywords and uncovered problems named.
Quality checklist	Mappings defensible from customer language; orphans flagged.
Failure conditions	Marks the mapping provisional on unvalidated problems.
Next prompt	PR-37

PR-37 Search surface map

AUDIENCE AND DEMAND · INTERMEDIATE

When to use	Quarterly, to know where each question is actually asked.
When NOT to use	Do not use it to justify presence on a surface your buyers do not use.
Required inputs	Question bank, observed buyer behaviour, surface evidence.
Databases used	Search Demand, Question Bank, AI Prompt Universe

COPY-PASTE PROMPT

ROLE: You are mapping questions to the surfaces where they are asked.

OBJECTIVE: Decide where coverage is needed, beyond classic search.

CONTEXT: Questions: [PASTE]. What we know about where buyers look: [PASTE evidence]. Surfaces we can realistically serve: [LIST].

CONSTRAINTS: Base surface assignment on evidence, not assumption. Mark inferred assignments. Do not recommend a surface we have no capacity to serve.

PROCESS: 1 Assign each question to the surfaces where it is evidenced. 2 Note which surfaces we currently serve. 3 Identify the highest-frequency questions on surfaces we ignore. 4 Assess capacity to serve each. 5 Recommend at most two new surfaces.

OUTPUT: A question-to-surface table, the coverage gaps, and at most two recommended surfaces with the capacity implication.

QUALITY CHECKS: Assignments cite evidence; recommendations respect capacity.

MISSING INFORMATION: Where surface evidence is absent, mark it unknown rather than assuming Google.

HUMAN REVIEW: The SEO lead approves any new surface commitment.

Expected output	A question-to-surface map with gaps and capacity-bounded recommendations.
Quality checklist	Evidence-based assignment; at most two new surfaces recommended.
Failure conditions	Marks surfaces unknown rather than defaulting to search.
Next prompt	PR-38

PR-38 Demand validation pack

AUDIENCE AND DEMAND · INTERMEDIATE

When to use	Before a topic enters the content build order.
When NOT to use	Do not use it to retro-justify a topic already commissioned.
Required inputs	The candidate topic, all supporting evidence, existing coverage.
Databases used	Customer Problems, Search Demand, Topics, Content

COPY-PASTE PROMPT

ROLE: You are validating whether a topic has real demand behind it.

OBJECTIVE: Produce a pass or fail verdict against six validation criteria.

CONTEXT: Candidate topic: [DESCRIBE]. Evidence: [PASTE all]. Existing coverage: [PASTE]. Our differentiating evidence: [DESCRIBE].

CONSTRAINTS: Six criteria – three or more distinct people, a commercial path, something specific to say, customer vocabulary, a known surface, and an extractable answer. All six must pass, or the verdict is fail.

PROCESS: 1 Test each criterion against the evidence. 2 State pass or fail per criterion with the evidence. 3 Give an overall verdict. 4 If it fails, say what evidence would change it.

OUTPUT: A six-criterion verdict table, an overall verdict, and the evidence needed to change a fail.

QUALITY CHECKS: All six tested; the overall verdict follows the criteria rather than enthusiasm.

MISSING INFORMATION: A criterion with no evidence is a fail, not a pass by default.

HUMAN REVIEW: The content lead may override, but must record why.

Expected output	A six-criterion verdict with the evidence needed to reverse a fail.
Quality checklist	All six criteria tested; missing evidence fails rather than passes.
Failure conditions	Treats absent evidence as failure, not neutrality.
Next prompt	PR-39

PR-39 Seed keyword expansion

KEYWORDS · FOUNDATIONAL

When to use	At the start of keyword research, after demand evidence exists.
When NOT to use	Do not use it as the starting point — seeds should come from customer language.
Required inputs	Seed terms drawn from the demand layer, our offer, competitor vocabulary.
Databases used	Search Demand, Keyword Master

COPY-PASTE PROMPT

ROLE: You are expanding seed keywords into a candidate set.

OBJECTIVE: Produce a broad candidate list organized by the sixteen query types.

CONTEXT: Seeds from customer language: [PASTE]. Our offer: [DESCRIBE]. Competitor vocabulary: [PASTE]. Markets: [LIST].

CONSTRAINTS: Expand from customer vocabulary first, internal naming last. Do not invent search volumes. Organize output by query type so gaps are visible.

PROCESS: 1 Expand each seed across the sixteen types. 2 Mark which came from customer language and which from our own. 3 Identify types with no candidates — those are coverage gaps. 4 Flag terms only we use.

OUTPUT: A candidate list grouped by query type, source-labelled, with the empty types named as gaps.

QUALITY CHECKS: Customer-language origin is labelled; no invented volumes; empty types surfaced.

MISSING INFORMATION: If no customer vocabulary was supplied, say the expansion will reflect internal language only.

HUMAN REVIEW: The SEO specialist prunes before scoring.

Expected output	A type-organized candidate list with source labels and coverage gaps.
Quality checklist	Customer language prioritized and labelled; no invented volumes.
Failure conditions	Warns when expansion reflects internal vocabulary only.
Next prompt	PR-40

PR-40 BOFU query mining

KEYWORDS · INTERMEDIATE

When to use	When the pipeline is full of traffic content and short of revenue-adjacent pages.
When NOT to use	Do not use it on a site with no commercial pages to point BOFU traffic at.
Required inputs	Product knowledge, pricing structure, competitor names, sales objections.
Databases used	Keyword Master, Content, Question Bank

COPY-PASTE PROMPT

ROLE: You are mining bottom-of-funnel queries.

OBJECTIVE: Find the queries closest to a purchase decision that we do not cover.

CONTEXT: What we sell: [DESCRIBE]. Pricing structure: [DESCRIBE]. Competitors: [LIST]. Sales objections: [PASTE]. Existing commercial pages: [LIST].

CONSTRAINTS: BOFU means the searcher is choosing, not learning. Do not include informational queries. Every suggestion must have a plausible commercial destination on our site or be flagged as needing one.

PROCESS: 1 Generate pricing, comparison, alternative, review, and integration queries. 2 Generate objection-shaped queries from the sales objections. 3 Map each to an existing page or flag it as a gap. 4 Rank by proximity to decision.

OUTPUT: A BOFU query table with intent, destination or gap, and decision proximity ranking.

QUALITY CHECKS: No informational queries included; every query has a destination or an explicit gap.

MISSING INFORMATION: If pricing structure is undisclosed, note which query types cannot be served honestly.

HUMAN REVIEW: The SEO lead confirms before commissioning commercial pages.

Expected output	A BOFU query table ranked by decision proximity with destinations or gaps.
Quality checklist	Only decision-stage queries; every query has a destination or a flagged gap.
Failure conditions	Notes which queries cannot be served without disclosed pricing.
Next prompt	PR-41

PR-41 Comparison and alternatives mining

KEYWORDS · INTERMEDIATE

When to use	When competitors are named in your sales calls or AI answers.
When NOT to use	Do not use it to produce misleading competitor claims — outputs feed honest comparisons only.
Required inputs	Competitor names, honest differentiators, known areas where they are stronger.
Databases used	Keyword Master, Content, AI Competitors

COPY-PASTE PROMPT

ROLE: You are mining comparison and alternative queries.

OBJECTIVE: Identify the vs and alternative queries worth covering honestly.

CONTEXT: Competitors: [LIST]. Our honest differentiators: [PASTE]. Where they are genuinely stronger: [PASTE]. Existing comparison pages: [LIST].

CONSTRAINTS: Every comparison page must be able to concede something true. Do not generate queries we could only answer by misrepresenting a competitor. Do not make factual claims about competitor products.

PROCESS: 1 Generate vs, alternative, and switching queries per competitor. 2 Assess whether we can answer each honestly. 3 Identify which need a concession to be credible. 4 Flag any query we should not target.

OUTPUT: A comparison query table with an honesty assessment, the required concession per page, and the queries to decline.

QUALITY CHECKS: Every targetable query can be answered honestly; concessions are identified.

MISSING INFORMATION: If we do not know where competitors are stronger, say the comparison would not be credible.

HUMAN REVIEW: Marketing approves every comparison claim before publication.

Expected output	A comparison query table with honesty assessments and required concessions.
Quality checklist	Every page can concede something true; no competitor product claims.
Failure conditions	Declines when competitor strengths are unknown.
Next prompt	PR-42

PR-42 Conversational query generation

KEYWORDS · INTERMEDIATE

When to use	When building the prompt universe or covering assistant-style questions.
When NOT to use	Do not use it to generate keyword-shaped strings — these must read as speech.
Required inputs	Question bank, buyer decisions, constraints buyers mention.
Databases used	Search Demand, AI Prompt Universe, Question Bank

COPY-PASTE PROMPT

ROLE: You are generating conversational queries as people ask assistants.

OBJECTIVE: Produce natural-language questions for coverage and prompt tracking.

CONTEXT: Question bank: [PASTE]. Buyer decisions: [LIST]. Constraints buyers mention: [PASTE, e.g. team size, budget, existing stack].

CONSTRAINTS: Phrase as a person speaks, including the constraint. No keyword strings. Each question must be answerable factually rather than promotionally.

PROCESS: 1 For each decision, write how someone would ask about it aloud. 2 Add constraint variations. 3 Add comparison and recommendation phrasings. 4 Mark which are suitable for the tracked prompt universe. 5 Note what a good answer contains.

OUTPUT: A conversational query set with constraint variants, prompt-universe candidates flagged, and good-answer definitions.

QUALITY CHECKS: Phrasing is speech-like; constraints included; good answers defined.

MISSING INFORMATION: Without a question bank, mark output as inferred and lower quality.

HUMAN REVIEW: Wording is fixed before any prompt enters the tracked set.

Expected output	A conversational query set with constraints and prompt-universe candidates.
Quality checklist	Speech-like phrasing; constraints present; good answers defined.
Failure conditions	Marks output inferred without a question bank.
Next prompt	PR-43

PR-43 Keyword-to-page mapping

KEYWORDS · FOUNDATIONAL

When to use	After scoring, before anything is commissioned.
When NOT to use	Do not use it to map several similar keywords to several similar pages.
Required inputs	Scored keyword list, existing page inventory with topics.
Databases used	Keyword Master, Content, Topic Clusters

COPY-PASTE PROMPT

ROLE: You are mapping scored keywords to owning pages.

OBJECTIVE: Assign exactly one owning page per keyword, or flag a gap.

CONTEXT: Scored keywords: [PASTE with intent and stage]. Existing pages: [PASTE URLs and topics]. Cluster plan: [PASTE].

CONSTRAINTS: One keyword, one owning page. A page may support several related variants but owns exactly one primary. Where two keywords want the same page, one must be reassigned or declined.

PROCESS: 1 Match each keyword to the page whose intent it fits. 2 Flag keywords with no suitable page as gaps. 3 Flag pages assigned more than one primary. 4 Resolve conflicts by intent, not by preference. 5 Recommend merge or differentiate where pages overlap.

OUTPUT: A mapping table, the gap list, the conflict list with resolutions, and merge recommendations.

QUALITY CHECKS: No page holds two primaries; conflicts resolved by intent.

MISSING INFORMATION: If page topics were not supplied, mapping is marked provisional.

HUMAN REVIEW: The content lead confirms before commissioning.

Expected output	A mapping table with gaps, conflicts, and merge recommendations.
Quality checklist	One primary per page; conflicts resolved on intent.
Failure conditions	Marks mapping provisional without page topic data.
Next prompt	PR-44

PR-44 Cannibalization diagnosis

KEYWORDS · ADVANCED

When to use	When ranking URLs for a query change back and forth over weeks.
When NOT to use	Do not use it to justify deleting a page before checking its links and traffic.
Required inputs	Query-level ranking URL history over 90 days, page inventory, internal links.
Databases used	Keyword Master, Content, Internal Link Opportunities

COPY-PASTE PROMPT

ROLE: You are diagnosing keyword cannibalization.

OBJECTIVE: Identify conflicts and recommend one resolution each.

CONTEXT: Ranking URL history by query, 90 days: [PASTE]. Page inventory: [PASTE]. Internal links to each competing page: [PASTE].

CONSTRAINTS: A conflict means two of our URLs alternate for one query, not merely that two pages mention it. Recommend one resolution per conflict: consolidate, differentiate, de-optimize, or redirect. Never recommend deletion without checking links and traffic.

PROCESS: 1 Identify queries with alternating URLs. 2 Determine which page should own each. 3 Check internal links – they often point at the wrong page. 4 Choose a resolution. 5 Specify the internal link changes required.

OUTPUT: A conflict table with the intended owner, the resolution, the internal link changes, and a 60-day recheck date.

QUALITY CHECKS: Alternation evidenced; one resolution each; link changes specified.

MISSING INFORMATION: Without ranking history, say the diagnosis is speculative.

HUMAN REVIEW: The SEO lead approves any redirect or removal.

Expected output	A conflict table with resolutions, link changes, and a recheck date.
Quality checklist	Alternation evidenced; one resolution per conflict; links addressed.
Failure conditions	Declares the diagnosis speculative without ranking history.
Next prompt	PR-45

PR-45 Ranking tracking rules

KEYWORDS · FOUNDATIONAL

When to use	Once at setup, and when the tracked set grows unmanageable.
When NOT to use	Do not use it to justify tracking everything — the discipline is what to exclude.
Required inputs	Keyword list with scores, business priorities, tool limits.
Databases used	Keyword Master, Reports

COPY-PASTE PROMPT

ROLE: You are defining what gets rank-tracked and how often.

OBJECTIVE: Produce a tracking set small enough to review and large enough to matter.

CONTEXT: Keyword list with priority scores: [PASTE]. Business priorities: [DESCRIBE]. Tool capacity: [DESCRIBE]. Markets and devices: [LIST].

CONSTRAINTS: Track what a decision depends on. Exclude anything you would not act on. Define frequency by volatility, not by habit. State the review cadence.

PROCESS: 1 Select the tracked set by score and business proximity. 2 Assign frequency per group. 3 Define the market and device split. 4 State what is deliberately untracked and why. 5 Define who reviews and when.

OUTPUT: A tracking specification with the set, frequency, splits, exclusions, and the review cadence.

QUALITY CHECKS: Exclusions are explicit; frequency justified by volatility.

MISSING INFORMATION: Without tool limits, produce the ideal set and note it may need trimming.

HUMAN REVIEW: The SEO lead approves the exclusion list.

Expected output	A tracking specification with an explicit exclusion list and cadence.
Quality checklist	Exclusions stated; frequency justified; a reviewer named.
Failure conditions	Produces an ideal set and flags trimming when limits are unknown.
Next prompt	PR-46

PR-46 Topic cluster design

CONTENT PLANNING · ADVANCED

When to use	When a topic has been prioritized and needs a build plan.
When NOT to use	Do not use it before keyword mapping — the cluster is built from mapped queries.
Required inputs	Topic record, mapped keywords, existing coverage, commercial destination.
Databases used	Topics, Topic Clusters, Keyword Master, Content

COPY-PASTE PROMPT

ROLE: You are designing a hub-and-spoke content cluster.

OBJECTIVE: Produce a complete build plan with interlinking specified.

CONTEXT: Topic: [DESCRIBE with our belief and evidence]. Mapped keywords: [PASTE]. Existing coverage: [PASTE]. Commercial destination: [URL].

CONSTRAINTS: One pillar, six to twelve spokes, and a commercial destination. Every spoke answers one question thoroughly. Do not plan a spoke for a query with no demand evidence.

PROCESS: 1 Define the pillar and what it must cover. 2 Assign each mapped keyword to a spoke or the pillar. 3 Identify spokes needing new content versus updates. 4 Specify interlinking in both directions. 5 Specify the path to the commercial page.

OUTPUT: A cluster plan with pillar scope, spoke list with keywords, new-versus-update marking, and the interlinking specification.

QUALITY CHECKS: Every spoke has demand evidence; interlinking is bidirectional; a commercial destination exists.

MISSING INFORMATION: If no commercial destination exists, flag the cluster as leading nowhere.

HUMAN REVIEW: The content lead approves before briefs are written.

Expected output	A cluster plan with pillar scope, spokes, and bidirectional interlinking.
Quality checklist	Every spoke evidenced; interlinking specified; commercial path present.
Failure conditions	Flags clusters with no commercial destination.
Next prompt	PR-47

PR-47 Content portfolio review

CONTENT PLANNING · ADVANCED

When to use	Quarterly, to see the shape of what you have rather than one page at a time.
When NOT to use	Do not use it on fewer than thirty pages — the patterns will not be meaningful.
Required inputs	Full content inventory with stage, format, performance, and age.
Databases used	Content, Topic Clusters, Keyword Master

COPY-PASTE PROMPT

ROLE: You are reviewing a content portfolio at the collection level.

OBJECTIVE: Identify the imbalances that individual page reviews cannot show.

CONTEXT: Full inventory: [PASTE with stage, format, cluster, performance, age]. Business priorities: [DESCRIBE].

CONSTRAINTS: Analyse distribution, not individual pages. Report counts and percentages. Do not recommend removals here — that is a per-page decision.

PROCESS: 1 Distribution by funnel stage. 2 Distribution by format. 3 Distribution by cluster and coverage completeness. 4 Age and freshness profile. 5 Performance distribution. 6 Identify the three most consequential imbalances.

OUTPUT: Five distribution tables and the three imbalances with their commercial implication.

QUALITY CHECKS: Counts and percentages given; imbalances tied to a commercial implication.

MISSING INFORMATION: If performance data is absent, produce structural distributions only and say so.

HUMAN REVIEW: The content lead decides what the imbalances mean for the plan.

Expected output	Five distribution tables and three consequential imbalances.
Quality checklist	Percentages given; imbalances tied to commercial implications.
Failure conditions	Produces structural analysis only without performance data.
Next prompt	PR-48

PR-48 Editorial calendar construction

CONTENT PLANNING · INTERMEDIATE

When to use	At the start of a quarter, from the build order and the constraint register.
When NOT to use	Do not use it to plan volume beyond the observed publishing capacity.
Required inputs	Build order, capacity, author availability, refresh queue.
Databases used	Content, Content Briefs, Content Refresh Queue

COPY-PASTE PROMPT

ROLE: You are building a quarterly editorial calendar.

OBJECTIVE: Schedule production within real capacity, refresh work included.

CONTEXT: Build order: [PASTE]. Refresh queue: [PASTE]. Observed capacity: [PIECES/MONTH]. Authors and availability: [DESCRIBE]. Cluster priorities: [LIST].

CONSTRAINTS: Refresh work occupies capacity – schedule it explicitly, not as overflow. Do not exceed observed capacity. Leave a fifth of capacity unscheduled for reactive work.

PROCESS: 1 Allocate capacity between new and refresh. 2 Sequence new pieces so pillars precede spokes. 3 Schedule refresh work by decline severity. 4 Assign authors by expertise. 5 Reserve reactive capacity.

OUTPUT: A twelve-week calendar with piece, type, author, due date, and the reserved capacity shown.

QUALITY CHECKS: Total scheduled is at or below 80% of capacity; refresh work is explicit; pillars precede spokes.

MISSING INFORMATION: If capacity is unknown, produce the sequence without dates.

HUMAN REVIEW: Authors confirm their assignments before the calendar is published.

Expected output	A twelve-week calendar with refresh work scheduled and capacity reserved.
Quality checklist	Under 80% of capacity; refresh explicit; pillars first.
Failure conditions	Produces sequence without dates when capacity is unknown.
Next prompt	PR-49

PR-49 Original research design

CONTENT PLANNING · ADVANCED

When to use	When you have first-party data and want a genuinely citable asset.
When NOT to use	Do not use it where the data cannot support a finding — a weak study damages credibility.
Required inputs	The available data, its provenance, permissions, and known limitations.
Databases used	Content, Digital PR Campaigns

COPY-PASTE PROMPT

ROLE: You are designing an original research asset from first-party data.

OBJECTIVE: Determine what the data can honestly support, and design around it.

CONTEXT: Data available: [DESCRIBE, including how it was collected]. Sample size: [N]. Permissions: [WHAT MAY BE PUBLISHED]. Known limitations: [DESCRIBE].

CONSTRAINTS: Never propose a finding the data cannot support. State sample size and methodology alongside every finding. Do not extrapolate to a population the sample does not represent. Anonymize completely.

PROCESS: 1 State what the data can and cannot support. 2 Identify the findings that are genuinely novel. 3 Design the methodology statement. 4 Draft the limitations section before the findings. 5 Identify the strongest citable claim.

OUTPUT: A research design with supportable findings, the methodology statement, the limitations section, and the citable claim.

QUALITY CHECKS: Limitations written before findings; no extrapolation beyond the sample; anonymization confirmed.

MISSING INFORMATION: If the sample cannot support a finding, say so and recommend not publishing.

HUMAN REVIEW: The data owner and legal confirm what may be published.

Expected output	A research design with supportable findings, methodology, and limitations.
Quality checklist	Limitations precede findings; no extrapolation; anonymization confirmed.
Failure conditions	Recommends not publishing when the sample cannot support a finding.
Next prompt	PR-50

PR-50 Case study production

CONTENT PLANNING · INTERMEDIATE

When to use	When a customer has a documented outcome and has agreed to be featured.
When NOT to use	Do not use it without written permission and verified figures from the customer.
Required inputs	Customer permission, verified outcome data, their own description of the change.
Databases used	Content, Voice of Customer

COPY-PASTE PROMPT

ROLE: You are structuring a customer case study.

OBJECTIVE: Produce a credible account built on the customer's own figures.

CONTEXT: Customer: [NAME, with permission status]. Their situation before: [PASTE their words]. What changed: [PASTE]. Their figures: [PASTE with source]. Permission scope: [WHAT MAY BE NAMED].

CONSTRAINTS: Only figures the customer supplied and approved. Never estimate an outcome. If permission is for anonymized use only, describe the profile rather than naming them. Include what did not go smoothly.

PROCESS: 1 Their situation in their words. 2 What they tried before. 3 What changed and how. 4 Their figures with their as-of dates. 5 What was harder than expected. 6 What they would tell a peer.

OUTPUT: The case study draft, the figure list with sources, and the permission scope restated at the top.

QUALITY CHECKS: Every figure customer-supplied; permission scope respected; a genuine difficulty included.

MISSING INFORMATION: Without written permission, produce the structure only and stop.

HUMAN REVIEW: The customer approves the final text before publication.

Expected output	A case study draft built on customer-supplied figures with permission stated.
Quality checklist	Customer figures only; permission respected; a real difficulty included.
Failure conditions	Stops at structure without written permission.
Next prompt	PR-51

PR-51 Glossary and definition system

CONTENT PLANNING · FOUNDATIONAL

When to use	When your category has terms buyers search and you want entity clarity.
When NOT to use	Do not use it to publish definitions that merely restate what already ranks.
Required inputs	Category terms, customer vocabulary, our own definitions.
Databases used	Content, Question Bank, Brand Facts

COPY-PASTE PROMPT

ROLE: You are designing a glossary for entity clarity and extractability.

OBJECTIVE: Produce definitions that are quotable and consistent with our positioning.

CONTEXT: Terms: [LIST]. Customer vocabulary: [PASTE]. Our category description: [PASTE from the brand fact sheet].

CONSTRAINTS: Each definition must be one sentence that stands alone accurately. No promotional language inside a definition. Definitions must be consistent with the brand fact sheet. Do not define a term we have nothing to add on.

PROCESS: 1 Write a one-sentence definition per term. 2 Add the practical explanation separately. 3 Check consistency against the fact sheet. 4 Identify terms where our definition differs from the common one, and say why. 5 Flag terms to omit.

OUTPUT: A glossary with one-sentence definitions, practical explanations, fact sheet consistency checks, and the omission list.

QUALITY CHECKS: Definitions stand alone; no promotional language; fact sheet consistency confirmed.

MISSING INFORMATION: Without a fact sheet, note that consistency is unverified.

HUMAN REVIEW: Marketing confirms definitions match positioning.

Expected output	A glossary with standalone definitions and fact-sheet consistency checks.
Quality checklist	One-sentence definitions; no promotion; consistency verified.
Failure conditions	Flags unverified consistency without a fact sheet.
Next prompt	PR-52

PR-52 Content consolidation decision

CONTENT PLANNING · ADVANCED

When to use	When two or more pages compete, overlap, or individually underperform.
When NOT to use	Do not use it on pages with meaningful backlinks until the link impact is assessed.
Required inputs	The candidate pages, their performance, backlinks, and internal links.
Databases used	Content, Backlinks, Internal Link Opportunities

COPY-PASTE PROMPT

ROLE: You are deciding whether to consolidate overlapping pages.
OBJECTIVE: Choose one action and specify the migration.
CONTEXT: Candidate pages: [PASTE with content summaries]. Performance each: [PASTE with dates]. Backlinks each: [PASTE]. Internal links each: [PASTE].
CONSTRAINTS: Assess link equity before recommending removal. The surviving page must genuinely cover what the merged pages covered. Specify redirects explicitly.
PROCESS: 1 Identify overlap precisely. 2 Determine which page should survive and why. 3 List what must be carried across. 4 Specify redirects for every retired URL. 5 Specify internal link updates. 6 Set a 60-day recheck.
OUTPUT: A consolidation decision with the surviving URL, the carry-across list, the redirect map, the link updates, and the recheck date.
QUALITY CHECKS: Link equity assessed; redirects specified per URL; carry-across list complete.
MISSING INFORMATION: Without backlink data, recommend gathering it before acting.
HUMAN REVIEW: The SEO lead approves every redirect.

Expected output	A consolidation decision with redirects, carry-across list, and recheck date.
Quality checklist	Link equity assessed; redirects per URL; nothing lost in the merge.
Failure conditions	Defers the decision until backlink data exists.
Next prompt	PR-53

PR-53 Publishing handoff

CONTENT PLANNING · FOUNDATIONAL

When to use	When an approved draft moves to publication.
When NOT to use	Do not use it before QA has passed — this is a handoff, not a review.
Required inputs	The approved draft, the brief, metadata, internal link plan.
Databases used	Content, Content Briefs, Internal Link Opportunities

COPY-PASTE PROMPT

ROLE: You are preparing a publishing handoff package.
OBJECTIVE: Ensure nothing in the brief is lost between approval and publication.
CONTEXT: Approved draft: [PASTE]. Brief: [PASTE]. Metadata: [PASTE]. Internal link plan: [PASTE]. Publishing platform requirements: [DESCRIBE].
CONSTRAINTS: Every brief requirement must be verifiably present. Every internal link in and out must be listed with its source page. Do not publish without an author with a visible credential.
PROCESS: 1 Verify each brief requirement is met. 2 Confirm metadata is final. 3 List links out to add on this page. 4 List links in to add on other pages. 5 Confirm author and credential. 6 Set the performance review date.
OUTPUT: A handoff checklist with pass or fail per item, the two link lists, and the review date.
QUALITY CHECKS: Every brief requirement verified; both link lists produced; author credential confirmed.
MISSING INFORMATION: A missing brief requirement is a fail, not a note.
HUMAN REVIEW: The editor signs the handoff before the page goes live.

Expected output	A pass/fail handoff checklist with link lists and a review date.
Quality checklist	Every brief requirement verified; both link directions listed.
Failure conditions	Fails on any unmet brief requirement.
Next prompt	PR-54

PR-54 Content distribution plan

CONTENT PLANNING · INTERMEDIATE

When to use	When a substantial piece publishes and deserves more than one post about it.
When NOT to use	Do not use it on routine pieces — distribution effort should follow asset value.
Required inputs	The published piece, the audience, available channels and their norms.
Databases used	Content, Authority Targets

COPY-PASTE PROMPT

ROLE: You are planning distribution for a published asset.

OBJECTIVE: Extract every legitimate distribution opportunity from one piece.

CONTEXT: The piece: [PASTE or summarize]. Audience: [DESCRIBE]. Channels available: [LIST with norms]. Communities where we have standing: [LIST].

CONSTRAINTS: Adapt to each channel rather than reposting. Do not propose posting in communities where we have no standing or where promotion is against the norms. No undisclosed promotion anywhere.

PROCESS: 1 Extract the standalone ideas. 2 Adapt each to the channels where it fits. 3 Identify communities where sharing would be welcome and disclosed. 4 Identify who should be told directly. 5 Sequence over four weeks.

OUTPUT: A distribution plan by channel with adapted assets, a four-week schedule, and the channels deliberately excluded.

QUALITY CHECKS: Assets adapted rather than copied; no undisclosed promotion; exclusions stated.

MISSING INFORMATION: Where channel norms are unknown, exclude the channel and say why.

HUMAN REVIEW: Marketing approves anything going to an external community.

Expected output	A channel-adapted distribution plan with a schedule and exclusions.
Quality checklist	Adaptation not reposting; disclosure everywhere; exclusions stated.
Failure conditions	Excludes channels whose norms are unknown.
Next prompt	PR-55

PR-55 Content measurement loop

CONTENT PLANNING · INTERMEDIATE

When to use	Sixty and ninety days after a piece publishes.
When NOT to use	Do not use it before sixty days — you will be reading noise.
Required inputs	The piece, its brief, its target keyword, and performance since publication.
Databases used	Content, Keyword Master, Reports

COPY-PASTE PROMPT

ROLE: You are reviewing whether a published piece did what the brief intended.

OBJECTIVE: Produce a verdict and one next action.

CONTEXT: The piece: [URL and summary]. Brief intent: [PASTE]. Target keyword: [PASTE]. Performance since publication: [PASTE with dates]. Internal links added: [LIST].

CONSTRAINTS: Compare against the brief's stated intent, not against a general expectation. Sixty days minimum. Distinguish a ranking problem from a conversion problem from a demand problem.

PROCESS: 1 State what the brief intended. 2 Compare against actual performance. 3 Classify any shortfall by type. 4 Check whether the planned internal links were actually added. 5 Recommend one action.

OUTPUT: A verdict against brief intent, the shortfall classification, the link verification, and one next action with an owner.

QUALITY CHECKS: Measured against the brief; shortfall type classified; links verified.

MISSING INFORMATION: Under sixty days, decline to judge and set a review date.

HUMAN REVIEW: The content lead decides the next action.

Expected output	A verdict against brief intent with a classified shortfall and one action.
Quality checklist	Brief-relative judgement; shortfall classified; links verified.
Failure conditions	Declines to judge before sixty days.
Next prompt	PR-56

PR-56 Product-led content design

CONTENT PLANNING · ADVANCED

When to use	When product knowledge is the differentiator and generic content will not compete.
When NOT to use	Do not use it to produce a feature tour disguised as a guide.
Required inputs	The workflow being explained, product capability, honest limitations.
Databases used	Content, Topics, Question Bank

COPY-PASTE PROMPT

ROLE: You are designing product-led content that teaches rather than sells.

OBJECTIVE: Explain how to do a job, using the product where it genuinely helps.

CONTEXT: The job: [DESCRIBE]. How our product helps: [DESCRIBE]. Where it does not: [DESCRIBE]. Alternatives readers might use: [LIST].

CONSTRAINTS: The piece must be useful to someone who never buys. Mention the product where it is genuinely the answer, and say where it is not. Acknowledge the alternatives honestly.

PROCESS: 1 Structure the job as steps someone could follow. 2 Identify where the product materially helps. 3 Identify where a general approach is sufficient. 4 State the limitation honestly. 5 Place the commercial link where it is genuinely useful.

OUTPUT: A content structure with product placement justified per section, the stated limitations, and the commercial link placement.

QUALITY CHECKS: Useful without purchase; limitations stated; product mentions justified.

MISSING INFORMATION: If the product does not materially help, recommend a non-product piece.

HUMAN REVIEW: The editor checks the balance before commissioning.

Expected output	A content structure with justified product placement and stated limitations.
Quality checklist	Useful without buying; limitations honest; placements justified.
Failure conditions	Recommends a non-product piece where the product does not help.
Next prompt	PR-57

PR-57 Expert interview design

CONTENT PLANNING · INTERMEDIATE

When to use	When the differentiator is expertise you have access to and competitors do not.
When NOT to use	Do not use it to fabricate expert commentary — the expert must be real and named.
Required inputs	The expert, their credentials, the topic, what only they can say.
Databases used	Content, Voice of Customer

COPY-PASTE PROMPT

ROLE: You are designing an expert interview to produce citable content.

OBJECTIVE: Extract what only this person can say about the topic.

CONTEXT: Expert: [NAME AND CREDENTIALS]. Topic: [DESCRIBE]. What they have done that others have not: [DESCRIBE]. Audience: [DESCRIBE].

CONSTRAINTS: Questions must produce specifics, not opinions. No question answerable by anyone in the field. The expert must be named with a verifiable credential. Do not put words in their mouth.

PROCESS: 1 Identify what they uniquely know. 2 Write questions that surface specifics — decisions, mistakes, numbers, situations. 3 Include one question inviting disagreement with common practice. 4 Note what a citable answer sounds like. 5 Plan the follow-up probes.

OUTPUT: A question set with probes, the citable-answer definition per question, and the credential to display.

QUALITY CHECKS: No question is answerable generically; the credential is verifiable; specifics are the target.

MISSING INFORMATION: Without a verifiable credential, note the content will not carry authority.

HUMAN REVIEW: The expert approves their quotes before publication.

Expected output	A question set with probes and a citable-answer definition for each.
Quality checklist	Questions produce specifics; credential verifiable; no fabrication.
Failure conditions	Flags absent credentials as an authority problem.
Next prompt	PR-58

PR-58 Heading architecture review

ON-PAGE · FOUNDATIONAL

When to use	On any draft or published page whose structure is unclear.
When NOT to use	Do not use it as a substitute for fixing an unclear argument — structure cannot rescue muddled thinking.
Required inputs	The full page text with its current headings.
Databases used	Content, Content Briefs

COPY-PASTE PROMPT

ROLE: You are reviewing the heading architecture of a page.

OBJECTIVE: Produce a hierarchy a reader could outline from and a system can parse.

CONTEXT: Page text with current headings: [PASTE]. Primary keyword: [KEYWORD]. Intent: [INTENT].

CONSTRAINTS: One H1. Headings describe content, not style. Do not add headings to break up text that is genuinely one idea. Do not stuff keywords into headings.

PROCESS: 1 Extract the current outline. 2 Identify where the hierarchy breaks. 3 Identify sections whose heading does not describe them. 4 Propose a corrected outline. 5 Flag sections that should merge or split.

OUTPUT: The current outline, the corrected outline, and the merge or split recommendations with reasons.

QUALITY CHECKS: One H1; every heading describes its section; no keyword stuffing.

MISSING INFORMATION: If intent was not supplied, state the assumption made.

HUMAN REVIEW: The editor approves structural changes to published pages.

Expected output	Current and corrected outlines with merge or split recommendations.
Quality checklist	One H1; descriptive headings; no stuffing.
Failure conditions	States its assumption when intent is missing.
Next prompt	PR-59

PR-59 Direct-answer block writing

ON-PAGE · FOUNDATIONAL

When to use	On any page targeting a question, before publication.
When NOT to use	Do not use it to add a summary box that repeats the intro without adding clarity.
Required inputs	The question, the correct answer, and any caveats.
Databases used	Content, Question Bank

COPY-PASTE PROMPT

ROLE: You are writing a direct-answer block for extractability.

OBJECTIVE: Produce a passage that answers the question accurately when quoted alone.

CONTEXT: The question: [QUESTION]. The correct answer: [PASTE]. Caveats: [PASTE]. Sources for any figure: [PASTE].

CONSTRAINTS: Under 60 words. Must remain accurate with no surrounding context. Every figure carries its unit, period, and source. No hedging that removes the answer. No promotional language.

PROCESS: 1 State the answer in one sentence. 2 Add the essential qualification. 3 Attach figures with units and sources. 4 Test the passage in isolation. 5 Rewrite anything that depends on prior context.

OUTPUT: The answer block, the isolation test result, and any caveat that could not be included without breaking the word limit.

QUALITY CHECKS: Accurate in isolation; under 60 words; figures fully attributed.

MISSING INFORMATION: If the answer requires a caveat that will not fit, say so rather than dropping it.

HUMAN REVIEW: The subject expert confirms accuracy.

Expected output	A sub-60-word answer block with an isolation test result.
Quality checklist	Accurate alone; figures attributed; no dropped caveats.
Failure conditions	Refuses to drop a caveat to meet the word limit.
Next prompt	PR-60

PR-60 Comparison table construction

ON-PAGE · INTERMEDIATE

When to use	On comparison and alternatives pages, where the table is what gets quoted.
When NOT to use	Do not use it to compare on criteria chosen because you win them.
Required inputs	The options, the buyer's stated criteria, and verifiable facts about each.
Databases used	Content, AI Competitors, Brand Facts

COPY-PASTE PROMPT

ROLE: You are building a comparison table that is safe to quote.
OBJECTIVE: Compare options on the buyer's criteria, honestly.
CONTEXT: Options: [LIST]. Buyer's criteria in their order: [PASTE]. Verifiable facts per option: [PASTE with sources]. Where others are stronger: [PASTE].
CONSTRAINTS: Use the buyer's criteria in the buyer's order. State the criterion before the verdict. Only verifiable facts about other products – no inference. Concede at least one criterion honestly.
PROCESS: 1 Set the criteria as rows in the buyer's order. 2 Fill only verifiable cells. 3 Mark unverifiable cells as not established rather than guessing. 4 Include the criterion where we are weaker. 5 Add the as-of date.
OUTPUT: The comparison table, the unverifiable cells listed, and the conceded criterion.
QUALITY CHECKS: Buyer's criteria and order preserved; a genuine concession present; no unverifiable claims about others.
MISSING INFORMATION: Unverifiable cells are marked, never filled.
HUMAN REVIEW: Marketing approves every competitor claim before publication.

Expected output	A comparison table with unverifiable cells marked and a genuine concession.
Quality checklist	Buyer criteria and order; concession present; no inference about others.
Failure conditions	Marks unverifiable cells rather than filling them.
Next prompt	PR-61

PR-61 FAQ block construction

ON-PAGE · FOUNDATIONAL

When to use	When real recurring questions exist in the question bank.
When NOT to use	Do not use it to invent questions for schema markup — that is a policy violation.
Required inputs	Questions from the question bank with frequency, and the correct answers.
Databases used	Question Bank, Content

COPY-PASTE PROMPT

ROLE: You are building an FAQ block from real questions.
OBJECTIVE: Answer recurring questions directly and extractably.
CONTEXT: Questions with frequency: [PASTE from the question bank]. Correct answers: [PASTE]. Page topic: [DESCRIBE].
CONSTRAINTS: Only questions people actually asked. Answers under 50 words, direct, and accurate in isolation. Do not include a question to justify markup. Every answer must be visible on the page, not only in the markup.
PROCESS: 1 Select questions by frequency and page relevance. 2 Write each answer directly, without preamble. 3 Test each answer in isolation. 4 Confirm each is visible on the page. 5 Note which need a source.
OUTPUT: The FAQ block with sources noted, the isolation test results, and any question excluded with the reason.
QUALITY CHECKS: Real questions only; answers visible on the page; accurate alone.
MISSING INFORMATION: Without a question bank, produce nothing and say why.
HUMAN REVIEW: The editor confirms every answer is visible before markup is added.

Expected output	An FAQ block from real questions with isolation test results.
Quality checklist	Real questions only; visible answers; accurate in isolation.
Failure conditions	Produces nothing without a real question source.
Next prompt	PR-62

PR-62 Evidence and citation pass

ON-PAGE · INTERMEDIATE

When to use	On any draft containing figures, claims, or assertions.
When NOT to use	Do not use it to add citations to claims that should simply be removed.
Required inputs	The draft and the sources available for each claim.
Databases used	Content, Brand Facts

COPY-PASTE PROMPT

ROLE: You are auditing the claims in a draft.

OBJECTIVE: Ensure every claim is sourced, or removed.

CONTEXT: Draft: [PASTE]. Available sources: [PASTE]. Approved brand facts: [PASTE].

CONSTRAINTS: Every non-obvious claim needs a source or removal – softening is not an option. Figures need unit, period, and origin. Brand facts must match the fact sheet exactly. Superlatives require external verification or deletion.

PROCESS: 1 Extract every claim. 2 Classify: sourced, sourceable, unsupportable. 3 Match brand facts against the sheet. 4 Flag superlatives. 5 Recommend source or removal per claim.

OUTPUT: A claim table with classification, source where available, and the removal list.

QUALITY CHECKS: Every claim classified; brand facts checked; superlatives flagged.

MISSING INFORMATION: Unsupportable claims are marked for removal, not softening.

HUMAN REVIEW: The editor applies removals before publication.

Expected output	A claim table with classifications, sources, and a removal list.
Quality checklist	Every claim classified; facts checked; superlatives flagged.
Failure conditions	Marks unsupportable claims for removal, never softening.
Next prompt	PR-63

PR-63 Author credibility review

ON-PAGE · FOUNDATIONAL

When to use	On any page where expertise affects whether the content is trusted.
When NOT to use	Do not use it to manufacture credentials — the credential must be real and verifiable.
Required inputs	The author, their actual credentials, and what the page claims.
Databases used	Content

COPY-PASTE PROMPT

ROLE: You are reviewing author credibility signals on a page.

OBJECTIVE: Ensure the person behind the content is identifiable and credible.

CONTEXT: Author: [NAME]. Their actual credentials: [PASTE]. Page topic: [DESCRIBE]. Current byline and bio: [PASTE].

CONSTRAINTS: Only real, verifiable credentials. The credential must be relevant to this topic specifically. Do not invent titles or affiliations. If the author has no relevant credential, say so.

PROCESS: 1 Assess whether the credential is relevant to this topic. 2 Check it is externally verifiable. 3 Draft a bio line that states it plainly. 4 Identify what else would strengthen credibility. 5 Flag if the wrong person is bylined.

OUTPUT: The credibility assessment, a drafted bio line, and what would strengthen it.

QUALITY CHECKS: Credential relevant and verifiable; no invention; a mismatch is flagged.

MISSING INFORMATION: If the credential is not verifiable externally, say the signal is weak.

HUMAN REVIEW: The author approves their own bio.

Expected output	A credibility assessment with a drafted bio line and strengthening options.
Quality checklist	Relevant, verifiable credentials; no invention; mismatches flagged.
Failure conditions	Reports a weak signal rather than inventing authority.
Next prompt	PR-64

PR-64 Crawl budget analysis

TECHNICAL SEO · ADVANCED

When to use	When index coverage or crawl frequency suggests waste.
When NOT to use	Do not use it without log data — everything else is inference.
Required inputs	Server logs filtered by crawler, plus the URL inventory.
Databases used	Technical Issues, Site Architecture

COPY-PASTE PROMPT

ROLE: You are analysing crawl budget distribution from log data.
OBJECTIVE: Find where crawl effort goes versus where it should.
CONTEXT: Log sample: [PASTE, filtered by user agent, with the period]. URL inventory with commercial importance: [PASTE]. Known parameter patterns: [LIST].
CONSTRAINTS: Log data only – do not infer crawl behaviour from index coverage. State the sample period. Do not recommend blocking anything without checking whether it is indexed and ranking.
PROCESS: 1 Group requests by URL pattern. 2 Compare share of crawl against commercial importance. 3 Identify patterns consuming disproportionate budget. 4 Check whether those URLs are indexed or ranking. 5 Recommend action per pattern.
OUTPUT: A crawl distribution table, the disproportionate patterns, their index status, and a recommended action each.
QUALITY CHECKS: Log-based only; sample period stated; index status checked before any blocking recommendation.
MISSING INFORMATION: Without logs, state that the analysis cannot be performed.
HUMAN REVIEW: Engineering confirms before any robots or parameter change.

Expected output	A crawl distribution table with actions per disproportionate pattern.
Quality checklist	Log-based; period stated; index status checked before blocking.
Failure conditions	Declines the analysis entirely without log data.
Next prompt	PR-65

PR-65 Rendering audit

TECHNICAL SEO · ADVANCED

When to use	When content may depend on JavaScript execution to appear.
When NOT to use	Do not use it on a summary of the page — it requires the raw and rendered HTML.
Required inputs	Raw HTML, rendered HTML, and the page as a user sees it.
Databases used	Technical Issues, Content

COPY-PASTE PROMPT

ROLE: You are auditing whether key content renders without JavaScript.
OBJECTIVE: Identify content that exists only after execution.
CONTEXT: Raw HTML: [PASTE]. Rendered HTML: [PASTE]. What the page should contain: [DESCRIBE]. Template: [NAME].
CONSTRAINTS: Compare raw against rendered directly. Do not assume behaviour from framework names. Identify which specific elements differ, not just that they do.
PROCESS: 1 Compare the two versions element by element. 2 List content present only in the rendered version. 3 Assess the SEO consequence of each. 4 Identify whether links are affected. 5 Assign severity.
OUTPUT: A difference table with element, consequence, and severity, plus whether navigation links are affected.
QUALITY CHECKS: Specific elements named; link impact assessed; severity assigned.
MISSING INFORMATION: Without both HTML versions, state the audit cannot proceed.
HUMAN REVIEW: Engineering confirms the diagnosis before ticketing.

Expected output	A raw-versus-rendered difference table with severity and link impact.
Quality checklist	Element-level specificity; link impact assessed.
Failure conditions	Cannot proceed without both HTML versions.
Next prompt	PR-66

PR-66 Migration redirect mapping

TECHNICAL SEO · ADVANCED

When to use	Before any URL structure change, however small.
When NOT to use	Do not use it after launch — a redirect map built retrospectively always loses the tail.
Required inputs	Full current URL inventory with traffic, rankings, and backlinks; the new structure.
Databases used	Technical Issues, Content, Backlinks

COPY-PASTE PROMPT

ROLE: You are building a migration redirect map.

OBJECTIVE: Ensure every current URL has a destination, one hop, before launch.

CONTEXT: Current URL inventory: [PASTE with traffic, rankings, backlinks]. New structure: [DESCRIBE]. URLs being removed entirely: [LIST].

CONSTRAINTS: Every URL gets a destination — no exceptions, including the long tail. One hop only. Redirect to the closest equivalent, never blanket to the homepage. Flag any URL with backlinks being removed.

PROCESS: 1 Map every URL to its new destination. 2 Identify URLs with no equivalent and choose the closest relevant page. 3 Flag chains. 4 Flag backlinked URLs being removed. 5 Produce the verification list for post-launch.

OUTPUT: The complete redirect map, the no-equivalent decisions with reasoning, the chain flags, and the post-launch verification list.

QUALITY CHECKS: Every URL mapped; one hop; no blanket homepage redirects; backlinked removals flagged.

MISSING INFORMATION: If the inventory is partial, state the map is incomplete and name what is missing.

HUMAN REVIEW: Engineering and SEO both sign the map before launch.

Expected output	A complete one-hop redirect map with verification list.
Quality checklist	Every URL mapped; no blanket redirects; backlinked removals flagged.
Failure conditions	Declares the map incomplete on a partial inventory.
Next prompt	PR-67

PR-67 Incident response triage

TECHNICAL SEO · ADVANCED

When to use	The moment a sitewide traffic or indexing anomaly appears.
When NOT to use	Do not use it for gradual declines — those are decay, not incidents.
Required inputs	What changed, when it started, recent releases, and current status codes.
Databases used	Technical Issues, Technical Releases & Verification

COPY-PASTE PROMPT

ROLE: You are triaging a suspected SEO incident.

OBJECTIVE: Determine what changed, when, and what to check first.

CONTEXT: What is observed: [DESCRIBE with dates]. Recent releases: [PASTE the release log]. Current robots.txt: [PASTE]. Current index coverage: [PASTE]. Status codes on key pages: [PASTE].

CONSTRAINTS: Check the release log before anything else — most incidents are a deploy. Do not speculate about algorithm updates until site-side causes are eliminated. Order checks by how quickly they can be ruled out.

PROCESS: 1 Correlate the start date against the release log. 2 Check robots and meta directives. 3 Check status codes and canonicals on affected templates. 4 Check index coverage change. 5 Only then consider external causes.

OUTPUT: A ranked check list with what each rules in or out, the most likely cause, and the immediate action.

QUALITY CHECKS: Release log checked first; site-side causes eliminated before external ones; checks ordered by speed.

MISSING INFORMATION: Without a release log, say that diagnosis will be substantially slower and start one.

HUMAN REVIEW: Engineering is engaged immediately on any critical finding.

Expected output	A ranked diagnostic checklist with the most likely cause and an action.
Quality checklist	Release log first; site-side before external; speed-ordered checks.
Failure conditions	Flags the absence of a release log as a diagnostic handicap.
Next prompt	PR-68

PR-68 Internal link opportunity generation

INTERNAL LINKING · FOUNDATIONAL

When to use	Monthly, and after any batch of publishing.
When NOT to use	Do not use it to link every page to every other — relevance below 3 is discarded.
Required inputs	The source page, candidate destinations with commercial importance, existing links.
Databases used	Internal Link Opportunities, Content

COPY-PASTE PROMPT

ROLE: You are generating internal link opportunities from one page.

OBJECTIVE: Suggest links a reader would genuinely want, prioritizing commercial destinations.

CONTEXT: Source page: [PASTE full text]. Candidate destinations: [LIST with topics and commercial importance]. Existing links on this page: [LIST].

CONSTRAINTS: Maximum six suggestions. Each needs a reader-benefit rationale. Relevance below 3 is discarded rather than included. Anchors must vary. Do not suggest a link purely because the destination is commercially important.

PROCESS: 1 Identify natural link moments in the text. 2 Match each to the most relevant destination. 3 Write a descriptive anchor. 4 Score target value, relevance, source strength. 5 Discard anything below 3 on relevance.

OUTPUT: A table with the sentence to link from, the anchor, the destination, the three scores, the priority product, and the reader rationale.

QUALITY CHECKS: Every link has a rationale; anchors vary; low relevance discarded.

MISSING INFORMATION: Without commercial importance, rank by relevance and say so.

HUMAN REVIEW: The editor approves before implementation.

Expected output	A scored link table with anchors, destinations, and reader rationales.
Quality checklist	Rationale per link; varied anchors; low-relevance discarded.
Failure conditions	Ranks on relevance alone when importance is unknown.
Next prompt	PR-69

PR-69 Orphan page resolution

INTERNAL LINKING · INTERMEDIATE

When to use	When a crawl identifies pages reachable only from the sitemap.
When NOT to use	Do not use it to link orphans indiscriminately — some should be removed instead.
Required inputs	The orphan list with traffic, backlinks, and topic; candidate source pages.
Databases used	Internal Link Opportunities, Content, Site Architecture

COPY-PASTE PROMPT

ROLE: You are resolving orphan pages.

OBJECTIVE: Decide, per orphan, whether to link, merge, redirect, or remove.

CONTEXT: Orphans: [PASTE with traffic, backlinks, topic, age]. Candidate source pages: [LIST]. Cluster structure: [PASTE].

CONSTRAINTS: An orphan with no traffic, no links, and no cluster fit is a removal candidate, not a linking candidate. Check backlinks before recommending removal. Every link recommendation needs a relevant source.

PROCESS: 1 Classify each orphan: valuable, marginal, or obsolete. 2 For valuable ones, identify relevant source pages. 3 For marginal ones, assess merge candidates. 4 For obsolete ones, specify a redirect. 5 Flag any with backlinks.

OUTPUT: A per-orphan decision table with the action, the source pages or redirect destination, and the backlink flags.

QUALITY CHECKS: Backlinks checked before removal; link recommendations have relevant sources; obsolete pages are not linked for the sake of it.

MISSING INFORMATION: Without backlink data, defer removal decisions.

HUMAN REVIEW: The SEO lead approves removals and redirects.

Expected output	A per-orphan decision table with actions, sources, and backlink flags.
Quality checklist	Backlinks checked; relevance required; removal is a valid outcome.
Failure conditions	Defers removals without backlink data.
Next prompt	PR-70

PR-70 Money page link support audit

INTERNAL LINKING · FOUNDATIONAL

When to use	Monthly, and immediately when inheriting a site.
When NOT to use	Do not use it to add links from irrelevant pages to commercial pages.
Required inputs	URL inventory with internal inbound link counts and commercial importance.
Databases used	Content, Internal Link Opportunities

COPY-PASTE PROMPT

ROLE: You are auditing internal link support for commercially important pages.

OBJECTIVE: Find the highest-return internal linking work available.

CONTEXT: URL inventory with inbound internal link counts: [PASTE]. Commercial importance rating per page: [PASTE]. Cluster structure: [PASTE].

CONSTRAINTS: Sort by importance descending and link count ascending – the top of that list is the work. Only recommend links from topically relevant sources. Do not recommend footer or sitewide links as a solution.

PROCESS: 1 Rank pages by importance and inverse link count. 2 For the top twenty, identify relevant source pages within the cluster. 3 Estimate the effort per page. 4 Flag any commercial page with zero inbound links. 5 Sequence the work.

OUTPUT: A ranked support table, the zero-link commercial pages flagged, relevant sources per page, and a sequenced implementation plan.

QUALITY CHECKS: Sources are topically relevant; no sitewide link recommendations; zero-link pages flagged prominently.

MISSING INFORMATION: Without importance ratings, ask for them – the audit is meaningless otherwise.

HUMAN REVIEW: The editor approves the implementation sequence.

Expected output	A ranked support table with zero-link pages flagged and a sequenced plan.
Quality checklist	Relevant sources only; no sitewide links; zero-link pages surfaced.
Failure conditions	Stops without commercial importance ratings.
Next prompt	PR-71

PR-71 Authority gap analysis

AUTHORITY AND PR · ADVANCED

When to use	Quarterly, and whenever competitors appear in sources you do not.
When NOT to use	Do not use it to justify buying links — gaps close through earning or not at all.
Required inputs	Your link and mention profile, competitor profiles, the source landscape.
Databases used	Backlinks, Authority Targets, AI Citations & Sources

COPY-PASTE PROMPT

ROLE: You are analysing the authority gap against competitors.

OBJECTIVE: Identify sources covering competitors and not us, ranked by closability.

CONTEXT: Our profile: [PASTE]. Competitor profiles: [PASTE]. Sources cited in AI answers: [PASTE]. Our assets available to offer: [LIST].

CONSTRAINTS: Never recommend paid, exchanged, or concealed placements. Rank by relevance and closability rather than by domain metric. Only include sources where we have something genuinely worth offering.

PROCESS: 1 Identify sources covering competitors but not us. 2 Assess topical relevance to our category. 3 Assess whether we have an asset worth their attention. 4 Rate closability. 5 Rank and propose a route each.

OUTPUT: A gap table with relevance, our asset, closability, and the proposed route, plus the sources deliberately excluded.

QUALITY CHECKS: Nothing recommended that requires concealment; every entry has an asset to offer.

MISSING INFORMATION: Without competitor profiles, produce a source landscape only.

HUMAN REVIEW: The outreach owner approves the target list.

Expected output	A ranked authority gap table with routes and an exclusion list.
Quality checklist	No paid or concealed tactics; every target has an asset offered.
Failure conditions	Produces a landscape only without competitor data.
Next prompt	PR-72

PR-72 Link reclamation

AUTHORITY AND PR · FOUNDATIONAL

When to use	Monthly, on unlinked mentions and lost links.
When NOT to use	Do not use it to chase every mention — prioritize by relevance and effort.
Required inputs	Unlinked mentions and lost links with sources, dates, and context.
Databases used	Backlinks, Authority Targets

COPY-PASTE PROMPT

ROLE: You are planning link reclamation outreach.
OBJECTIVE: Convert unlinked mentions and recover lost links, politely.
CONTEXT: Unlinked mentions: [PASTE with source, date, context]. Lost links: [PASTE with last seen date]. Our relevant pages: [LIST].
CONSTRAINTS: One request per contact, one follow-up maximum. Never imply an obligation. For lost links, establish whether removal was deliberate before asking. Do not chase mentions on irrelevant sites.
PROCESS: 1 Prioritize by relevance and likely ease. 2 For mentions, identify the correct destination URL. 3 For lost links, check whether the page still exists. 4 Draft one message per situation type. 5 Set the follow-up rule.
OUTPUT: A prioritized reclamation list with the destination URL, a drafted message per type, and the follow-up rule.
QUALITY CHECKS: One follow-up maximum; no implied obligation; deliberate removals respected.
MISSING INFORMATION: Where context is missing, exclude the mention rather than guessing.
HUMAN REVIEW: The outreach owner sends; nothing is automated.

Expected output	A prioritized reclamation list with drafted messages and follow-up rules.
Quality checklist	One follow-up; no obligation implied; deliberate removals respected.
Failure conditions	Excludes mentions where context is missing.
Next prompt	PR-73

PR-73 Journalist research brief

AUTHORITY AND PR · INTERMEDIATE

When to use	Before any editorial pitch, to know whose beat it actually is.
When NOT to use	Do not use it to build a list from a database without reading recent work.
Required inputs	The story, target publication types, and recent coverage you have read.
Databases used	Authority Targets, Digital PR Campaigns

COPY-PASTE PROMPT

ROLE: You are researching journalists for an editorial pitch.
OBJECTIVE: Match the story to people whose beat it genuinely is.
CONTEXT: Our story: [DESCRIBE]. Our data and its limits: [PASTE]. Publication types: [LIST]. Recent coverage read: [PASTE what and by whom].
CONSTRAINTS: Match on beat evidenced by recent work, not on publication name. Use public professional contact routes only. Do not record personal information. One pitch per person.
PROCESS: 1 Identify beats the story fits. 2 For each journalist, note the specific recent piece that shows the fit. 3 Draft the angle for their audience, not ours. 4 Identify what they would need from us. 5 Rank by fit.
OUTPUT: A journalist table with the evidencing piece, the audience-specific angle, what they would need, and the fit ranking.
QUALITY CHECKS: Beat fit evidenced by a named piece; angle framed for their audience; public routes only.
MISSING INFORMATION: Without recent coverage read, the fit cannot be assessed — say so.
HUMAN REVIEW: The PR owner approves the list and sends personally.

Expected output	A journalist table with evidenced fit and audience-specific angles.
Quality checklist	Fit evidenced by a named piece; public routes; one pitch each.
Failure conditions	Refuses to assess fit without recent coverage.
Next prompt	PR-74

PR-74 Outreach personalization

AUTHORITY AND PR · FOUNDATIONAL

When to use	On every outreach message, before sending.
When NOT to use	Do not use it to generate a personalization line that would fit anyone.
Required inputs	The specific target, what they published, and our asset.
Databases used	Authority Targets, Content

COPY-PASTE PROMPT

ROLE: You are personalizing an outreach message.

OBJECTIVE: Produce an opening that could only have been written for this person.

CONTEXT: Target: [NAME AND ROLE]. What they published: [PASTE the specific piece]. Our asset: [DESCRIBE]. Why it improves their work: [DESCRIBE].

CONSTRAINTS: Under 120 words total. Reference the specific piece, not the publication. No flattery openers. Do not claim to have read something unless it was supplied. Make declining easy.

PROCESS: 1 Identify the specific point in their piece our asset relates to. 2 Open on that. 3 State what we have in one sentence. 4 Make one clear, small ask. 5 Make declining costless.

OUTPUT: The message, the specific reference used, and the test result for whether it would work for another target.

QUALITY CHECKS: Reference is specific; under 120 words; one ask; declining is easy.

MISSING INFORMATION: Without their actual published work, produce nothing.

HUMAN REVIEW: The outreach owner sends; nothing is automated.

Expected output	A sub-120-word message with a specific reference and a portability test.
Quality checklist	Specific reference; one ask; declining made easy.
Failure conditions	Produces nothing without the target's actual work.
Next prompt	PR-75

PR-75 Toxic link review

AUTHORITY AND PR · ADVANCED

When to use	Only when a manual action or a clear manipulative pattern exists.
When NOT to use	Do not use it routinely — most sites need this far less often than tools suggest.
Required inputs	The link profile, the suspected pattern, and any manual action notice.
Databases used	Backlinks

COPY-PASTE PROMPT

ROLE: You are reviewing whether links warrant disavowal.

OBJECTIVE: Distinguish genuinely manipulative links from merely low-metric ones.

CONTEXT: Link profile: [PASTE]. Suspected pattern: [DESCRIBE]. Manual action notice if any: [PASTE]. Links we built ourselves or inherited: [DESCRIBE].

CONSTRAINTS: A low metric is not toxicity. Require a pattern, not isolated instances. Attempt removal before disavowal and document it. Do not disavow without a plausible connection to an actual problem.

PROCESS: 1 Identify the pattern and its scale. 2 Determine origin — built, inherited, or organic. 3 Check for a manual action or a clear ranking event. 4 Recommend removal attempts first. 5 Only then consider disavowal, with the reasoning recorded.

OUTPUT: A pattern assessment, an origin determination, a removal plan, and a disavowal recommendation with its reasoning — or an explicit recommendation to take no action.

QUALITY CHECKS: Pattern evidenced; origin determined; removal attempted first; no action is a valid outcome.

MISSING INFORMATION: Without evidence of a problem, recommend no action.

HUMAN REVIEW: The SEO lead approves any disavowal, and the reasoning is logged.

Expected output	A pattern assessment with a removal plan and a documented recommendation.
Quality checklist	Pattern required; removal first; no action is valid.
Failure conditions	Recommends no action absent evidence of an actual problem.
Next prompt	PR-76

PR-76 Google Business Profile audit

LOCAL SEO · FOUNDATIONAL

When to use	Monthly, per location.
When NOT to use	Do not use it on a profile you have not claimed — claim it first.
Required inputs	The canonical NAP record and the current profile contents.
Databases used	Local Locations, Local Citations

COPY-PASTE PROMPT

ROLE: You are auditing a business profile against the canonical record.
OBJECTIVE: Find every discrepancy and completeness gap.
CONTEXT: Canonical NAP: [PASTE exactly]. Current profile: [PASTE every field]. Primary and secondary categories: [PASTE]. Services and products listed: [PASTE].
CONSTRAINTS: Compare character by character. Do not estimate map positions or local traffic. Assess category choice against what customers actually search, not what sounds accurate.
PROCESS: 1 Compare NAP field by field. 2 Assess category selection. 3 Check completeness: hours, services, products, attributes, photos. 4 Check description against the brand fact sheet. 5 Rank fixes by likely effect.
OUTPUT: A discrepancy table with exact differences, a completeness result, and a ranked fix list.
QUALITY CHECKS: Differences quoted exactly; no estimated local metrics; category assessed against search behaviour.
MISSING INFORMATION: Fields not supplied are marked unchecked, not assumed correct.
HUMAN REVIEW: The location owner confirms the canonical NAP before corrections.

Expected output	A discrepancy table with exact differences and a ranked fix list.
Quality checklist	Character-level comparison; no estimated metrics; unchecked fields marked.
Failure conditions	Marks unsupplied fields unchecked rather than assuming.
Next prompt	PR-77

PR-77 Review response drafting

LOCAL SEO · FOUNDATIONAL

When to use	On every review, within the agreed response window.
When NOT to use	Do not use it to dispute facts publicly or to respond to every review identically.
Required inputs	The review text, the customer context if known, and the response policy.
Databases used	Local Locations

COPY-PASTE PROMPT

ROLE: You are drafting a response to a customer review.
OBJECTIVE: Respond in a way that is human, specific, and safe.
CONTEXT: Review: [PASTE]. Rating: [N]. What we know about the situation: [PASTE or state unknown]. Our response policy: [PASTE].
CONSTRAINTS: Reference what they specifically said. Never dispute facts publicly. Never disclose customer details. For negative reviews, acknowledge and move offline. No keywords, no sales language. Vary phrasing across responses.
PROCESS: 1 Identify what they actually said. 2 Choose the response type. 3 Acknowledge specifically. 4 State what happens next, if anything. 5 Offer a direct route where appropriate.
OUTPUT: The response under 80 words, the response type, and a note on anything that should be escalated internally.
QUALITY CHECKS: Specific to the review; no public dispute; no customer details; under 80 words.
MISSING INFORMATION: Where the situation is unknown, acknowledge without conceding or disputing facts.
HUMAN REVIEW: The location owner sends; nothing is automated.

Expected output	A sub-80-word response with a type classification and escalation note.
Quality checklist	Specific; no public dispute; no customer details disclosed.
Failure conditions	Acknowledges without conceding when the situation is unknown.
Next prompt	PR-78

PR-78 Multi-location content system

LOCAL SEO · ADVANCED

When to use	When more than three locations need pages that are genuinely different.
When NOT to use	Do not use it to spin near-identical pages with the place name swapped.
Required inputs	The locations, what genuinely differs between them, and local evidence.
Databases used	Local Locations, Content

COPY-PASTE PROMPT

ROLE: You are designing a multi-location content system.

OBJECTIVE: Produce location pages that are genuinely distinct.

CONTEXT: Locations: [LIST]. What actually differs: [PASTE – services, team, conditions, regulations, customers]. Local evidence per location: [PASTE].

CONSTRAINTS: A page with only the place name changed is duplication. Each page needs at least three genuinely local elements. Do not fabricate local detail. If a location has no distinct content, say the page should not exist yet.

PROCESS: 1 Identify what genuinely differs per location. 2 Define the shared structure. 3 Define the three or more required local elements per page. 4 Flag locations with insufficient distinct content. 5 Specify local schema per page.

OUTPUT: A page structure, the required local elements per location, the flagged locations, and the schema specification.

QUALITY CHECKS: Three or more genuine local elements per page; no fabricated detail; insufficient locations flagged.

MISSING INFORMATION: Without local evidence, recommend gathering it before building pages.

HUMAN REVIEW: The location owner verifies every local claim.

Expected output	A page structure with required local elements and flagged locations.
Quality checklist	Three genuine local elements minimum; nothing fabricated.
Failure conditions	Recommends gathering evidence before building thin location pages.
Next prompt	PR-79

PR-79 AI answer capture and classification

LLM AND AI SEO · FOUNDATIONAL

When to use	Every month, on every prompt in the tracked universe.
When NOT to use	Do not use it with edited prompt wording — comparability is destroyed.
Required inputs	The exact prompt, the assistant, the date, and the verbatim answer.
Databases used	AI Prompt Runs, Brand Facts, AI Competitors

COPY-PASTE PROMPT

ROLE: You are classifying a captured AI assistant answer.

OBJECTIVE: Record mention, citation, sentiment, errors, and sources consistently.

CONTEXT: Prompt exactly as run: [PASTE]. Assistant: [NAME]. Date: [DATE]. Answer verbatim: [PASTE]. Our brand fact sheet: [PASTE]. Competitor list: [LIST].

CONSTRAINTS: Mention and citation are separate – record both. Compare every statement about us against the fact sheet. Do not speculate about why the system answered this way. Quote errors exactly.

PROCESS: 1 Record whether we were mentioned. 2 Record whether we were cited with attribution. 3 Note position in the answer if named. 4 Classify sentiment. 5 Check every claim about us against the fact sheet. 6 List competitors named and sources cited.

OUTPUT: A classification record with all fields, the factual errors quoted, and the source list.

QUALITY CHECKS: Mention and citation separate; errors quoted exactly; no speculation about system behaviour.

MISSING INFORMATION: If the prompt wording differs from the universe record, void the run.

HUMAN REVIEW: Factual errors escalate to the brand facts owner immediately.

Expected output	A full classification record with quoted errors and a source list.
Quality checklist	Mention and citation separated; errors quoted; no speculation.
Failure conditions	voids the run if prompt wording drifted.
Next prompt	PR-80

PR-80 Share of AI voice calculation

LLM AND AI SEO · INTERMEDIATE

When to use	Monthly, once every core prompt has been run.
When NOT to use	Do not use it to compare against another company or a published figure.
Required inputs	All runs for the month with mentions and competitors named.
Databases used	AI Prompt Runs, Historical Scorecards, AI Competitors

COPY-PASTE PROMPT

ROLE: You are calculating the monthly AI visibility measures.

OBJECTIVE: Produce five comparable measures with their limitations stated.

CONTEXT: All runs this month: [PASTE]. Prompt set size: [N]. Assistants: [LIST]. Previous month's scorecard: [PASTE].

CONSTRAINTS: State the prompt set size and assistant list with every figure. If the set changed, flag the month as not comparable. Never present share of voice as market share. Report method alongside every number.

PROCESS: 1 Calculate mention rate, citation rate, answer coverage, share of AI voice, accuracy rate. 2 Compare with the previous month. 3 Check whether the set changed. 4 State the limitation of each measure. 5 Note hypotheses for movement, labelled as hypotheses.

OUTPUT: The five measures with method, the month-on-month comparison, the comparability flag, and hypotheses clearly labelled.

QUALITY CHECKS: Set size stated; comparability flagged; hypotheses labelled as such; no market-share framing.

MISSING INFORMATION: If any core prompt was not run, the scorecard is provisional.

HUMAN REVIEW: The SEO lead approves before it enters leadership reporting.

Expected output	Five measures with method, comparison, and labelled hypotheses.
Quality checklist	Set size stated; comparability flagged; no market-share claims.
Failure conditions	Marks the scorecard provisional if a core prompt was skipped.
Next prompt	PR-81

PR-81 Citation gap prioritization

LLM AND AI SEO · ADVANCED

When to use	Monthly, after the tracking run, to convert gaps into work.
When NOT to use	Do not use it to target sources with no realistic route to inclusion.
Required inputs	Sources cited across runs, whether they mention us, and our available assets.
Databases used	AI Citations & Sources, Authority Targets, Content

COPY-PASTE PROMPT

ROLE: You are prioritizing citation gaps into actionable work.

OBJECTIVE: Rank the sources worth pursuing and specify a route for each.

CONTEXT: Sources cited this month: [PASTE with frequency]. Which mention us: [PASTE]. Which mention competitors: [PASTE]. Our assets: [LIST]. Our capacity: [DESCRIBE].

CONSTRAINTS: Rank by citation frequency and closability, not by domain metric. Every entry needs a specific route. Exclude sources with no legitimate route. Never propose payment or concealment.

PROCESS: 1 Identify sources covering competitors and not us. 2 Rank by how often they appear across the prompt set. 3 Assess closability. 4 Specify a route per closable gap. 5 Estimate effort and sequence.

OUTPUT: A ranked gap table with frequency, closability, route, effort, and the excluded sources with reasons.

QUALITY CHECKS: Ranked on frequency and closability; every entry has a route; no payment or concealment.

MISSING INFORMATION: Without frequency data, rank on relevance and say so.

HUMAN REVIEW: The outreach owner takes the closable gaps into the target list.

Expected output	A ranked citation gap table with routes, effort, and exclusions.
Quality checklist	Frequency and closability ranking; routes specified; nothing paid or hidden.
Failure conditions	Ranks on relevance alone without frequency data.
Next prompt	PR-82

PR-82 Factual error correction plan

LLM AND AI SEO · ADVANCED

When to use	Immediately, whenever a tracked answer states something incorrect about you.
When NOT to use	Do not use it to attempt to influence a model directly — corrections happen at source.
Required inputs	The error, the correct fact, and the sources that may be causing it.
Databases used	Brand Facts, AI Prompt Runs, Local Citations, Content

COPY-PASTE PROMPT

ROLE: You are planning a correction for a factual error in an AI answer.
OBJECTIVE: Find the likely source of the error and correct it there.
CONTEXT: The error, quoted: [PASTE]. The correct fact from our sheet: [PASTE]. Where this fact appears publicly: [PASTE all sources]. Assistant and date: [PASTE].
CONSTRAINTS: Corrections happen at source. Do not propose prompt-based attempts to influence a model. Identify every public source stating the wrong version. Correct owned properties first.
PROCESS: 1 Search our own properties for the incorrect version. 2 Search third-party sources for it. 3 Rank by likely influence. 4 Correct owned sources immediately. 5 Draft correction requests for third parties. 6 Set a recheck date.
OUTPUT: A source list ranked by likely influence, the owned-source corrections, drafted third-party requests, and the recheck date.
QUALITY CHECKS: Owned sources corrected first; no attempt to influence the model directly; recheck date set.
MISSING INFORMATION: If the source cannot be identified, correct owned properties and monitor.
HUMAN REVIEW: The brand facts owner approves and sends correction requests.

Expected output	A ranked source list with corrections, requests, and a recheck date.
Quality checklist	Source-level correction only; owned properties first; recheck scheduled.
Failure conditions	Corrects owned sources and monitors when the origin is unclear.
Next prompt	PR-83

PR-83 AI-ready content brief

LLM AND AI SEO · INTERMEDIATE

When to use	When a citation gap indicates content is the missing piece.
When NOT to use	Do not use it to produce thin pages optimized only for extraction.
Required inputs	The gap, the question being answered, and our evidence.
Databases used	AI Visibility Opportunities, Content Briefs, Question Bank

COPY-PASTE PROMPT

ROLE: You are briefing content designed to be citable.
OBJECTIVE: Produce a brief where extractability is designed in, not retrofitted.
CONTEXT: The gap: [DESCRIBE with the prompts that revealed it]. The question: [PASTE]. Our evidence: [PASTE]. Brand facts relevant: [PASTE].
CONSTRAINTS: The piece must be genuinely useful to a reader first. Specify which passages must stand alone. Every fact must match the brand fact sheet. Do not brief a page whose only purpose is extraction.
PROCESS: 1 State the question and the correct answer. 2 Specify the direct-answer block. 3 Specify the supporting depth a reader needs. 4 Identify facts requiring sheet consistency. 5 Specify sources for every figure.
OUTPUT: The brief with the direct-answer specification, the extractable passage list, the fact consistency requirements, and the source requirements.
QUALITY CHECKS: Reader value first; extractable passages specified; fact consistency required.
MISSING INFORMATION: Without our own evidence, note the piece would be a summary of what already ranks.
HUMAN REVIEW: The editor approves before commissioning.

Expected output	A brief with direct-answer specification and fact consistency requirements.
Quality checklist	Reader value first; passages specified; facts checked against the sheet.
Failure conditions	Flags briefs that would only summarize existing content.
Next prompt	PR-84

PR-84 Competitor recommendation analysis

LLM AND AI SEO · ADVANCED

When to use	Quarterly, on recommendation prompts where competitors appear and you do not.
When NOT to use	Do not use it to speculate about ranking factors inside these systems.
Required inputs	Captured recommendation answers, competitor public presence, our own presence.
Databases used	AI Prompt Runs, AI Competitors, AI Visibility Opportunities

COPY-PASTE PROMPT

ROLE: You are analysing why competitors are recommended and we are not.
OBJECTIVE: Identify observable differences we could legitimately close.
CONTEXT: Recommendation answers: [PASTE verbatim with assistant and date].
 Competitors named: [LIST]. Their public presence: [PASTE what you observed]. Our equivalent: [PASTE].
CONSTRAINTS: Describe observable differences only. Do not speculate about system internals or ranking factors. Do not conclude a single factor caused a recommendation. Report reasons only where the answer states them.
PROCESS: 1 Count which brands are named across the set. 2 Record any stated reasons. 3 Compare public presence – coverage, listings, reviews, content depth, pricing transparency. 4 Identify closable differences. 5 Rank by effort and frequency.
OUTPUT: A presence comparison table, the stated reasons, and a ranked list of closable differences with effort.
QUALITY CHECKS: Observable differences only; no causal claims; gaps are specific.
MISSING INFORMATION: Exclude competitors whose presence data was not supplied.
HUMAN REVIEW: Strategy decides which differences to close.

Expected output	A presence comparison with stated reasons and ranked closable differences.
Quality checklist	Observable only; no causal or ranking-factor speculation.
Failure conditions	Excludes competitors without supplied presence data.
Next prompt	PR-85

PR-85 Monthly performance analysis

MEASUREMENT · INTERMEDIATE

When to use	Monthly, before the report is written.
When NOT to use	Do not use it to explain movement the data cannot account for.
Required inputs	All metrics with sources and as-of dates, plus the previous month.
Databases used	Reports, Business Goals, Content, Keyword Master

COPY-PASTE PROMPT

ROLE: You are analysing a month of SEO performance.
OBJECTIVE: Separate what the data shows from what is being inferred.
CONTEXT: This month's metrics: [PASTE with sources and as-of dates]. Previous month: [PASTE]. What we changed and when: [PASTE]. Quarterly goals: [PASTE].
CONSTRAINTS: Every figure needs a source and a date. Do not attribute a change to a cause the data cannot support – list hypotheses instead. Report sample sizes with every rate. Do not compare to industry benchmarks.
PROCESS: 1 Report progress against each goal. 2 Report month-on-month change per metric. 3 Identify what changed in the same period. 4 Separate observation from inference. 5 List hypotheses for unexplained movement.
OUTPUT: A goal progress table, a metric change table with sample sizes, the observations, and the hypotheses labelled as hypotheses.
QUALITY CHECKS: Sources and dates on every figure; hypotheses labelled; no benchmark comparison.
MISSING INFORMATION: Mark unavailable figures explicitly rather than dropping the row.
HUMAN REVIEW: The SEO lead interprets before it becomes a report.

Expected output	Goal and metric tables with observations and labelled hypotheses.
Quality checklist	Sourced and dated; hypotheses labelled; no benchmarks.
Failure conditions	Marks unavailable figures rather than dropping rows.
Next prompt	PR-86

PR-86 Executive SEO summary

REPORTING · INTERMEDIATE

When to use	Monthly or quarterly, for an audience with no SEO context.
When NOT to use	Do not use it to present activity volume as achievement.
Required inputs	The performance analysis, the goals, and the decisions required.
Databases used	Reports, Business Goals, Decisions & Change Log

COPY-PASTE PROMPT

ROLE: You are writing an executive summary of SEO performance.

OBJECTIVE: Tell a business audience what changed, what it means, and what you need.

CONTEXT: Performance analysis: [PASTE]. Business goals: [PASTE]. Decisions required: [LIST]. Audience: [ROLES and what they care about].

CONSTRAINTS: Under 300 words. No SEO jargon without a plain-language equivalent. Do not report activity volume. Do not forecast. Every figure carries its as-of date. End with what you need from them.

PROCESS: 1 State progress against the business goal in their terms. 2 State what changed and the leading hypothesis. 3 State the risk that matters most. 4 State the decisions required and their consequence. 5 State what you need.

OUTPUT: The summary under 300 words, and a one-slide version of the same content.

QUALITY CHECKS: No jargon; no activity metrics; no forecasts; ends in a specific ask.

MISSING INFORMATION: Where a business metric is unavailable, say so plainly rather than substituting an SEO proxy.

HUMAN REVIEW: The SEO lead owns every number presented.

Expected output	A sub-300-word summary and a one-slide version.
Quality checklist	Plain language; no activity metrics or forecasts; ends in an ask.
Failure conditions	States unavailable business metrics rather than substituting proxies.
Next prompt	PR-87

PR-87 System governance audit

QA AND GOVERNANCE · ADVANCED

When to use	Quarterly, on the workspace itself rather than on the site.
When NOT to use	Do not use it as a performance review of people — it audits the system.
Required inputs	Database usage data, SOP run history, prompt usage, and ownership records.
Databases used	SOP Library, Prompt Library, all databases

COPY-PASTE PROMPT

ROLE: You are auditing whether the SEO operating system is actually being used.

OBJECTIVE: Find the parts that are decorative and the parts that are load-bearing.

CONTEXT: Database record counts and last-modified dates: [PASTE]. SOP last-run dates: [PASTE]. Prompt usage counts: [PASTE]. Ownership assignments: [PASTE].

CONSTRAINTS: Judge the system, not the people. A database nobody updates is a design problem before it is a discipline problem. Recommend removal as readily as enforcement.

PROCESS: 1 Identify databases with no activity in 60 days. 2 Identify SOPs never run. 3 Identify prompts never used. 4 Identify records with no owner. 5 For each, decide: enforce, simplify, or remove.

OUTPUT: An activity audit, the unowned records, and a decision per unused asset – enforce, simplify, or remove.

QUALITY CHECKS: Removal is offered as readily as enforcement; the audit targets the system rather than individuals.

MISSING INFORMATION: Where usage data is unavailable, note which assets cannot be assessed.

HUMAN REVIEW: The SEO lead decides what to remove; removals are logged.

Expected output	An activity audit with an enforce, simplify, or remove decision per asset.
Quality checklist	System-focused; removal offered readily; unowned records surfaced.
Failure conditions	Names assets that cannot be assessed without usage data.
Next prompt	—

WHICH TWELVE TO SET UP FIRST

PR-01 strategy, PR-27 goal mapping, PR-31 call mining, PR-39 seed expansion, PR-43 keyword mapping, PR-46 cluster design, PR-06 content brief, PR-59 direct-answer blocks, PR-68 internal links, PR-79 AI answer capture, PR-80 share of AI voice, and PR-85 monthly analysis. Those twelve cover one complete cycle of the system. The other seventy-five exist for when a specific job comes up, and several will never apply to your business at all.

The SOP library, continued

SOP-01 to SOP-16 appear later in this document in their original form. SOP-17 to SOP-42 complete the library, covering keywords, clusters, authority, local, AI visibility, reporting, and governance.

Forty-two procedures in total. Each carries an ID, an owner, a trigger, inputs, tools, preconditions, steps, an output, a quality gate, failure conditions, an escalation route, a handoff, a cadence, a related database, a related prompt, and an example run. Build them as one SOP database with these fields as properties.

SOP-17

Keyword scoring pass

QUARTERLY

Purpose	Score every candidate against the ten criteria consistently.
Owner	SEO specialist
Trigger	Quarterly, or when 20+ new candidates accumulate
Inputs	Candidate keywords, demand evidence, current rankings, the scoring model
Tools	Keyword tool, Notion, SERP checks
Preconditions	The scoring weights are current and versioned
STEPS	
1	Confirm the weight version has not changed since the last pass.
2	Score S1 to S10 for each candidate, leaving nothing blank.
3	Let the formula calculate the total and the band.
4	Review the distribution — if everything scores 60–70, the model is not discriminating.
5	Spot-check ten scores against a colleague's independent scoring.
Output	A fully scored candidate set with bands assigned
Quality gate	No candidate has an unscored criterion; the distribution spreads across bands
Failure conditions	BLOCKED if weights changed mid-pass; INCOMPLETE if any criterion is blank
Escalation	If two people score the same keyword more than two points apart, revise the criterion definition
Handoff	To SOP-18 for mapping
Related database	Keyword Master
Related prompt	PR-15, PR-27
EXAMPLE RUN	A pass of 84 candidates produced 11 in Build now, 23 this quarter, and 14 declines; the distribution check flagged S6 as never varying, and its definition was rewritten.

SOP-18 Keyword-to-page mapping

QUARTERLY

Purpose	Assign one owning page per keyword, or record a gap.
Owner	SEO specialist
Trigger	After every scoring pass
Inputs	Scored keywords, page inventory with topics, cluster plans
Tools	Notion, crawl export
Preconditions	Scoring is complete and the page inventory is current

STEPS

- 1 Match each scored keyword to the page whose intent it fits.
- 2 Flag keywords with no suitable page as content gaps.
- 3 Flag any page holding more than one primary keyword.
- 4 Resolve conflicts on intent, never on preference.
- 5 Record the mapping and update the cluster relation.

Output	A mapped keyword list with a gap list and resolved conflicts
Quality gate	No page holds two primary keywords
Failure conditions	BLOCKED if the page inventory is more than 30 days old
Escalation	Unresolvable conflicts go to the SEO lead with both options stated
Handoff	To SOP-19 if conflicts exist, otherwise to SOP-10
Related database	Keyword Master, Content
Related prompt	PR-43

EXAMPLE RUN Mapping 84 keywords surfaced 6 pages holding two primaries each; four were resolved by differentiating intent and two by merging the pages.

SOP-19 Cannibalization review

MONTHLY

Purpose	Resolve queries where two of our URLs alternate.
Owner	SEO specialist
Trigger	Monthly, and after any mapping pass
Inputs	90 days of query-level ranking URL history, internal link counts
Tools	Search Console, crawl export, Notion
Preconditions	Ranking history covers at least 90 days

STEPS

- 1 Export query-level ranking URLs for the period.
- 2 Identify queries where the ranking URL alternates.
- 3 Determine the intended owner by intent and conversion path.
- 4 Check internal links — they usually point at the wrong page.
- 5 Choose one resolution: consolidate, differentiate, de-optimize, or redirect.
- 6 Set a 60-day recheck and record it.

Output	A conflict list with one resolution and a recheck date each
Quality gate	Every conflict has a single chosen resolution and a recheck date
Failure conditions	INSUFFICIENT EVIDENCE below 90 days of history — record a baseline instead
Escalation	Redirects and removals require SEO lead approval
Handoff	To SOP-22 for the internal link changes
Related database	Keyword Master, Content
Related prompt	PR-44

EXAMPLE RUN Nine alternating queries; six resolved by de-optimizing the weaker page, two by merging, one deferred pending backlink review.

SOP-20 Topic cluster creation

PER TOPIC

Purpose	Turn a prioritized topic into a build plan.
Owner	Content lead
Trigger	When a topic reaches This quarter priority
Inputs	Topic record, mapped keywords, existing coverage, commercial destination
Tools	Notion
Preconditions	Keyword mapping for the topic is complete
STEPS	
1	Define the pillar scope and what it must cover.
2	Assign every mapped keyword to the pillar or a spoke.
3	Mark each spoke as new content or an update.
4	Specify bidirectional interlinking and the commercial path.
5	Create the cluster record with planned spoke count.
Output	A cluster record with a spoke list and interlinking specification
Quality gate	Every spoke has demand evidence and a commercial path exists
Failure conditions	BLOCKED without a commercial destination — the cluster would lead nowhere
Escalation	A cluster with no commercial destination goes to the SEO lead for a strategy decision
Handoff	To SOP-10 for briefing
Related database	Topic Clusters, Keyword Master
Related prompt	PR-46
EXAMPLE RUN	A nine-spoke cluster planned; four spokes were updates to existing pages rather than new commissions, which halved the estimated build time.

SOP-21 Content consolidation

AS NEEDED

Purpose	Merge overlapping pages without losing equity.
Owner	Content lead
Trigger	When the cannibalization review or portfolio audit flags overlap
Inputs	Candidate pages, performance, backlinks, internal links
Tools	Crawl export, backlink tool, CMS
Preconditions	Backlink data is available for every candidate
STEPS	
1	Identify the overlap precisely, section by section.
2	Choose the surviving URL on intent and equity.
3	List everything that must be carried across.
4	Specify redirects for every retired URL, one hop.
5	Update internal links to the surviving URL.
6	Recheck rankings and traffic after 60 days.
Output	A consolidated page, redirects live, and a recheck scheduled
Quality gate	Nothing in the retired pages was lost; every URL redirects one hop
Failure conditions	BLOCKED without backlink data on every candidate
Escalation	Any page with significant backlinks being retired goes to the SEO lead
Handoff	To SOP-22 for internal link updates
Related database	Content, Backlinks
Related prompt	PR-52
EXAMPLE RUN	Three thin pages merged into one; the surviving page recovered its previous combined traffic within seven weeks, verified at the 60-day recheck.

THE SOP LIBRARY

SOP-22

Internal linking implementation

WEEKLY

Purpose	Turn scored opportunities into live, verified links.
Owner	Editor
Trigger	Weekly, from the approved queue
Inputs	Approved link opportunities with anchors and rationales
Tools	CMS, Notion
Preconditions	Opportunities have passed relevance scoring
STEPS	
1	Take the approved queue in priority order.
2	Implement each link with the specified anchor.
3	Confirm the link is live and points at the correct URL.
4	Update the opportunity record to Verified.
5	Recheck the target page inbound count.
Output	Live, verified internal links with updated records
Quality gate	Every implemented link is confirmed live before the record is closed
Failure conditions	INCOMPLETE if a link is marked implemented without a live check
Escalation	If a CMS limitation prevents implementation, log it as a technical issue
Handoff	To SOP-39 for reporting
Related database	Internal Link Opportunities, Content
Related prompt	PR-68
EXAMPLE RUN	Twenty-two links implemented across nine source pages; the pricing page moved from 3 to 14 inbound internal links.

SOP-23

Orphan page resolution

MONTHLY

Purpose	Decide and act on pages reachable only from the sitemap.
Owner	SEO specialist
Trigger	Monthly, from the crawl
Inputs	Orphan list with traffic, backlinks, topic, and age
Tools	Crawler, backlink tool, CMS
Preconditions	A full crawl has completed within 7 days
STEPS	
1	Classify each orphan as valuable, marginal, or obsolete.
2	For valuable pages, identify relevant source pages and queue the links.
3	For marginal pages, assess merge candidates.
4	For obsolete pages, specify a redirect destination.
5	Check backlinks before any removal.
Output	A decision and action per orphan page
Quality gate	No page is removed without a backlink check and a redirect
Failure conditions	BLOCKED on removal decisions without backlink data
Escalation	Removals of pages with backlinks require SEO lead approval
Handoff	To SOP-22 or SOP-21 depending on the decision
Related database	Content, Internal Link Opportunities
Related prompt	PR-69
EXAMPLE RUN	Thirty-one orphans found; 12 linked from cluster pages, 7 merged, 9 redirected, 3 held pending a backlink review.

SOP-24 Linkable asset selection

QUARTERLY

Purpose	Choose what is worth building a campaign around.
Owner	Marketing owner
Trigger	Quarterly, before campaign planning
Inputs	Available data, expertise, product knowledge, and permissions
Tools	Notion
Preconditions	Data permissions have been checked

STEPS

- 1 List what we genuinely hold that others do not.
- 2 Assess each for durability and effort.
- 3 Confirm what may be published, with the data owner.
- 4 Select one asset per quarter, not three.
- 5 Define the methodology and limitations before building.

Output	One selected asset with methodology and limitations defined
Quality gate	Permissions confirmed in writing before any build begins
Failure conditions	BLOCKED if permissions are unclear
Escalation	Unclear permissions go to whoever owns the data
Handoff	To SOP-25 for campaign design
Related database	Content, Digital PR Campaigns
Related prompt	PR-49

EXAMPLE RUN Of four candidates, the anonymized usage benchmark was selected; two were dropped for insufficient sample and one for permission constraints.

SOP-25 Digital PR campaign run

PER CAMPAIGN

Purpose	Take an asset to market without overclaiming.
Owner	Marketing or PR owner
Trigger	When an asset is published with its methodology
Inputs	The asset, methodology, target research, spokesperson
Tools	Email, Notion
Preconditions	Methodology and limitations are published alongside the asset

STEPS

- 1 Confirm the asset and methodology are live.
- 2 Build the target list by beat, from recent work read.
- 3 Personalize each message with a specific reference.
- 4 Send, follow up once, then stop.
- 5 Log every outcome, including refusals and silence.
- 6 Report outcomes against effort at close.

Output	A completed campaign with every outcome logged
Quality gate	No pitch claim exceeds what the published methodology supports
Failure conditions	BLOCKED if the methodology is not published
Escalation	Requests for paid placement are declined and logged
Handoff	To SOP-27 for link monitoring
Related database	Digital PR Campaigns, Authority Targets
Related prompt	PR-73, PR-74

EXAMPLE RUN Thirty-four targets contacted, six outcomes earned, twenty-two no responses logged; the refusal reasons showed the angle was too broad for trade press.

SOP-26 Link reclamation run

MONTHLY

Purpose	Convert unlinked mentions and recover lost links.
Owner	Outreach owner
Trigger	Monthly
Inputs	Unlinked mentions, lost links, and the correct destination URLs
Tools	Mention monitoring, email, Notion
Preconditions	Mention data covers the last 90 days

STEPS

- 1 Review new unlinked mentions and lost links.
- 2 Prioritize by relevance and likely ease.
- 3 Check whether lost links were deliberately removed.
- 4 Send one request per contact, one follow-up maximum.
- 5 Log every outcome.

Output	Reclamation outcomes logged against every attempt
Quality gate	One follow-up maximum; deliberate removals are respected
Failure conditions	INCOMPLETE if outcomes are not logged for unsuccessful attempts
Escalation	Repeated refusals from one publisher stop further approaches
Handoff	To SOP-27
Related database	Backlinks
Related prompt	PR-72

EXAMPLE RUN Fourteen unlinked mentions actioned; five converted, two declined, seven no response — all seven logged rather than left blank.

SOP-27 Backlink monitoring

WEEKLY

Purpose	Notice gained and lost links while they are still actionable.
Owner	Outreach owner
Trigger	Weekly
Inputs	Backlink export, previous week's snapshot
Tools	Backlink tool, Notion
Preconditions	A previous snapshot exists for comparison

STEPS

- 1 Export the current profile and compare against last week.
- 2 Record new links with source, target, and anchor.
- 3 Record lost links with a last-seen date.
- 4 Assess lost links for reclamation.
- 5 Flag any unusual acquisition pattern for review.

Output	An updated link profile with new and lost links recorded
Quality gate	Every change is dated; nothing is recorded without a source
Failure conditions	INCOMPLETE without a previous snapshot — record a baseline instead
Escalation	An unusual pattern of new links goes to the SEO lead before any action
Handoff	To SOP-26 for reclamation
Related database	Backlinks
Related prompt	PR-72

EXAMPLE RUN Nine new links, three lost; two of the lost were site redesigns and were reclaimed, one was a deliberate editorial removal and was left alone.

SOP-28 Toxic link review

AS NEEDED

Purpose	Assess whether disavowal is genuinely warranted.
Owner	SEO lead
Trigger	Only on a manual action or a clear manipulative pattern
Inputs	Link profile, the suspected pattern, any manual action notice
Tools	Backlink tool, Search Console
Preconditions	A specific problem has been identified

STEPS

- 1 Establish the pattern and its scale.
- 2 Determine origin — built, inherited, or organic.
- 3 Attempt removal first and document each attempt.
- 4 Only then consider disavowal.
- 5 Record the reasoning before filing anything.

Output	A documented decision, which may be to take no action
Quality gate	Removal attempted and documented before any disavowal
Failure conditions	BLOCKED without evidence of an actual ranking or manual action problem
Escalation	Disavowal requires SEO lead approval and a logged decision
Handoff	To SOP-42 for the decision log
Related database	Backlinks, Decisions & Change Log
Related prompt	PR-75

EXAMPLE RUN A pattern of 40 low-quality links was assessed; origin was an inherited agency campaign, removal succeeded on 11, and the decision was recorded not to disavow the remainder.

SOP-29 Google Business Profile maintenance

MONTHLY

Purpose	Keep every profile complete and consistent.
Owner	Local SEO owner
Trigger	Monthly, per location
Inputs	Canonical NAP, current profile contents, category research
Tools	Profile dashboard, Notion
Preconditions	The canonical NAP record exists and is confirmed

STEPS

- 1 Compare every profile field against the canonical record.
- 2 Assess category selection against how customers search.
- 3 Complete services, products, hours, and attributes.
- 4 Check the description against the brand fact sheet.
- 5 Record discrepancies and assign corrections.

Output	A corrected, complete profile and a discrepancy log
Quality gate	Every difference is quoted exactly with its field named
Failure conditions	BLOCKED if the canonical NAP has not been confirmed by the location owner
Escalation	A category change affecting visibility goes to the SEO lead first
Handoff	To SOP-30 for citations
Related database	Local Locations
Related prompt	PR-76

EXAMPLE RUN Two discrepancies found on one location — an abbreviated street type and a missing service category; both corrected within the week.

SOP-30 Local citation audit

MONTHLY

Purpose	Find and correct NAP drift across directories.
Owner	Local SEO owner
Trigger	Monthly, per location
Inputs	Canonical NAP, the citation list, directory access
Tools	Citation tool or manual checks, Notion
Preconditions	The canonical NAP is confirmed

STEPS

- 1 Check each listing character by character against the canonical record.
- 2 Quote every difference exactly, with its source.
- 3 Claim any unclaimed listing where possible.
- 4 Request corrections on listings we do not control.
- 5 Identify and merge duplicates.
- 6 Update the consistency rollup.

Output	A corrected citation set and a tracked correction queue
Quality gate	Consistency is measured character by character, not by impression
Failure conditions	INCOMPLETE if any listing is marked correct without an exact comparison
Escalation	Directories refusing correction are logged and noted in the AI visibility review
Handoff	To SOP-39
Related database	Local Citations, Local Locations
Related prompt	PR-76

EXAMPLE RUN Of 22 listings, 19 matched; two had an abbreviated street type and one carried a phone number from before a 2024 change.

SOP-31 Review response

CONTINUOUS

Purpose	Respond to every review within the agreed window.
Owner	Local SEO owner
Trigger	Within 48 hours of any new review
Inputs	The review, the response policy, and internal context where available
Tools	Profile dashboard
Preconditions	A response policy exists

STEPS

- 1 Read the review and identify what was specifically said.
- 2 Choose the response type from the policy.
- 3 Acknowledge specifically, never generically.
- 4 Move any dispute offline.
- 5 Escalate anything indicating a service failure.

Output	A sent response and an internal escalation where warranted
Quality gate	No two responses use identical wording; no facts disputed publicly
Failure conditions	NEEDS HUMAN REVIEW on any review alleging a serious service failure
Escalation	Serious allegations go to the location owner immediately, before responding
Handoff	To SOP-39 for reporting
Related database	Local Locations
Related prompt	PR-77

EXAMPLE RUN A negative review about response times was acknowledged publicly, moved to a direct conversation, and logged internally as a support theme.

SOP-32 Local landing page build

PER LOCATION

Purpose	Create location pages that are genuinely distinct.
Owner	Content lead
Trigger	When a location is added or a page audit finds duplication
Inputs	Canonical NAP, local evidence, services available at the location
Tools	CMS, Notion
Preconditions	At least three genuinely local elements exist

STEPS

- 1 Confirm three or more genuinely local elements exist.
- 2 Build the page from the location template.
- 3 Match the address and hours to the canonical record exactly.
- 4 Add local schema matching the profile.
- 5 Add internal links from the site to the page.
- 6 Verify the page renders and the schema validates.

Output	A live, distinct location page with valid schema
Quality gate	Three or more genuine local elements; NAP matches exactly
Failure conditions	BLOCKED if only the place name differs from another page
Escalation	A location with no distinct content is escalated rather than published thin
Handoff	To SOP-22 for internal links
Related database	Content, Local Locations
Related prompt	PR-78

EXAMPLE RUN One planned location page was held because the only difference from an existing page was the city name; it launched a month later with local team and service detail.

SOP-33 Prompt universe creation

ONCE, REVIEWED QUARTERLY

Purpose	Build the fixed tracking set, once.
Owner	SEO lead
Trigger	Once at setup, reviewed quarterly
Inputs	ICP decisions, question bank, competitor list
Tools	Notion, assistant access
Preconditions	The question bank contains real customer questions

STEPS

- 1 List the decisions where we should be a candidate.
- 2 Write conversational prompts for each, in speech form.
- 3 Add category, brand, competitor, comparison, and problem prompts.
- 4 Cap the set at 40 unless capacity justifies more.
- 5 Define what a good result looks like per prompt.
- 6 Lock the wording and record the version.

Output	A locked, typed prompt set with good-result definitions
Quality gate	All six prompt families represented; wording locked before the first run
Failure conditions	BLOCKED without a real question bank — problem prompts would be invented
Escalation	Any wording change creates a new record at version+1, never an edit
Handoff	To SOP-34 for the first run
Related database	AI Prompt Universe
Related prompt	PR-23, PR-42

EXAMPLE RUN A 38-prompt set was built: 12 category, 6 brand, 8 competitor, 6 comparison, 6 problem; wording was locked and versioned before the first run.

SOP-34 Monthly AI visibility run

MONTHLY

Purpose	Execute the tracking set identically every month.
Owner	AI visibility owner
Trigger	First week of each month
Inputs	The locked prompt universe, the brand fact sheet, assistant access
Tools	Assistant access, Notion
Preconditions	The prompt set has not changed since the last run

STEPS

- 1 Confirm no prompt wording has changed.
- 2 Run every core prompt on every tracked assistant.
- 3 Paste each answer verbatim into a run record.
- 4 Record mention and citation separately.
- 5 Compare every statement about us against the fact sheet.
- 6 Log factual errors as the highest-priority findings.

Output	A complete dated run set with classifications
Quality gate	Prompt wording matches the universe records exactly
Failure conditions	VOID any run where the wording differed; INCOMPLETE if a core prompt was skipped
Escalation	Factual errors escalate to the brand facts owner the same day
Handoff	To SOP-35 for scoring
Related database	AI Prompt Runs, Brand Facts
Related prompt	PR-79

EXAMPLE RUN Thirty-eight prompts across three assistants produced 114 run records; two factual errors were found and escalated within the hour.

SOP-35 AI scorecard calculation

MONTHLY

Purpose	Turn the run set into five comparable measures.
Owner	AI visibility owner
Trigger	Immediately after the monthly run
Inputs	All run records for the month, the previous scorecard
Tools	Notion
Preconditions	Every core prompt has been run

STEPS

- 1 Confirm every core prompt was run.
- 2 Calculate the five measures.
- 3 Compare against the previous month.
- 4 Check whether the prompt set or assistant list changed.
- 5 Flag the month as comparable or not.
- 6 Record hypotheses for movement, labelled as hypotheses.

Output	A dated scorecard with five measures and a comparability flag
Quality gate	Set size and assistant list stated alongside every figure
Failure conditions	PROVISIONAL if any core prompt was skipped
Escalation	A set change is flagged to leadership before the figures are presented
Handoff	To SOP-40 for reporting
Related database	Historical Scorecards
Related prompt	PR-80

EXAMPLE RUN Mention rate rose from 34% to 38%; the notes recorded the extractability pass as a hypothesis rather than a cause.

SOP-36 Brand fact correction

AS NEEDED

Purpose	Correct an incorrect fact at every source that carries it.
Owner	Marketing owner
Trigger	Immediately on any factual error in a tracked answer
Inputs	The error quoted, the canonical fact, and every public source
Tools	CMS, profile dashboards, email
Preconditions	The canonical fact sheet exists
STEPS	
1	Quote the error exactly and identify the correct fact.
2	Search our own properties for the incorrect version.
3	Search third-party sources for it.
4	Correct owned properties immediately.
5	Request corrections from third parties, in writing.
6	Set a recheck date on the next tracking run.
Output	Corrected owned sources and a tracked third-party request list
Quality gate	Owned properties corrected before any external request is sent
Failure conditions	BLOCKED from attempting to influence a model directly — corrections happen at source
Escalation	A third party refusing correction is logged and noted in reporting
Handoff	To SOP-34 for the recheck
Related database	Brand Facts, AI Prompt Runs
Related prompt	PR-82
EXAMPLE RUN	An outdated category description was found on our own about page and two directories; all three were corrected and the next run showed the correct description.

SOP-37 AI citation gap analysis

MONTHLY

Purpose	Convert observed source gaps into outreach work.
Owner	SEO lead
Trigger	Monthly, after the run
Inputs	Sources cited across runs, our presence in each, our assets
Tools	Notion
Preconditions	The monthly run is complete
STEPS	
1	List every source cited across the run set with its frequency.
2	Identify which mention us and which mention competitors only.
3	Assess closability for each gap.
4	Specify a route per closable gap.
5	Rank by frequency and effort.
6	Create authority targets for the top gaps.
Output	A ranked gap list with routes and created outreach targets
Quality gate	Every entry has a specific route; nothing requiring payment is included
Failure conditions	INSUFFICIENT EVIDENCE without frequency data — rank on relevance and say so
Escalation	A gap requiring a paid placement is declined and recorded
Handoff	To SOP-25 for outreach
Related database	AI Citations & Sources, Authority Targets
Related prompt	PR-81
EXAMPLE RUN	Eleven sources appeared more than twice; four covered competitors and not us, of which two had a realistic editorial route.

SOP-38 Competitor recommendation review

QUARTERLY

Purpose	Understand observable differences, quarterly.
Owner	SEO lead
Trigger	Quarterly, on recommendation prompts
Inputs	Captured recommendation answers, competitor presence data
Tools	Notion
Preconditions	At least three months of run data exists

STEPS

- 1 Count which brands are named across recommendation prompts.
- 2 Record any reasons the answers state.
- 3 Compare public presence across observable dimensions.
- 4 Identify differences we could legitimately close.
- 5 Rank by effort and frequency.

Output	A presence comparison and a ranked list of closable differences
Quality gate	Only observable differences recorded; no ranking-factor speculation
Failure conditions	INSUFFICIENT EVIDENCE below three months of runs
Escalation	Strategy decides which differences to close; nothing is actioned automatically
Handoff	To SOP-01 for the strategy review
Related database	AI Competitors, AI Visibility Opportunities
Related prompt	PR-84

EXAMPLE RUN Two competitors appeared in 60% of recommendation prompts; both had published pricing structures and we did not, which became a Q4 content decision.

SOP-39 Weekly reporting

WEEKLY

Purpose	Turn the week into three decisions.
Owner	SEO lead
Trigger	Every Friday
Inputs	All performance data with sources and dates, last week's report
Tools	Notion, analytics, Search Console
Preconditions	Last week's report exists

STEPS

- 1 Assemble every figure with its source and as-of date.
- 2 Compare against last week and the quarterly goal.
- 3 Separate observation from inference.
- 4 State the three decisions the week requires.
- 5 Record the system changes each decision implies.

Output	A report under 500 words ending in three owned decisions
Quality gate	Every figure sourced and dated; the report ends in decisions
Failure conditions	INCOMPLETE if any figure lacks a source or an as-of date
Escalation	A decision with no owner is reassigned before the report is filed
Handoff	To SOP-42 for the decision log
Related database	Reports, Decisions & Change Log
Related prompt	PR-85, PR-86

EXAMPLE RUN A report noted a 12% impression rise on comparison pages; the three decisions were to accelerate two briefs, defer a technical ticket, and re-run one experiment.

THE SOP LIBRARY

SOP-40

Monthly reporting

MONTHLY

Purpose	Report to a business audience in their terms.
Owner	SEO lead
Trigger	First week of each month
Inputs	The monthly analysis, goals, AI scorecard, and decisions
Tools	Notion
Preconditions	The monthly analysis and AI scorecard are complete
STEPS	
1	Assemble the analysis, goals, and scorecard.
2	Write the executive summary under 300 words.
3	Include the AI visibility trend with its method.
4	State the three decisions and their consequence.
5	End with what is needed from the audience.
Output	A business-language report with three decisions and a specific ask
Quality gate	No jargon, no activity metrics, no forecasts
Failure conditions	INCOMPLETE if a business metric is substituted with an SEO proxy
Escalation	Any figure the SEO lead cannot personally defend is removed
Handoff	To SOP-01 quarterly
Related database	Reports, Historical Scorecards
Related prompt	PR-86
EXAMPLE RUN	The monthly summary reported trial starts from organic against the baseline, with the AI mention trend shown alongside its method and set size.

SOP-41

Experiment design and evaluation

PER EXPERIMENT

Purpose	Change one thing and know whether it worked.
Owner	SEO lead
Trigger	When the team disagrees about what would work
Inputs	The hypothesis, a dated baseline, available volume
Tools	Notion, analytics
Preconditions	A baseline has been captured before any change
STEPS	
1	Record the baseline with source and date, before anything changes.
2	Define exactly one variable and what stays constant.
3	Set the minimum sample and the evaluation date.
4	Write the decision rule before starting.
5	Run without interference until the evaluation date.
6	Apply the rule and record the outcome either way.
Output	A concluded experiment with an adopt, reject, or extend decision
Quality gate	The decision rule predates the result; one variable only
Failure conditions	BLOCKED without a baseline; INSUFFICIENT EVIDENCE if volume cannot support a conclusion
Escalation	If volume is too low, run structured interviews instead and record that choice
Handoff	To SOP-42 for the decision log
Related database	Experiments, Decisions & Change Log
Related prompt	PR-99 equivalent — see Section 12
EXAMPLE RUN	An extractability experiment on twelve pages ran three monthly cycles; the rule was met and the change was adopted across the cluster.

SOP-42 Decision log update

CONTINUOUS

Purpose	Record what was decided and why, while the reasoning is fresh.
Owner	SEO lead
Trigger	On every material decision
Inputs	The decision, its context, evidence, and expected outcome
Tools	Notion
Preconditions	The decision has actually been made
STEPS	
1	Record the decision and the date.
2	Record the context as it was at the time.
3	Record the evidence and the options rejected.
4	Write a falsifiable expected outcome.
5	Set a review date.
6	On the review date, record the actual outcome and a verdict.
Output	A logged decision with a falsifiable prediction and a review date
Quality gate	The prediction is specific enough to be shown wrong
Failure conditions	INCOMPLETE without evidence or a review date
Escalation	A decision nobody will own is escalated rather than logged
Handoff	Feeds SOP-01 at the quarterly review
Related database	Decisions & Change Log
Related prompt	PR-87
EXAMPLE RUN	A decision to stop pursuing free-tier keywords was logged with the segment analysis as evidence and a 90-day review date; the verdict at review was Correct.

SECTION 14 · TEMPLATES

The template library

Fifty-five reusable documents. Nobody in the team should write anything from a blank page — every document type in the system has a starting structure with a completion criterion.

Each template below names the fields it contains, the criterion that says it is finished, the SOP and prompt that produce it, and the database it lands in. Build them as Notion database templates so a new record starts complete rather than empty.

TEMPLATE	FIELDS IT CONTAINS	COMPLETION CRITERION	SOP / PROMPT	DESTINATION DB
SEO strategy one-pager	Goal, audience, where we compete and do not, three bets with signals, constraints	Three bets each with a measurable signal and an owner	SOP-01 / PR-01	SEO Strategy
Stakeholder interview guide	Per-role question set, metric mapping, conflict questions	Every question maps to something the role is measured on	SOP-01 / PR-26	Business Goals
Business goal mapping	Objective, SEO outcome, primary metric, leading indicator, baseline	Every metric names its data source	SOP-01 / PR-27	Business Goals
SEO operating charter	Decision types, single decider each, consulted parties, escalation route	One decider per decision type	SOP-01 / PR-28	SEO Strategy
Constraint register	Capacity, publishing rate, dev availability, approval duration, budget	Values observed rather than aspirational	SOP-01 / PR-29	SEO Strategy
Maturity assessment	Ten dimensions scored 1–5 with the honest question for each	The lowest dimension is named, not averaged away	SOP-01	SEO Strategy
Customer problem record	Problem, verbatim quotes, source, date, distinct people, stage	At least one verbatim quote with a source and date	SOP-03 / PR-31	Customer Problems
Voice of customer record	Quote verbatim, theme, role, segment, source, date, permission	A recorded permission decision before publication	SOP-03 / PR-31	Voice of Customer
Demand validation record	Six criteria with pass or fail and the evidence for each	All six criteria tested; missing evidence fails	SOP-05 / PR-38	Search Demand
Question bank entry	Question as asked, asker role, stage, frequency, answered on, extractable	Extractable only ticked after the checklist is run	SOP-03 / PR-31	Question Bank
Search demand record	Query, demand type, evidence, volume with source, intent, surface	No volume without a named source and date	SOP-05 / PR-35	Search Demand
Keyword record	Query, type, intent, stage, ten criterion scores, mapping, owner	All ten criteria scored before mapping	SOP-17 / PR-15	Keyword Master
Keyword scoring sheet	The ten criteria with anchors, weights, and the band thresholds	Weight version recorded with the scoring date	SOP-17	Keyword Master
Declined keyword record	Keyword, reason, detail, decider, date, revisit condition	A revisit condition or an explicit permanent decline	SOP-17	Declined Keywords
Topic record	Topic, our belief, audience, business outcome, our evidence, priority	A stated belief and evidence only we hold	SOP-20 / PR-46	Topics
Topic cluster plan	Pillar scope, spoke list with keywords, commercial destination, interlinking	A commercial destination and bidirectional interlinking	SOP-20 / PR-46	Topic Clusters
Content brief	Keyword, reader, decision, angle, evidence, coverage, extractable passages, links	A differentiated angle and a specified commercial link	SOP-10 / PR-06	Content Briefs
Article outline	H2 and H3 structure, coverage per section, extractable passage markers	Every brief coverage question has a section	SOP-10 / PR-58	Content
Comparison page template	Buyer criteria in their order, verifiable cells, a genuine concession	A conceded criterion and no unverifiable competitor claims	SOP-10 / PR-60	Content
Alternatives page template	Honest positioning, who each option suits, switching considerations	States who the alternative genuinely suits better	SOP-10 / PR-41	Content
Product-led article	The job in steps, where the product helps, where it does not	Useful to someone who never buys	SOP-10 / PR-56	Content
Expert interview guide	Questions producing specifics, probes, citable-answer definitions	No question answerable generically	SOP-10 / PR-57	Content

TEMPLATE	FIELDS IT CONTAINS	COMPLETION CRITERION	SOP / PROMPT	DESTINATION DB
Case study template	Their situation, prior attempts, change, their figures, a real difficulty	Written permission and customer-supplied figures only	SOP-10 / PR-50	Content
Original research design	Data, methodology, limitations written first, supportable findings	Limitations drafted before findings	SOP-24 / PR-49	Content
Glossary entry	One-sentence definition, practical explanation, fact sheet consistency	Definition stands alone accurately	SOP-10 / PR-51	Content
Content QA checklist	On-page, extractability, claims, links, author credential	Unsourced claims are an automatic return	SOP-13 / PR-62	Content
Content refresh record	Decline evidence, cause type, decision, specific changes, recheck	A cause type before an action is chosen	SOP-14 / PR-16	Content Refresh Queue
Consolidation decision	Overlap, surviving URL, carry-across list, redirect map, link updates	Backlink equity assessed before removal	SOP-21 / PR-52	Content
Publishing handoff	Brief verification, metadata, links in and out, author, review date	Every brief requirement verified as present	SOP-13 / PR-53	Content
Technical audit sheet	Twenty modules with checks, evidence required, and severity	Every finding has evidence attached	SOP-06 / PR-64	Technical Issues
Technical issue record	Issue, module, evidence, five scoring inputs, owner, ticket	Confidence above 3, or a diagnostic instead	SOP-06 / PR-13	Technical Issues
Developer ticket	Problem, evidence, expected behaviour, testable criteria, rollback	Criteria testable by a non-SEO	SOP-07 / PR-14	Technical Issues
Migration map	URL inventory, destinations, one hop, backlinked removals flagged	Every current URL has a destination	SOP-19 / PR-66	Technical Issues
Release verification	Pre-deploy check, post-deploy crawl, 7-day and 30-day index checks	Both index checks completed before closure	SOP-20	Technical Releases & Verification
Incident response sheet	Release correlation, robots, status codes, index coverage, escalation	Release log checked before external causes	SOP-06 / PR-67	Technical Issues
Internal link opportunity	Source, target, anchor, rationale, three scores, priority	A reader-benefit rationale per link	SOP-22 / PR-68	Internal Link Opportunities
Site architecture map	Nodes, intended depth, actual depth, parent, navigation status	Intended depth recorded for every commercial node	SOP-22	Site Architecture
Authority target record	Target, type, relevance, specific reason, asset, contact route	A specific reason and an asset before contact	SOP-25 / PR-73	Authority Targets
Outreach email	Specific reference, one-sentence offer, one small ask, easy decline	Would not work for a different target	SOP-25 / PR-74	Authority Targets
Digital PR campaign plan	Asset, story, timing hook, methodology, limits, targets, spokesperson	Permission cleared and methodology published	SOP-25 / PR-18	Digital PR Campaigns
Linkable asset scorecard	Durability, effort, uniqueness, permission status, citation potential	Permissions confirmed in writing	SOP-24 / PR-49	Content
Location record	Canonical NAP, categories, profile status, landing page, audit date	Canonical NAP confirmed by the location owner	SOP-29 / PR-76	Local Locations
GBP audit sheet	Field-by-field comparison, category assessment, completeness result	Differences quoted exactly	SOP-29 / PR-76	Local Locations
Local landing page	Local services, local context, team, exact NAP, local schema, real proof	Three or more genuinely local elements	SOP-32 / PR-78	Content
Citation record	Directory, NAP as listed verbatim, match status, difference, correction date	Character-level comparison performed	SOP-30 / PR-76	Local Citations
Brand fact sheet	Every canonical fact, must-match surfaces, verification status, version	One canonical wording per fact	SOP-36 / PR-21	Brand Facts
AI prompt record	Exact wording, ID, type, intent, locale, version, good-result definition	Wording locked before the first run	SOP-33 / PR-23	AI Prompt Universe
AI answer run record	Prompt, assistant, date, mention, citation, sentiment, verbatim, errors	Verbatim answer pasted; mention and citation separate	SOP-34 / PR-79	AI Prompt Runs
AI citation record	Source, type, run, mentions us, mentions competitors, gap, route	Closability assessed before it becomes a target	SOP-37 / PR-81	AI Citations & Sources
AI visibility opportunity	Gap type, evidence runs, frequency, proposed action, area, owner	Linked to the runs that evidenced it	SOP-37 / PR-81	AI Visibility Opportunities

TEMPLATE	FIELDS IT CONTAINS	COMPLETION CRITERION	SOP / PROMPT	DESTINATION DB
AI scorecard	Five measures with method, set size, assistant list, comparability flag	Set size stated alongside every figure	SOP-35 / PR-80	Historical Scorecards
Monthly report	Metrics with sources and dates, changes, hypotheses, three decisions	Ends in three owned decisions	SOP-40 / PR-86	Reports
Quarterly review	Stated versus actual reconstruction, bet verdicts, new bets, not-doing list	The stated-versus-actual gap is named	SOP-01 / PR-30	SEO Strategy
Experiment record	Hypothesis, baseline, one variable, decision rule, evaluation date	Decision rule predates the result	SOP-41	Experiments
Decision log entry	Decision, context, evidence, options rejected, prediction, review date	The prediction is falsifiable	SOP-42 / PR-87	Decisions & Change Log

BUILD TEMPLATES AS DATABASE TEMPLATES, NOT PAGES

A template stored as a page gets duplicated, edited, and forked. A Notion database template applies automatically when a record is created, carries its own checklist inside the record, and cannot drift out of sync with the database it belongs to. The five-minute difference in setup is the difference between a system people use and a folder of documents they forget.

Northstar Analytics: one opportunity, end to end

A fictional B2B analytics SaaS company, from a market signal to a quarterly strategy change. Every value shown is invented to demonstrate the shape of a completed record.

NORTHSTAR ANALYTICS IS FICTIONAL

Northstar Analytics does not exist. It is an invented B2B analytics SaaS company selling to data teams at mid-market software businesses. Every figure, URL, keyword, score, and outcome shown for it is a placeholder written to show what a completed record looks like — none of it is a claim about anything real, and none of it should be read as a benchmark.

#	STAGE	DATABASE	THE RECORD (FICTIONAL)	QUALITY GATE	HUMAN DECISION	H A N D O F F
0 1	Business goal	Business Goals	Grow qualified trial starts from organic. Baseline 38/month as of 2026-07-01. Owner: Head of Growth. Primary metric: trial starts from organic. Leading indicator: impressions on BOFU pages.	A metric and a dated baseline exist	Head of Growth approves	G o a →S t r a t e g y
0 2	Customer problem	Customer Problems	"Monthly reporting takes three days." Seven distinct people, sourced from discovery calls between 2026-03 and 2026-05. Evidence strength: Strong. Stage: problem-aware.	Verbatim quotes present with sources and dates	Content lead validates	P r o b l e →m D e m a n d
0 3	Demand evidence	Search Demand	"analytics tool for data teams" — established demand, commercial intent, surfaces: Google and AI answers. Volume 390 per Tool X, 2026-06. Linked to the reporting problem.	Evidence names a source; volume carries its tool and date	SEO specialist validates	D e m a n →d K e y w o r d

#	STAGE	DATABASE	THE RECORD (FICTIONAL)	QUALITY GATE	HUMAN DECISION	H A N D O F F
04	Keyword record	Keyword Master	Scored 87/105 — Build now. Business relevance 5, intent match 5, conversion potential 5, difficulty 2, content gap 4. Mapped to the comparison hub. Cluster: tool selection.	All ten criteria scored; one owning page assigned	SEO lead confirms mapping	K e y w o r d C l u s t e r
05	Topic cluster	Topic Clusters	Tool selection cluster. Pillar: "How to choose an analytics tool". Nine spokes planned, four published — 44% coverage. Commercial destination: pricing page.	A commercial destination exists and interlinking is specified	Content lead approves	C l u s t e r B r i e f
06	Content brief	Content Briefs	Angle: selection should start from the reporting job, not the feature grid. Evidence: anonymized usage patterns. Will not claim: any time-saving percentage. Five coverage questions. Extractable: the definition and the criteria list.	A differentiated angle and a stated will-not-claim list	Editor approves	B r i e f C o n t e n t
07	On-page QA	Content	Direct answer at 62 words in the opening. Criteria list quotable in isolation. Two unsourced claims removed rather than softened. Author: Head of Data, credential visible.	No unsourced claims; key passages accurate alone	Editor signs the QA	C o n t e n t P u b l i s h i n g a n d o f f

#	STAGE	DATABASE	THE RECORD (FICTIONAL)	QUALITY GATE	HUMAN DECISION	H A N D O F F
08	Technical issue	Technical Issues	Facet URLs indexable. Module: faceted navigation. 2,400 URLs. Score 96, action Fix. Ticketed to ENG-4412, deployed 2026-07-18, verified 2026-07-20.	Evidence attached, confidence 4, verified in production	Engineering owner verifies	I s s u e → R e l e a s e d v e r i f i c a t i o n
09	Internal link opportunity	Internal Link Opportunities	Reporting guide → pricing. Anchor: "how our pricing works". Target value 5, relevance 4, source strength 3 — priority 60. Rationale: readers ask about cost at this point.	A reader-benefit rationale recorded	Editor approves before implementation	L i n k → V e r i f i e d
10	Authority campaign	Digital PR Campaigns	Reporting benchmark 2026. Methodology published, limits stated (self-reported, 140 responses, one sector). 34 targets contacted, 6 outcomes earned, 22 no responses logged.	Permission cleared and methodology published before outreach	Data owner clears publication	C a m p a i g → B a c k l i n k s
11	Prompt universe record	AI Prompt Universe	PU-014, category prompt: "What is the best analytics tool for a small data team?" Good result: named among 3–5 options with a link. Wording locked at version 1.	Wording locked before the first run	SEO lead fixes the wording	P r o m p t → R u n s

#	STAGE	DATABASE	THE RECORD (FICTIONAL)	QUALITY GATE	HUMAN DECISION	H A N D O F F
1 2	AI answer result	AI Prompt Runs	PU-014, Assistant A, 2026-08-04. Mentioned: no. Cited: no. Competitors named: A and B. Sources drawn on: two publishers and one directory. Verbatim answer recorded.	Verbatim answer pasted; mention and citation recorded separately	AI visibility owner classifies	R u →n C i t a t i o n s
1 3	Citation gap	AI Citations & Sources	Publisher covering Competitor A and not us, appearing in 4 of 38 prompts. Gap: Yes. Closable: Yes. Route: editorial pitch with the benchmark study.	Closability assessed and a route stated	Outreach owner accepts the target	G a →p A u t h o r i t y a r g e t
1 4	Opportunity	AI Visibility Opportunities	Publish a pricing structure page. Gap type: Absence. Evidence: PU-014, PU-019, PU-027. Frequency 3. Area: Content. Effort: Days. Priority: Now.	Linked to the runs that evidenced it	Content lead schedules it	O p p o r t u n i t →y C o n t e n t r i e f
1 5	Monthly report	Reports	August: mention rate 38% (up from 34%), citation rate 14%, coverage 52%, accuracy 92%. Set size 38 prompts, 3 assistants, unchanged. Extractability pass noted as a hypothesis, not a cause. Three decisions assigned.	Every figure sourced and dated; set size stated	SEO lead assigns the decisions	R e p o r →t D e c i s i o n s

#	STAGE	DATABASE	THE RECORD (FICTIONAL)	QUALITY GATE	HUMAN DECISION	HAND OFF
16	Decision log entry	Decisions & Change Log	Publish the pricing structure page in Q4. Evidence: absent from 3 recommendation prompts; both named competitors publish pricing. Prediction: named in at least one of the three by the January run. Review: 2027-01-31.	A falsifiable prediction and a review date	SEO lead records it	Decision strategy

WHAT THE SIXTEEN STEPS SHOW

Every record traces backwards to a customer statement and forwards to a commercial decision. The keyword exists because seven people described a problem. The AI opportunity exists because three tracked prompts showed absence. The quarterly decision exists because the monthly report ended in decisions rather than observations. That traceability is the point of the whole system — not the individual records.

SECTION 16 · WORKFLOWS

Ten cross-functional workflows

Each one crosses several operating areas. These are what the databases exist to support — the individual records are only useful because these sequences run through them.

WORKFLOW 01 Customer problem to published content

W-01

Trigger	A validated problem reaches three distinct people
Databases touched	Customer Problems → Search Demand → Keyword Master → Topic Clusters → Content Briefs → Content
Status changes	New → Validated → Mapped → Brief → Drafting → Review → Published
SOPs used	SOP-03, SOP-17, SOP-18, SOP-20, SOP-10, SOP-13
Prompts used	PR-31, PR-36, PR-43, PR-46, PR-06, PR-62
Human owners	Content lead owns the sequence; the editor owns the QA gate
Quality gates	Demand validation pack, keyword mapping, brief approval, content QA
Final output	A published page answering an evidenced problem
Handoff	To internal linking (W-05) and AI visibility tracking
Learning recorded	The angle that worked and the coverage questions readers still asked

WORKFLOW 02 Keyword opportunity to mapped page

W-02

Trigger	A keyword scores above 60 and has no owning page
Databases touched	Keyword Master → Content → Topic Clusters → Content Briefs
Status changes	Unmapped → Mapped → Brief → Published
SOPs used	SOP-17, SOP-18, SOP-19, SOP-10
Prompts used	PR-15, PR-43, PR-44, PR-06
Human owners	SEO specialist maps; content lead commissions
Quality gates	All ten criteria scored; no page holds two primaries
Final output	A keyword with exactly one owning page
Handoff	To content production (W-01)
Learning recorded	Whether the mapping held or produced cannibalization at 60 days

WORKFLOW 03 Declining page to recovered asset

W-03

Trigger	The decay monitor flags a page declining over 90 days
Databases touched	Content → Content Refresh Queue → Content Briefs → Content
Status changes	Declining → Diagnosed → Scheduled → Shipped → Verified
SOPs used	SOP-09, SOP-14, SOP-13
Prompts used	PR-16, PR-55, PR-62
Human owners	Content lead diagnoses; the writer executes; the editor QAs
Quality gates	Cause type identified before an action; 60-day recovery verification
Final output	An updated page with a verified recovery outcome
Handoff	To reporting (W-10) with the recovery result
Learning recorded	Which cause types actually recover and which do not

WORKFLOW 04 Technical issue to verified fix

W-04

Trigger	A crawl, alert, or log review surfaces an issue
Databases touched	Technical Issues → Technical Releases & Verification → Site Architecture
Status changes	Found → Diagnosed → Ticketed → In progress → Deployed → Verified
SOPs used	SOP-06, SOP-07, SOP-20
Prompts used	PR-13, PR-14, PR-64, PR-65
Human owners	Technical SEO owner diagnoses; engineering owner ships; both verify
Quality gates	Confidence above 3 before ticketing; production verification before closure
Final output	A verified fix with a recorded date
Handoff	To release verification and the 30-day index check
Learning recorded	Whether the issue recurs after the next release

WORKFLOW 05 Orphan money page to supported architecture

W-05

Trigger	A crawl shows a commercial page with fewer than five inbound internal links
Databases touched	Content → Internal Link Opportunities → Site Architecture
Status changes	Identified → Approved → Implemented → Verified
SOPs used	SOP-22, SOP-23
Prompts used	PR-68, PR-69, PR-70
Human owners	SEO specialist identifies; the editor implements and verifies
Quality gates	Relevance above 3; every link carries a reader rationale
Final output	A commercial page with meaningful internal support
Handoff	To reporting, with the inbound count before and after
Learning recorded	Which source pages actually drive the most qualified onward clicks

WORKFLOW 06 Original data to earned citations

W-06

Trigger	First-party data exists with publication permission
Databases touched	Content → Digital PR Campaigns → Authority Targets → Backlinks → AI Citations & Sources
Status changes	Planning → Live → Closed
SOPs used	SOP-24, SOP-25, SOP-27
Prompts used	PR-49, PR-73, PR-74
Human owners	Marketing owner runs the campaign; the data owner clears publication
Quality gates	Methodology published; no claim exceeds the data; permissions in writing
Final output	Earned coverage, links, and citations logged against effort
Handoff	To AI visibility, where cited sources are tracked
Learning recorded	Which asset types and angles actually earn coverage

WORKFLOW 07 Local inconsistency to corrected presence

W-07

Trigger	A monthly audit finds a listing that does not match the canonical record
Databases touched	Local Locations → Local Citations → Content
Status changes	Found → Inconsistent → Correction requested → Correct
SOPs used	SOP-29, SOP-30, SOP-32
Prompts used	PR-76, PR-78
Human owners	Local SEO owner audits and requests; the location owner confirms the canonical NAP
Quality gates	Character-level comparison; differences quoted exactly
Final output	A consistent NAP across every listing
Handoff	To reporting, with the consistency percentage
Learning recorded	Which directories drift most often and how quickly they respond

WORKFLOW 08 AI prompt gap to content opportunity

W-08

Trigger	A tracked prompt shows absence across three consecutive runs
Databases touched	AI Prompt Runs → AI Visibility Opportunities → Content Briefs → Content
Status changes	Identified → Scheduled → In progress → Done
SOPs used	SOP-34, SOP-37, SOP-10
Prompts used	PR-79, PR-81, PR-83
Human owners	AI visibility owner identifies; content lead schedules and delivers
Quality gates	Linked to the evidencing runs; reader value before extractability
Final output	Published content designed to be both useful and citable
Handoff	To the next monthly run, where the gap is rechecked
Learning recorded	Whether the content changed the mention or citation outcome

WORKFLOW 09 Factual AI error to corrected source ecosystem

W-09

Trigger	A tracked answer states something incorrect about the brand
Databases touched	AI Prompt Runs → Brand Facts → Content → Local Citations → AI Prompt Runs
Status changes	Error found → Sources identified → Owned corrected → Requests sent → Rechecked
SOPs used	SOP-34, SOP-36, SOP-30
Prompts used	PR-79, PR-82, PR-21
Human owners	Marketing owner corrects at source; AI visibility owner rechecks
Quality gates	Owned properties corrected first; no attempt to influence the model directly
Final output	A corrected public source ecosystem and a rechecked answer
Handoff	To the next monthly run and the accuracy measure
Learning recorded	How long correction takes to appear, per source type

WORKFLOW 10 Monthly performance data to quarterly strategy change

W-10

Trigger	The monthly reporting cycle completes
Databases touched	Reports → Decisions & Change Log → SEO Strategy → Business Goals
Status changes	Analysed → Decisions assigned → Strategy revised → Bets set
SOPs used	SOP-39, SOP-40, SOP-01, SOP-42
Prompts used	PR-85, PR-86, PR-30, PR-87
Human owners	SEO lead analyses and proposes; the business owner approves the strategy
Quality gates	Every figure sourced and dated; the stated-versus-actual gap named
Final output	A versioned strategy with three new bets and a not-doing list
Handoff	To the next quarter, where the cycle restarts at W-01
Learning recorded	Whether the previous quarter's predictions proved correct

INSTRUMENT THREE, NOT TEN

W-01, W-04, and W-10 together cover demand becoming content, findings becoming shipped fixes, and data becoming strategy. If those three run reliably, the other seven are variations on patterns the team already knows. Trying to instrument all ten in the first month produces ten half-followed processes.

How to build it in Notion

Ten steps, roughly nine hours for the full system. The order is a dependency chain, not a preference — building relations before properties is the most common reason people abandon the build.

STEP	WHAT YOU DO	TIME	DEPENDS ON	EXPECTED RESULT	QUALITY CHECK	COMMON SETUP ERROR
STEP 1	Create the workspace structure	45 minutes	None	Seventeen top-level pages, empty, in the order given in Section 02	Every page exists and is named exactly as specified	Creating sub-pages before the top level, which makes the sidebar unreadable
STEP 2	Create the databases in dependency order	2–3 hours	Step 1	Thirty-two databases with their properties, and no relations yet	Every property has the correct Notion type and default; no relation properties exist yet	Adding relations while building, which creates half-linked databases you cannot untangle
STEP 3	Add relations	60 minutes	Step 2 complete for all thirty-two	Every relation from the specifications, created once from the source side	Every relation appears on both databases and points at the right one	Creating a relation from both sides, which produces two separate relation properties
STEP 4	Add rollups	40 minutes	Step 3	Fourteen rollups calculating from their relations	Every rollup returns a number rather than an error	Building a rollup on a relation with no records yet, then assuming it is broken
STEP 5	Add formulas	60 minutes	Steps 2 and 4	Twenty-four formulas, tested on one record each	Each formula returns the expected value on a test record	Copying a formula without renaming the properties it references
STEP 6	Create database templates	90 minutes	Step 5	One template per database, with pre-filled sections and a completion checklist	Creating a new record produces a structured, prompted starting point	Building templates as separate pages rather than as database templates
STEP 7	Create dashboard views	90 minutes	Steps 3 to 5	The Command Center with six blocks and its linked views	Every block returns rows and answers the question it is named for	Building views before there is data, then filtering on properties that are all empty
STEP 8	Import or create example records	45 minutes	Step 6	Three records per flagship database so views and formulas can be checked	Every view shows something; every formula calculates	Skipping this, then discovering a formula error three weeks into real use
STEP 9	Assign owners	30 minutes	Step 8	A named person on every database and every SOP	No database and no SOP has an empty owner	Assigning a team rather than a person, which means nobody
STEP 10	Schedule review cadences	30 minutes	Step 9	Every cadence in a shared calendar with the owner invited	The weekly rhythm and every monthly and quarterly review has a calendar entry	Relying on the Notion cadence property alone, which nobody looks at

The database build order

Ten waves. Each wave relates only to databases already created.

WAVE	DATABASES	WHY THIS ORDER
First	Business Goals, SEO Strategy	Nothing else has anything to point at without these
Second	Customer Problems, Voice of Customer, Search Demand, Question Bank	The demand layer feeds keywords and content
Third	Keyword Master, Query Patterns, Declined Keywords	Keywords relate to demand records already created
Fourth	Topics, Topic Clusters, Content, Content Briefs, Content Refresh Queue	Content relates to keywords and clusters, which must exist first
Fifth	Site Architecture, Technical Issues, Technical Releases & Verification, Internal Link Opportunities	These relate to Content, which is now in place
Sixth	Backlinks, Authority Targets, Digital PR Campaigns	Authority relates to Content assets
Seventh	Local Locations, Local Citations	Independent of the rest; build only if local applies
Eighth	Brand Facts, AI Prompt Universe, AI Prompt Runs, AI Citations & Sources, AI Competitors, AI Visibility Opportunities, Historical Scorecards	The AI layer relates to Content and Brand Facts, both now in place
Ninth	Experiments, Reports, Decisions & Change Log	These relate to everything, so they come last
Tenth	Prompt Library, SOP Library, Template Library	Reference databases with no dependencies — build them whenever

Three deployment modes

The same architecture, scaled to who is operating it.

Solo consultant or founder

Build	About 90 minutes. Build nine databases: Business Goals, Customer Problems, Keyword Master, Topics, Content, Technical Issues, Internal Link Opportunities, Brand Facts, and Decisions & Change Log. Skip clusters, briefs, and the full AI layer until there is enough activity to justify them.
Run	Run SOP-02 weekly and SOP-39 on Fridays. Everything else waits.

COMMON MISTAKE

The temptation is to build all thirty-two and use four. Build nine and use nine.

Small marketing team

Build	About five hours across two sessions. Build all thirty-two databases but populate the demand, keyword, and content layers first. Assign owners at the database level rather than the record level initially.
Run	Run the full weekly rhythm with a named owner per day. Add the AI layer in month two, once the content layer is producing.

COMMON MISTAKE

Distributing the message approval gate across the team, after which quality drifts within a fortnight.

Agency or multi-client team

Build	About nine hours plus per-client setup. Build the master workspace once, add a Client relation to every client-specific database, and keep the SOP, Prompt, and Template libraries shared and unfiltered.
Run	Run per-client weekly rhythms offset across the week. Client-specific Brand Facts and voice files are non-negotiable.

COMMON MISTAKE

Sharing a brand fact sheet or voice file across clients, after which every client's output starts sounding identical.

SECTION 19 · MULTI-CLIENT

The agency and multi-project model

One master workspace with a Client relation, shared libraries, and strict separation of anything brand-specific. The alternative — one workspace per client — is defensible but does not scale past four.

DECISION	ONE WORKSPACE PER CLIENT	MASTER WORKSPACE WITH A CLIENT RELATION
Setup time	Multiplied by the number of clients	Once, plus about 20 minutes per client
SOP and prompt updates	Repeated in every workspace, and they drift	Updated once, applied everywhere
Cross-client learning	Effectively impossible	A Client filter removed shows patterns across the portfolio
Data separation	Absolute by construction	Depends entirely on discipline with the Client property
Client access	Simple — share the workspace	Requires per-database permissions or a filtered client-facing page
Offboarding	Archive or transfer the whole workspace	Filter, export, and archive by Client — more work
Best for	Three or fewer clients, or where the contract requires isolation	Four or more clients, and any team that wants shared learning

What must be client-specific, and what must not

Always client-specific

- Brand Facts — a shared fact sheet is a data leak waiting to happen
- Voice and style — sharing one is why agency output sounds identical
- Business Goals and SEO Strategy
- Every demand, keyword, content, and technical record
- AI prompt universes and all run data
- Reports, decisions, and experiment records

Always shared

- The SOP Library — the procedures are your product
- The Prompt Library — with client context pasted at run time
- The Template Library
- Query Patterns and the scoring model definitions
- The technical module list and severity definitions
- The workflow definitions in Section 16

Preventing data leakage between clients

Six rules. The first one prevents most of what goes wrong.

- ☐ Every client-specific database has a required Client relation — no record without one
- ☐ Every default view filters on Client; the unfiltered view is not the default
- ☐ Prompt runs paste the client context explicitly; nothing relies on remembered context
- ☐ Brand Facts and voice files are never shared, referenced, or copied between clients
- ☐ Client-facing pages are separate filtered pages, never the working database
- ☐ Offboarding filters by Client, exports, archives, then removes access — in that order


ON CLIENT CONFIDENTIALITY

Nothing in this playbook assumes you may use one client's data to inform another's work. Cross-client learning means procedural learning — which SOPs work, which prompt phrasings produce usable output, which technical modules generate the most issues. It does not mean sharing keyword lists, performance figures, brand facts, or research between accounts. Check your contracts before assuming otherwise.

The SEO operating dashboard

Fifteen connected databases and eleven views. This is the Notion workspace the previous ten sections have been feeding – the place the system actually lives.

THE SEO BRAIN WORKSPACE, IN FIVE LAYERS



Fifteen databases, five layers, one workspace. Build the demand and production layers first — the rest connect to them.

The fifteen databases

Nine are specified in full earlier in this playbook. The remaining six are below.

DATABASE	LAYER	SPECIFIED IN	REVIEW CADENCE
Customer Problems	Demand	Section 02	Quarterly
Search Demand	Demand	Section 02	Monthly
Question Bank	Demand	Section 02	Monthly
Keywords	Demand	Section 03	Weekly status, quarterly scoring
Content	Production	Section 04	Weekly
Technical Issues	Infrastructure	Section 06	Weekly
Internal Links	Production	Section 07	Weekly
Backlinks and Authority	Infrastructure	Section 08	Weekly
Local SEO	Infrastructure	Section 09	Monthly
AI Visibility	Infrastructure	Section 10	Monthly
Strategy	Strategy	Below	Quarterly
Topics	Demand	Below	Monthly
Competitors	Strategy	Below	Quarterly
Experiments	Strategy	Below	Weekly
Opportunities	Strategy	Below	Weekly

CONTENT, KEYWORDS, AI VISIBILITY Topics Database

PROPERTY	TYPE	WHAT IT HOLDS
Topic	Title	The cluster name, in customer language
Pillar page	Relation	Content database — the anchor
Supporting content	Relation	Content database — the spokes
Keywords	Relation	Keyword Master Database
Commercial destination	Relation	The page the cluster points at
Coverage	Formula	Published pieces / planned pieces
Authority status	Select	None / building / established
Owner	Person	Who owns this cluster
Priority	Select	This quarter / next / later
PURPOSE The clusters that organize everything else.		REVIEW Prevent content from being a list of posts rather than a body of work.
		RELATED Monthly

TOPICS, AI VISIBILITY, CONTENT Competitors Database

PROPERTY	TYPE	WHAT IT HOLDS
Competitor	Title	Name
Type	Select	Direct / adjacent / publisher / aggregator
Surfaces	Multi-select	Where they compete with you
Topics they own	Relation	Topics database
Their positioning	Text	In their words, dated
Named in AI answers	Number	Count from the AI Visibility database
Last reviewed	Date	Drives the review cadence
Notes	Text	Fact and inference labelled separately
PURPOSE Who you are actually competing with on each surface.		REVIEW Keep competitive understanding current and evidence-based.
		RELATED Quarterly

CONTENT, TECHNICAL ISSUES, REPORTING Experiments Database

PROPERTY	TYPE	WHAT IT HOLDS
Experiment	Title	One variable, stated
Hypothesis	Text	What you believe and why
Baseline	Text	The figure, its source, and the as-of date
Variable changed	Text	Exactly one
Held constant	Text	Everything else
Start / evaluation date	Date ×2	Committed in advance
Decision rule	Text	Written before the experiment starts
Result	Text	The figure with its sample context
Decision	Select	Adopt / reject / extend / inconclusive
Owner	Person	Who runs it

PURPOSE What you changed, why, and what happened.

REVIEW Make change attributable instead of coincidental.

RELATED Weekly

ALL DATABASES Opportunities Database

PROPERTY	TYPE	WHAT IT HOLDS
Opportunity	Title	What could be done
Source	Select	Audit / research / AI visibility / competitor / idea
Area	Select	Which of the ten operating areas
Effort	Select	Hours / days / weeks
Expected effect	Text	Described, not estimated numerically
Evidence	Text	What suggests this is worth doing
Priority	Select	Now / this quarter / later / declined
Owner	Person	Who would do it
Status	Select	Identified / scheduled / in progress / done / declined

PURPOSE Everything worth doing that is not yet scheduled.

REVIEW Stop good ideas from living in people's heads and Slack threads.

RELATED Weekly

ALL DATABASES Reporting Database

PROPERTY	TYPE	WHAT IT HOLDS
Period	Title	Week or month
Metrics	Text	With sources and as-of dates
What changed	Text	Versus the previous period
Hypotheses	Text	What we think caused it, labelled as hypotheses
Decisions	Text	Three, each with an owner
System changes	Text	What was updated as a result
Author	Person	Who wrote it

PURPOSE The record of what was reported, decided, and changed.

REVIEW Turn reporting into a learning record rather than a status update.

RELATED Weekly

The eleven dashboard views

Build these as filtered views over the databases above. Each answers one question.

VIEW	THE QUESTION IT ANSWERS	FILTER
Executive overview	Are we making progress on the quarter's goals?	Strategy bets, with their success signals
This week	What is happening in the next seven days?	Due date within 7 days, across all databases
High-priority actions	What matters most right now?	Priority = now, sorted by area
Content pipeline	What is being written and where is it stuck?	Content where status is brief, drafting, or review
Technical backlog	What is broken and who is fixing it?	Technical issues where status is not verified, sorted by priority score
AI visibility opportunities	Where are we absent from AI answers?	AI Visibility where mentioned is false, grouped by prompt type
Quick wins	What could we do this week for a real effect?	Effort = hours and priority = now or this quarter

VIEW	THE QUESTION IT ANSWERS	FILTER
At risk	What is slipping?	Due date passed and status not done
Waiting for review	What is blocked on someone?	Status = review, grouped by reviewer
Completed	What shipped, and when?	Status = done or verified, sorted by date descending
Monthly review	What should we change about the system?	Reporting entries plus experiments concluded this month

BUILD ORDER, IF YOU ARE STARTING FROM NOTHING

Keywords and Content first — those two alone replace most spreadsheets. Then Technical Issues and Internal Links, because that is where the fastest wins usually are. Then AI Visibility, which needs a month of data before it says anything. Strategy, Experiments, and Reporting last, once there is enough activity to review. Building all fifteen in week one produces fourteen empty databases and one you actually use.

The weekly SEO operating rhythm

A system that is not on a calendar is a document. This is the cadence that turns the previous eleven sections into something that runs.

THE WEEKLY SEO OPERATING RHYTHM

MON	TUE	WED	THU	FRI
Review & decide - Performance and rankings - AI visibility check - Technical alerts - Confirm the week's priorities	Research - Keyword and topic research - Competitor analysis - Content briefs	Produce - Content production - On-page optimization - Internal linking	Build authority - Technical fixes shipped - Digital PR and outreach - Local SEO - AI visibility actions	Report & learn - Reporting - Learning review - Update dashboards - Archive completed work

Batch by activity type, not by page. Friday is the step that makes next week better than this one.

The week in detail

DAY	WHAT HAPPENS	OUTPUT BY END OF DAY	TIME
Monday	Review performance, rankings, AI visibility, technical alerts. Confirm priorities.	The week's three priorities, assigned	60-90 min
Tuesday	Keyword and topic research, competitor analysis, content briefs.	Briefs ready for anything being written this week	Half day
Wednesday	Content production, on-page optimization, internal linking.	Drafts written, pages optimized, links implemented	Half day
Thursday	Technical fixes shipped, authority outreach, local SEO, AI visibility actions.	Tickets progressed, outreach sent, actions logged	Half day
Friday	Reporting, learning review, dashboard updates, archiving.	The weekly report and its three decisions	60-90 min

The longer cadences

Weekly work compounds only if these happen on schedule rather than when someone remembers.

CADENCE	WHAT HAPPENS	OUTPUT
Monthly review	Content decay audit, AI visibility run, internal link audit, backlink review	An updated opportunity list and a refresh queue
Quarterly strategy review	Re-run the maturity assessment, revise the strategy one-pager, re-score the keyword list	A versioned strategy and a revised build order
Content refresh cycle	Every published page reviewed at least twice a year, high-value pages quarterly	Update, merge, redirect, or remove decisions
Technical crawl	Full crawl monthly; after every release touching templates; immediately on any alert	New issues scored and added to the backlog
AI prompt tracking	The full prompt universe re-run monthly, same wording, same assistants	A comparable data point added to the trend

CADENCE	WHAT HAPPENS	OUTPUT
Log file review	Quarterly, or when crawl or indexing behaviour changes unexpectedly	Crawl budget findings and orphan discovery

IF YOU ONLY HAVE TWO HOURS A WEEK

Do Monday and Friday. Thirty minutes on Monday to look at what changed and pick one thing; ninety minutes on Friday to do that one thing, record what happened, and pick the next. A two-hour rhythm sustained for a year beats a full programme abandoned in March, and it is the version most solo founders and small teams actually manage.

SECTION 13 · PROMPTS

The Claude prompt library

Twenty-five prompts covering the full operating system. Each one carries its required inputs, its constraints, its output format, and what to do when information is missing.

These are not one-line prompts. Each is structured so the output is consistent enough to file in the databases from Section 11, and constrained enough that it declines to invent what it does not have. The bracketed variables are where your business goes — a prompt run without filling them produces exactly the generic output people complain about.

THE RULE THAT APPLIES TO ALL TWENTY-FIVE

Every output is a draft. Nothing here goes live, gets sent, or changes a page without a person reviewing it. Where a prompt says it will name a gap rather than fill it, that behaviour is the point — an answer that invents a benchmark is worse than no answer, because it looks like research.

PROMPT 01 SEO strategy one-pager

STRATEGY

Required inputs Business goal, audience, current performance, capacity, competitors.

COPY-PASTE PROMPT

ROLE: You are a senior SEO strategist advising an in-house team. You are paid to be useful rather than agreeable.

OBJECTIVE: Produce a one-page SEO strategy that connects a business goal to a small number of defensible bets.

CONTEXT:

Business goal this quarter: [GOAL]

What we sell and to whom: [OFFER AND ICP]

Current organic performance: [PASTE — traffic, rankings, conversions, dates]

Capacity: [people, publishing rate, dev capacity]

Named competitors: [LIST]

Our unfair advantage: [first-party data, expertise, product, reach]

CONSTRAINTS: Use only what I supplied. Do not invent benchmarks, industry averages, or traffic estimates. Do not forecast rankings or traffic. Where a recommendation depends on information I did not give you, name the gap instead of filling it.

PROCESS:

1. Restate the business goal in SEO terms.
2. Identify where we could realistically compete, given capacity and authority.
3. Identify where we should not compete, and say why.
4. Propose three bets with a measurable success signal for each.
5. Name the constraint most likely to break the plan.

OUTPUT FORMAT: The one-page strategy in the template structure, then a table of the three bets with signal, owner, and first action.

QUALITY CHECKS: Every bet traces to the stated business goal. Nothing is forecast. The not-doing list is specific rather than generic.

MISSING INFORMATION: List what you would need and how it would change the recommendation. Do not proceed on assumption.

HUMAN REVIEW: A named owner approves before this drives any work.

Output format A one-page strategy plus a three-bet table with owners and signals.

Quality checks No forecasts, no invented benchmarks, a specific not-doing list.

If information is missing Names the gap and stops rather than assuming.

Human review The strategy owner approves and versions it with a date.

PROMPT 02 Audience and problem research

RESEARCH

Required inputs Transcripts, reviews, tickets, or interview notes — labelled by source and date.

COPY-PASTE PROMPT

ROLE: You are a voice-of-customer analyst.

OBJECTIVE: Extract the problems customers actually describe, grouped by underlying problem rather than by wording.

CONTEXT:
Who we sell to: [ICP]
Material: [PASTE transcripts, reviews, tickets, notes – label the source and date of each item]

CONSTRAINTS: Quote exactly when quoting; never smooth a quote into better prose. Group by underlying problem, not by similar words. Mark any group supported by one person as a single data point, not a pattern. Do not add problems you think the audience probably has.

PROCESS:

1. Extract every problem statement verbatim, with its source.
2. Group by underlying problem.
3. Count distinct people per group, not mentions.
4. Note the trigger that preceded each search where it is stated.
5. Identify vocabulary customers use that our material does not.

OUTPUT FORMAT: A frequency-ranked problem table with distinct-person counts, a verbatim quote bank grouped by theme, and the vocabulary gap list.

QUALITY CHECKS: Every group cites its evidence. Single mentions are flagged. No group merges two different underlying problems.

MISSING INFORMATION: If the material is paraphrased rather than verbatim, say so and stop – a language bank built on paraphrase is worthless.

HUMAN REVIEW: Marketing confirms phrases before they enter customer-facing copy.

Output format	A ranked problem table, a verbatim quote bank, and a vocabulary gap list.
Quality checks	Verbatim quotes only; distinct-person counts; single mentions flagged.
If information is missing	Stops if the input is paraphrased rather than verbatim.
Human review	Marketing approves phrases before they appear in copy.

PROMPT 03 Keyword clustering

KEYWORDS

Required inputs A raw keyword export with volume and difficulty where available.

COPY-PASTE PROMPT

ROLE: You are an SEO analyst clustering keywords into topics.

OBJECTIVE: Turn a flat keyword list into clusters that each map to one page.

CONTEXT:

Keyword list: [PASTE with volume and difficulty where you have them]

What we sell: [OFFER]

Existing pages: [PASTE URLs and their primary topics]

CONSTRAINTS: Cluster by search intent and by whether one page could reasonably satisfy all of them – not by string similarity. Do not invent volumes. Where a keyword could belong to two clusters, place it in one and note the ambiguity.

PROCESS:

1. Group keywords that one page could satisfy together.
2. Name each cluster in the language customers use.
3. Identify the head term and the supporting long-tail for each.
4. Map each cluster to an existing page, or mark it as a gap.
5. Flag clusters that overlap existing pages – these are cannibalization risks.

OUTPUT FORMAT: A table of clusters with head term, supporting keywords, mapped page or GAP, and the cannibalization flags.

QUALITY CHECKS: No keyword appears in two clusters. Every cluster could plausibly be one page. Ambiguous placements are noted.

MISSING INFORMATION: If existing pages were not supplied, produce the clusters and mark all mapping as UNKNOWN rather than guessing.

HUMAN REVIEW: The SEO owner confirms the mapping before content is commissioned.

Output format	A cluster table with head terms, supporting keywords, mapping, and flags.
Quality checks	One keyword, one cluster; every cluster is plausibly one page.
If information is missing	Marks mapping UNKNOWN rather than guessing at existing pages.
Human review	The SEO owner confirms mapping before commissioning.

PROMPT 04 Search intent classification

KEYWORDS

Required inputs A keyword list, plus what currently ranks for each if available.

COPY-PASTE PROMPT

ROLE: You are classifying search intent for a keyword set.

OBJECTIVE: Assign each keyword an intent and a buyer stage so pages can be matched to them correctly.

CONTEXT:
Keywords: [PASTE]
What currently ranks: [PASTE the top three result types per keyword, if you have them – this is the strongest available evidence of intent]
Our page types: [LIST what we can offer: guide, comparison, product, pricing]

CONSTRAINTS: Where SERP evidence is supplied, it outranks your assumption about the intent. Where it is not supplied, mark the classification as inferred. Do not classify commercial intent onto a query simply because we would prefer it.

PROCESS:
1. Classify intent: informational, comparison, commercial, navigational.
2. Assign a buyer stage: problem-aware, solution-aware, vendor-aware, customer.
3. Recommend the page type that matches.
4. Flag mismatches where our existing page type conflicts with the intent.

OUTPUT FORMAT: A table with keyword, intent, stage, recommended page type, evidence used, and a mismatch flag.

QUALITY CHECKS: Every classification states whether it was evidenced or inferred. Mismatches are flagged rather than smoothed over.

MISSING INFORMATION: If no SERP evidence is supplied, say that all classifications are inferred and lower-confidence.

HUMAN REVIEW: Spot-check ten classifications against live results before acting.

Output format	An intent and stage table with recommended page types and mismatch flags.
Quality checks	SERP evidence outranks assumption; inferred classifications are labelled.
If information is missing	Declares the whole set inferred if no SERP evidence exists.
Human review	Ten classifications spot-checked against live results.

PROMPT 05 Topic prioritization

STRATEGY

Required inputs Candidate topics, the scoring criteria, and honest capacity.

COPY-PASTE PROMPT

ROLE: You are prioritizing a content backlog for a team with limited capacity.

OBJECTIVE: Produce a ranked build order using the content scoring model.

CONTEXT:

Candidate topics: [LIST with demand evidence for each]

Our evidence and expertise: [what we hold that others do not]

Existing coverage: [PASTE what is already published]

Capacity: [pieces per month, realistically]

Business goal: [GOAL]

CONSTRAINTS: Score against the six criteria in the content scoring model – business proximity, evidence we hold, demand evidence, competitive gap, cluster contribution, citation potential. Do not score a topic highly for demand without evidence supplied. Cut the list at stated capacity.

PROCESS:

1. Score each candidate on the six criteria.
2. Apply the weights and rank.
3. Cut at capacity and place the remainder on an explicit not-now list.
4. For the top five, state the angle that makes ours worth reading.
5. Identify any candidate better served by updating an existing page.

OUTPUT FORMAT: The scoring table, the capacity-cut build order, the not-now list, and the update-instead recommendations.

QUALITY CHECKS: The cut matches stated capacity. Every high score cites its evidence. Update-instead options are surfaced before new commissions.

MISSING INFORMATION: Where demand evidence is absent for a topic, score it as unevidenced rather than assuming.

HUMAN REVIEW: The content owner approves the build order.

Output format	A scored, capacity-cut build order with a not-now list.
Quality checks	Cut matches capacity; scores cite evidence; refresh options surfaced first.
If information is missing	Scores unevidenced rather than assuming demand.
Human review	The content owner approves the order.

PROMPT 06 Content brief generation

CONTENT

Required inputs The keyword, the intent, our evidence, and the current top results.

COPY-PASTE PROMPT

ROLE: You are a content strategist writing a brief a writer could work from without asking questions.

OBJECTIVE: Produce a complete content brief.

CONTEXT:

Primary keyword: [KEYWORD] Intent: [INTENT] Stage: [STAGE]

Reader: [who, and what they already know]

Current top results: [PASTE what they cover, or describe them]

Our evidence: [first-party data, customer quotes, product experience]

Cluster: [which pillar this supports]

Commercial destination: [the page this must link to]

CONSTRAINTS: The angle must be something the current results do not offer. Do not specify a word count as the primary target – specify required coverage. Every claim the piece will make must be traceable to the evidence I supplied; where it is not, write [NEEDS SOURCE].

PROCESS:

1. State the decision this page supports.
2. Define the angle and why it is defensible.
3. List required coverage as questions, in order.
4. Identify which passages must be self-contained and extractable.
5. Specify internal links in and out, including the commercial page.

OUTPUT FORMAT: The brief in the content brief template structure, then a list of every [NEEDS SOURCE] marker.

QUALITY CHECKS: The angle is genuinely differentiated. Extractable passages are identified. The commercial link is specified.

MISSING INFORMATION: If no differentiating evidence exists, say so and recommend not commissioning the piece.

HUMAN REVIEW: The editor approves the brief before writing starts.

Output format	A complete brief plus a list of unsourced claims to resolve.
Quality checks	A differentiated angle, extractable passages named, commercial link specified.
If information is missing	Recommends not commissioning where no differentiation exists.
Human review	The editor approves before writing starts.

PROMPT 07 SERP analysis

RESEARCH

Required inputs The query and the current top results, pasted or described.

COPY-PASTE PROMPT

ROLE: You are analysing a results page to understand what the query rewards.

OBJECTIVE: Determine what kind of page could realistically win this query.

CONTEXT:

Query: [QUERY]

Top results: [PASTE titles, URLs, page types, and what each covers]

Our current position and page: [if any]

CONSTRAINTS: Analyse only what I supplied. Do not estimate traffic, authority scores, or ranking probability. Do not assume ranking factors – describe patterns you can observe in the results themselves.

PROCESS:

1. Identify the dominant page type and format.
2. Identify what every result covers – the table stakes.
3. Identify what none of them cover – the opening.
4. Assess how much depth the query appears to reward.
5. Judge whether our current page type matches the intent shown.

OUTPUT FORMAT: A pattern summary, the table-stakes list, the gap, and a verdict on whether to target this query with the page type we have.

QUALITY CHECKS: Observations are grounded in the supplied results. No estimation of metrics you cannot see.

MISSING INFORMATION: If fewer than five results were supplied, note that the pattern read is low-confidence.

HUMAN REVIEW: An SEO reviews before the finding changes a content plan.

Output format A pattern read, the table stakes, the gap, and a target-or-not verdict.

Quality checks Grounded in supplied results only; no estimated metrics.

If information is missing Flags low confidence below five supplied results.

Human review An SEO reviews before it changes a plan.

PROMPT 08 Competitor content analysis

RESEARCH

Required inputs Competitor names and their content you have gathered.

COPY-PASTE PROMPT

ROLE: You are a competitive content analyst.

OBJECTIVE: Identify where competitors are strong, where they are thin, and where a real opening exists.

CONTEXT:

Competitors: [LIST]

Their content: [PASTE URLs, titles, and what each covers]

Our content: [PASTE the same]

Our differentiators: [what is true of us and not of them]

CONSTRAINTS: Work only from supplied material. Do not estimate their traffic, authority, or performance. Separate what you observed from what you inferred, and label both. Do not recommend copying a competitor's approach.

PROCESS:

1. Map coverage by topic for each competitor and for us.
2. Identify topics everyone covers identically – the crowded ground.
3. Identify topics nobody covers well – the openings.
4. Identify where our differentiator lets us cover something they cannot.
5. Flag anything they do better that we should honestly acknowledge.

OUTPUT FORMAT: A coverage matrix, the crowded list, the opening list, and the honest acknowledgement list.

QUALITY CHECKS: Observation and inference are labelled separately. No estimated performance metrics. The acknowledgement list is not empty.

MISSING INFORMATION: Name the competitors whose content was not supplied and exclude them from conclusions.

HUMAN REVIEW: Strategy reviews before it changes the roadmap.

Output format	A coverage matrix with crowded ground, openings, and honest acknowledgements.
Quality checks	Observation separated from inference; no estimated competitor metrics.
If information is missing	Excludes competitors whose content was not supplied.
Human review	Strategy reviews before the roadmap changes.

PROMPT 09 Content gap analysis

CONTENT

Required inputs Your published content, the question bank, and the cluster definitions.

COPY-PASTE PROMPT

ROLE: You are auditing content coverage against known demand.

OBJECTIVE: Find the questions your audience asks that your site does not answer.

CONTEXT:

Our published content: [PASTE URLs and topics]

Question bank: [PASTE the questions with their sources]

Clusters: [PASTE cluster definitions]

Business priorities: [what matters commercially this quarter]

CONSTRAINTS: A question is only covered if a specific page answers it directly – not if a page mentions the topic. Do not count partial coverage as coverage. Rank gaps by commercial proximity, not by volume.

PROCESS:

1. Match each question to a page that answers it directly, or mark it a gap.
2. Mark partial coverage separately from no coverage.
3. Rank the gaps by commercial proximity and demand evidence.
4. Identify gaps better closed by expanding an existing page.
5. Identify pages that cover nothing anyone asked about.

OUTPUT FORMAT: A coverage table, a ranked gap list, an expand-instead list, and a list of content covering no known demand.

QUALITY CHECKS: Partial coverage is distinguished from full. Ranking follows commercial proximity. Expand-instead options come before new commissions.

MISSING INFORMATION: If the question bank is thin, say so – the analysis is only as good as the demand evidence behind it.

HUMAN REVIEW: The content owner decides what to build.

Output format	A coverage table, a ranked gap list, and expand-instead recommendations.
Quality checks	Partial coverage distinguished; ranked by commercial proximity.
If information is missing	Flags a thin question bank as limiting the analysis.
Human review	The content owner decides what gets built.

PROMPT 10 On-page optimization review

ON-PAGE

Required inputs The page content, the primary keyword, and the intent.

COPY-PASTE PROMPT

ROLE: You are reviewing a page against an on-page standard.

OBJECTIVE: Identify what would make this page clearer for readers, crawlers, and models – in priority order.

CONTEXT:
Page content: [PASTE the full text, headings, title, and meta]
Primary keyword: [KEYWORD] Intent: [INTENT]
Commercial destination: [where this page should send people]

CONSTRAINTS: Do not recommend keyword density targets or word count minimums. Do not suggest adding content that says nothing. Every recommendation must improve clarity for a reader first.

PROCESS:
1. Check intent match between page type and query.
2. Review title, meta, URL, H1, and heading hierarchy.
3. Check whether the direct answer appears in the first 100 words.
4. Assess extractability – could each key passage stand alone?
5. Check internal links, especially to the commercial destination.
6. Identify unsupported claims to remove or source.

OUTPUT FORMAT: A prioritized recommendation list with the specific change for each, then rewritten title and meta options.

QUALITY CHECKS: Every recommendation improves reader clarity. No density or word-count advice. Unsupported claims are flagged.

MISSING INFORMATION: If intent was not supplied, state the assumption you made and flag it.

HUMAN REVIEW: The editor approves changes to published pages.

Output format	A prioritized recommendation list plus rewritten title and meta options.
Quality checks	Reader clarity first; no density advice; unsupported claims flagged.
If information is missing	States and flags any assumption about intent.
Human review	The editor approves changes to published pages.

PROMPT 11 Metadata generation

ON-PAGE

Required inputs The page content and the primary keyword.

COPY-PASTE PROMPT

ROLE: You are writing titles and meta descriptions.

OBJECTIVE: Produce accurate, click-worthy metadata that describes the page.

CONTEXT:

Page content: [PASTE or summarize accurately]

Primary keyword: [KEYWORD]

What the reader gets: [the specific value of this page]

Brand name: [BRAND]

CONSTRAINTS: The title must describe what is actually on the page. Roughly 60 characters for titles, 155 for descriptions – approximate, since display varies. No clickbait, no false urgency, no promises the page does not keep. Do not front-load the brand name unless the query is branded.

PROCESS:

1. Identify the single most compelling accurate thing about the page.
2. Write five title options taking different angles.
3. Write three meta descriptions that expand on the title honestly.
4. Flag any option that overstates what the page delivers.

OUTPUT FORMAT: Five titles with character counts, three descriptions with character counts, and a recommendation with reasoning.

QUALITY CHECKS: Every option is accurate to the page. No clickbait. Character counts included.

MISSING INFORMATION: If the page content was summarized rather than pasted, note that accuracy could not be fully verified.

HUMAN REVIEW: Whoever owns the page confirms accuracy before publishing.

Output format Five titles and three descriptions with character counts and a recommendation.

Quality checks Accurate to the page; no clickbait; counts included.

If information is missing Notes when accuracy could not be verified from a summary.

Human review The page owner confirms accuracy before publishing.

PROMPT 12 Internal link suggestions

Required inputs The page, a list of candidate destinations, and their commercial importance.

COPY-PASTE PROMPT

ROLE: You are identifying internal linking opportunities.

OBJECTIVE: Suggest links a reader would genuinely want, prioritizing commercially important destinations.

CONTEXT:
Source page: [PASTE the full content]
Candidate destinations: [LIST URLs with their topics and commercial importance]
Existing internal links on this page: [LIST]

CONSTRAINTS: Only suggest a link where a reader at that point in the page would genuinely want it. Do not suggest a link purely because the destination matters commercially. Suggest a maximum of six new links. Do not suggest exact-match anchors repeatedly.

PROCESS:

1. Read the page and identify natural link moments.
2. Match each to the most relevant candidate destination.
3. Write a descriptive anchor for each.
4. Score each on target value, relevance, and reader benefit.
5. Discard anything scoring below 3 on relevance.

OUTPUT FORMAT: A table with the sentence to link from, the anchor text, the destination, the score, and the reason a reader would click.

QUALITY CHECKS: Each link has a reader-benefit rationale. Anchors vary. Fewer than six suggestions is an acceptable answer.

MISSING INFORMATION: If commercial importance was not supplied, rank by topical relevance and say so.

HUMAN REVIEW: The editor approves before implementation.

Output format	A link table with anchors, destinations, scores, and reader-benefit reasons.
Quality checks	Reader benefit required; anchors vary; low-relevance links discarded.
If information is missing	Ranks by relevance alone if commercial importance is missing.
Human review	The editor approves before implementation.

PROMPT 13 Technical issue explanation

TECHNICAL

Required inputs	The issue, the evidence, and the affected URLs.
COPY-PASTE PROMPT ROLE: You are explaining a technical SEO issue to a non-specialist stakeholder. OBJECTIVE: Explain what is wrong, why it matters commercially, and what should happen – without jargon and without exaggeration. CONTEXT: The issue: [DESCRIBE what you observed] Evidence: [PASTE crawl data, log sample, or screenshots described] Affected URLs: [count and examples] Commercial importance of those pages: [DESCRIBE] CONSTRAINTS: Do not estimate traffic loss or recovery. Do not claim certainty about cause where the evidence is correlational. Explain in plain language a non-technical stakeholder could repeat accurately. PROCESS: 1. State what is happening, factually. 2. State what should be happening instead. 3. Explain the commercial consequence in terms of the affected pages. 4. State the confidence level in the diagnosis and what would raise it. 5. Recommend the next action – a fix, or a diagnostic if confidence is low. OUTPUT FORMAT: A plain-language explanation under 200 words, then the technical detail separately for engineering. QUALITY CHECKS: No traffic estimates. Confidence stated explicitly. Plain language genuinely plain. MISSING INFORMATION: If evidence is thin, recommend a diagnostic rather than a fix, and say what to gather. HUMAN REVIEW: An engineer validates the diagnosis before a ticket is raised.	
Output format	A plain-language explanation plus separate technical detail.
Quality checks	No traffic estimates; confidence stated; language genuinely plain.
If information is missing	Recommends a diagnostic rather than a fix when evidence is thin.
Human review	An engineer validates before ticketing.

PROMPT 14 Dev ticket drafting

TECHNICAL

Required inputs The diagnosed issue, its evidence, and the priority score inputs.

COPY-PASTE PROMPT

ROLE: You are writing an engineering ticket for an SEO fix.

OBJECTIVE: Produce a ticket an engineer could pick up without asking questions.

CONTEXT:

Issue: [DESCRIBE] Evidence: [PASTE or reference attachments]

Affected URLs: [count, with a sample of ten]

Expected behaviour: [what should happen instead]

Priority inputs: impact [n], pages [n], business [n], confidence [n], effort [n]

CONSTRAINTS: Acceptance criteria must be testable by someone who is not an SEO. Do not specify implementation unless the method genuinely matters. State what is out of scope explicitly.

PROCESS:

1. Write the problem statement with evidence.
2. Write the expected behaviour.
3. Write three to five testable acceptance criteria.
4. Calculate the priority score from the inputs.
5. Define verification steps and a rollback plan.

OUTPUT FORMAT: The ticket in the dev-ticket template structure.

QUALITY CHECKS: Every acceptance criterion is testable. Out of scope is stated. Verification and rollback are defined.

MISSING INFORMATION: If any priority input is missing, leave the score uncalculated rather than assuming a value.

HUMAN REVIEW: Engineering confirms feasibility and sizing before scheduling.

Output format	A complete dev ticket with testable acceptance criteria.
Quality checks	Criteria testable by a non-SEO; scope, verification, and rollback defined.
If information is missing	Leaves the score uncalculated rather than assuming inputs.
Human review	Engineering confirms feasibility before scheduling.

PROMPT 15 Structured data suggestions

TECHNICAL

Required inputs The page content and the page type.

COPY-PASTE PROMPT

ROLE: You are recommending structured data for a page.

OBJECTIVE: Suggest markup that accurately describes what is visibly on the page.

CONTEXT:
Page type: [article / product / FAQ / how-to / local business / organization]
Page content: [PASTE, including anything visible to a reader]
Existing markup: [PASTE if any]

CONSTRAINTS: Only recommend markup for content that is visibly present on the page. Do not recommend FAQ markup for invented questions. Do not recommend review markup for reviews you do not genuinely have. Markup that describes something a reader cannot see is a policy violation and a credibility risk.

PROCESS:

1. Identify which schema types genuinely apply.
2. List the required and recommended properties for each.
3. Map each property to the visible content that supports it.
4. Flag any property that cannot be filled from visible content.
5. Note anything currently marked up that should be removed.

OUTPUT FORMAT: The recommended schema with each property mapped to its visible source, plus the removal list.

QUALITY CHECKS: Every property maps to visible content. No invented FAQs or reviews. Removals identified.

MISSING INFORMATION: Where a property cannot be supported, leave it out rather than suggesting placeholder content.

HUMAN REVIEW: Validate the output in a structured data testing tool before deploy.

Output format	Schema recommendations with every property mapped to visible content.
Quality checks	Nothing marked up that a reader cannot see; no invented FAQs or reviews.
If information is missing	Omits unsupportable properties rather than inventing content.
Human review	Validated in a testing tool before deployment.

PROMPT 16 Content refresh analysis

CONTENT

Required inputs The declining page, its performance history, and what now outranks it.

COPY-PASTE PROMPT

ROLE: You are diagnosing why a page has declined and what to do about it.

OBJECTIVE: Produce a specific refresh plan, or a recommendation to merge, redirect, or remove.

CONTEXT:

The page: [PASTE full content]

Performance history: [PASTE traffic, rankings, and dates]

What now outranks it: [PASTE the current top results and what they cover]

Business importance: [why this page matters, or does not]

CONSTRAINTS: Distinguish a content problem from a technical problem from a competitive problem – they need different responses. Do not recommend a rewrite where the issue is intent mismatch or cannibalization.

PROCESS:

1. Identify when the decline started and what changed around then.
2. Compare our coverage with what now outranks us.
3. Check for intent drift – has the query's meaning changed?
4. Check for cannibalization against our own pages.
5. Recommend one action: update, merge, redirect, remove, or investigate.

OUTPUT FORMAT: The diagnosis with evidence, the single recommended action, and – if updating – the specific changes in priority order.

QUALITY CHECKS: The diagnosis distinguishes problem types. One recommendation, not a menu. Changes are specific rather than "improve the content".

MISSING INFORMATION: Without performance history, say the diagnosis is speculative and recommend gathering it first.

HUMAN REVIEW: The content owner approves before republishing.

Output format	A diagnosis, one recommended action, and specific prioritized changes.
Quality checks	Problem types distinguished; one recommendation; changes specific.
If information is missing	Declares the diagnosis speculative without performance history.
Human review	The content owner approves before republishing.

PROMPT 17 Local SEO audit

LOCAL

Required inputs The Business Profile details, listings, and the location page.

COPY-PASTE PROMPT

ROLE: You are auditing local visibility for one location.

OBJECTIVE: Find inconsistencies and gaps that are limiting local visibility.

CONTEXT:

Canonical NAP: [the single source of truth – name, address, phone]

Business Profile: [PASTE current details, categories, hours, services]

Listings found: [PASTE each with the NAP exactly as it appears there]

Location page: [PASTE content]

Local competitors: [LIST]

CONSTRAINTS: Compare NAP character by character – "St" and "Street" are different. Do not estimate map pack positions or local traffic. Flag inconsistency rather than guessing which version is correct.

PROCESS:

1. Compare every listing NAP against the canonical record.
2. Check profile completeness and category accuracy.
3. Review the location page for unique, genuinely local content.
4. Check schema against the profile.
5. Identify missing citations relative to competitors.

OUTPUT FORMAT: An inconsistency table with the exact difference and the source, a completeness checklist result, and a prioritized fix list.

QUALITY CHECKS: Differences are quoted exactly. No estimated local metrics. Fixes are ordered by likely effect.

MISSING INFORMATION: List the directories not supplied and exclude them.

HUMAN REVIEW: The location owner confirms the canonical NAP before corrections.

Output format	An inconsistency table, a completeness result, and a prioritized fix list.
Quality checks	Character-level NAP comparison; no estimated local metrics.
If information is missing	Excludes directories that were not supplied.
Human review	The location owner confirms the canonical NAP first.

PROMPT 18 Digital PR angle development

AUTHORITY

Required inputs Your data or expertise, and the publications you want to reach.

COPY-PASTE PROMPT

ROLE: You are developing digital PR angles from an asset a company holds.

OBJECTIVE: Find the stories a journalist would actually want, from what we genuinely have.

CONTEXT:

What we have: [first-party data, expertise, unusual perspective, access]

What we can publish: [what is permitted – check before offering]

Target publications: [type and examples]

Our spokesperson: [named, with credentials]

CONSTRAINTS: Do not invent findings, statistics, or data points. Do not propose angles requiring data we do not hold. Every angle must be true and supportable. State the methodology limits before pitching, not after.

PROCESS:

1. Identify what is genuinely newsworthy in what we hold.
2. For each, write the headline a journalist would write.
3. Identify the timing hook that makes it relevant now.
4. State the methodology and its limits.
5. Flag any angle where the data cannot support the headline.

OUTPUT FORMAT: Angles ranked by newsworthiness, each with the headline, the hook, the supporting data, and the stated limits.

QUALITY CHECKS: Nothing invented. Every headline is supported by the data. Limits stated explicitly.

MISSING INFORMATION: If the data cannot support a genuine story, say so – a weak angle wastes relationships that take years to build.

HUMAN REVIEW: Legal or data owner confirms what may be published.

Output format	Ranked angles with headlines, hooks, supporting data, and stated limits.
Quality checks	Nothing invented; every headline supported; limits stated.
If information is missing	Says plainly when the data cannot support a story.
Human review	The data owner confirms what may be published.

PROMPT 19 Backlink opportunity analysis

AUTHORITY

Required inputs Your linkable asset and a list of candidate sites you have researched.

COPY-PASTE PROMPT

ROLE: You are qualifying backlink opportunities.

OBJECTIVE: Rank opportunities by relevance and realistic likelihood, and discard the ones not worth pursuing.

CONTEXT:

Our asset: [what we are offering and why it is useful]

Candidate sites: [PASTE with what each publishes]

Our topic and audience: [DESCRIBE]

CONSTRAINTS: Do not recommend paid links, exchanges, or any arrangement that would need to be hidden. Do not estimate domain authority – use whatever metric I supplied and name the tool. Relevance outranks metrics.

PROCESS:

1. Assess topical relevance for each candidate.
2. Assess whether our asset genuinely improves their page.
3. Identify the specific reason each would consider it.
4. Rank by relevance first, then by likelihood.
5. Discard candidates where the ask would be unreasonable.

OUTPUT FORMAT: A ranked table with relevance, the specific reason, the suggested approach, and the discard list with reasons.

QUALITY CHECKS: Relevance leads. Every entry has a specific reason. Nothing recommended that requires concealment.

MISSING INFORMATION: If what a site publishes was not supplied, exclude it rather than assuming.

HUMAN REVIEW: The outreach owner approves the target list.

Output format	A ranked opportunity table with specific reasons and a discard list.
Quality checks	Relevance leads; specific reasons required; nothing concealable.
If information is missing	Excludes candidates whose content was not supplied.
Human review	The outreach owner approves the list.

PROMPT 20 AI visibility audit

AI SEO

Required inputs Your prompt universe and the answers you captured, verbatim.

COPY-PASTE PROMPT

ROLE: You are analysing AI assistant answers for brand visibility.

OBJECTIVE: Determine where we appear, how we are described, and where the gaps are.

CONTEXT:

Our brand and category: [DESCRIBE]

Prompt universe and captured answers: [PASTE each prompt, the assistant used, the date, and the answer verbatim]

Our competitors: [LIST]

Our brand fact sheet: [PASTE]

CONSTRAINTS: Analyse only the answers I captured. Do not speculate about why a system produced an answer – describe what it said. Flag every factual error against the fact sheet as the highest-priority finding.

PROCESS:

1. Record mention and citation separately for each prompt.
2. Record which competitors appeared and how often.
3. Compare every statement about us against the fact sheet.
4. Identify which of our pages were cited, and what they have in common.
5. Identify the sources answers drew on where they are named.

OUTPUT FORMAT: A per-prompt table, the five scorecard measures, the factual error list, and the source list with citation gaps.

QUALITY CHECKS: Mention and citation counted separately. Errors listed first. No speculation about system behaviour.

MISSING INFORMATION: If prompt wording varied between runs, note that the comparison is unreliable.

HUMAN REVIEW: Factual errors are escalated for correction at source.

Output format	A per-prompt table, five scorecard measures, an error list, and source gaps.
Quality checks	Mention and citation separated; errors prioritized; no speculation.
If information is missing	Flags unreliable comparison if prompt wording drifted.
Human review	Factual errors escalated for correction at source.

PROMPT 21 Brand fact extraction

AI SEO

Required inputs	Your site, profiles, and listings as they currently read.
COPY-PASTE PROMPT ROLE: You are auditing brand fact consistency across public sources. OBJECTIVE: Find every place our own facts contradict each other. CONTEXT: Canonical fact sheet: [PASTE, or state that none exists] Sources: [PASTE how the brand is described on the homepage, about page, LinkedIn, directories, review platforms, and any coverage] CONSTRAINTS: Compare exact wording. Note differences in category description, founding year, location, product naming, and what the company does. Do not decide which version is correct – surface the conflict for a human. PROCESS: 1. Extract each stated fact from each source. 2. Group by fact type and compare. 3. Flag every inconsistency with both versions quoted. 4. Identify facts stated in some places and absent in others. 5. Note anything outdated relative to current positioning. OUTPUT FORMAT: A conflict table with the fact, each version, and its source, then a completeness gap list. QUALITY CHECKS: Both versions quoted exactly. No version chosen for me. Outdated descriptions identified. MISSING INFORMATION: List sources not supplied – the audit is only as complete as the inputs. HUMAN REVIEW: A named owner decides the canonical version and corrects at source.	
Output format	A conflict table with quoted versions and sources, plus a gap list.
Quality checks	Exact quotes; no version chosen automatically; outdated entries flagged.
If information is missing	Lists unsupplied sources as limits on completeness.
Human review	A named owner decides the canonical version.

PROMPT 22 AI citation readiness review

AI SEO

Required inputs The page content, and the fact sheet it must be consistent with.

COPY-PASTE PROMPT

ROLE: You are assessing whether a page could be safely quoted by an AI system.

OBJECTIVE: Identify what prevents this page from being extracted and cited accurately.

CONTEXT:
Page content: [PASTE in full]
Brand fact sheet: [PASTE]
The question this page should answer: [QUESTION]

CONSTRAINTS: Test each key passage in isolation. A passage that requires surrounding context to be accurate is not citable. Do not recommend adding summary boxes that duplicate content without adding clarity.

PROCESS:
1. Identify the passage that answers the question directly.
2. Test whether it makes sense and stays accurate quoted alone.
3. Check every fact against the fact sheet.
4. Check numbers carry unit, period, and source.
5. Identify unsupported claims and superlatives to remove.

OUTPUT FORMAT: A pass or fail verdict per key passage, the specific rewrite for each failure, and the list of claims to remove or source.

QUALITY CHECKS: Passages tested in isolation. Verdicts binary. Rewrites specific rather than directional.

MISSING INFORMATION: If no fact sheet exists, flag that consistency cannot be verified.

HUMAN REVIEW: The page owner approves rewrites before publishing.

Output format	Pass or fail per passage, specific rewrites, and a claims-to-remove list.
Quality checks	Passages tested alone; binary verdicts; rewrites specific.
If information is missing	Flags that consistency is unverifiable without a fact sheet.
Human review	The page owner approves rewrites.

PROMPT 23 Prompt universe creation

AI SEO

Required inputs Your ICP, the decisions they make, and your question bank.**COPY-PASTE PROMPT****ROLE:** You are building a prompt set for tracking AI visibility.**OBJECTIVE:** Produce a fixed, re-runnable set of prompts covering the decisions where we should be a candidate.**CONTEXT:**

Our brand and category: [DESCRIBE]

Our ICP: [DESCRIBE]

Decisions they make where we should appear: [LIST]

Question bank: [PASTE customer questions in their own words]

Competitors: [LIST]

CONSTRAINTS: Prompts must be phrased as a real person would ask, not as keywords. Cap the set at 40 unless I state otherwise – a set that is too large will not be re-run. Cover category, brand, competitor, comparison, and problem prompts.**PROCESS:**

1. Write category prompts where recommendations happen.
2. Write brand prompts testing understanding and accuracy.
3. Write competitor and alternative prompts.
4. Write problem prompts in customer language.
5. Label each prompt by type and note what a good result looks like.

OUTPUT FORMAT: The prompt set in a table with type, the decision it tests, and the definition of a good result.**QUALITY CHECKS:** Phrasing is conversational. All five types are represented. The set is small enough to re-run monthly.**MISSING INFORMATION:** If the question bank was not supplied, note that problem prompts are inferred and lower quality.**HUMAN REVIEW:** Fix the wording before the first run – it must never change.**Output format** A capped, typed prompt set with a good-result definition for each.**Quality checks** Conversational phrasing; all five types present; re-runnable size.**If information is missing** Flags problem prompts as inferred without a question bank.**Human review** Wording is fixed before the first run and never changed.

PROMPT 24 Competitor recommendation analysis

AI SEO

Required inputs Captured answers to recommendation prompts across assistants.

COPY-PASTE PROMPT

ROLE: You are analysing why competitors are recommended by AI systems and we are not.

OBJECTIVE: Identify the observable differences between recommended brands and ours.

CONTEXT:
Recommendation prompts and answers: [PASTE verbatim, with assistant and date]
Competitors named: [LIST]
What we know about their public presence: [PASTE – coverage, directories, reviews, published content]
Our equivalent presence: [PASTE]

CONSTRAINTS: Describe observable differences only. Do not speculate about ranking factors or system internals. Do not conclude that any single factor caused a recommendation.

PROCESS:
1. List which brands are named and how often across the set.
2. Record the reasons the answers give, where they give any.
3. Compare public presence – coverage, listings, reviews, content depth.
4. Identify differences we could legitimately close.
5. Rank those by effort and by how often they appear across answers.

OUTPUT FORMAT: A presence comparison table, the stated reasons where given, and a ranked list of closable gaps.

QUALITY CHECKS: Observable differences only. No causal claims. Gaps are specific and actionable.

MISSING INFORMATION: If presence data for a competitor was not supplied, exclude them from the comparison.

HUMAN REVIEW: Strategy decides which gaps to close.

Output format	A presence comparison, the stated reasons, and a ranked closable-gap list.
Quality checks	Observable differences only; no causal claims; gaps specific.
If information is missing	Excludes competitors with no supplied presence data.
Human review	Strategy decides which gaps to close.

PROMPT 25 Weekly SEO report

REPORTING

Required inputs This week's data with sources and dates, plus last week's report.

COPY-PASTE PROMPT

ROLE: You are writing the weekly SEO report for a business audience.

OBJECTIVE: Report what changed, what it means, and what we are deciding – not what we did.

CONTEXT:

This week's data: [PASTE traffic, rankings, conversions, technical status, AI visibility results – each with its source and as-of date]

Last week's report: [PASTE]

This quarter's goals: [STATE]

CONSTRAINTS: Every figure needs a source and an as-of date. Do not forecast. Do not explain a change with a cause the data cannot support – list hypotheses instead. Report quality alongside volume. Do not compare to industry benchmarks.

PROCESS:

1. Report progress against each quarterly goal.
2. Report what moved versus last week, with sample context.
3. Separate what the data shows from what we are inferring.
4. State the three decisions required this week.
5. Note what should change in the system as a result.

OUTPUT FORMAT: The report under 500 words, a metric table with week-on-week change, three decisions with owners, and the system changes implied.

QUALITY CHECKS: Every figure sourced and dated. No forecasts. The report ends in decisions rather than observations.

MISSING INFORMATION: Where a figure is unavailable, mark it unavailable rather than omitting the row silently.

HUMAN REVIEW: The SEO owner assigns the three decisions.

Output format A report under 500 words with a change table, three decisions, and system changes.

Quality checks Sourced and dated figures; no forecasts; ends in decisions.

If information is missing Marks unavailable figures explicitly rather than dropping rows.

Human review The SEO owner assigns the decisions.

SEO templates and SOPs

Sixteen standard operating procedures. Each has an owner, a trigger, a quality gate, and an escalation rule — because a procedure without those is a suggestion.

Build these as a single SOP database in Notion, one entry per procedure, with the fields below as properties. Filter by owner for a personal view and by cadence for a calendar view. The templates these procedures reference live in the sections where they were introduced — the one-page strategy in Section 01, the content brief in Section 04, the dev ticket in Section 06, the local landing page in Section 09, and the brand fact sheet in Section 10.

SOP 01

Quarterly SEO strategy review

SOP

Purpose	Keep the strategy connected to the business, and versioned.
Owner	SEO lead
Trigger	Two weeks before each quarter starts
Inputs	Last quarter's reports, maturity assessment, performance data, business goals
STEPS	
1	Re-run the maturity assessment and note the lowest dimension.
2	Review last quarter's three bets against their success signals.
3	Reconstruct where effort actually went versus where the strategy said it should.
4	Revise the strategy one-pager, including the not-doing list.
5	Set three bets with a measurable signal and an owner each.
6	Version the document with a date and circulate it.
Output	A versioned one-page strategy with three owned bets
Quality gate	The gap between stated and actual effort is named explicitly
Escalation	If the lowest maturity dimension has not moved in two quarters, escalate to the business owner
Review cadence	Quarterly

SOP 02

Weekly performance review

SOP

Purpose	Notice what changed while it is still cheap to respond to.
Owner	SEO lead
Trigger	Every Monday morning
Inputs	Search Console, analytics, rank tracking, technical alerts, AI visibility log
STEPS	
1	Check technical alerts first — a crawl or index problem outranks everything else.
2	Review rankings and clicks against the previous four weeks, not the previous week.
3	Check conversions from organic against the quarter's goal.
4	Scan the At Risk and Waiting For Review dashboard views.
5	Pick three priorities for the week and assign each one.
Output	Three assigned priorities in the Opportunities database
Quality gate	Every priority names an owner and a day it will happen
Escalation	Any sitewide indexing or traffic anomaly escalates immediately, outside this cadence
Review cadence	Weekly

SOP 03 Keyword research and scoring

SOP

Purpose	Maintain one prioritized keyword list everyone works from.
Owner	SEO specialist
Trigger	Quarterly, or when entering a new topic or market
Inputs	Demand databases, keyword tool exports, current rankings, business priorities
STEPS	
1	Start from the Customer Problem and Question Bank databases, not a tool export.
2	Expand with tool data and label which keywords came from where.
3	Classify intent and stage, using SERP evidence where available.
4	Score against the ten criteria and apply the weights.
5	Map each keyword to exactly one owning page, or mark it a gap.
6	Run the cannibalization audit on anything already published.
Output	An updated Keyword Master Database with scores and mapping
Quality gate	No keyword maps to two pages; every score has its criteria recorded
Escalation	If more than a quarter of the list is unmappable, the problem is strategy, not keywords
Review cadence	Quarterly

SOP 04 Content brief and commission

SOP

Purpose	Ensure nothing is written without a defined purpose and angle.
Owner	Content lead
Trigger	When a topic reaches the top of the build order
Inputs	Keyword, intent, cluster, our evidence, current top results, commercial destination
STEPS	
1	Confirm the topic is not better served by updating an existing page.
2	Run the SERP analysis to establish what the query rewards.
3	Define the angle — what we say that current results do not.
4	Write the brief, including extractable passages and internal links.
5	Name the author and the reviewer, and set the due date.
Output	A complete brief in the Content database, status = brief
Quality gate	The angle is genuinely differentiated and the commercial link is specified
Escalation	If no differentiating angle exists, decline the commission and record why
Review cadence	Per piece

SOP 05 Article QA before publication

SOP

Purpose	Nothing goes live that would embarrass you if it were quoted.
Owner	Editor
Trigger	When a draft reaches review status
Inputs	The draft, the brief, the on-page checklist, the extractability checklist

STEPS

- 1 Check the draft answers the brief's required coverage in order.
- 2 Run the on-page checklist and fix every failure.
- 3 Run the extractability checklist on each key passage in isolation.
- 4 Verify every claim has a source, or remove it.
- 5 Confirm internal links in and out, including the commercial destination.
- 6 Confirm the author is named with a relevant credential.

Output	A publish-or-return verdict with the specific fixes required
Quality gate	Unsourced claims are an automatic return, not a note for later
Escalation	Anything involving legal, medical, or financial claims goes to a qualified reviewer
Review cadence	Per piece

SOP 06 Monthly technical audit

SOP

Purpose	Catch technical decay before it becomes a traffic event.
Owner	Technical SEO owner
Trigger	First week of each month, and after template releases
Inputs	Crawler, Search Console, log files, the previous audit

STEPS

- 1 Run a full crawl and compare against the previous month.
- 2 Check index coverage and investigate any change over five percent.
- 3 Review Core Web Vitals field data on key templates.
- 4 Check orphan pages, redirect chains, and broken links.
- 5 Score every new issue with the prioritization formula.
- 6 Create tickets for anything scoring above your threshold.

Output	Scored issues in the Technical Issues database, with tickets raised
Quality gate	Every issue has evidence attached and confidence above 3, or a diagnostic instead
Escalation	Critical severity issues bypass the cadence and go straight to engineering
Review cadence	Monthly

SOP 07 Technical issue to shipped fix

SOP

Purpose	Make sure findings become deployed changes rather than a backlog.
Owner	Technical SEO owner with an engineering owner
Trigger	When an issue is ticketed
Inputs	The scored issue, its evidence, the affected URL list

STEPS

- 1 Write the ticket using the dev-ticket template.
- 2 Confirm acceptance criteria are testable by a non-SEO.
- 3 Agree the sprint and the engineering owner.
- 4 Verify on staging before deploy where possible.
- 5 Verify in production within 48 hours of deploy.
- 6 Record the verified date and close the issue.

Output	A verified fix with the date recorded
Quality gate	The issue is not closed until verified in production
Escalation	If a fix is deferred twice, escalate with the business impact restated
Review cadence	Per issue

SOP 08 Internal linking pass

SOP

Purpose	Direct authority and readers towards commercially important pages.
Owner	SEO specialist
Trigger	Monthly, and after any batch of publishing
Inputs	Content database, internal link counts, commercial importance ratings

STEPS

- 1 Export every URL with its internal inbound link count.
- 2 Sort by commercial importance descending, link count ascending.
- 3 For the top twenty, identify relevant source pages.
- 4 Score each opportunity and discard anything below 3 on relevance.
- 5 Write descriptive anchors and implement in priority order.
- 6 Verify implementation and update the Internal Links database.

Output	Implemented links recorded with anchors and destinations
Quality gate	Every implemented link has a reader-benefit rationale recorded
Escalation	If a money page still has no internal links after two passes, raise it in the weekly review
Review cadence	Monthly

SOP 09 Content decay review

SOP

Purpose	Recover declining pages before writing new ones.
Owner	Content lead
Trigger	Monthly
Inputs	Performance history by URL, current rankings, the Content database
STEPS	
1	Identify every page whose traffic or rankings declined over 90 days.
2	For each, check what now outranks it and what changed.
3	Distinguish content, technical, intent, and cannibalization causes.
4	Decide: update, merge, redirect, remove, or investigate.
5	Add updates to the build order ahead of new commissions.
Output	A refresh queue with one decided action per declining page
Quality gate	Cause type is identified before an action is chosen
Escalation	A sitewide decline pattern is a technical or algorithmic question, not a content one
Review cadence	Monthly

SOP 10 Digital PR campaign

SOP

Purpose	Earn coverage and citations from something you genuinely have.
Owner	Marketing or PR owner
Trigger	When a linkable asset is ready
Inputs	The asset, its methodology, permissions, target publications, spokesperson
STEPS	
1	Confirm what may be published, with the data owner.
2	Develop angles and select the strongest supportable one.
3	Publish the asset with its methodology visible.
4	Build the target list by beat, not by domain metric.
5	Run outreach — three touches maximum per contact.
6	Log every outcome, including the refusals and the silence.
Output	Coverage, citations, and mentions logged in the Backlinks database
Quality gate	No claim in the pitch exceeds what the data supports
Escalation	Any request to pay for placement is declined and recorded
Review cadence	Per campaign

SOP 11 Backlink opportunity outreach

SOP

Purpose	Pursue links without doing anything you would have to hide.
Owner	SEO or outreach owner
Trigger	When an asset and a target list exist
Inputs	The asset, qualified targets, the outreach templates
STEPS	
1	Qualify each target on topical relevance first.
2	Confirm the asset genuinely improves their page.
3	Personalize each message with a specific reason.
4	Send, then follow up once, then stop.
5	Log the outcome regardless of what it was.
Output	Outcomes recorded against every contacted target
Quality gate	Nothing sent that would require concealment if published
Escalation	Decline any exchange, payment, or undisclosed arrangement, and record the decline
Review cadence	Per campaign

SOP 12 Monthly local SEO audit

SOP

Purpose	Keep every location consistent everywhere it appears.
Owner	Local SEO owner
Trigger	Monthly, per location
Inputs	Canonical NAP record, Business Profile, citation list, location page
STEPS	
1	Compare every listing NAP against the canonical record, character by character.
2	Check profile completeness, categories, hours, and services.
3	Review new reviews and respond within the agreed window.
4	Check the location page for unique local content and correct schema.
5	Identify missing citations relative to local competitors.
Output	An inconsistency list with corrections assigned
Quality gate	Every difference is quoted exactly, with its source
Escalation	A NAP change anywhere triggers an immediate full-consistency pass
Review cadence	Monthly

SOP 13 Monthly AI visibility run

SOP

Purpose	Track AI visibility as a trend rather than an impression.
Owner	SEO lead
Trigger	First week of each month
Inputs	The fixed prompt universe, the brand fact sheet, the AI Visibility database
STEPS	
1	Run every prompt with identical wording to previous months.
2	Record mention, citation, page cited, competitors named, and the verbatim passage.
3	Compare every statement about the brand against the fact sheet.
4	Log factual errors as the highest-priority findings.
5	Calculate the five scorecard measures and add to the trend.
6	Convert citation gaps into entries in the Opportunities database.
Output	A dated set of results and five updated scorecard measures
Quality gate	Prompt wording is unchanged from previous runs
Escalation	A factual error about the brand triggers immediate correction at source
Review cadence	Monthly

SOP 14 Brand fact consistency correction

SOP

Purpose	Keep the facts about your brand identical everywhere.
Owner	Marketing owner
Trigger	Quarterly, or whenever an inconsistency is found
Inputs	The canonical fact sheet and every public source describing the brand
STEPS	
1	Audit every public source against the fact sheet.
2	Log each inconsistency with both versions quoted.
3	Decide the canonical version where the fact sheet is silent.
4	Correct owned properties first, then request third-party corrections.
5	Update the fact sheet and version it.
Output	A corrected source list and an updated, versioned fact sheet
Quality gate	One canonical wording per fact, with no third version introduced
Escalation	If a third party declines to correct a material error, record it and note it in the AI visibility log
Review cadence	Quarterly

SOP 15 Monthly SEO report

SOP

Purpose	Report what changed and what is being decided, not what was done.
Owner	SEO lead
Trigger	First week of each month
Inputs	All metrics with sources and as-of dates, last month's report, quarterly goals

STEPS

- 1 Assemble every figure with its source and as-of date.
- 2 Compare against last month and against the quarterly goal.
- 3 Separate what the data shows from what is being inferred.
- 4 State the three decisions the month requires.
- 5 Record what changes in the system as a result.

Output	A report under 500 words with three owned decisions
Quality gate	Every figure sourced and dated; nothing forecast
Escalation	A decision with no owner is not a decision — reassign before the report is filed
Review cadence	Monthly

SOP 16 Experiment design and evaluation

SOP

Purpose	Make change attributable rather than coincidental.
Owner	SEO lead
Trigger	When the team disagrees about what would work
Inputs	The hypothesis, the baseline, the available volume

STEPS

- 1 Record the baseline with its source and date before changing anything.
- 2 Change exactly one variable and record what is held constant.
- 3 Set the minimum sample and the evaluation date in advance.
- 4 Write the decision rule before the experiment starts.
- 5 Evaluate on the date, apply the rule, record the outcome either way.

Output	A completed experiment record with an adopt, reject, or extend decision
Quality gate	One variable only; the decision rule predates the result
Escalation	If volume cannot support a conclusion, run structured interviews instead and say so
Review cadence	Per experiment

SEO scorecards

Nine scorecards using transparent criteria. Each is scored 0, 1, or 2 — coarse deliberately, because false precision in SEO scoring is worse than none at all.

HOW TO READ EVERY SCORECARD IN THIS SECTION

These are diagnostic instruments, not measurements. They tell you where to look and make judgement consistent between people and across quarters. They are not comparable to anyone else's score, they do not predict performance, and the criteria are judgement expressed as a scale. Score honestly — a flattering score is a wasted afternoon.

Website SEO health

Ten criteria covering the fundamentals. Run quarterly.

CRITERION	WHAT A 0 LOOKS LIKE	WHAT A 2 LOOKS LIKE
Indexation matches intent	Index count unknown or unexplained	Indexed pages match what we intend, and we can explain the difference
Crawl efficiency	No log file review has ever been done	Crawl budget goes to important pages, verified in logs
Speed on key templates	No field data collected	Field data collected and key templates within targets
Mobile experience	Never tested on a real device	Verified on real devices, no usability errors
Architecture depth	Money pages more than four clicks deep	Every important page within three clicks
Duplicate handling	Duplicates and parameters uncontrolled	Canonicals deliberate, parameters documented, facets controlled
Structured data	Absent or invalid	Present, valid, and matching visible content
Internal linking	Money pages under-linked	Commercial pages are among the best-linked pages
Content freshness	No refresh process exists	Every page reviewed at least twice yearly
Measurement	Cannot connect SEO to a business metric	At least one SEO metric maps to a business objective
SCORE	READ IT AS	NEXT ACTION
0-6	Foundational gaps	Fix these before anything else can compound
7-12	Functional, with real weaknesses	Pick the two lowest criteria this quarter
13-16	Solid	Maintain, and move to the next scorecard on your list
17-20	Strong	Protect it — review quarterly rather than rebuilding

Content quality

Ten criteria applied to a representative sample of ten pages.

CRITERION	WHAT A 0 LOOKS LIKE	WHAT A 2 LOOKS LIKE
Intent match	Page type conflicts with the query	Page type matches what the query rewards
Answer position	The answer arrives after several paragraphs	The direct answer is in the first 100 words
Original substance	Summarizes what already ranks	Contains evidence or experience nobody else has
Author credibility	No author, or a generic byline	Named author with a relevant, verifiable credential
Claims sourced	Statistics with no origin	Every non-obvious claim links to its source
Extractability	Passages need context to make sense	Key passages stand alone accurately
Structure	Headings used for styling	A hierarchy a reader could outline from
Conversion path	No next step, or four competing ones	One clear next step matched to the stage
Internal links	No links, or generic anchors	Descriptive anchors pointing at relevant commercial pages
Freshness	Undated, or visibly out of date	Published and updated dates visible and current
SCORE	READ IT AS	NEXT ACTION
0–6	Foundational gaps	Fix these before anything else can compound
7–12	Functional, with real weaknesses	Pick the two lowest criteria this quarter
13–16	Solid	Maintain, and move to the next scorecard on your list
17–20	Strong	Protect it — review quarterly rather than rebuilding

AI visibility

Ten criteria. Run after your first two monthly AI visibility passes, not before.

CRITERION	WHAT A 0 LOOKS LIKE	WHAT A 2 LOOKS LIKE
Entity consistency	Facts differ between our own sources	One canonical version of every fact, everywhere
Third-party corroboration	Only we describe us	Independent sources describe us consistently
Prompt universe	No defined prompt set exists	A fixed set, re-run monthly with unchanged wording
Mention rate trend	Not measured	Measured monthly with a visible trend
Citation rate	Never cited with attribution	Cited with attribution on a meaningful share of prompts
Answer accuracy	Answers about us contain errors	Answers about us are factually correct
Extractable content	Nothing structured for quoting	Key pages contain self-contained, quotable passages
Source presence	Absent from the sources answers draw on	Present in the sources these systems actually use
Category coverage	Absent from category and comparison prompts	Appear where we are a legitimate answer
Correction process	No process when an error appears	A named owner corrects errors at source within a defined window
SCORE	READ IT AS	NEXT ACTION
0–6	Foundational gaps	Fix these before anything else can compound
7–12	Functional, with real weaknesses	Pick the two lowest criteria this quarter
13–16	Solid	Maintain, and move to the next scorecard on your list
17–20	Strong	Protect it — review quarterly rather than rebuilding

The remaining six scorecards

Same 0-2 structure, same bands. Build them as one scorecard database in Notion with a type property.

SCORECARD	THE TEN CRITERIA, IN SHORT
Topic authority	Cluster completeness · pillar strength · supporting depth · interlinking within the cluster · coverage of the question bank · original evidence · competitor comparison · rankings across the topic rather than one query · citations earned on the topic · refresh currency
Technical health	Crawlability · indexability · render completeness · Core Web Vitals field data · mobile usability · redirect hygiene · status code correctness · duplicate control · structured data validity · monitoring and alerting
Internal linking	Money page link counts · orphan pages · click depth · anchor descriptiveness · blog-to-commercial links · cluster interlinking · link decay monitoring · navigation stability · breadcrumb presence · a tracked opportunity backlog
Authority	Linkable assets held · original data published · editorial coverage earned · unlinked mentions converted · directory and profile completeness · expert visibility · community presence · review platform standing · citation authority · outreach hygiene
Local visibility	Profile completeness · category accuracy · NAP consistency · citation coverage · review volume and recency · review response rate · location page uniqueness · local schema · locally relevant content · multi-location consistency
Brand recommendation readiness	Findability across surfaces · entity clarity · fact consistency · third-party validation · review presence · comparison content published · pricing transparency · extractable answers · category association · absence of unsupported claims

SCORE TWO, NOT NINE

Running all nine in one sitting produces a nine-page document nobody acts on. Run Website SEO Health first, because it gates everything else. Then run the one covering whichever area your maturity assessment flagged as weakest. Return to the others next quarter.

SECTION 16 · ROADMAP

The 30-day implementation roadmap

Four weeks to a working system. This is deliberately smaller than you want it to be, because the version you finish beats the version you abandon in week three.

WEEK 1 Audit, goals, baseline, workspace

WEEK 1

By Friday you know where you stand and you have somewhere to put things.

DAY	WHAT HAPPENS
Day 1	Run the maturity assessment. Write down the lowest dimension.
Day 2	Full technical crawl. Record index coverage, orphan pages, and critical issues.
Day 3	Baseline everything — rankings, traffic, conversions, with dates. This is what you will compare against for the rest of the year.
Day 4	Build the Keywords and Content databases in Notion. Import what already exists.
Day 5	Write the one-page strategy, including the not-doing list. Version it.

QUALITY GATE

A maturity score, a dated baseline, two live databases, and a versioned strategy.

WEEK 2 Demand, keywords, content mapping, technical backlog

WEEK 2

By Friday you know what to work on and in what order.

DAY	WHAT HAPPENS
Day 6	Mine customer language from calls, reviews, and tickets. Build the Question Bank.
Day 7	Keyword research from the demand evidence, not from a tool export.
Day 8	Score and prioritize the keyword list. Map each keyword to one owning page.
Day 9	Run the cannibalization audit and the content gap analysis.
Day 10	Score every technical issue and raise tickets for anything above threshold.

QUALITY GATE

A scored keyword list, a mapped content plan, and a prioritized technical backlog.

WEEK 3 Optimization, internal linking, authority, AI visibility

WEEK 3

By Friday the first improvements have shipped and AI tracking has begun.

DAY	WHAT HAPPENS
Day 11	Internal linking pass. Start with money pages that have the fewest links.
Day 12	On-page and extractability pass on your five most commercially important pages.
Day 13	Build the brand fact sheet. Audit your own sources against it and correct what you own.
Day 14	Build the prompt universe. Cap it at 40 prompts and fix the wording.
Day 15	Run the first AI visibility pass. Record everything, including the errors.

QUALITY GATE

Links implemented, key pages optimized, a fact sheet, and the first AI visibility data point.

WEEK 4 Reporting, learning, prioritization, next quarter

WEEK 4

By Friday the system runs without you having to remember it.

DAY	WHAT HAPPENS
Day 16	Build the Reporting and Opportunities databases. Move everything loose into them.
Day 17	Write the first weekly report against the week-one baseline.
Day 18	Set up the SOP database. Add the six SOPs you will actually run.
Day 19	Schedule the weekly rhythm in a calendar, with owners attached.
Day 20	Write the next-quarter roadmap from what the month revealed.

QUALITY GATE

A live reporting loop, six running SOPs, a scheduled rhythm, and a next-quarter plan.

IF YOU HAVE LESS THAN A DAY A WEEK

Do week one in full — it takes about six hours and it is the only part where skipping steps invalidates everything after it. Then run week two across a fortnight and stop. A baseline, a scored keyword list, and a mapped content plan is already a real system. Weeks three and four can wait for the following quarter.

SECTION 17 · VERIFICATION

Final readiness checklist

Twelve areas. If you can tick most of these, you have an SEO operating system rather than a list of tactics.

Strategy

- ☐ A one-page strategy exists, versioned, with a not-doing list
- ☐ Three bets have measurable signals and named owners
- ☐ The maturity assessment has been run and the weakest dimension named

Content

- ☐ Every page has a stated purpose and a commercial destination
- ☐ A brief standard exists and is used before commissioning
- ☐ A decay review runs monthly and produces refresh decisions

Internal linking

- ☐ Money pages are among the best-linked pages on the site
- ☐ Orphan pages are identified and resolved
- ☐ A link opportunity backlog is tracked and worked

Local SEO

- ☐ A canonical NAP record exists per location
- ☐ Listings are audited monthly against it
- ☐ Reviews are responded to within a defined window

Measurement

- ☐ A baseline exists, with dates, for everything you track
- ☐ Every reported figure has a source and an as-of date
- ☐ At least one SEO metric maps to a business objective

Human review

- ☐ Nothing publishes without a named reviewer
- ☐ AI-drafted output is reviewed before it goes anywhere
- ☐ Claims are sourced or removed, never softened

Keywords

- ☐ One keyword list exists that everyone works from
- ☐ Every keyword is scored with its criteria recorded
- ☐ Each keyword maps to exactly one owning page

Technical SEO

- ☐ A scored technical backlog exists with engineering owners
- ☐ A crawl runs monthly and after template releases
- ☐ Critical issues have a defined escalation route

Authority

- ☐ At least one linkable asset exists that others would reference
- ☐ Outreach is logged, including refusals and silence
- ☐ Nothing in the programme would need to be hidden

AI visibility

- ☐ A fixed prompt universe exists and is re-run monthly
- ☐ A brand fact sheet exists and public sources match it
- ☐ Factual errors trigger correction at source

Governance

- ☐ Every database has an owner and a review cadence
- ☐ The weekly rhythm is in a calendar, not in someone's memory
- ☐ Decisions are recorded together with their reasoning

Data quality

- ☐ Demand evidence is verbatim, not paraphrased
- ☐ Unknown values are recorded as unknown rather than estimated
- ☐ Where a tool supplies a metric, the tool is named

Build your SEO brain.

Modern SEO is no longer one channel, one ranking report, or one content calendar. It is a connected system that helps your brand become easier to find, understand, trust, cite, and recommend.

Everything in this playbook is designed to live in one workspace and run on a weekly rhythm. Not because systems are inherently virtuous, but because the alternative — scattered tactics, decisions reopened every month, and a strategy that exists only in one person’s head — is what makes SEO feel unpredictable when it is mostly just undocumented.

Start with the parts that compound. The baseline, so you can tell whether anything is working. The keyword list, so the same decisions stop being reopened. The internal linking pass, because it is free and fast. The brand fact sheet, because entity consistency is the cheapest AI visibility work available to anyone. Then let the weekly rhythm carry the rest.

1 WORKSPACE	10 OPERATING AREAS	5 HOURS A WEEK	30 DAYS TO RUNNING
----------------	-----------------------	-------------------	-----------------------

The same shift is happening in how buyers discover brands.

An increasing share of evaluation now happens before anyone visits a website. A buyer asks an assistant who the credible options are, and a shortlist forms without your involvement. Being in that answer is not a different discipline from SEO — it is the same discipline, measured on a new surface.

Klaym AI helps B2B brands track, analyze, and improve their visibility across AI search and answer engines including ChatGPT, Claude, Gemini, Perplexity, and Google AI Overviews — where you appear, why competitors get recommended instead, and what would change that.

BECOME THE BRAND AI RECOMMENDS

<https://klaym.ai>

Join the waitlist. · Comment "SEO" to get free access.

This playbook is an educational resource. Nothing in it guarantees rankings, traffic, leads, citations, revenue, or AI recommendations — search systems and answer engines are opaque, change without notice, and do not publish their weighting. Every framework, score, and formula here is a way of making judgement consistent and reviewable, not a measurement of objective truth. All names, figures, and examples shown in brackets are placeholders. Every output produced with the prompts in Section 13 is a draft requiring human review before it is published, sent, or acted on.